

MCAST

Data Driven Automated Algorithmic Trading

by

Gabriel Gauci Maistre

A thesis submitted in partial fulfillment for the
degree of Bachelor of Science

in the
Information and Communications Technology
Malta College of Art, Science, and Technology

June 2017

Declaration of Authorship

I, Gabriel Gauci Maistre, declare that this thesis titled, ‘Data Driven Automated Algorithmic Trading’ and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.
- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.
- Where I have consulted the published work of others, this is always clearly attributed.
- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this thesis is entirely my own work.
- I have acknowledged all main sources of help.
- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

Date:

“If past history was all there was to the game, the richest people would be librarians.”

Warren Buffett

MCAST

Abstract

Information and Communications Technology
Malta College of Art, Science, and Technology

Bachelor of Science

by Gabriel Gauci Maistre

Various existing stock market price forecasting methods were analysed in this report. Three methods were applied towards the problem making use of Technical Analysis, these were Time Series Analysis, Machine Learning, and Bayesian Statistics. Through the results of this report, it was found that the Efficient Market Hypothesis remains true, that past data does not contain enough useful information to forecast future prices and gain an advantage over the market. However, the results proved that Technical Analysis and Machine Learning could still be used to guide an investors decision. It was also found that the Random Walk Hypothesis was not necessarily true, as some stocks showed signs of auto and partial correlation. A common application of technical analysis was demonstrated and shown to produce limited useful information in beating the market. Based on the findings, a number of automated trading algorithms were developed using machine learning and backtested to determine their effectiveness.

Acknowledgements

I would like to express my special thanks of gratitude to my supervisor, Alan Gatt, for the patient guidance, encouragement, and advice he has provided throughout my time as his student.

I would also like to thank Luke Vella Critien, for guiding me towards the right path in the early stages of my research and for recommending Alan Gatt as my tutor.

My gratitude is also extended to Emma Galea and Miguel Attard for their valuable input while carrying out my research.

Completing this work would have been all the more difficult were it not for the support and friendship provided by the other members of the Malta College of Arts, Sciences, and Technology, and the institute of Information and Technology. I am indebted to them for their help.

Finally, I would like to thank my family who have supported me all throughout the final year of my bachelor's.

Contents

Declaration of Authorship	i
Abstract	iii
Acknowledgements	iv
List of Figures	ix
List of Tables	xiv
Abbreviations	xv
1 Introduction	1
1.1 Project goals and scope	1
1.1.1 Efficient Market Hypothesis	1
1.1.2 Random Walk Hypothesis	2
2 Background Theory	3
2.1 Time Series Analysis	3
2.1.1 Auto Regressive Moving Average (ARMA)	4
2.1.2 Auto Regressive Integrated Moving Average (ARIMA)	5
2.2 Statistical Machine Learning	6
2.2.1 Classification	6
2.2.2 Support Vector Machines	7
2.2.3 Regression	7
2.2.4 Decision Trees	8
2.3 Bayesian Statistics	8
2.3.1 Markov Chain Monte Carlo (MCMC)	9
3 Experimental Setup	11
3.1 Stock Selection	11
3.2 Time Series Analysis	11
3.2.1 Random Walk	12
3.2.2 Ordinary Least Squares (OLS)	12
3.2.3 Auto Regressive (AR)	12

3.2.4	Moving Average (MA)	12
3.2.5	Auto Regressive Moving Average (ARMA)	12
3.2.6	Auto Regressive Integrated Moving Average (ARIMA)	13
3.3	Machine Learning	13
3.3.1	Classification	13
3.3.1.1	Decision Tree	13
3.3.1.2	Boosted Decision Tree	14
3.3.1.3	Support Vector Machine (SVM)	14
3.3.1.4	Random Forest	14
3.3.1.5	K-Nearest Neighbour	15
3.3.1.6	Logistic Regression	15
3.3.1.7	Bernoulli Naive Bayes	15
3.3.1.8	Gaussian Naive Bayes	16
3.3.1.9	Neural Network	16
3.3.1.10	Stochastic Gradient Descent	16
3.3.2	Regression	17
3.3.2.1	Decision Tree	17
3.3.2.2	Boosted Decision Tree	18
3.3.2.3	Random Forest	18
3.3.2.4	Linear Regression	18
3.3.2.5	Neural Network	18
3.3.2.6	Stochastic Gradient Descent	18
3.4	Bayesian Statistics	19
3.4.1	No-U-Turn Sampler (NUTS)	19
3.4.2	Metropolis-Hastings	19
3.5	Strategy	19
3.5.1	Classification	19
3.5.2	Regression	20
4	Research Findings	21
4.1	Time Series Analysis	22
4.1.1	Random Walk	22
4.1.2	Ordinary Least Squares (OLS)	23
4.1.3	Auto Regressive (AR)	25
4.1.4	Moving Average (MA)	27
4.1.5	Auto Regressive Moving Average (ARMA)	30
4.1.6	Auto Regressive Integrated Moving Average (ARIMA)	32
4.2	Machine Learning	35
4.2.1	Classification	35
4.2.1.1	Decision Tree	35
4.2.1.2	Boosted Decision Tree	35
4.2.1.3	Support Vector Machine (SVM)	35
4.2.1.4	Random Forest	36
4.2.1.5	K-Nearest Neighbour	36
4.2.1.6	Logistic Regression	36
4.2.1.7	Gaussian Naive Bayes	37
4.2.1.8	Bernoulli Naive Bayes	37

4.2.1.9	Neural Network	37
4.2.1.10	Stochastic Gradient Descent	38
4.2.2	Regression	38
4.2.2.1	Decision Tree	38
4.2.2.2	Boosted Decision Tree	39
4.2.2.3	K-Nearest Neighbour	41
4.2.2.4	Random Forest	42
4.2.2.5	Linear Regression	44
4.2.2.6	Neural Network	45
4.2.2.7	Stochastic Gradient Descent	47
4.3	Bayesian Statistics	48
4.3.1	Metropolis-Hastings	48
4.3.2	No-U-Turn Sampler (NUTS)	50
4.4	Strategy	51
4.4.1	Classification	51
4.4.2	Regression	54
5	Discussion and Suggestions for Future Research	58
5.1	Time Series Analysis	58
5.1.1	Random Walk	58
5.1.2	Ordinary Least Squares (OLS)	58
5.1.3	Auto Regressive (AR)	59
5.1.4	Moving Average (MA)	59
5.1.5	Auto Regressive Moving Average (ARMA)	59
5.1.6	Auto Regressive Integrated Moving Average (ARIMA)	59
5.2	Machine Learning	60
5.2.1	Classification	60
5.2.1.1	Decision Tree	60
5.2.1.2	Boosted Decision Tree	60
5.2.1.3	Support Vector Machine (SVM)	60
5.2.1.4	Random Forest	60
5.2.1.5	K-Nearest Neighbour	60
5.2.1.6	Logistic Regression	60
5.2.1.7	Bernoulli Naive Bayes	61
5.2.1.8	Gaussian Naive Bayes	61
5.2.1.9	Neural Network	61
5.2.1.10	Stochastic Gradient Descent	61
5.2.2	Regression	61
5.2.2.1	Decision Tree	61
5.2.2.2	Boosted Decision Tree	61
5.2.2.3	K-Nearest Neighbour	62
5.2.2.4	Random Forest	62
5.2.2.5	Linear Regression	62
5.2.2.6	Neural Network	62
5.2.2.7	Stochastic Gradient Descent	62
5.3	Bayesian Statistics	63
5.3.0.1	Metropolis-Hastings	63

5.3.0.2	No-U-Turn-Sampler (NUTS)	63
5.4	Strategy	63
5.4.1	Classification	63
5.4.2	Regression	63
5.5	Future Research	63
6	Conclusion	65
A	An Appendix	66
A.1	Time Series Analysis	66
A.1.1	Random Walk	66
A.1.2	Ordinary Least Squares (OLS)	70
A.1.3	Ordinary Least Squares (OLS)	76
A.1.4	Auto Regressive (AR)	82
A.1.5	Moving Average (MA)	92
A.1.6	Auto Regressive Moving Average (ARMA)	102
A.1.7	Auto Regressive Integrated Moving Average (ARIMA)	112
A.2	Machine Learning	122
A.2.0.1	Decision Tree	122
A.2.0.2	Boosted Decision Tree	128
A.2.0.3	K-Nearest Neighbour	134
A.2.0.4	Random Forest	140
A.2.0.5	Linear Regression	146
A.2.0.6	Neural Network	152
A.2.0.7	Stochastic Gradient Descent	158
A.3	Bayesian Statistics	164
A.3.1	Metropolis-Hastings	164
A.3.2	No-U-Turn Sampler (NUTS)	170
Bibliography		177

List of Figures

4.1	Basket of stocks	21
4.2	HRG time series analysis	22
4.3	HRG histogram of returns	23
4.4	MSFT OLS in-sample prediction	24
4.5	100 day MSFT OLS in-sample forecast	24
4.6	MSFT AR time series analysis	25
4.7	MSFT AR histogram of returns	26
4.8	MSFT AR in-sample returns prediction	26
4.9	100 day MSFT AR in-sample returns forecast	27
4.10	HRG MA time series analysis	28
4.11	HRG MA histogram of returns	28
4.12	HRG MA in-sample returns prediction	29
4.13	100 day HRG MA in-sample returns forecast	29
4.14	CDE ARMA time series analysis	30
4.15	CDE ARMA histogram of returns	31
4.16	CDE ARMA in-sample returns prediction	31
4.17	100 day CDE ARMA in-sample returns forecast	32
4.18	MSFT ARIMA time series analysis	33
4.19	MSFT ARIMA histogram of returns	33
4.20	MSFT ARIMA in-sample returns prediction	34
4.21	100 day MSFT ARIMA in-sample returns forecast	34
4.22	MSFT Decision Trees in-sample prediction	38
4.23	100 day MSFT Decision Trees out-of-sample forecast	39
4.24	MSFT Boosted Decision Trees in-sample prediction	40
4.25	100 day MSFT Boosted Decision Trees out-of-sample forecast	40
4.26	MSFT K-Nearest Neighbour in-sample prediction	41
4.27	100 day MSFT K-Nearest Neighbour out-of-sample forecast	42
4.28	MSFT Random Forest in-sample prediction	43
4.29	100 day MSFT Random Forest out-of-sample forecast	43
4.30	NAV B Linear Regression in-sample prediction	44
4.31	100 day NAV B Linear Regression out-of-sample forecast	45
4.32	CDE Neural Network in-sample prediction	46
4.33	100 day CDE Neural Network out-of-sample forecast	46
4.34	CDE SGD in-sample prediction	47
4.35	100 day CDE SGD out-of-sample forecast	48
4.36	MSFT Metropolis-Hastings in-sample prediction	49
4.37	100 day MSFT Metropolis-Hastings out-of-sample forecast	49

4.38	HL NUTS in-sample prediction	50
4.39	100 day HL NUTS out-of-sample forecast	51
4.40	Machine Learning Classifier strategy with only upwards forecasts	52
4.41	Machine Learning Classifier strategy with upwards and downwards forecasts	53
4.42	Machine Learning Classifier strategy with stop loss	54
4.43	Machine Learning Regression strategy with only upwards forecasts	55
4.44	Machine Learning Regression strategy with upwards and downwards forecasts	56
4.45	Machine Learning Regression strategy with stop loss	57
A.1	MSFT time series analysis	66
A.2	MSFT histogram of returns	67
A.3	CDE time series analysis	67
A.4	CDE histogram of returns	68
A.5	NAV B time series analysis	68
A.6	NAV B histogram of returns	69
A.7	HL time series analysis	69
A.8	HL histogram of returns	70
A.9	CDE OLS in-sample prediction	71
A.10	100 Day CDE OLS out of sample forecast	71
A.11	NAV B OLS in-sample prediction	72
A.12	100 day NAV B OLS in-sample forecast	73
A.13	HRG OLS in-sample prediction	74
A.14	100 day HRG OLS in-sample forecast	74
A.15	HL OLS in-sample prediction	75
A.16	100 day HL OLS in-sample forecast	76
A.17	CDE OLS in-sample prediction	77
A.18	100 Day CDE OLS out of sample forecast	77
A.19	NAV B OLS in-sample prediction	78
A.20	100 day NAV B OLS in-sample forecast	79
A.21	HRG OLS in-sample prediction	80
A.22	100 day HRG OLS in-sample forecast	80
A.23	HL OLS in-sample prediction	81
A.24	100 day HL OLS in-sample forecast	82
A.25	CDE AR time series analysis	83
A.26	CDE AR histogram of returns	83
A.27	CDE AR in-sample returns prediction	84
A.28	100 day CDE AR in-sample returns forecast	84
A.29	NAV B AR time series analysis	85
A.30	NAV B AR histogram of returns	86
A.31	NAV B AR in-sample returns prediction	86
A.32	100 day NAV B AR in-sample returns forecast	87
A.33	HRG AR time series analysis	88
A.34	HRG AR histogram of returns	88
A.35	HL AR in-sample returns prediction	89
A.36	100 day HRG AR in-sample returns forecast	89
A.37	HL AR time series analysis	90

A.38 HL AR histogram of returns	91
A.39 HL AR in-sample returns prediction	91
A.40 100 day HL AR in-sample returns forecast	92
A.41 CDE MA time series analysis	93
A.42 CDE MA histogram of returns	93
A.43 CDE MA in-sample returns prediction	94
A.44 100 day CDE MA in-sample returns forecast	94
A.45 NAVB MA time series analysis	95
A.46 NAVB MA histogram of returns	96
A.47 NAVB MA in-sample returns prediction	96
A.48 100 day NAVB MA in-sample returns forecast	97
A.49 HRG MA time series analysis	98
A.50 HRG MA histogram of returns	98
A.51 HRG MA in-sample returns prediction	99
A.52 100 day HRG MA in-sample returns forecast	99
A.53 HL MA time series analysis	100
A.54 HL MA histogram of returns	101
A.55 HL MA in-sample returns prediction	101
A.56 100 day HL MA in-sample returns forecast	102
A.57 MSFT ARMA time series analysis	103
A.58 MSFT ARMA histogram of returns	103
A.59 MSFT ARMA in-sample returns prediction	104
A.60 100 day MSFT ARMA in-sample returns forecast	104
A.61 NAVB ARMA time series analysis	105
A.62 NAVB ARMA histogram of returns	106
A.63 NAVB ARMA in-sample returns prediction	106
A.64 100 day NAVB ARMA in-sample returns forecast	107
A.65 HRG ARMA time series analysis	108
A.66 HRG ARMA histogram of returns	108
A.67 HRG ARMA in-sample returns prediction	109
A.68 100 day HRG ARMA in-sample returns forecast	109
A.69 HL ARMA time series analysis	110
A.70 HL ARMA histogram of returns	111
A.71 HL ARMA in-sample returns prediction	111
A.72 100 day HL ARMA in-sample returns forecast	112
A.73 CDE ARIMA time series analysis	113
A.74 CDE ARIMA histogram of returns	113
A.75 CDE ARIMA in-sample returns prediction	114
A.76 100 day CDE ARIMA in-sample returns forecast	114
A.77 NAVB ARIMA time series analysis	115
A.78 NAVB ARIMA histogram of returns	116
A.79 NAVB ARIMA in-sample returns prediction	116
A.80 100 day NAVB ARIMA in-sample returns forecast	117
A.81 HRG ARIMA time series analysis	118
A.82 HRG ARIMA histogram of returns	118
A.83 HRG ARIMA in-sample returns prediction	119
A.84 100 day HRG ARIMA in-sample returns forecast	119

A.85 HL ARIMA time series analysis	120
A.86 HL ARIMA histogram of returns	121
A.87 HL ARIMA in-sample returns prediction	121
A.88 100 day HL ARIMA in-sample returns forecast	122
A.89 CDE Decision Trees in-sample prediction	123
A.90 100 day CDE Decision Trees out-of-sample forecast	123
A.91 NAVB Decision Trees in-sample prediction	124
A.92 100 day NAVB Decision Trees out-of-sample forecast	125
A.93 HRG Decision Trees in-sample prediction	126
A.94 100 day HRG Decision Trees out-of-sample forecast	126
A.95 HL Decision Trees in-sample prediction	127
A.96 100 day HL Decision Trees out-of-sample forecast	128
A.97 CDE Boosted Decision Trees in-sample prediction	129
A.98 100 day CDE Boosted Decision Trees out-of-sample forecast	129
A.99 NAVB Boosted Decision Trees in-sample prediction	130
A.100 100 day NAVB Boosted Decision Trees out-of-sample forecast	131
A.101 HRG Boosted Decision Trees in-sample prediction	132
A.102 100 day HRG Boosted Decision Trees out-of-sample forecast	132
A.103 HL Boosted Decision Trees in-sample prediction	133
A.104 100 day HL Boosted Decision Trees out-of-sample forecast	134
A.105 CDE K-Nearest Neighbour in-sample prediction	135
A.106 100 day CDE K-Nearest Neighbour out-of-sample forecast	135
A.107 NAVB K-Nearest Neighbour in-sample prediction	136
A.108 100 day NAVB K-Nearest Neighbour out-of-sample forecast	137
A.109 HRG K-Nearest Neighbour in-sample prediction	138
A.110 100 day HRG K-Nearest Neighbour out-of-sample forecast	138
A.111 HL K-Nearest Neighbour in-sample prediction	139
A.112 100 day HL K-Nearest Neighbour out-of-sample forecast	140
A.113 CDE Random Forest in-sample prediction	141
A.114 100 day CDE Random Forest out-of-sample forecast	141
A.115 NAVB Random Forest in-sample prediction	142
A.116 100 day NAVB Random Forest out-of-sample forecast	143
A.117 HRG Random Forest in-sample prediction	144
A.118 100 day HRG Random Forest out-of-sample forecast	144
A.119 HL Random Forest in-sample prediction	145
A.120 100 day HL Random Forest out-of-sample forecast	146
A.121 MSFT Linear Regression in-sample prediction	147
A.122 100 day MSFT Linear Regression out-of-sample forecast	147
A.123 CDE Linear Regression in-sample prediction	148
A.124 100 day CDE Linear Regression out-of-sample forecast	149
A.125 HRG Linear Regression in-sample prediction	150
A.126 100 day HRG Linear Regression out-of-sample forecast	150
A.127 HL Linear Regression in-sample prediction	151
A.128 100 day HL Linear Regression out-of-sample forecast	152
A.129 MSFT Neural Network in-sample prediction	153
A.130 100 day MSFT Neural Network out-of-sample forecast	153
A.131 NAVB Neural Network in-sample prediction	154

A.132	100 day NAVB Neural Network out-of-sample forecast	155
A.133	HRG Neural Network in-sample prediction	156
A.134	100 day HRG Neural Network out-of-sample forecast	156
A.135	HL Neural Network in-sample prediction	157
A.136	100 day HL Neural Network out-of-sample forecast	158
A.137	MSFT SGD in-sample prediction	159
A.138	100 day MSFT SGD out-of-sample forecast	159
A.139	NAV B SGD in-sample prediction	160
A.140	100 day NAVB SGD out-of-sample forecast	161
A.141	HRG SGD in-sample prediction	162
A.142	100 day HRG SGD out-of-sample forecast	162
A.143	HL SGD in-sample prediction	163
A.144	100 day HL SGD out-of-sample forecast	164
A.145	CDE Metropolis-Hastings in-sample prediction	165
A.146	100 day CDE Metropolis-Hastings out-of-sample forecast	165
A.147	NAV B Metropolis-Hastings in-sample prediction	166
A.148	100 day NAVB Metropolis-Hastings out-of-sample forecast	167
A.149	HRG Metropolis-Hastings in-sample prediction	168
A.150	100 day HRG Metropolis-Hastings out-of-sample forecast	168
A.151	HL Metropolis-Hastings in-sample prediction	169
A.152	100 day HL Metropolis-Hastings out-of-sample forecast	170
A.153	MSFT NUTS in-sample prediction	171
A.154	100 day MSFT NUTS out-of-sample forecast	171
A.155	CDE NUTS in-sample prediction	172
A.156	100 day CDE NUTS out-of-sample forecast	173
A.157	NAV B NUTS in-sample prediction	174
A.158	100 day NAVB NUTS out-of-sample forecast	174
A.159	HRG NUTS in-sample prediction	175
A.160	100 day HRG NUTS out-of-sample forecast	176

List of Tables

4.1	Equities Descriptive Statistics	22
4.2	Decision Tree results	35
4.3	Boosted Decision Tree results	35
4.4	Support Vector Machine results	35
4.5	Random Forest results	36
4.6	K-Nearest Neighbour results	36
4.7	Logistic Regression results	36
4.8	Gaussian Naive Bayes results	37
4.9	Bernoulli Naive Bayes results	37
4.10	Neural Network results	37
4.11	Stochastic Gradient Descent results	38
4.12	Machine Learning Classifier strategy with only upwards forecasts	51
4.13	Machine Learning Classifier strategy with upwards and downwards forecasts	52
4.14	Machine Learning Classifier strategy with stop loss	53
4.15	Machine Learning Regression strategy with only upwards forecasts	54
4.16	Machine Learning Regression strategy with upwards and downwards forecasts	55
4.17	Machine Learning Regression strategy with stop loss	56

Abbreviations

EMH	Efficient Market Hypothesis
RWH	Random Walk Hypothesis
OLS	Ordinary Least Squares
AR	Auto Regressive
MA	Moving Average
ARMA	Auto Regressive Moving Average
ARIMA	Auto Regressive Integrated Moving Average
GARCH	Generalised Auto Regressive Conditional Heteroskedasticity
EGARCH	Exponential Generalised Auto Regressive Conditional Heteroskedasticity
ADF	Augmented Dickey Fuller
SVM	Support Vector Machine
ADT	Alternating Decision Tree
SGD	Stochastic Gradient Descent
SMA	Simple Moving Average
AIC	Alkaline Information Criterion
IC	Information Criterion
NUTS	No-U-Turn Sampler
QQ	Quantile-Quantile
ETF	Exchange Traded Fund
S&P	Standards & Poor's
CAPE	Cyclically Adjusted Price-to-Earnings
ANN	Artificial Neural Network

To my parents, Keith & Christine Gauci Maistre. Without them, and their unconditional love and support, none of this would have been possible.

Chapter 1

Introduction

The stock market retains its status as a prime location for investors to invest in the market and earn a profit, however this is not always easy due to the constantly thriving and changing nature which follows the stock market. Investors are constantly presented with numerous profit potential opportunities, however without intensive planning and analysis, these opportunities could easily turn into losses. This means that it is crucial for every investor to carry out stock market analysis prior to any investment by monitoring past price movements in order to forecast future trends. Even though past data is not a clear indication of future movement, it is still proven to provide some useful insight.

1.1 Project goals and scope

This report aims to disprove two hypotheses, the Efficient Market Hypotheses (EMH), and the Random Walk Hypothesis (RWH).

1.1.1 Efficient Market Hypothesis

By definition of the EMH, it is impossible for any investor to beat the market as the markets are efficient and all stock prices always incorporate and reflect all relevant information. This essentially means that stocks are always traded at their fair value, and their prices only fluctuate when new information is released to the public. This essentially makes it impossible for any investors to exploit any inefficiencies in the market and earn a profit. Believers of the EMH claim that it is pointless to try and search for undervalued stocks, or predict market trends.

Although a vast number of academics find the EMH to be true, there are still a number of investors who have consistently beaten the market over long periods of time, something which is impossible by definition according to the EMH. Stock market crashes, such as that of the financial crisis of 1987, in which the market fell 28.3% in a single day, proving that stock prices can seriously deviate from their fair values.

It is the firm belief of proponents of the EMH that investors are better off investing in low risk index funds. It is a known fact that only a few number of active managers successfully manage to outperform passive funds.

1.1.2 Random Walk Hypothesis

According to the RWH, it is not possible to use the past movement or trend of a stock price to predict its future movement. The main idea of the RWH is all stocks take a random and unpredictable path. Strong believers of the RWH believe that it is not possible to beat the market without taking on more risk.

There are however a vast number of critics of this theory, who state that stocks do indeed follow trends over time, making it possible to carry out stock market analysis and carefully select entry and exit points for stock investments.

Chapter 2

Background Theory

The field of computational finance consists of one of the many branches of applied computer science, where the problems dealt with are in the interest of finance. Some slightly different definitions are the study of data and algorithms currently used in finance and the mathematics of computer programs that realize financial models or systems. Using computational finance in order to allocate assets in a portfolio is not at all unheard of and was first documented 1952.^[1] Markowitz first introduced the concept of portfolio selection as an exercise in mean-variance optimisation. This required more computer power than was available at the time, so he worked on useful algorithms for approximate solutions. He theorised that investors looking for low risk investments are able to form portfolios that are able to optimise or maximise expected return based on a given level of market risk, emphasising that a higher reward cannot be achieved without taking on more risk. Following his theory, it is possible to construct portfolios which are optimised to achieve the maximum possible expected return for a given level of risk.

2.1 Time Series Analysis

A time series is a series of data points which may be indexed, listed, or graphed, in a time order. Most commonly, a time series is a sequence taken at successive equally spaced points in time. Thus it is a sequence of discrete-time data. Examples of time series are daily temperature fluctuations, the number of cars on our roads, and the daily closing price of the S&P500. Brockwell et al. provide a formal description of time series as having a set of observations x_t , each one being recorded at a specific time t . A discrete-time time series is one in which the set T_0 of times at which observations are made is a discrete set, as is the case, for example, when observations are made at

fixed time intervals. Continuous time series are obtained when observations are recorded continuously over some time interval, e.g., when $T_0 = [0, 1]$.^[2]

Noteworthy "time series momentum" has likewise been achieved in equity index, currency, commodity, and bond futures for each of the 58 liquid instruments which were considered.^[3] Moskowitz et al. discover persistence in returns ranging from 1 to 12 months that partially reverses over longer horizons, consistent with sentiment theories of initial under-reaction and delayed over-reaction. Their study brought forward evidence that time series momentum strategies were found to deliver substantial abnormal returns using a diversified portfolio consisting of all asset classes having little exposure to standard asset pricing factors. Moskowitz et al. also back their claims by testing their strategies during extreme markets, which they describe to perform best in. Moskowitz et al. also studied the trading patterns of speculators and hedgers, in which they concluded that speculators tended to make profit mostly when making use of time series momentum strategies, all at the expense of hedgers.

2.1.1 Auto Regressive Moving Average (ARMA)

In the statistical analysis of time series, ARMA models provide a parsimonious description of a weakly stationary stochastic process in terms of two polynomials, one for the autoregression and the second for the moving average. The notation ARMA(p, q) refers to the model with p autoregressive terms and q moving-average terms. This model contains the AR(p) and MA(q) models,

$$x_t = c + \varepsilon_t + \sum_{i=1}^q \varphi_i X_{t-i} \theta_i \varepsilon_{t-i}$$

There is a growing demand for forecasting interest rates, as financial researchers, economists, and players in the fixed income markets seek to find the best method to get ahead of the market. A financial model was developed to forecast short-term interest rates, implicit yield on 91 day treasury bill, overnight MIBOR rate, and call money rate.^[4] Radha et al. made use of univariate models to forecast the short-term interest rates. These models include Random Walk, ARIMA, ARMA-GARCH, and ARMA-EGARCH. They selected the best performing model by considering a six year period starting from 1999. Radha et al. showed evidence that GARCH models are best suited for forecastint when applied towards time series having volatility clustering effects. It was their firm belief that ARIMA-EGARCH is the most appropriate forecasting model for these circumstances.

2.1.2 Auto Regressive Integrated Moving Average (ARIMA)

In time series analysis, an autoregressive integrated moving average (ARIMA) model is quite similar to an ARMA model as it is simply a generalisation of it. Given a time series of data X_t where t is an integer index and the X_t are real numbers, an ARMA(p, q) model is given by

$$x_t - \alpha_1 X_{t-1} - \dots - \alpha_p X_{t-p} = \varepsilon_t + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q}$$

The aforementioned model is fitted to time series data either to better understand the data or to predict future points in the series, also referred to as forecasting. ARIMA models are best suited for data showing non-stationarity, and require data to be normalised in most cases. These models provide an initial differencing step in order to reduce the stationarity found in data.

Abrosimova et al. examined the existence of weak-form efficiency in the Russian stock market for the period 1st September 1995 to 1st May 2001, while making use of daily, weekly, and monthly Russian Trading System index data.^[5] They made use of a number of approaches in order to determine the predictability of the RTS index time series. Abrosimova et al. tested the null hypothesis of a random walk model using unit root, autocorrelation, and variance ratio tests. It was evident in their results that the null hypothesis was only true for the monthly data, making them focus more on the daily and weekly data. Using linear and non-linear modelling, they fit ARIMA and GARCH models to the data in the in-sample tests over the period 1st September 1995 to 1st January 2001. Abrosimova et al. carried out forecast by selecting the models which achieved the best fit in the out-of-sample tests over the period 2nd January 2001 to 1st May 2001. Their results clearly indicated that there was no clear winner from the models, and the most accurate forecasts were only obtained in the out-of-sample tests. Although their research shed some light on the predictability of the RTS, there was not enough evidence to suggest that it would lead to a profitable trading rule, once transaction costs and risk were taken into account.

Darrat et al. set out to investigate with the use of new daily data, whether prices in the two Chinese stock exchanges (Shanghai and Shenzhen) follow a random walk process as required by market efficiency.^[6] Two different approaches were applied, the standard variance-ratio test, and a model-comparison test that compares the ex post forecasts from a naive model with those obtained from several alternative models such as ARIMA, GARCH, and ANNs. To evaluate ex post forecasts, Darrat et al. made use of several procedures including root-mean-square error (RMSE), mean absolute error (MAE), uncertainty coefficient, and encompassing tests. It was concluded that the model-comparison approach yielded results which were quite strongly rejected the RWH

in both Chinese stock markets when compared with the variance-ratio test. Darret et al. recommended the use of ANNs, as their results showed strong support for the model as a potentially useful factor for forecasting stock prices in emerging markets.

2.2 Statistical Machine Learning

A machine learning algorithm is an algorithm that is able to learn through examples from data. To understand what an algorithm is, Cormen et al. informally describe algorithms as "any well-defined computational procedures which takes some value, or set of values, as input and produce some value, or set of values, as output. An algorithm is thus a sequence of computational steps that transform the input into the output."^[7] In simple terms, it is possible to say that an algorithm is a sequence of steps which allow to solve a certain task. Similarly to a normal algorithm, a machine learning algorithm as defined formally by Tom M. Mitchell, states that "A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E."^[8]

2.2.1 Classification

Classification is a form of machine learning which is used to categorise data, such as predicting whether a tumor is benign or malignant. Murugesan et al. used this form of machine learning to categorise a stock as manipulated or non-manipulated depending on certain features in the data that characterises the particular stock.^[9] The model they selected was able to identify particular stocks which were prone to any kind of potential manipulation, including action-based, information-based, and trade-based. Murugesan et al. aided investigators to narrow down a list of stocks being potentially manipulated requiring further investigation down the line to determine the root of the manipulation. This model was of great use for markets like India, where over 5000 stocks are traded every day on the major exchanges, making it cumbersome to track all stocks for potential manipulation. Murugesan et al. carried out their research by testing their model on data from the Indian capital market. It was made evident to them from the start of their research that the data does not comply with the assumptions that govern the use of the linear classification function, which goes against academics who have earlier used discriminant analysis have used the Linear classification function. Murugesan et al. claimed that this was because academics failed to validate the assumption that governs the model in their studies. It was for these reasons that Murugesan et al. decided to use the Quadratic classification function, being the appropriate method for instances

where the data does not meet the assumptions, to categorise stocks into two categories, being manipulated and non-manipulated.

2.2.2 Support Vector Machines

Support vector machines (SVM), are a class of machine learning algorithms that have become incredibly popular in the past few years. SVMs are very similar to classifiers in the sense that they also classify data by drawing a line, called a decision boundary, to separate them. However, SVMs go a step further by calculating a vector from the data point with the smallest margin to the decision boundary. This is called a support vector. A vast majority of academics tend to predict the price of stocks in financial markets, however most models used are flawed and only focus on the accurate forecasting of the levels of the underlying stock index. There is a lack of studies examining the predictability of the direction of stock index movement. Given the notion that a prediction with little forecast error does not necessarily translate into capital gain, the authors of this research attempt to predict the direction of the S&P CNX NIFTY Market Index of the National Stock Exchange, one of the fastest growing financial exchanges in developing Asian countries.^[10] Machine learning models such as random forest and SVMs, differ widely from other models, and are making strides in predicting the financial markets. Kumar et al. tested classification models to predict the direction of the markets, by applying models such as linear discriminant analysis, logistic regression, ANNs, random forest, and SVM. Their evidence shows that SVMs outperform the other classification methods in terms of predicting the direction of the stock market movement, and that the random forest model outperforms other models such as ANNs, discriminant analysis, and logistic regression.

2.2.3 Regression

Regression is another form of machine learning which is used to predict data, allowing researchers to correctly identify the relationship that lies between the independent dependent variables. These models make it easier to understand how the typical value of the dependent variable changes when any one of the independent variables is varied, while the other independent variables are held fixed. Kakushadze et al. provide a systematic quantitative framework in what is intended to be a pedagogical fashion for discussing mean-reversion and optimisation.^[11] In their paper, they start off their research with pair trading and add complexity by following the sequence mean-reversion via demeaning, regression, weighted regression, (constrained) optimisation, factor models. Using this approach, Kakushadze et al. studied mean reversion models in further

detail, going into the common pitfalls found in practical applications, including the difference between maximising the Sharpe ratio and minimising an objective function when trading costs are included. Kakushadze et al. also discuss explicit algorithms for optimization with linear costs, constraints and bounds, and also illustrate their discussion on an explicit intraday mean-reversion alpha.

2.2.4 Decision Trees

Decision trees are a class of machine learning models which make use of a decision tree as a predictive model which creates a map of observations about an item, represents them in the form branches, to conclude about the item's target value represented in the form of leaves. There are two types of decision trees, classification trees, which its target variable can take a finite set of values, and regression trees, in which the target variable can take continuous values. In classification trees, each leaf represents a class label, and each branch represents conjunctions of features that lead to those class labels. In their research, Creamer et al. developed an automated trading algorithm making use of multiple stocks relying on a layered structure consisting of a machine learning algorithm, an online learning utility, and a risk management overlay.[\[12\]](#) The machine learning algorithm which they made use of was an Alternating Decision Tree (ADT) implemented with Logitboost. Their algorithm was able to select the best combination of rules derived from well known technical analysis indicators, and the best parameters of the indicators in question. Additionally, their online learning layer was also able to combine the output of several ADTs, suggesting a short or long position. Finally, the risk management layer in which they implemented, was able to validate the trading signal once it exceeds a specified non-zero threshold and limit the application of their trading strategy when it is not profitable. They tested the expert weighting algorithm with data of 100 randomly selected companies of the S&P 500 index during the period 20032005. They found that their algorithm generated abnormal returns during the test period. Their experiments show that the boosting approach was able to improve the predictive capacity when indicators were combined and aggregated as a single predictor. Furthermore, their results indicated that the combination of indicators of different stocks were adequate in order to reduce the use of computational resources, while still maintaining an adequate predictive capacity.

2.3 Bayesian Statistics

Bayesian Statistics, a form of probabilistic programming, describes probabilistic models and then performs inference in those models. Probabilistic reasoning is a foundational

technology of machine learning and has been used by companies such as Google, Amazon, and Microsoft. Probabilistic reasoning has been used for predicting stock prices, recommending movies, diagnosing computers, detecting cyber intrusions, and image detection. Gelman et al. defines bayesian inference as the process of fitting a probability model to a set of data and summarising the result by a probability distribution on the parameters of the model and on unobserved quantities such as predictions for new observations.[\[13\]](#)

2.3.1 Markov Chain Monte Carlo (MCMC)

MCMC models are a family of algorithms for sampling from a probability distribution based on constructing a Markov Chain that has the desired distribution of its equilibrium distribution. The state of the chain after a number of steps is then used as a sample of the desired distribution. The quality of the sample improves as a function of the number of steps. At the time of conducting their research, bond yields were well below and stock market valuations were well above their historical average.[\[14\]](#) Blanchett et al. claim this to be a unique opportunity as this has never happened before in the United States. It is also known for portfolio returns to have an outsize impact on retirement income strategies in the first decade. When conducting their research, academics do not normally incorporate market valuations in their study. This is of course a flawed approach as a big piece of the puzzle is being left out. Blanchett et al. set out to simulate how a low return environment for both stocks and bonds would affect retirees. In their study, they used current valuations to predict the impact on retirement portfolios. Blanchett et al. used an autoregressive model to estimate bond returns, making use of an initial bond yield value where yields drift in the future. CAPE ratios were also used in order to estimate the market valuation and predict short-run stock performance. It was evident from the number of simulations that they ran, that the risk factor of withdrawal strategies was directly affected by the initial bond yield and the CAPE value at retirement. The relative impact was also found to vary based on the portfolio stock allocation. Blanchett et al. found the probability of success to sit at approximately 48% when having a 40% stock allocation with a 4% initial withdrawal rate over a 30 year period by using the valuation measures from April 15, 2013. They concluded that the success rate in question is far lower than previous studies, which points to having serious implications on the likelihood of success for retirees today, as well as how much those near retirement may need to save to ensure a successful retirement.

In their paper, Hoffman et al. introduced the NUTS model, an improvement to the HMC model. The NUTS model is a model of MCMC family of algorithms, in which its main focus is to avoid the random walk behaviour and sensitivity to correlated parameters

affecting many other MCMC models. Hoffman et al. achieved this by taking a series of steps informed by first-order gradient information.[\[15\]](#) In their study, they claim HMC to be flawed as its performance is highly sensitive to two user-specified parameters, the step size, and the desired number of steps. NUTS is an improvement from HMC as it eliminates the need for the researcher to set the number of steps, and works by building a set of likely candidate points which span a wide swath of the target distribution, stopping automatically when it starts to double back and retrace its steps. They achieved all of this by making use of a recursive algorithm.

Chapter 3

Experimental Setup

A data set containing end of day stock prices, dividends, and splits for 3,000 US companies, curated by the Quandl community[16], and released into the public domain, was used. The date column in the CSV file was loaded into memory, and said column was converted to a date data type. The data frame was then sorted using the date column, starting from the oldest date, ending with the latest. The date column was also set to the index. The data frame was split into two, the training data set consisting of 80%, and the test data set consisting of 20% of the original data frame.

3.1 Stock Selection

The pairwise Pearson correlation coefficient of all the columns in the data frame was computed and stored in a new data frame. A list of stock pairs with low correlation were extracted.

3.2 Time Series Analysis

The selected stocks was extracted from the data set and stored in a data frame. The log returns of the stocks were calculated by calculating the logarithm of the stock's adjusted close price divided by the following day's adjusted close price. The resulting values from the calculation were then stored in a new column in the data frame and all infinite values were dropped from the series.

3.2.1 Random Walk

The first difference of the stocks were calculated and stored in a new column in the data frame and all infinite valued were dropped.

3.2.2 Ordinary Least Squares (OLS)

The series was fitted to an OLS model. The 15 day SMA was used as a nobs x k array where nobs is the number of observations and k is the number of regressors, to train the model, while the adjusted close price of the stock was used as the dependent variable for the model to predict. The mean absolute error, mean squared error, median absolute error, and r2 score were used as metrics in order to rank the performance of the model's prediction capabilities in in-sample testing.

3.2.3 Auto Regressive (AR)

The series was fitted to an AR(p) model with a maximum lag value of 30. The IC was used to fit the AR model in order to select the optimal lag length. No constants were passed when fitting the AR model to the series. The optimal lag for the fit of the AR model was then calculated using the same parameters passed when fitting the AR model.

3.2.4 Moving Average (MA)

The series was fitted to an MA(p, q) model with an order selected based on the lowest AIC. A maximum lag of 30 was once again passed to the MA model with no constant. The exact loglikelihood for the fit of the MA model was maximized via the Kalman Filter.

3.2.5 Auto Regressive Moving Average (ARMA)

The series was fitted to an ARMA(p, q, r) model with an order selected based on the lowest AIC. No constants were passed to the ARMA model. The exact loglikelihood for the fit of the ARMA model was maximized via the Kalman Filter.

3.2.6 Auto Regressive Integrated Moving Average (ARIMA)

The series was fitted to an ARIMA(p, q, r) model with an order selected based on the lowest AIC. No constants were passed to the ARIMA model. The exact loglikelihood for the fit of the ARIMA model was maximized via the Kalman Filter.

3.3 Machine Learning

The data set was split for training and testing purposes when fitting and predicting data. The training data set consisted of 80% of the whole data set, while the test data set consisted of 20%.

3.3.1 Classification

2, 3, 4, 5, and 6 day SMAs were used as the features to train the models, while a binary value used to determine whether the current day's adjusted close has risen or not from the previous day was used as the target valued for the model to predict. In-sample testing was carried out using the models predict function, passing the test data set's features in order to predict the output. The classification report and confusion matrix were used as metrics in order to rank the performance of the model's prediction capabilities in in-sample testing. An algorithm was developed for out-of-sample testing. The algorithm iterates for a number of n steps, increasing the index by 1 each step, forecasting the next day's outcome, and calculating the SMAs.

3.3.1.1 Decision Tree

The decision tree classifier was fit using the Gini impurity criterion to measure the quality of the split. The model was given the liberty to select the best strategy in order to split the tree at each node. The maximum allowed depth of the tree was left unrestricted, which was the same as for the maximum leaf nodes. A threshold of 1e-7 was used to terminate the tree growth to determine if a node is a leaf, if the impurity of a node is below the threshold, the node is a leaf. The minimum number of samples required to be at a leaf node was set to 1, while the minimum number of samples required to split an internal node was set to 2. The data fitted to the model was not pre-sorted, and the random number generator used by the model was that of Numpy's RandomState.

3.3.1.2 Boosted Decision Tree

The boosted decision tree was fit using the 'SAME.R' real algorithm which converges at a faster rate, achieving a lower test error with fewer boosting iterations. The maximum number of estimators at which boosting is terminated was set to 50; in case of perfect fit, the learning procedure is stopped early. The learning rate which is the contribution of each classifier of the model was shrunk by 1. The best estimator used to fit the data to the model was a Decision Tree Classifier, and the random number generator used by the model was that of Numpy's RandomState.

3.3.1.3 Support Vector Machine (SVM)

The C-Support Vector Classification implementation was used for the SVM model with a 0 penalty parameter of the error term. A kernal type of 'rbf' was used when fitted the model to the data, along with a polynomial kernel function degree of 3. The gama 'rbf' Kernel coefficient was calculated by dividing the number of features by 1, while no probability estimates were used when fitting. A shrinking heuristic was used and a tolerance of 1e-3 was used for stopping criterion. A cache size of 200MB was used for the kernel when fitting the model, and all classes were assigned a weight of one. Verbose output was not used, and the random number generator used by the model was that of Numpy's RandomState. No limits were set on the iterations within the solver, and a one-vs-rest decision function of shape (number of samples, number of classes) was returned.

3.3.1.4 Random Forest

The random forest was fit with bootstrap samples when building the trees, while all the weights associated were set to 1. The 'gini' function to measure the quality of a split, and no maximum depth of the tree was set, allowing the nodes to expand until all leaves are pure or until all leaves contain less than the minimum split samples. The number of features to consider when looking for the best split was the square root of the number of passed, and no limit on the maximum leaf nodes for growing trees was set. A threshold of 1e-7 was used to terminate the tree growth to determine if a node is a leaf, if the impurity of a node is below the threshold, the node is a leaf. The minimum number of samples required to be at a leaf node was set to 1, and the minimum number of samples required to split an internal node was set to 2. The minimum weighted fraction of the sum total of weights (of all the input samples) required to be at a leaf node was set to 0, and the number of trees in the forest was set to 10. The number of jobs to use for the

computation was set to 1, making use of only 1 CPU core, and out-of-bag samples to estimate the generalization accuracy were not used. The verbosity of the tree building process was not controlled, and the random number generator used by the model was that of Numpy’s RandomState. The model was built using a cold start by not making use of the previous call to fit and add more estimators to the ensemble, meaning a whole new forest was fit instead.

3.3.1.5 K-Nearest Neighbour

The k-nearest neighbour was fit with a number of 5 neighbors to use for k-neighbours queries, with a uniform weight making all points in each neighbourhood weighted equally when carrying out predictions. The model was left at liberty to select the most appropriate algorithm based on the values passed when fitting the model to the data. The lead size of the model was set to 30, and a minkowski distance metric with p equal to 2 which equivalent to the standard Euclidean metric. The number of jobs to use for the computation was set to 1, making use of only 1 CPU core.

3.3.1.6 Logistic Regression

The logistic regression was fit with an inverse of regularization of strength 0 specifying stronger regularisation, while all the weights associated were set to 1. Dual formulation was used in conjunction with the l2 penalty with a liblinear solver. A constant (bias or intercept) was added to the decision function, and the intercept scalar was set to 1. The maximum number of iterations taken for the solvers to converge were set to 100, while the multiclass option used was ovr, thus fitting binary problem for each label. The number of jobs to use for the computation was set to 1, making use of only 1 CPU core, and the random number generator used by the model was that of Numpy’s RandomState. The tolerance for stopping criteria was set to 0.0001, and verbosity was used with a value of 0 for the liblinear solver. The model was built using a cold start by not making use of the previous call to fit and add more estimators to the ensemble, meaning a whole new forest was fit instead.

3.3.1.7 Bernoulli Naive Bayes

The model was fit with an additive (Laplace/Lidstone) smoothing parameter of 0 for no smoothing, and the threshold for binarising (mapping to booleans) of sample features was set to 0, and was presumed to already consist of binary vectors. The model was set

to learn class prior probabilities, which means that the priors were adjusted according to the data.

3.3.1.8 Gaussian Naive Bayes

The model was fit to the data with no previous prior probabilities of the classes, thus the priors were not adjusted according to the data.

3.3.1.9 Neural Network

The model was fit to the data with 3 hidden layer of 100 neurons each using the rectified linear unit activation function for the hidden layer. The L2 penalty parameter used, which is the regularisation term, was set to 0.0001, and the size of minibatches for stochastic optimisers was set to the minimum of the two arguments being the number 200 and the number of samples. The solver used for weight optimisation was the 'adam' solver, which refers to a stochastic gradient-based optimiser. Since the 'adam' solver was used, the exponential decay rate for estimates of first moment vector was set to 0.9, while the second moment vector was set to 0.999. Early stopping was not used to terminate training when validation score is not improving whilst fitting the data to the model, and the value for numerical stability in the 'adam' solver was set to 1e-08. A constant learning rate was used for weight updates in conjunction with the initial learning rate, which controls the step-size in updating the weights, and was set to 0.001. The maximum number of iterations which the solver iterates until convergence was set to 200, and the momentum for gradient descent update was set to 0.9. Nesterov's momentum was used, and the exponent for the inverse scaling learning rate, which is used in updating the effective learning rate, was set to 0.5. A random number generator was not used, and the samples were shuffled in each iteration. The tolerance for optimisation was set to 0.0001 when the loss or score is not improving by at least the tolerance score set for two consecutive iterations in which convergence is considered to be reached and training is stopped. The validation fraction, which is the portion of training data to set aside as validation set for early stopping, was set to 0.1. The model was built using a cold start by not making use of the previous call to fit as initialisation, meaning the previous solution was erased.

3.3.1.10 Stochastic Gradient Descent

The model was fit to the data using the ordinary least squares fit squared epsilon insensitive loss function, which ignores errors less than epsilon and is linear past that; this is

the loss function used in SVR. The penalty used, also referred to regularisation term, was 'l2' which is the standard regularizer for linear SVM models, and the constant used to multiply the regularisation term was 0.0001. The elastic net mixing parameter was set to 0.15, while the fit intercept was estimated, meaning the data was not assumed to be already centered. A total of 5 passes over the training data, also known as epochs, were used, and the training data was shuffled after each epoch. No seed from the pseudo random number generated was used while shuffling the data, and a level 0 verbosity was used. The epsilon threshold in the epsilon-insensitive loss functions was set to 0.1, and any differences between the current prediction and the correct label were ignored if they were less than the threshold. The learning rate schedule used was 'invscaling', and the initial learning rate was set to 0.01. The exponent for inverse scaling learning rate was set to 0.25, and the model was built using a cold start by not making use of the previous call to fit as initialisation. The SGD weights of the model were not averaged.

3.3.2 Regression

15 and 50 day SMAs were used as the features to train the models, while the adjusted close price of the stock was used as the target value for the model to predict. In-sample testing was carried out using the models predict function, passing the test data set's features in order to predict the output. The mean absolute error, mean squared error, median absolute error, and R^2 (coefficient of determination) score were used as metrics in order to rank the performance of the model's prediction capabilities in in-sample testing. An algorithm was developed for out-of-sample testing. The algorithm iterates for a number of n steps, increasing the index by 1 each step, forecasting the next day's outcome, and calculating the SMAs.

3.3.2.1 Decision Tree

The decision tree model was fit with a mean squared error, which is equal to variance reduction as feature selection criterion, and mae for the mean absolute error. The maximum allowed depth of the tree was left unrestricted, which was the same as for the maximum leaf nodes. A threshold of 1e-7 was used to terminate the tree growth to determine if a node is a leaf, if the impurity of a node is below the threshold, the node is a leaf. The minimum number of samples required to be at a leaf node was set to 1, while the minimum number of samples required to split an internal node was set to 2. The data fitted to the model was not pre-sorted, and the random number generator used by the model was that of Numpy's RandomState. The model was given the liberty to select the best strategy in order to split the tree at each node.

3.3.2.2 Boosted Decision Tree

The boosted decision tree was fit with with a decision tree regressor as the base estimator from which the boosted ensemble is built. The maximum number of estimators at which the boosting is terminated was set to 50, and the learning rate which shrinks the contribution of each regressor was set to 1.0. The loss function used when updating the weights after each boosting iteration was set to linear, and the random number generator used by the model was that of Numpy's RandomState.

3.3.2.3 Random Forest

The random forest was fit with a mean absolute error, and bootstrap samples were used when building trees. The maximum features to consider when looking for the best split were left to the model to select the best number, and the number of jobs to run in parallel for both the fitting of the model and its predictions. The model was built using a cold start by not making use of the previous call to fit and add more estimators to the ensemble, meaning a whole new forest was fit instead. A threshold of 1e-7 was used to terminate the tree growth to determine if a node is a leaf, if the impurity of a node is below the threshold, the node is a leaf. The minimum number of samples required to be at a leaf node was set to 1, while the minimum number of samples required to split an internal node was set to 2.

3.3.2.4 Linear Regression

The linear regression was fit with no normalised regressors, note that this makes the hyperparameters learnt more robust and almost independent of the number of samples. The number of jobs to use for the computation was set to 1, making use of only 1 CPU core. The intercept of the model was calculated which centres the data being fit to the model.

3.3.2.5 Neural Network

The model was fit to the data with the same paramaters as the Neural Network classifier.

3.3.2.6 Stochastic Gradient Descent

The model was fit to the data with the same paramaters as the SGD classifier.

3.4 Bayesian Statistics

The last 500 rows of the selected stocks were extracted from the data set and stored into a data frame. The log returns were calculated by dividing each day's adjusted close with the adjusted close of the following day, in logarithmic form. The resulting values from the said calculation were then stored in a new column in the data frame and all infinite values were dropped from the series. The sharpe ratio of the original and predicted price returns was calculated to serve as an accuracy score in the in-sample tests.

3.4.1 No-U-Turn Sampler (NUTS)

The model returns were modeled with a Student-t distribution with an unknown degrees of freedom paramater, and a scale paramater determined by a latent process.

3.4.2 Metropolis-Hastings

This model was fit to the data with the same paramaters as the NUTS algorithm except that the model was optimised using the L-BFGS-B algorithm.

3.5 Strategy

The automated algorithmic strategies were run over the period of 01/05/2014. An ordered dictionary was used to store a series of data frames container the price data of the stocks selected. The data was tidied where the columns 'open', 'high', 'low', 'close', 'ex-dividend', and 'split_ratio' were dropped from each data frame and the columns 'ticker', 'adj_open', 'adj_high', 'adj_low', and 'adj_close' were renamed to 'sid', 'open', 'high', 'low', and 'close' respectively. The ordered dictionary was converted into a panel and passed to the trading algorithm. The allocation amount of stock to invest in was determined based on the size of the portfolio.

3.5.1 Classification

An ordered dictionary was used to store the day's open and close price for each stock as the backtester simulated each trading day. The algorithm was only allowed to run once there was enough price data, the amount chosen was 6. The 2, 3, 4, 5, and 6 day SMAs were calculated and each stored in an array. The six arrays were converted into an array of tuples, where the i-th tuple contains the i-th element from each of the

argument sequences or iterables. An array was created to store a 1 if the close price increased from the open price, while a 0 if it decreased. The independent variables, the SMAs, and the dependant variables, the binary outcome for the particular stock, were passed to the machine learning algorithm to predict the next day's close price. An order was placed on a stock if the algorithm predicted an increase in the following day's price. The stock was shorted if the algorithm predicted a decrease in the following day's price. A stop loss of 80% was placed on each position.

3.5.2 Regression

An ordered dictionary was used to store the day's close price for each stock as the back-tester simulated each trading day. The algorithm was only allowed to run once there was enough price data, the amount chosen was 50. The 15 and 50 day SMAs were calculated and each stored in an array. The two arrays were converted into an array of tuples, where the i -th tuple contains the i -th element from each of the argument sequences or iterables. The independent variables, the SMAs, and the dependant variables, the close prices for the particular stock, were passed to the machine learning algorithm to predict the next day's close price. An order was placed on a stock if the predicted price was higher than the current day's price. The stock was shorted if the predicted price was lower than the current day's price. A stop loss of 80% was placed on each position.

Chapter 4

Research Findings

For the purpose of benchmarking the performance of the algorithms, a total of five stocks from the basket of uncorrelated stocks were selected. These were MSFT, CDE, NAVB, HRG, and HL.

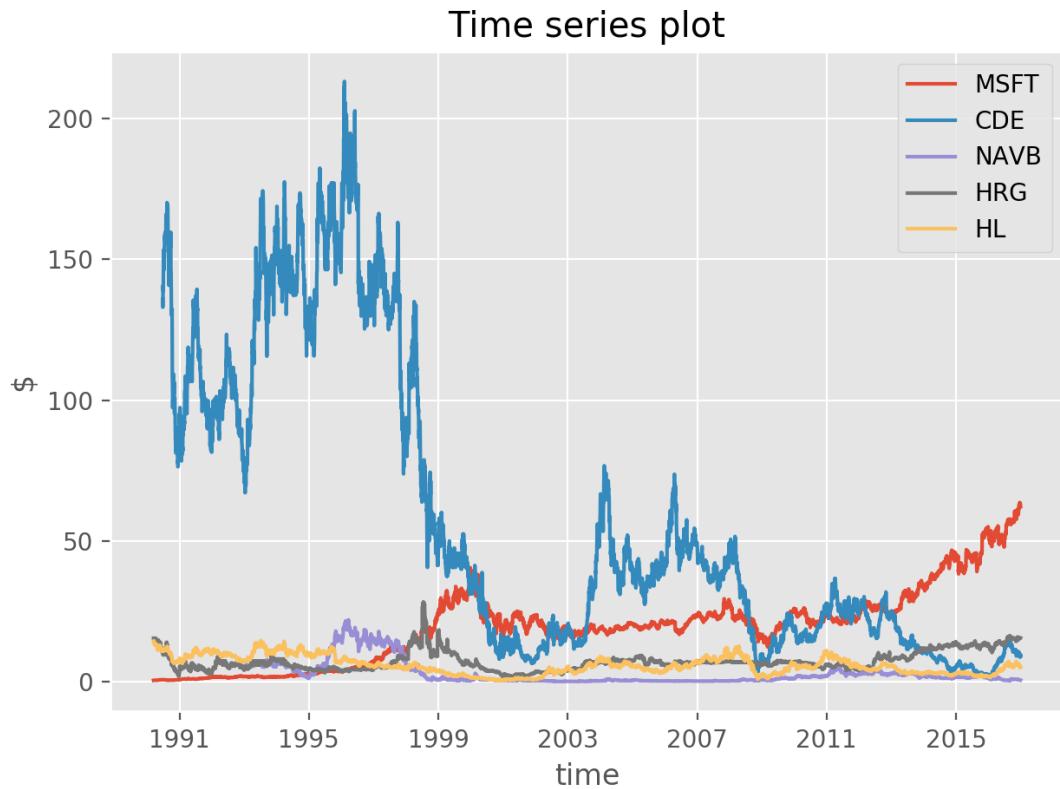


FIGURE 4.1: Basket of stocks

	MSFT	CDE	NAV	HRG	HL
Mean	16.983	57.716	3.000	13.328	7.154
Median	18.656	36.700	1.250	7.190	6.196
Maximum	63.620	213.178	22.000	89.340	23.830
Minimum	0.061	1.730	0.080	1.725	0.486
Variance	204.707	2807.852	18.445	249.105	18.445
Standard Deviation	14.308	52.989	4.29	15.783	4.295
Skewness	0.716	0.918	2.213	2.708	0.651
Kurtosis	0.250	-0.550	4.193	6.865	-0.081

TABLE 4.1: Equities Descriptive Statistics

4.1 Time Series Analysis

4.1.1 Random Walk

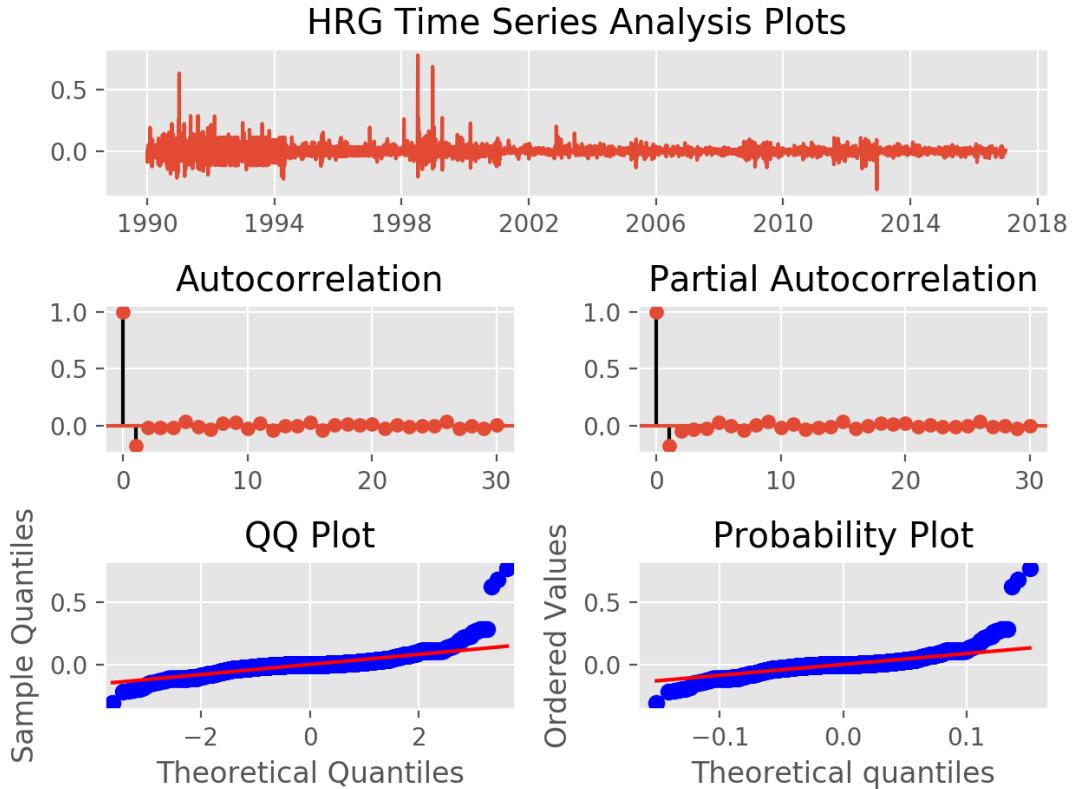


FIGURE 4.2: HRG time series analysis

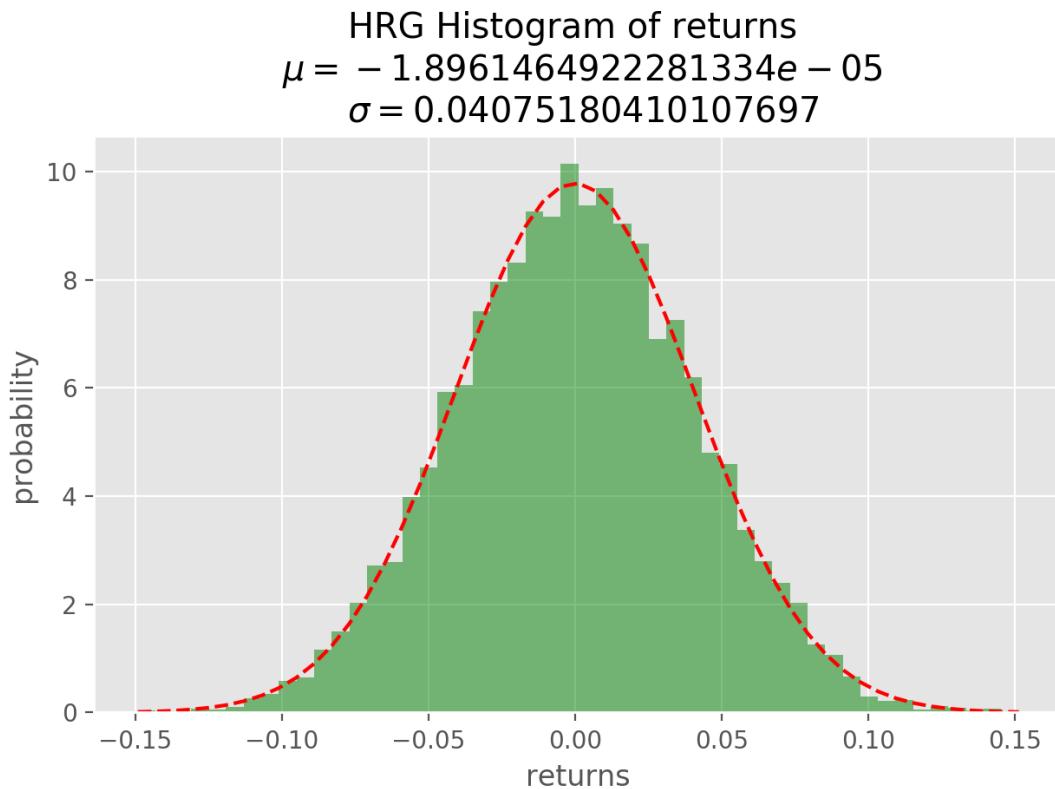


FIGURE 4.3: HRG histogram of returns

4.1.2 Ordinary Least Squares (OLS)

MSFT scored a mean absolute error regression loss of 0.810, and a coefficient of determination of 0.991.

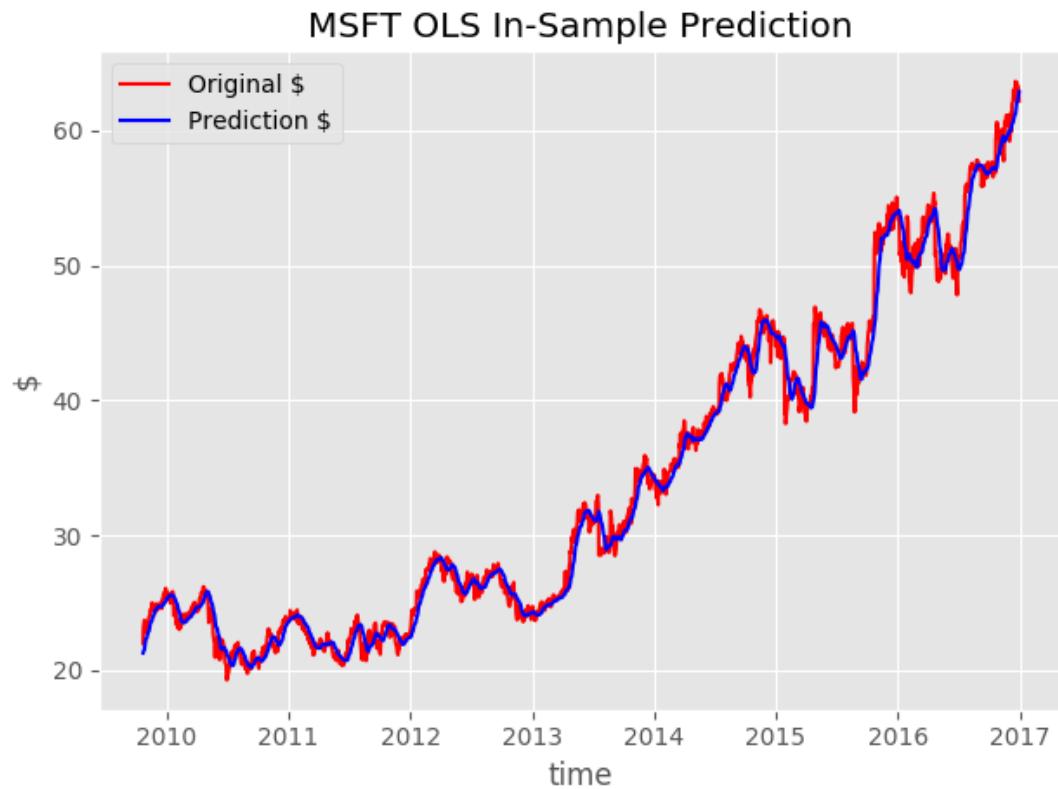


FIGURE 4.4: MSFT OLS in-sample prediction

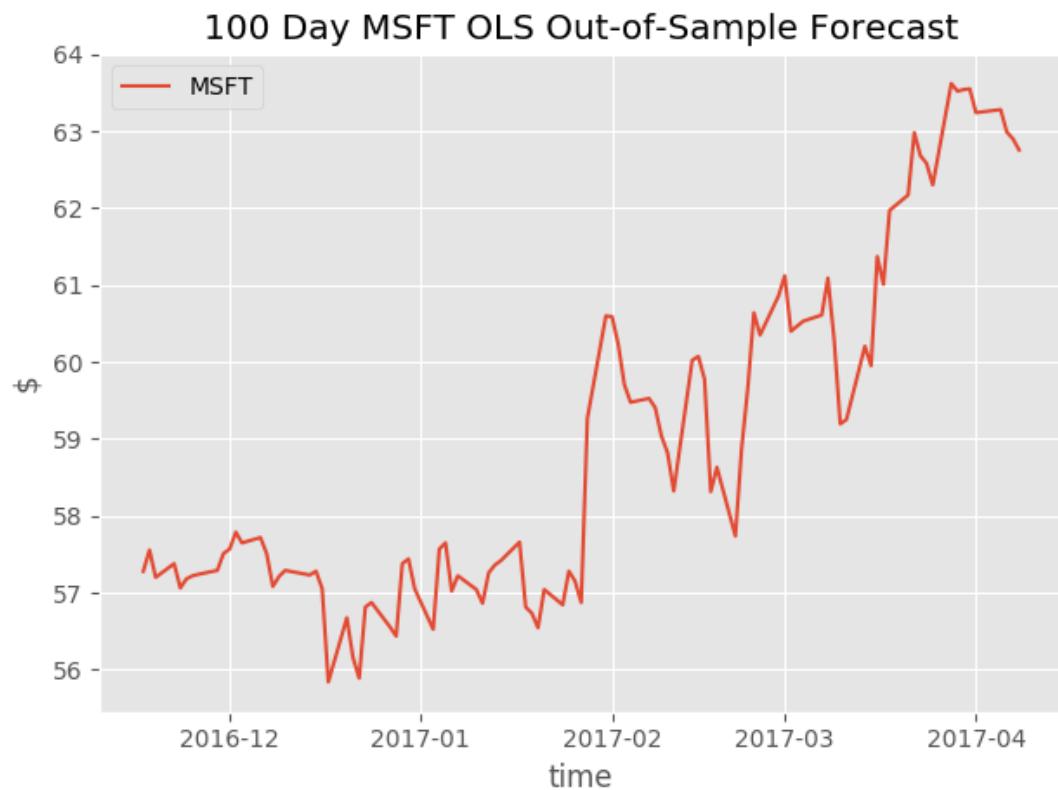


FIGURE 4.5: 100 day MSFT OLS in-sample forecast

4.1.3 Auto Regressive (AR)

MSFT scored sharpe ratios of 1.152 for the original returns, and -34.641 for the predicted returns in the in-sample test.

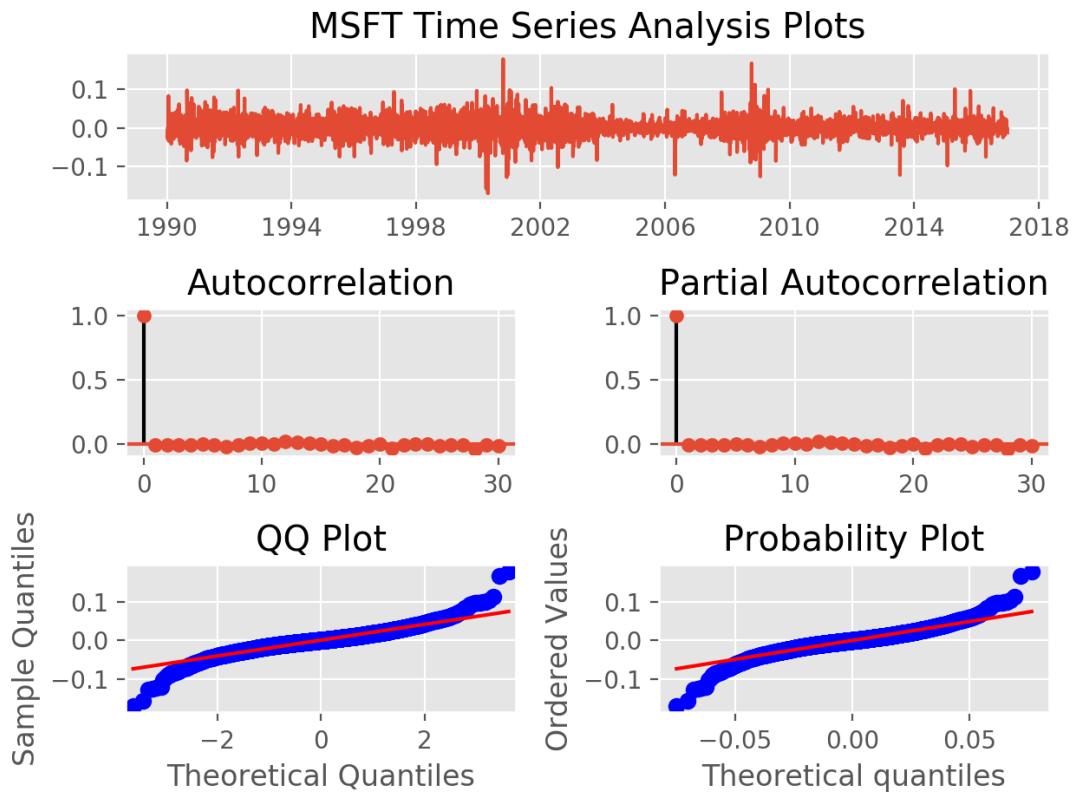


FIGURE 4.6: MSFT AR time series analysis

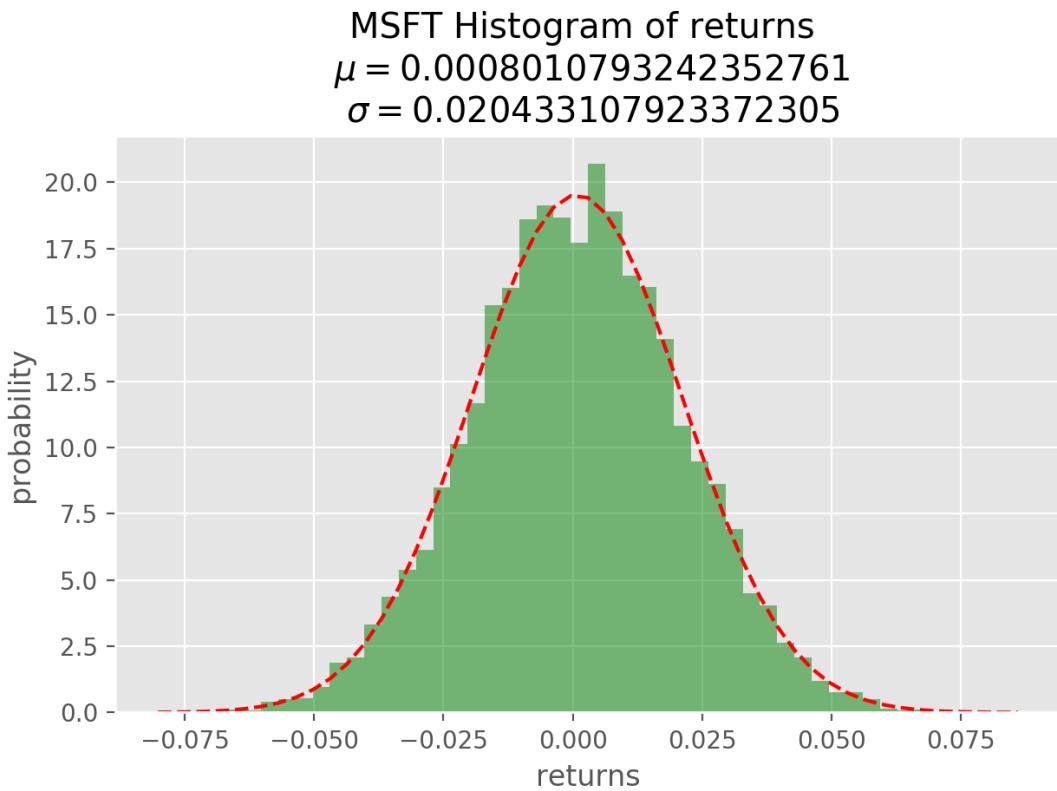


FIGURE 4.7: MSFT AR histogram of returns

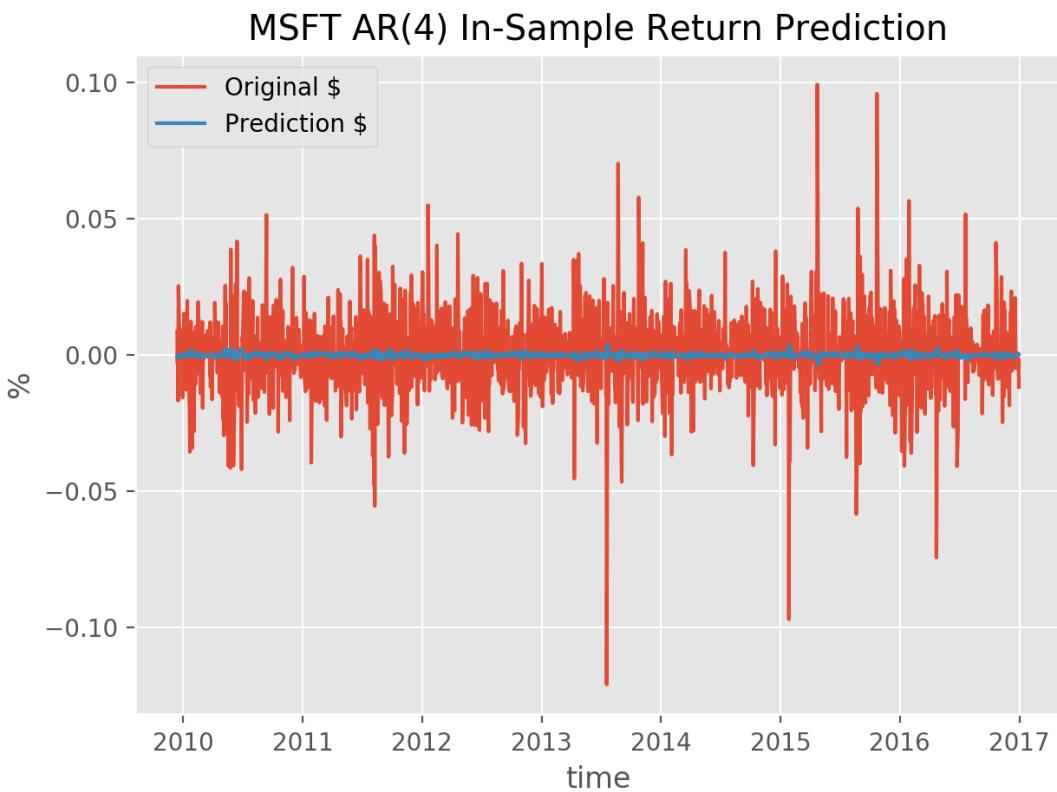


FIGURE 4.8: MSFT AR in-sample returns prediction

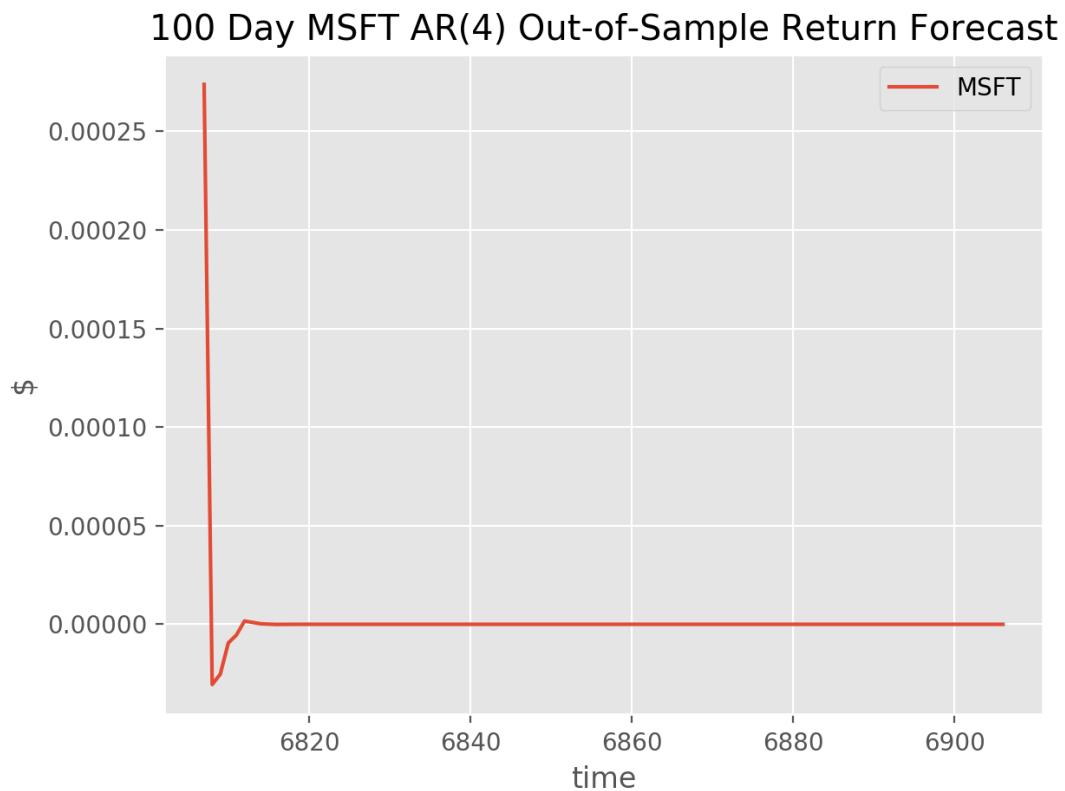


FIGURE 4.9: 100 day MSFT AR in-sample returns forecast

4.1.4 Moving Average (MA)

HRG scored sharpe ratios of 0.084 for the original returns, and -1.020 for the predicted returns in the in-sample test.

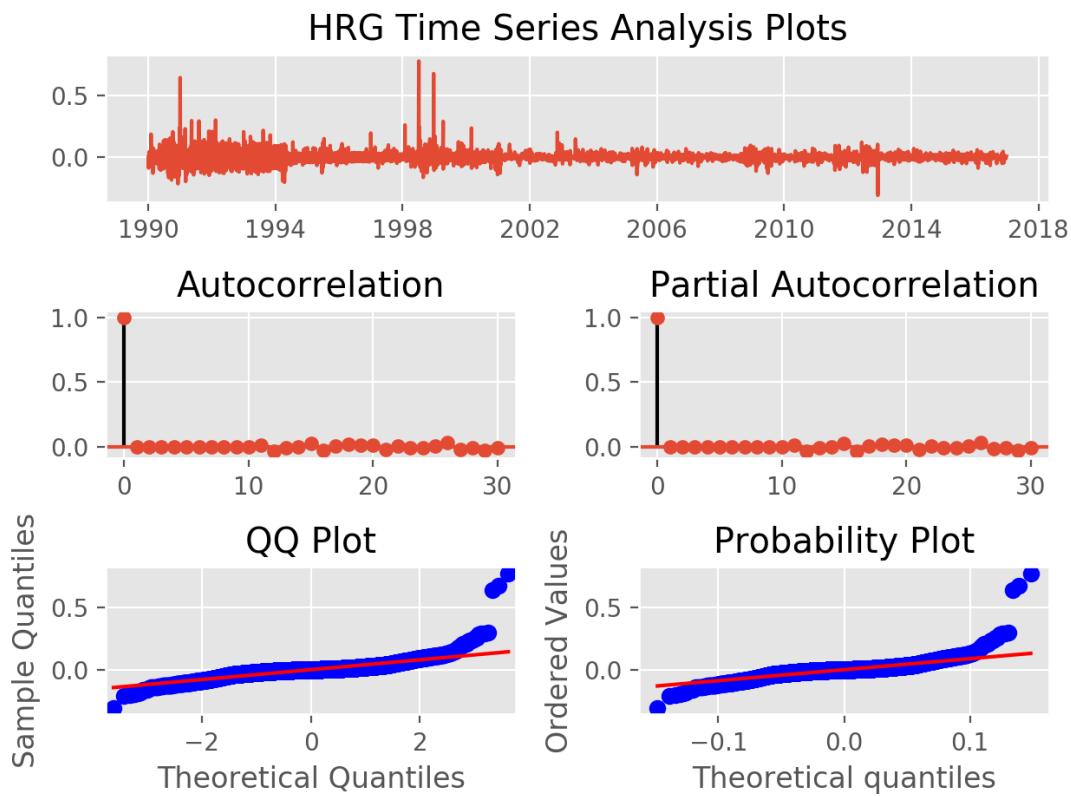


FIGURE 4.10: HRG MA time series analysis

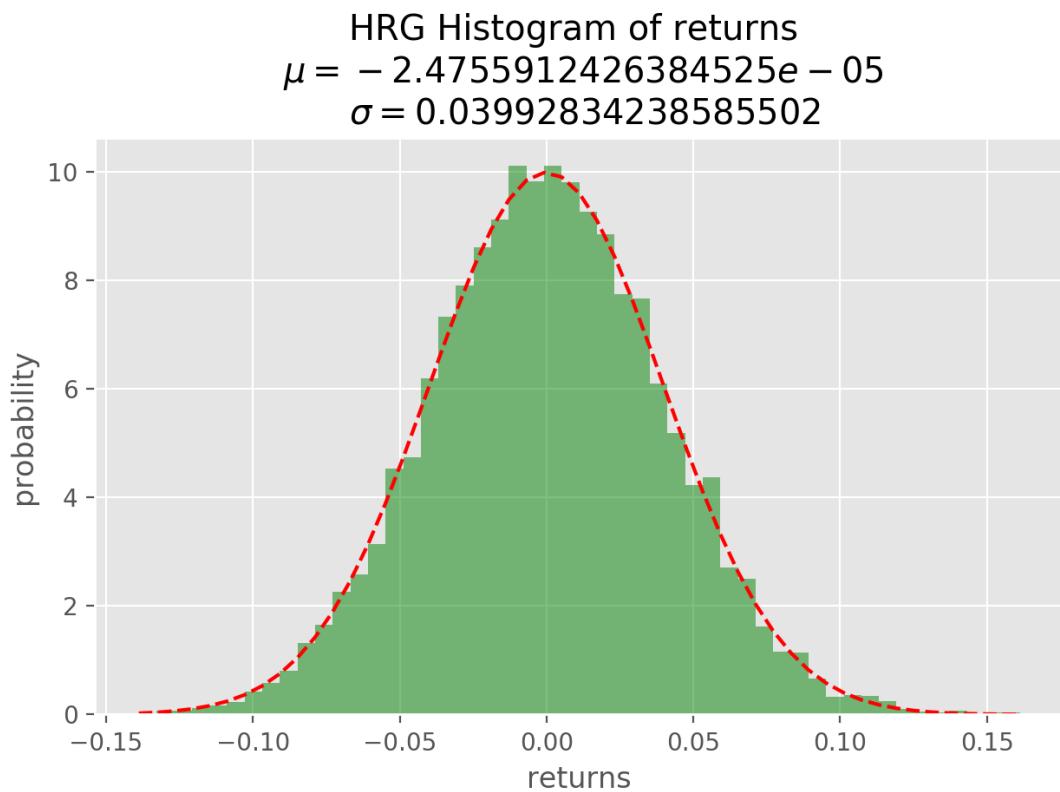


FIGURE 4.11: HRG MA histogram of returns

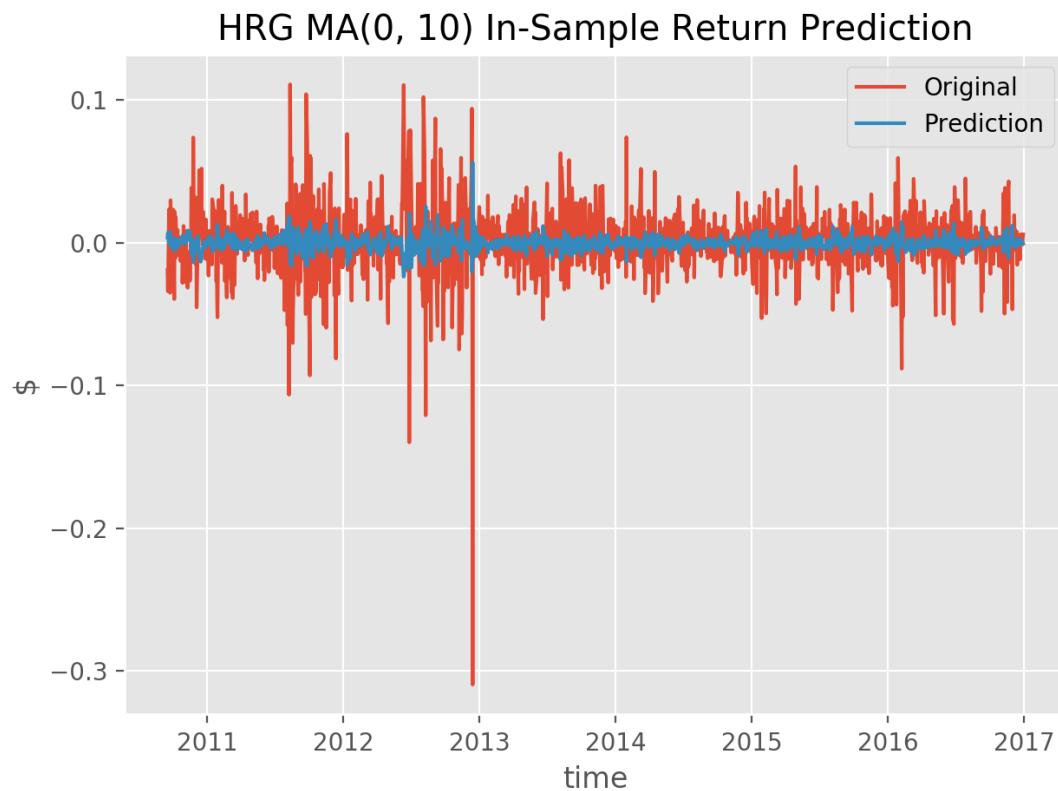


FIGURE 4.12: HRG MA in-sample returns prediction

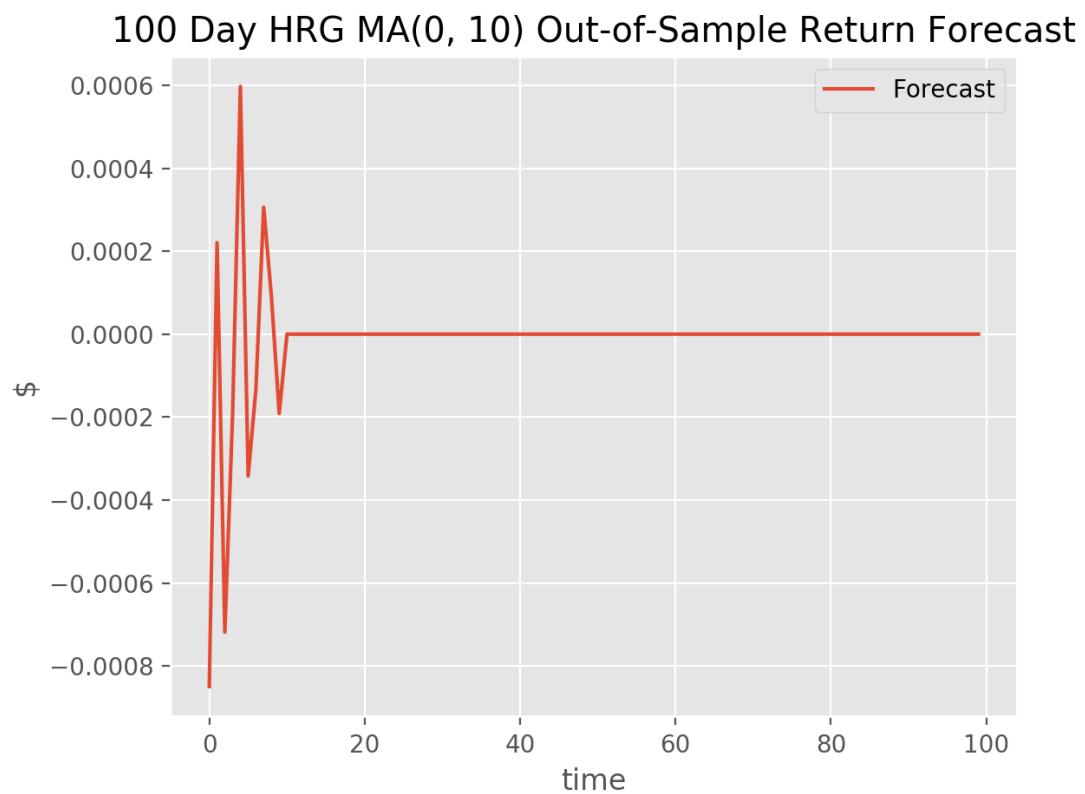


FIGURE 4.13: 100 day HRG MA in-sample returns forecast

4.1.5 Auto Regressive Moving Average (ARMA)

CDE scored sharpe ratios of -0.128 for the original returns, and -0.977 for the predicted returns in the in-sample test.

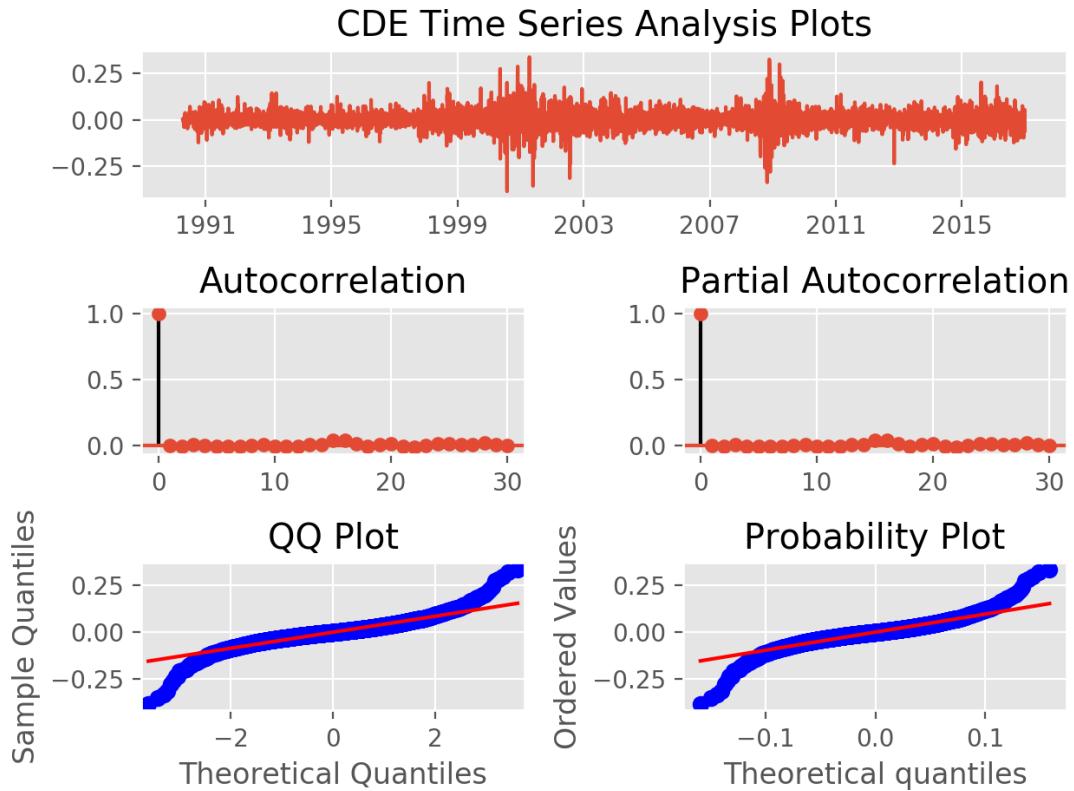


FIGURE 4.14: CDE ARMA time series analysis

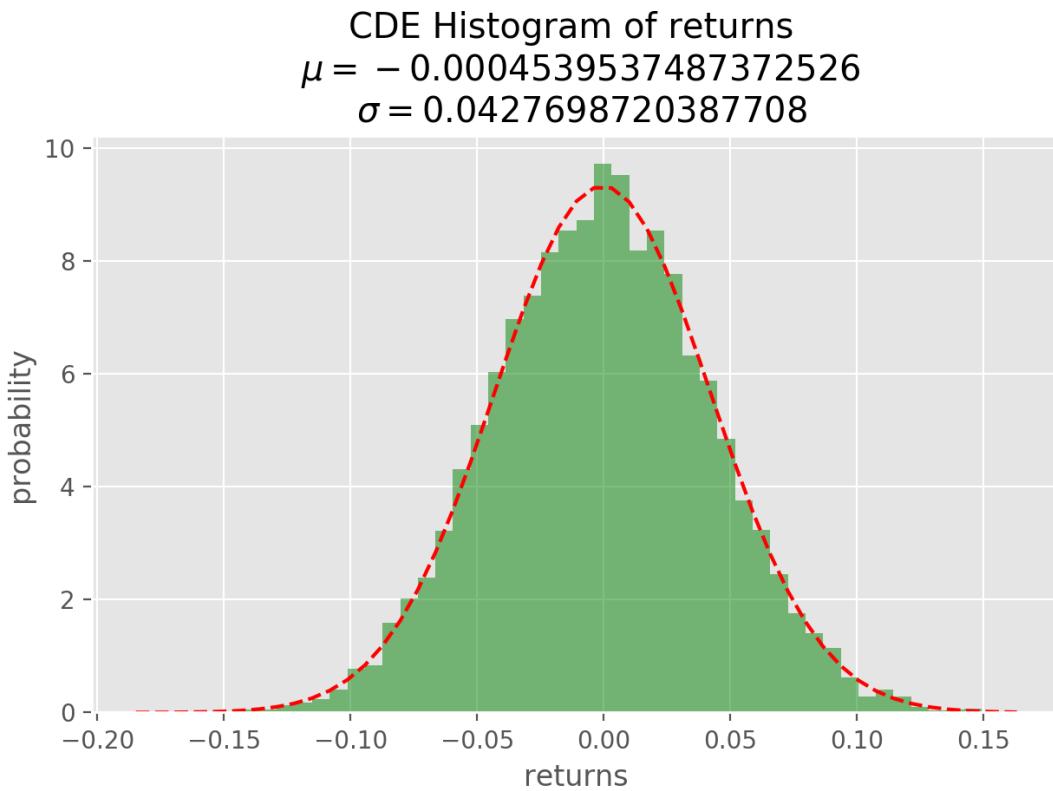


FIGURE 4.15: CDE ARMA histogram of returns

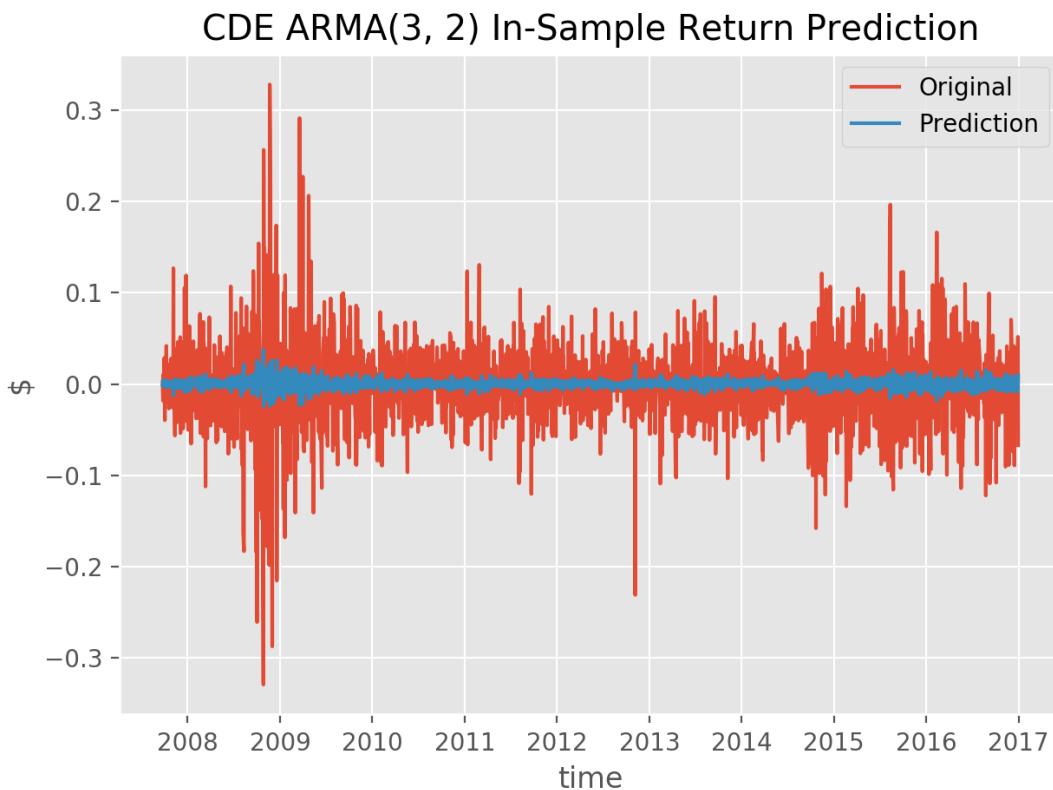


FIGURE 4.16: CDE ARMA in-sample returns prediction

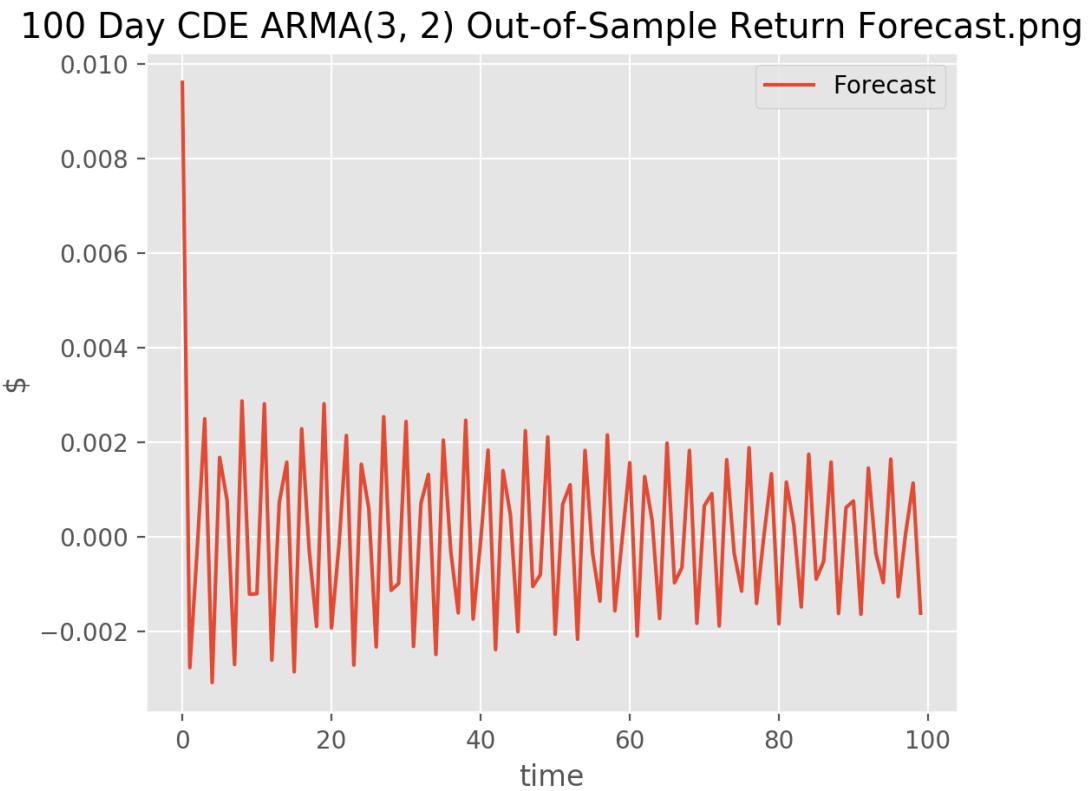


FIGURE 4.17: 100 day CDE ARMA in-sample returns forecast

4.1.6 Auto Regressive Integrated Moving Average (ARIMA)

MSFT scored sharpe ratios of 0.123 for the original returns, and -4.904 for the predicted returns in the in-sample test.

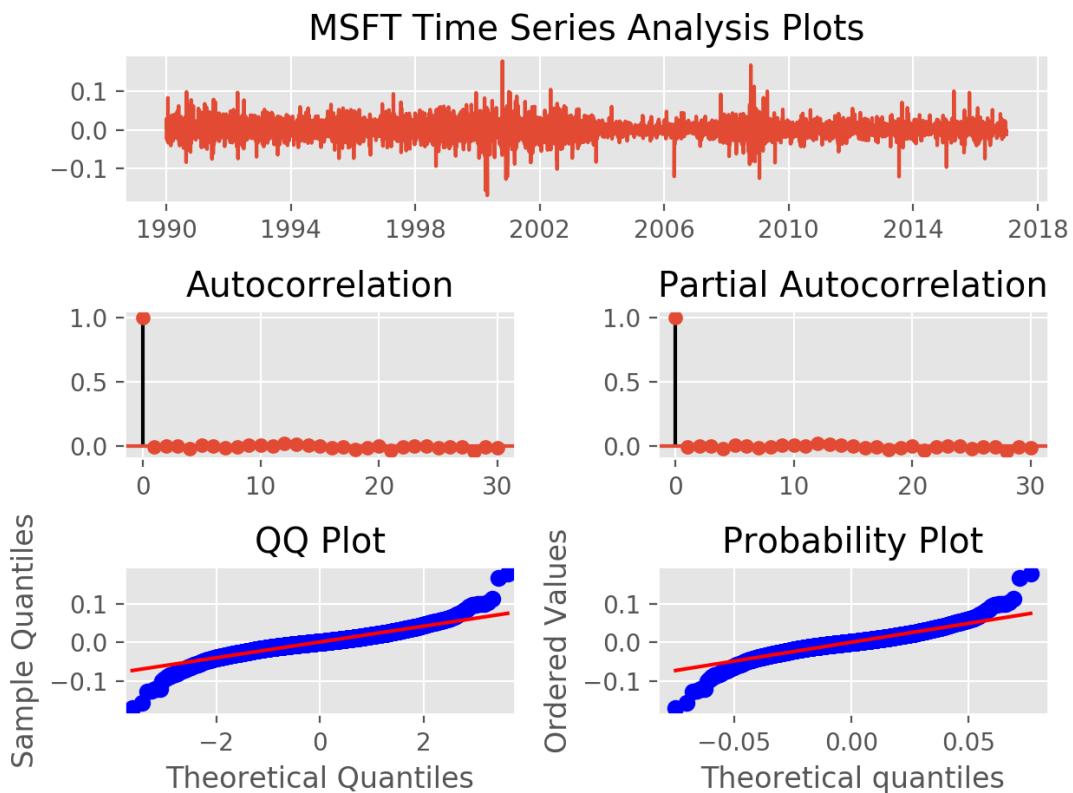


FIGURE 4.18: MSFT ARIMA time series analysis

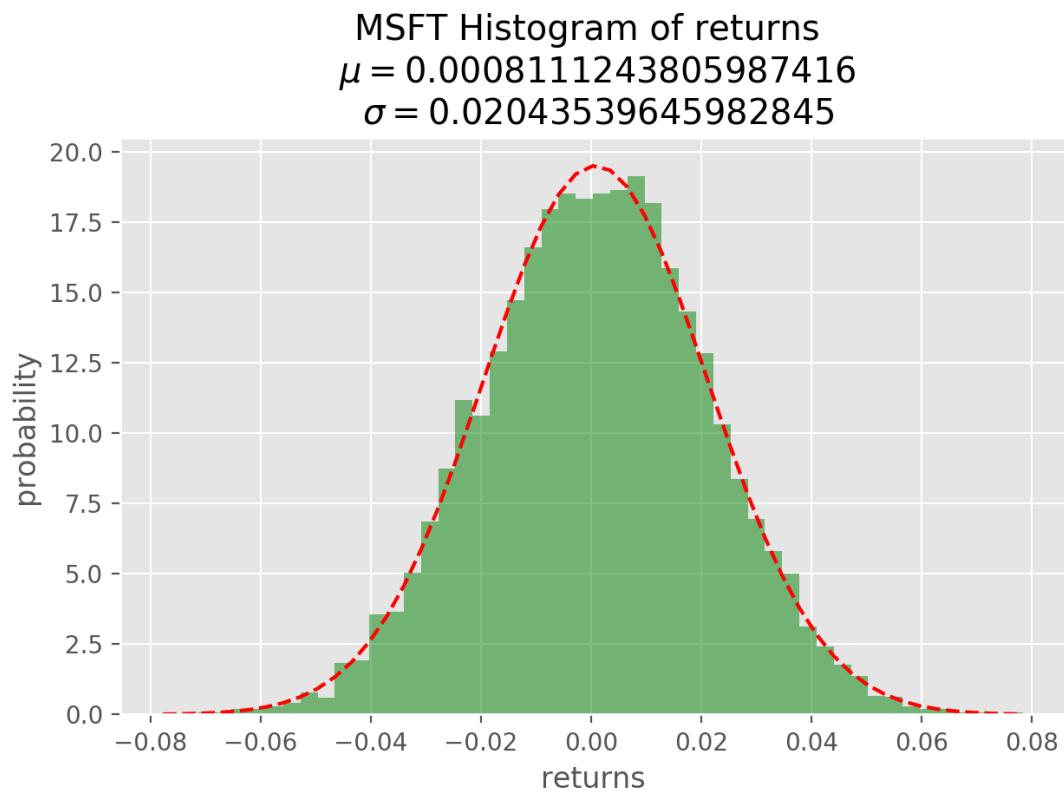


FIGURE 4.19: MSFT ARIMA histogram of returns

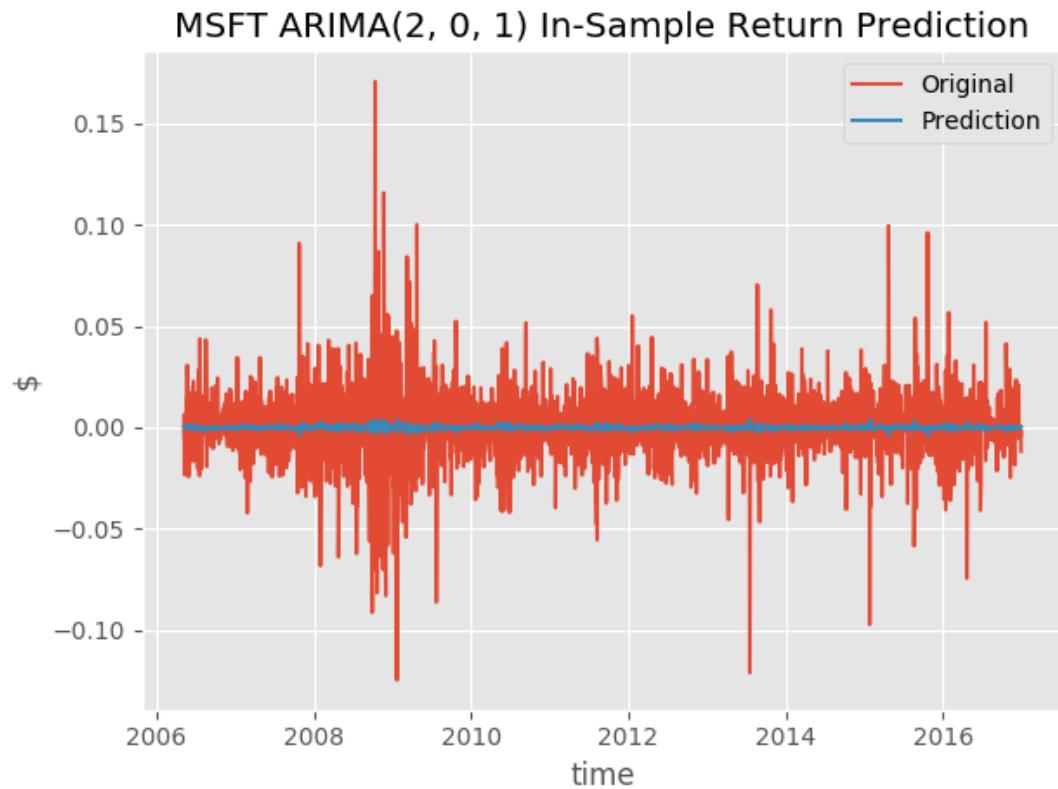


FIGURE 4.20: MSFT ARIMA in-sample returns prediction

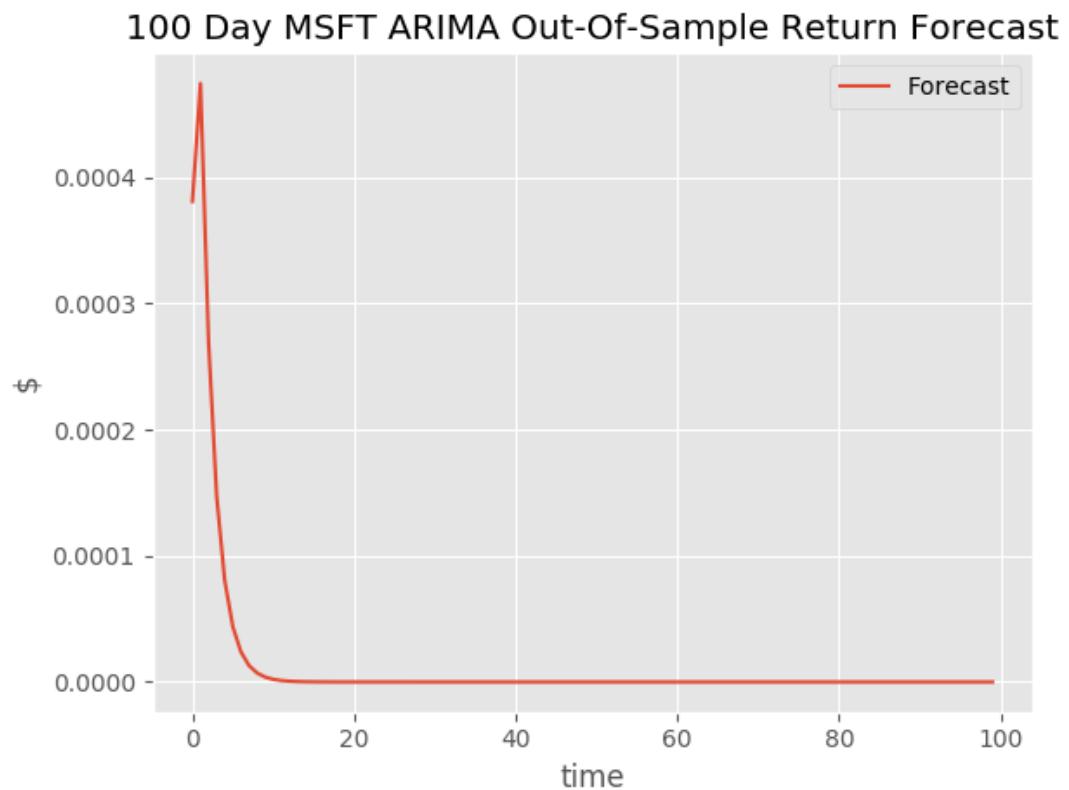


FIGURE 4.21: 100 day MSFT ARIMA in-sample returns forecast

4.2 Machine Learning

4.2.1 Classification

4.2.1.1 Decision Tree

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.77	493	131	547	190
CDE	0.79	565	144	504	133
NAV	0.76	592	165	331	126
HRG	0.75	553	210	462	135
HL	0.79	603	135	475	147

TABLE 4.2: Decision Tree results

4.2.1.2 Boosted Decision Tree

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.77	493	130	548	190
CDE	0.79	564	143	505	134
NAV	0.76	588	164	332	130
HRG	0.75	542	191	481	146
HL	0.79	602	132	478	149

TABLE 4.3: Boosted Decision Tree results

4.2.1.3 Support Vector Machine (SVM)

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.77	493	130	548	190
CDE	0.79	564	143	505	134
NAV	0.76	593	169	327	125
HRG	0.75	542	191	481	146
HL	0.81	579	98	512	171

TABLE 4.4: Support Vector Machine results

4.2.1.4 Random Forest

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.77	493	131	547	190
CDE	0.79	566	145	503	132
NAV	0.76	590	164	332	128
HRG	0.75	554	210	462	134
HL	0.81	581	101	509	169

TABLE 4.5: Random Forest results

4.2.1.5 K-Nearest Neighbour

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.76	489	130	548	194
CDE	0.80	574	147	501	126
NAV	0.75	573	161	335	145
HRG	0.74	560	233	439	128
HL	0.81	581	101	509	169

TABLE 4.6: K-Nearest Neighbour results

4.2.1.6 Logistic Regression

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.77	493	130	548	190
CDE	0.79	564	143	505	134
NAV	0.76	588	164	332	130
HRG	0.75	542	191	481	146
HL	0.79	601	132	478	149

TABLE 4.7: Logistic Regression results

4.2.1.7 Gaussian Naive Bayes

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.73	478	168	510	205
CDE	0.76	555	187	461	143
NAV	0.74	553	153	343	165
HRG	0.73	491	170	502	197
HL	0.77	590	157	453	160

TABLE 4.8: Gaussian Naive Bayes results

4.2.1.8 Bernoulli Naive Bayes

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.73	479	169	509	204
CDE	0.76	558	187	461	140
NAV	0.74	553	153	343	165
HRG	0.73	492	170	502	196
HL	0.77	590	158	452	160

TABLE 4.9: Bernoulli Naive Bayes results

4.2.1.9 Neural Network

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.77	685	181	787	258
CDE	0.79	616	156	542	152
NAV	0.76	691	183	403	153
HRG	0.74	690	269	542	164
HL	0.77	1059	270	848	164

TABLE 4.10: Neural Network results

4.2.1.10 Stochastic Gradient Descent

Ticker	Precision	True Negatives	False Negatives	True Positives	False Positives
MSFT	0.77	493	130	548	190
CDE	0.76	462	106	542	236
NAV	0.59	1032	707	118	80
HRG	0.72	761	225	1059	518
HL	0.80	535	83	527	215

TABLE 4.11: Stochastic Gradient Descent results

4.2.2 Regression

4.2.2.1 Decision Tree

MSFT scored a mean absolute error regression loss of 5.606, and a coefficient of determination of 0.458.

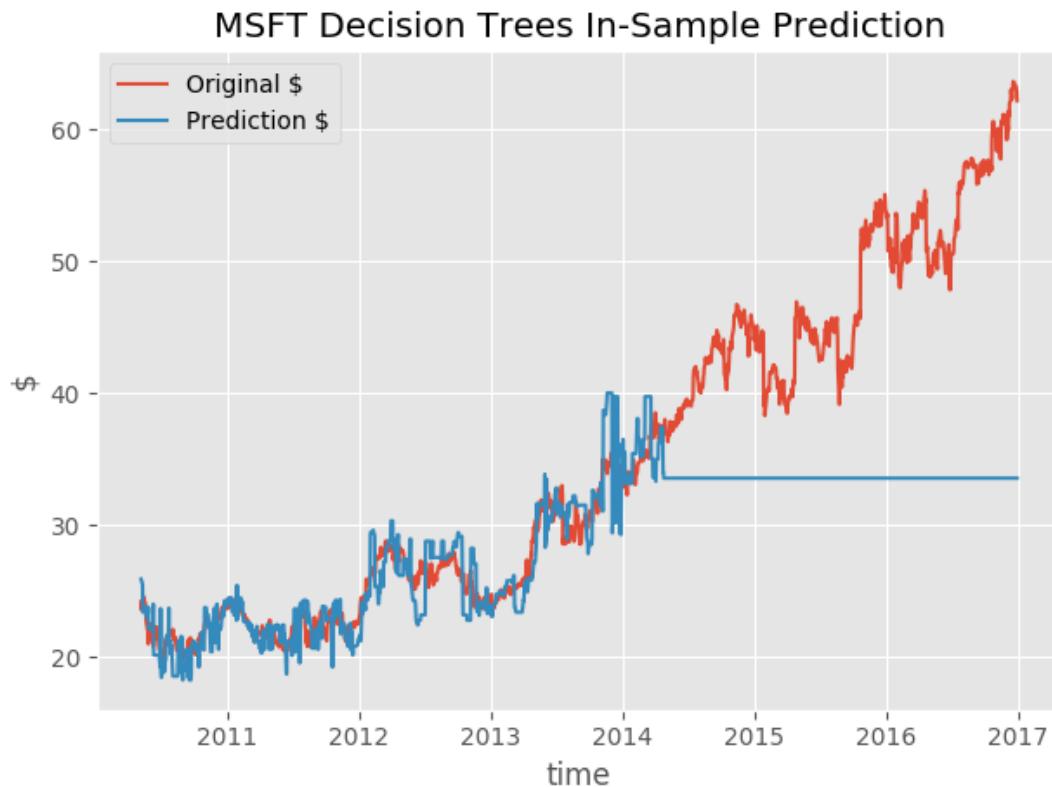


FIGURE 4.22: MSFT Decision Trees in-sample prediction

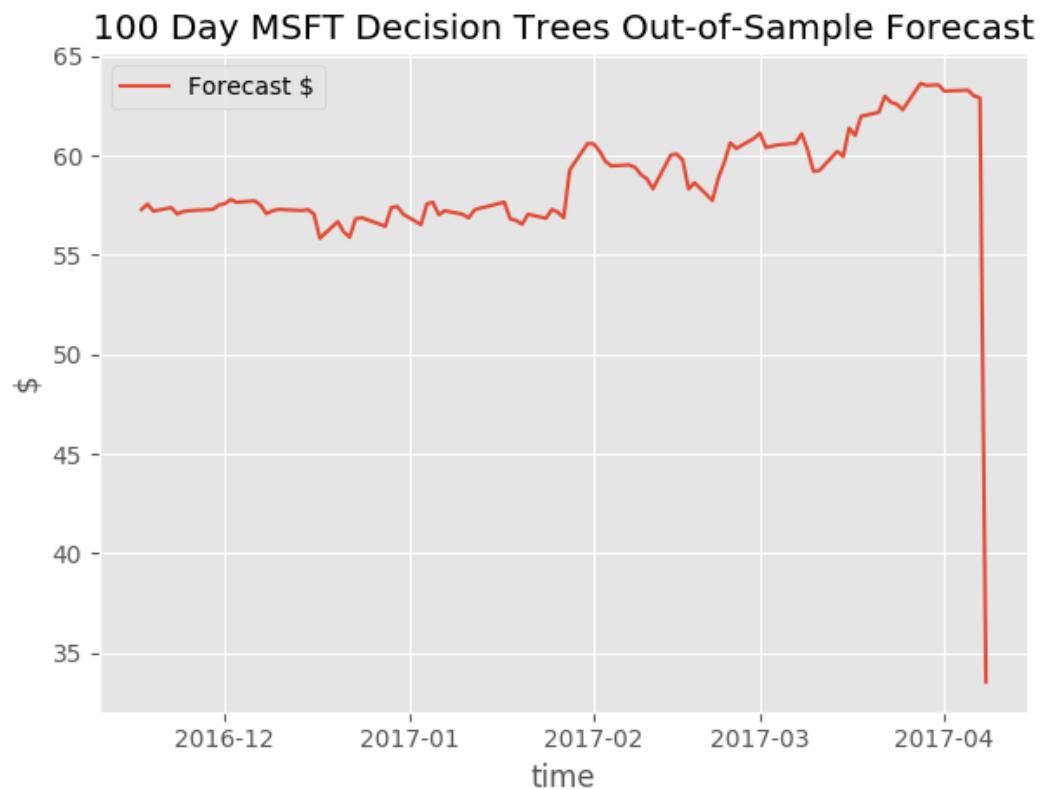


FIGURE 4.23: 100 day MSFT Decision Trees out-of-sample forecast

4.2.2.2 Boosted Decision Tree

MSFT scored a mean absolute error regression loss of 4.426, and a coefficient of determination of 0.580.

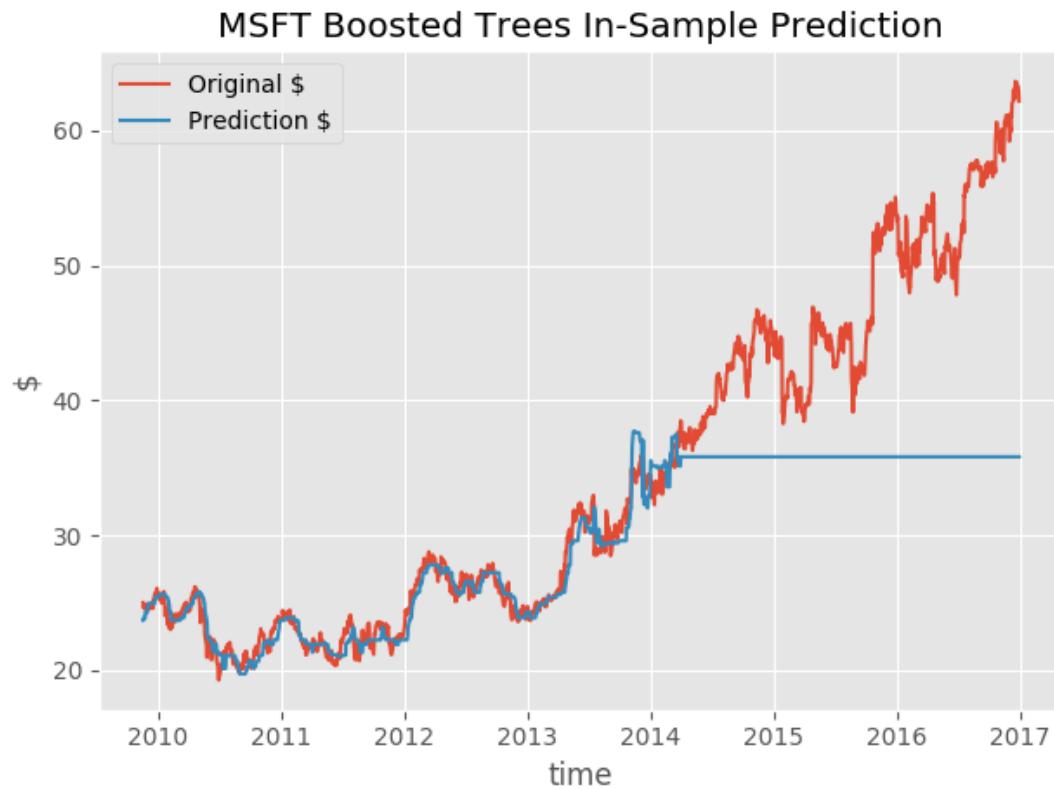


FIGURE 4.24: MSFT Boosted Decision Trees in-sample prediction

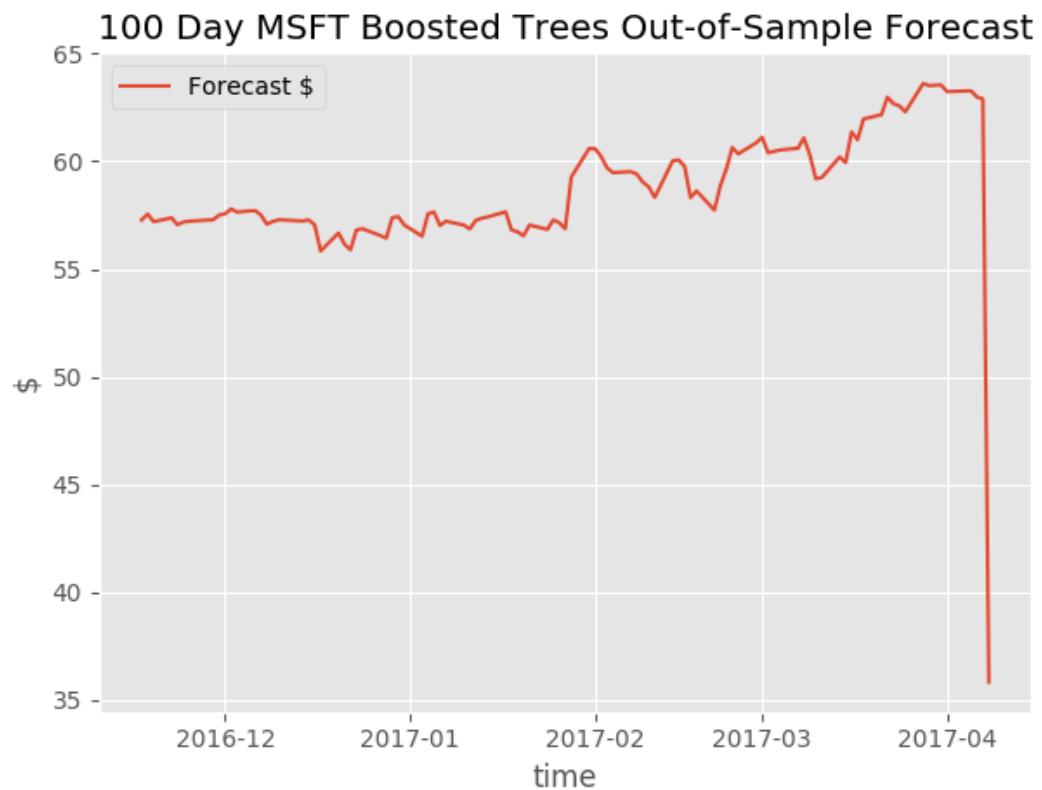


FIGURE 4.25: 100 day MSFT Boosted Decision Trees out-of-sample forecast

4.2.2.3 K-Nearest Neighbour

MSFT scored a mean absolute error regression loss of 4.426, and a coefficient of determination of 0.580.

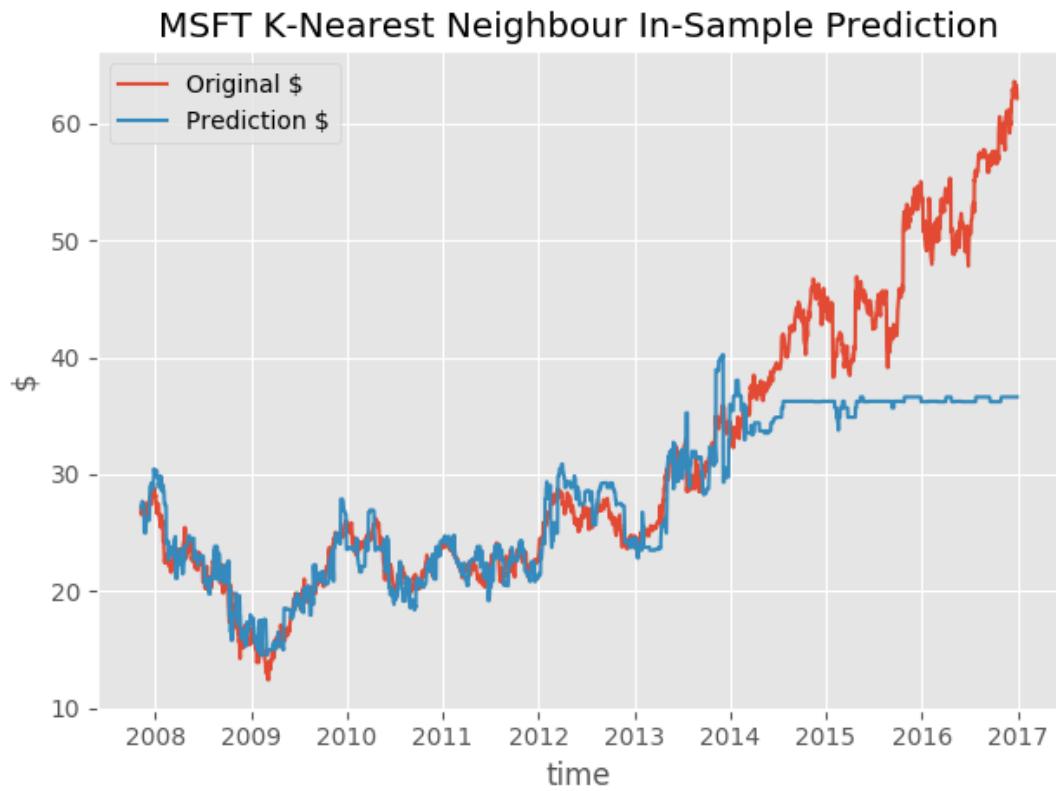


FIGURE 4.26: MSFT K-Nearest Neighbour in-sample prediction

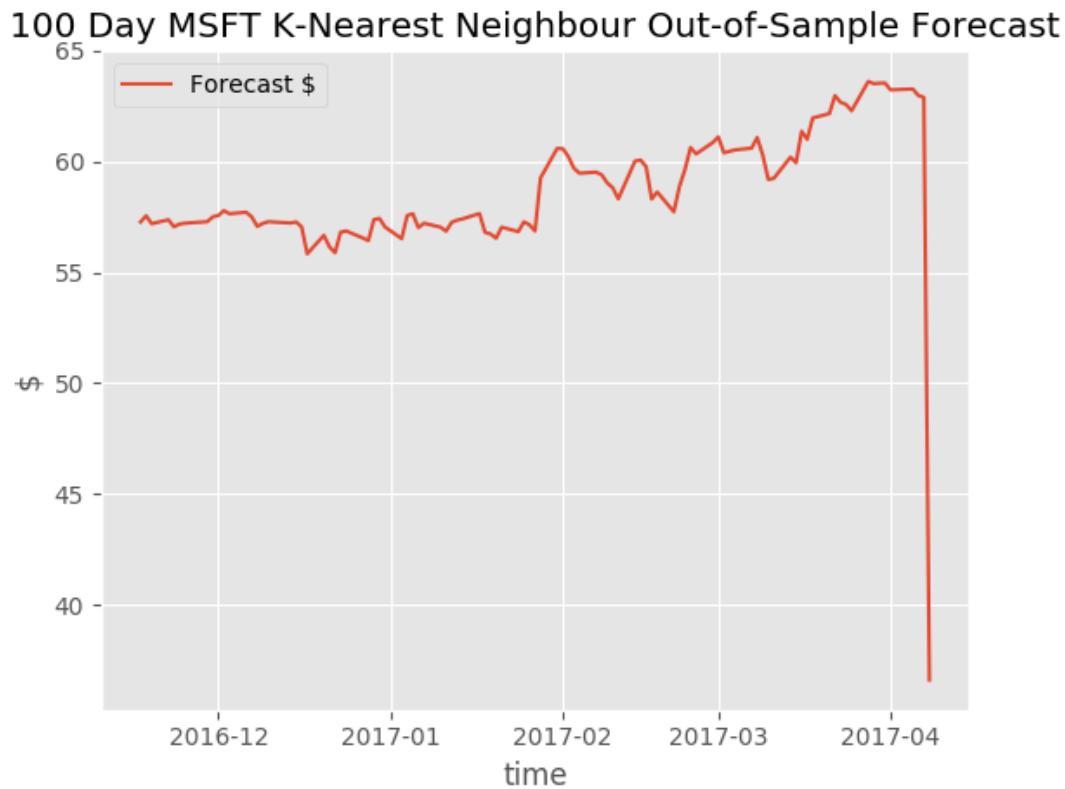


FIGURE 4.27: 100 day MSFT K-Nearest Neighbour out-of-sample forecast

4.2.2.4 Random Forest

MSFT scored a mean absolute error regression loss of 5.199, and a coefficient of determination of 0.483.

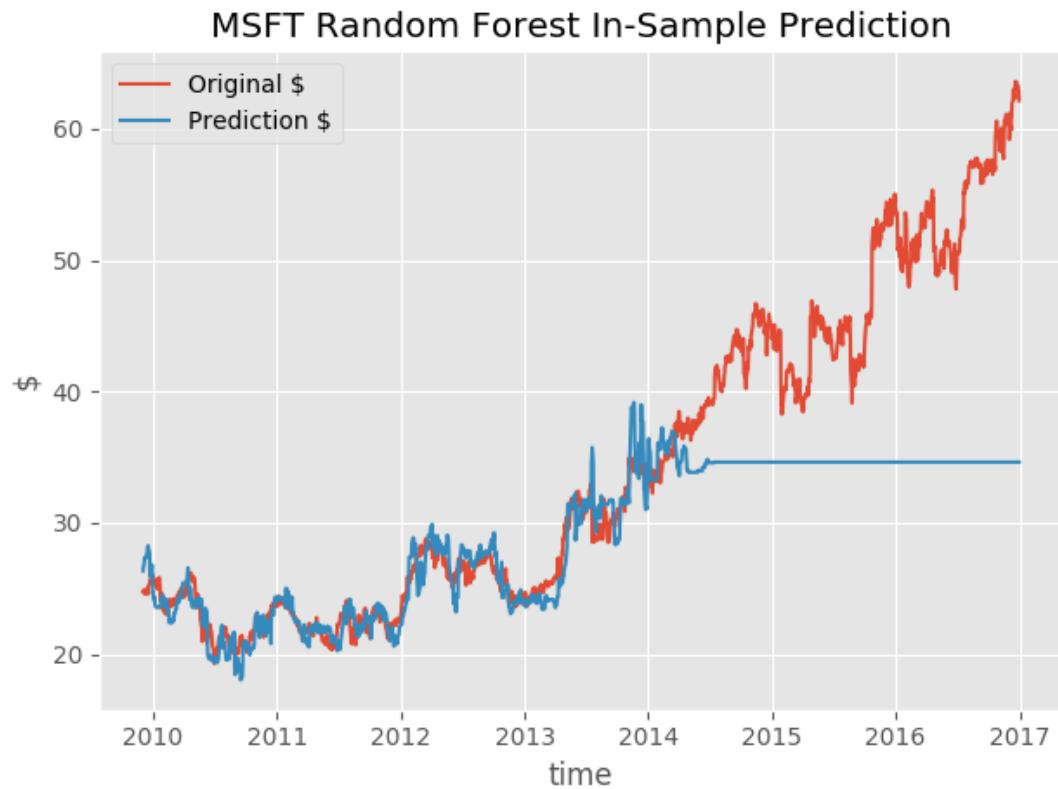


FIGURE 4.28: MSFT Random Forest in-sample prediction

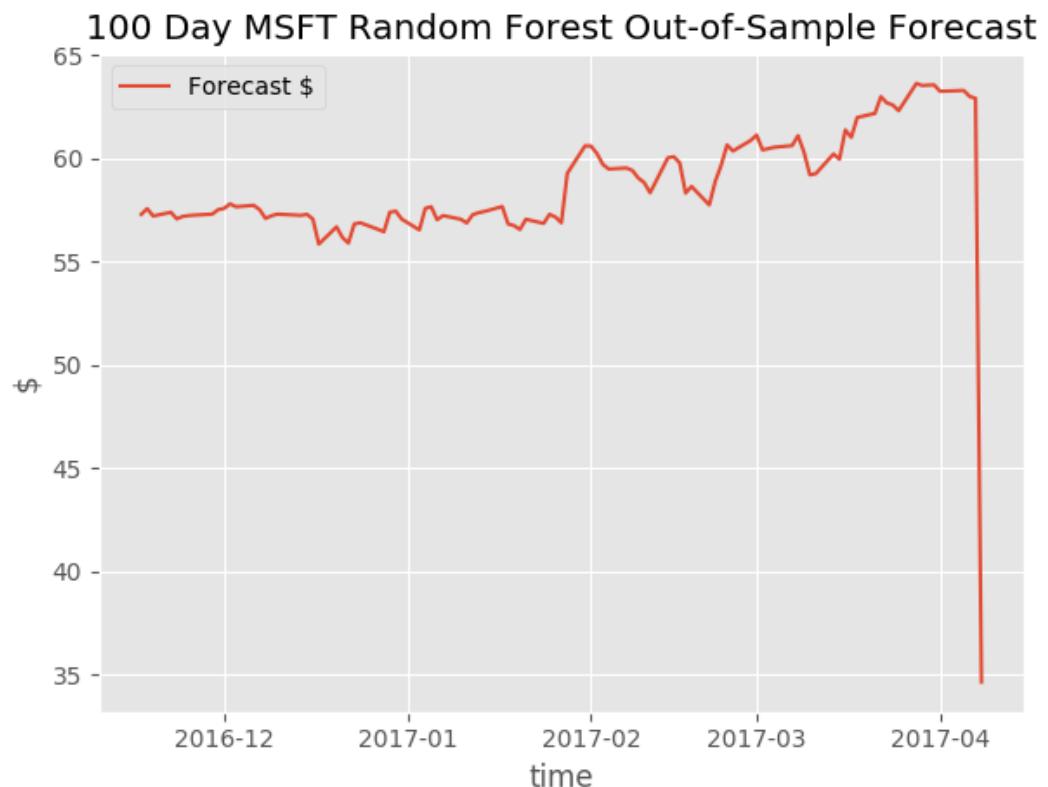


FIGURE 4.29: 100 day MSFT Random Forest out-of-sample forecast

4.2.2.5 Linear Regression

NAVB scored a mean absolute error regression loss of 0.128, and a coefficient of determination of 0.962.

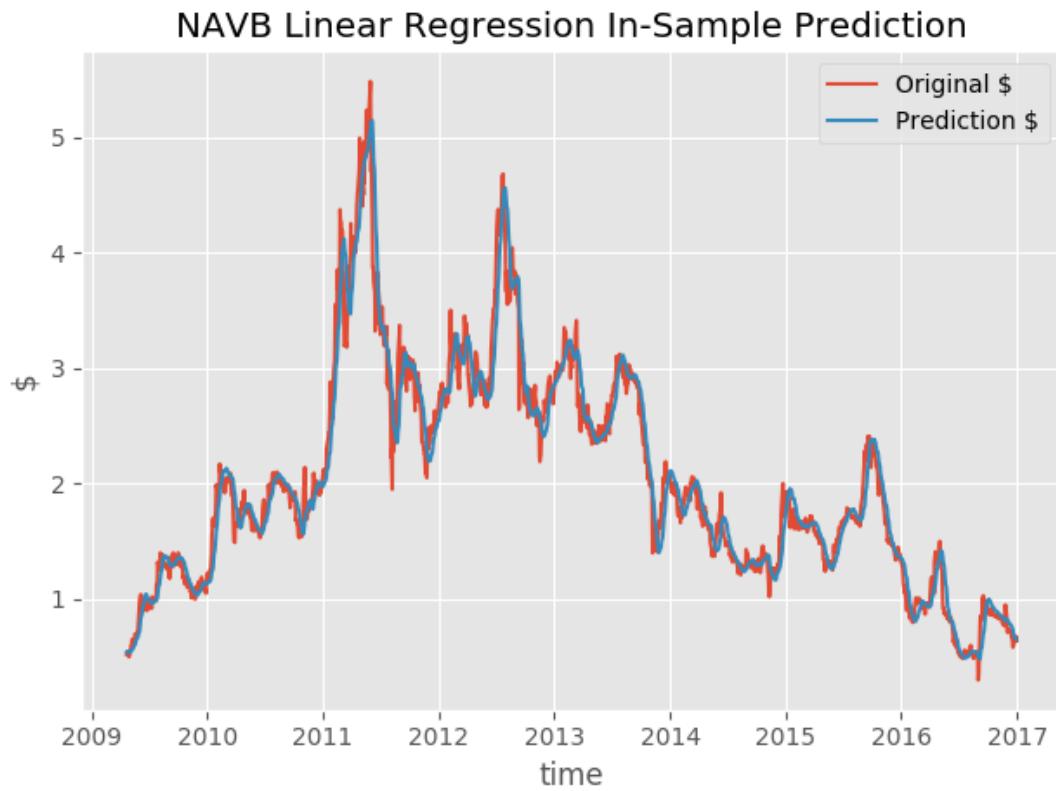


FIGURE 4.30: NAVB Linear Regression in-sample prediction

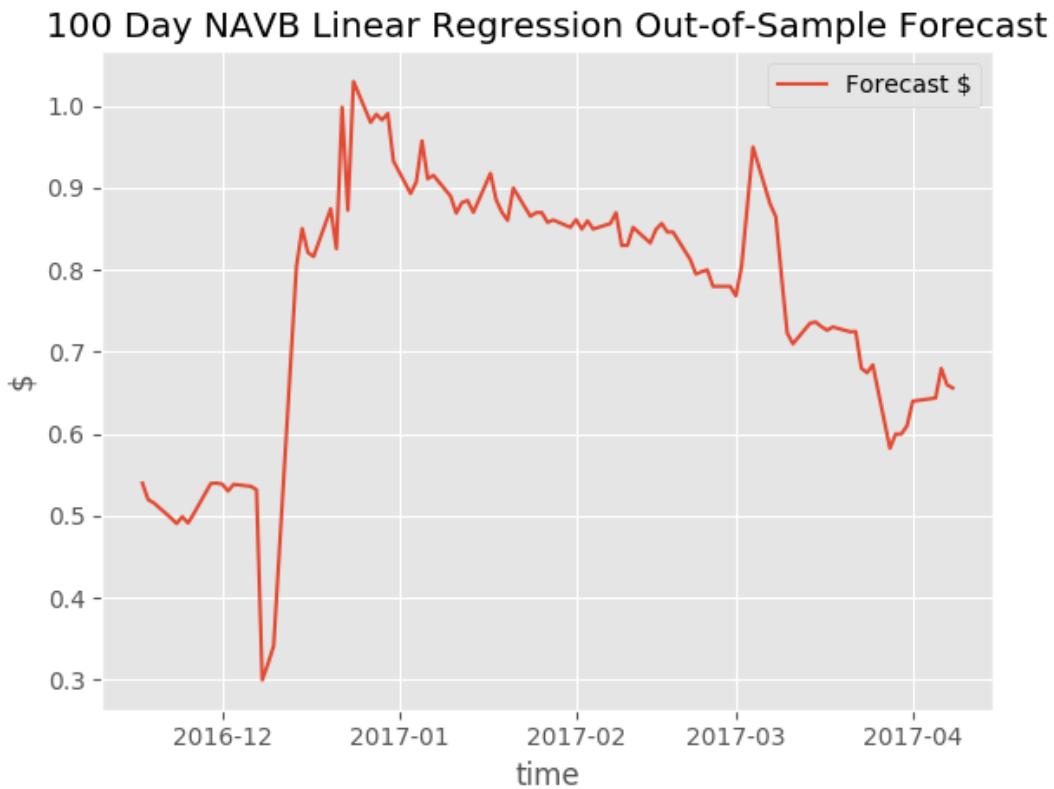


FIGURE 4.31: 100 day NAVB Linear Regression out-of-sample forecast

4.2.2.6 Neural Network

CDE scored a mean absolute error regression loss of 2.906, and a coefficient of determination of 0.816.

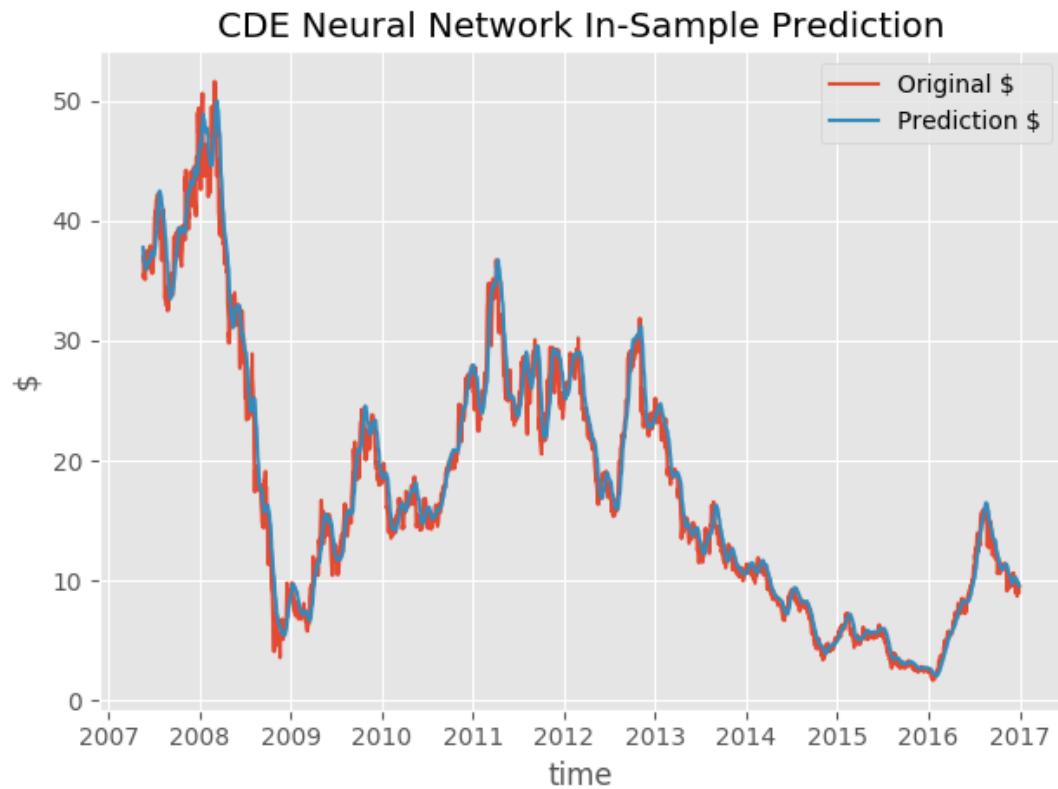


FIGURE 4.32: CDE Neural Network in-sample prediction

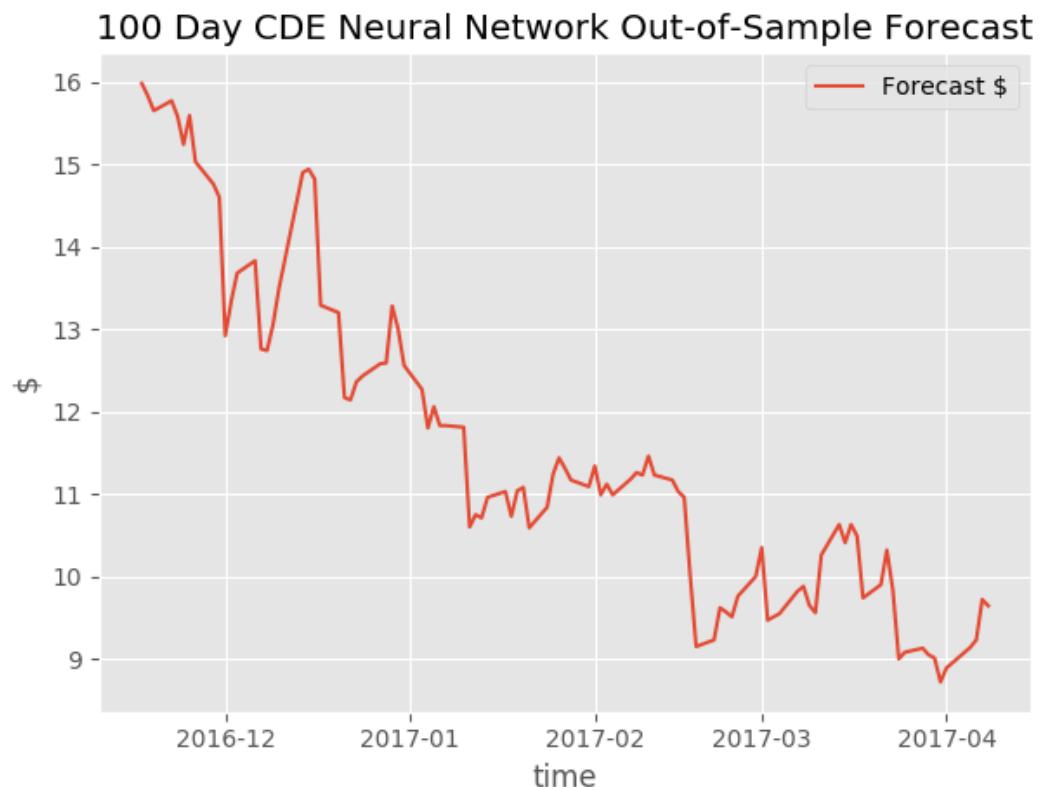


FIGURE 4.33: 100 day CDE Neural Network out-of-sample forecast

4.2.2.7 Stochastic Gradient Descent

CDE scored a mean absolute error regression loss of 4.724, and a coefficient of determination of 0.788.

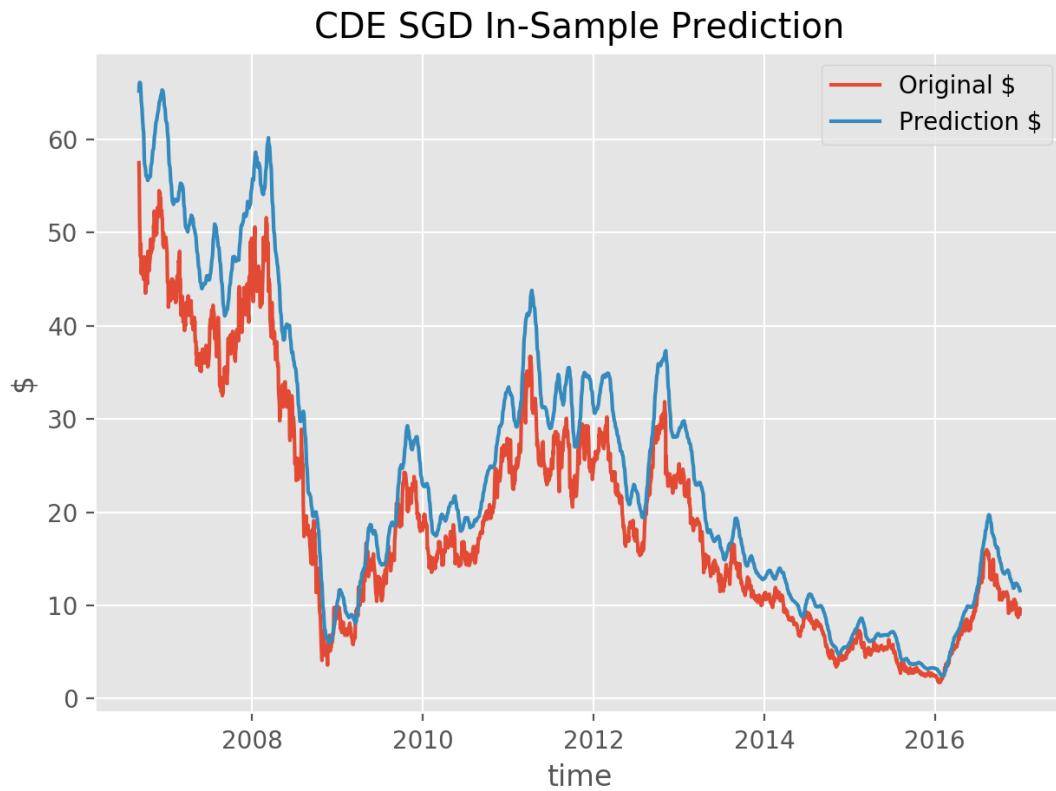


FIGURE 4.34: CDE SGD in-sample prediction

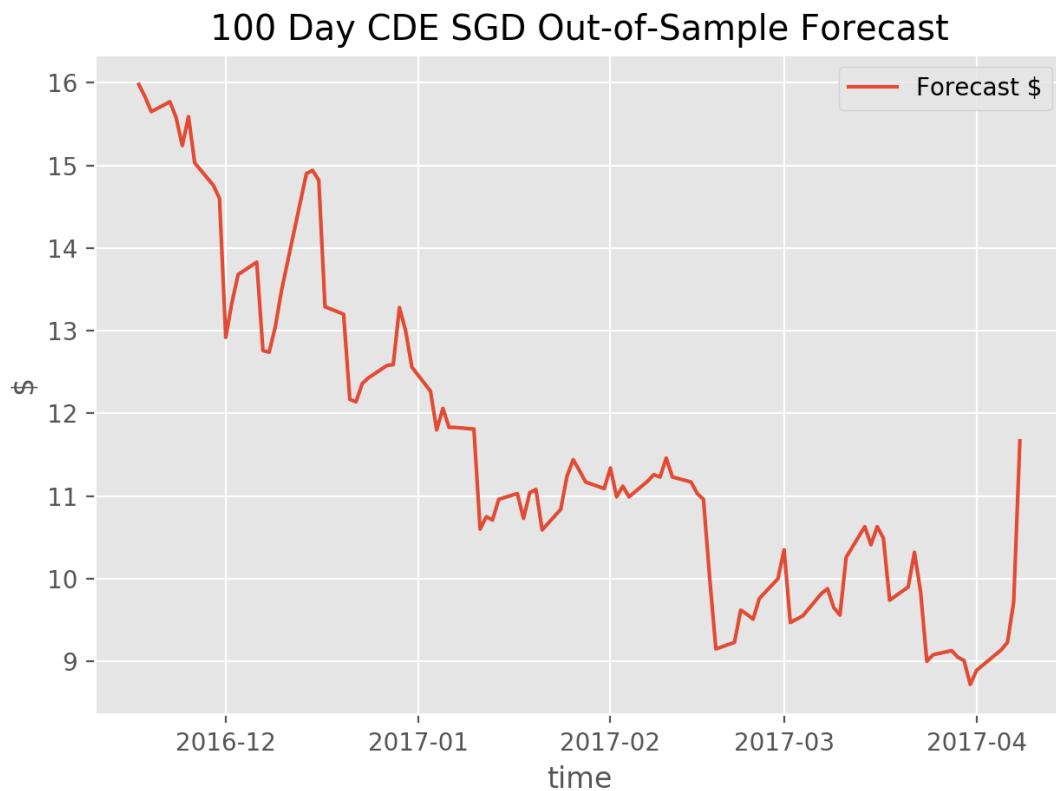


FIGURE 4.35: 100 day CDE SGD out-of-sample forecast

4.3 Bayesian Statistics

4.3.1 Metropolis-Hastings

MSFT scored sharpe ratios of 0.439 for the original returns, and -4.796 for the predicted returns in the in-sample test.

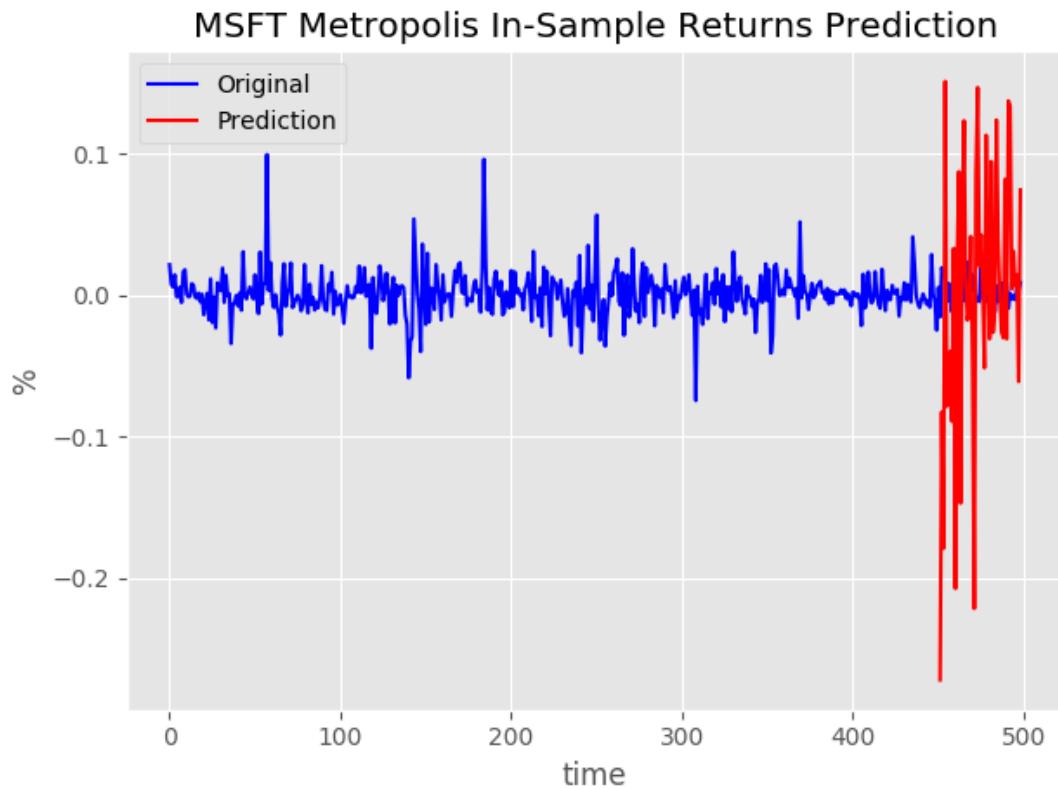


FIGURE 4.36: MSFT Metropolis-Hastings in-sample prediction

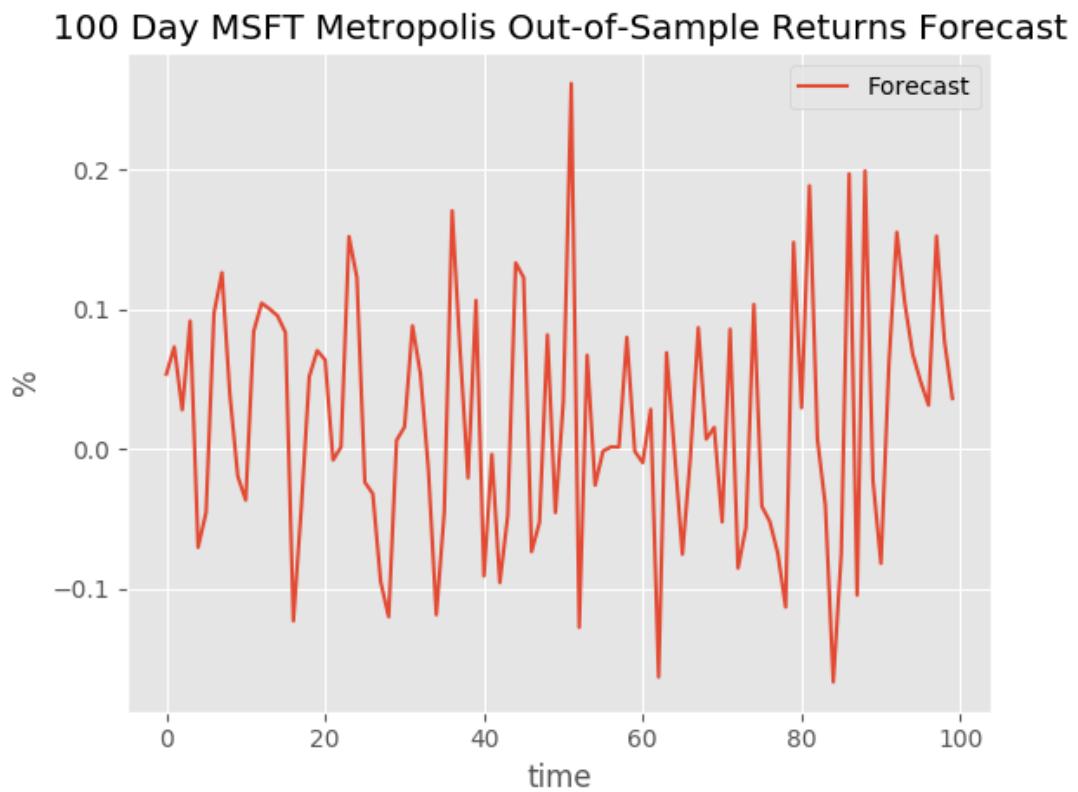


FIGURE 4.37: 100 day MSFT Metropolis-Hastings out-of-sample forecast

4.3.2 No-U-Turn Sampler (NUTS)

HL scored sharpe ratios of 0.506 for the original returns, and 0.923 for the predicted returns in the in-sample test.

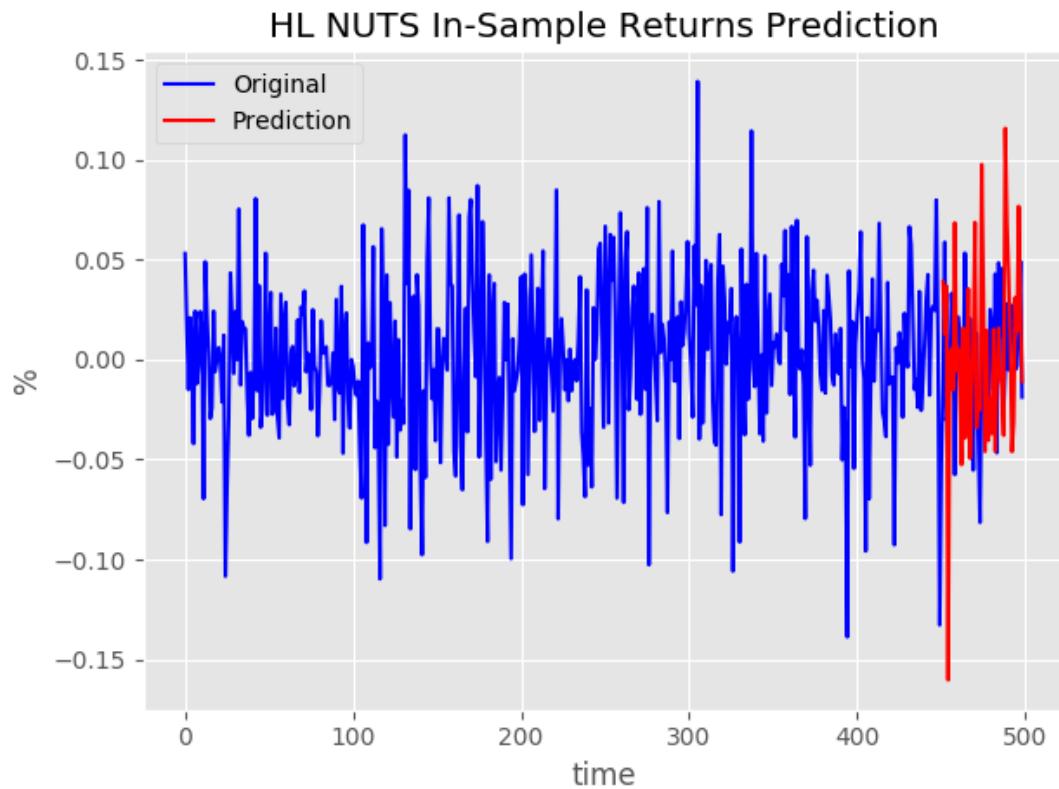


FIGURE 4.38: HL NUTS in-sample prediction

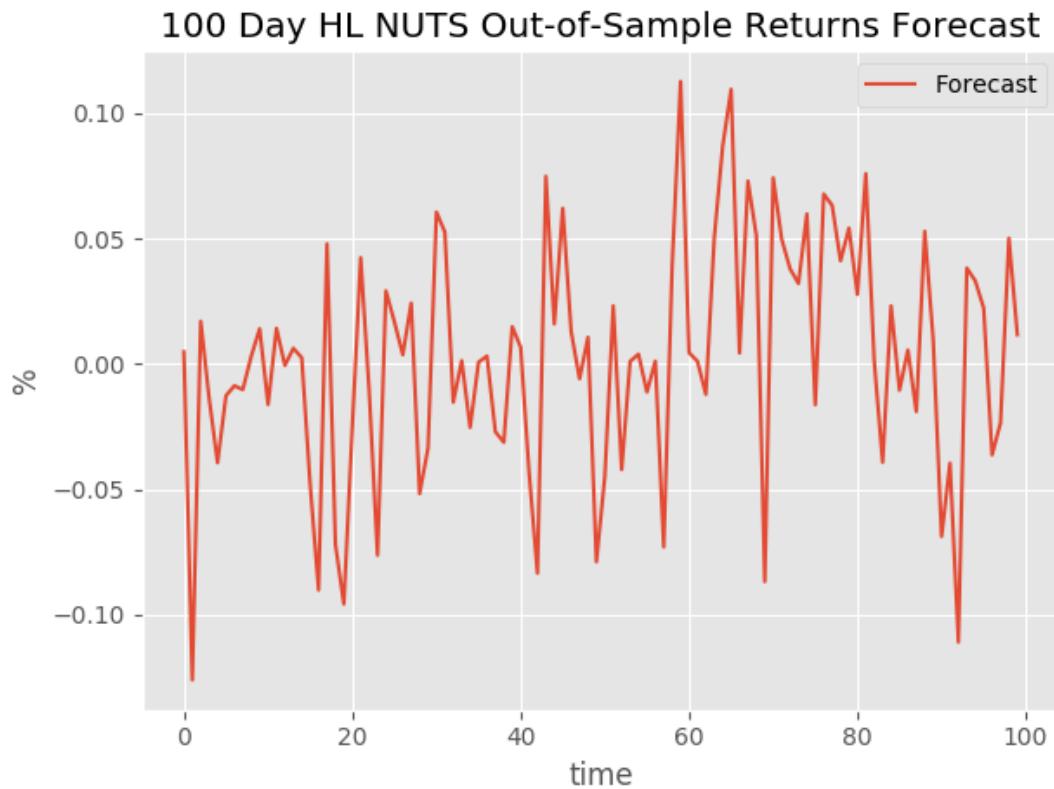


FIGURE 4.39: 100 day HL NUTS out-of-sample forecast

4.4 Strategy

4.4.1 Classification

Starting Capital	\$100,000
Total Capital Used	\$103,110.65
Sharpe Ratio	0.346
Portfolio Value	\$149,126.306
Algorithm Period Return	0.491
Benchmark Period Return	1.008
Algorithm Volatility	0.272
Benchmark Volatility	0.156

TABLE 4.12: Machine Learning Classifier strategy with only upwards forecasts

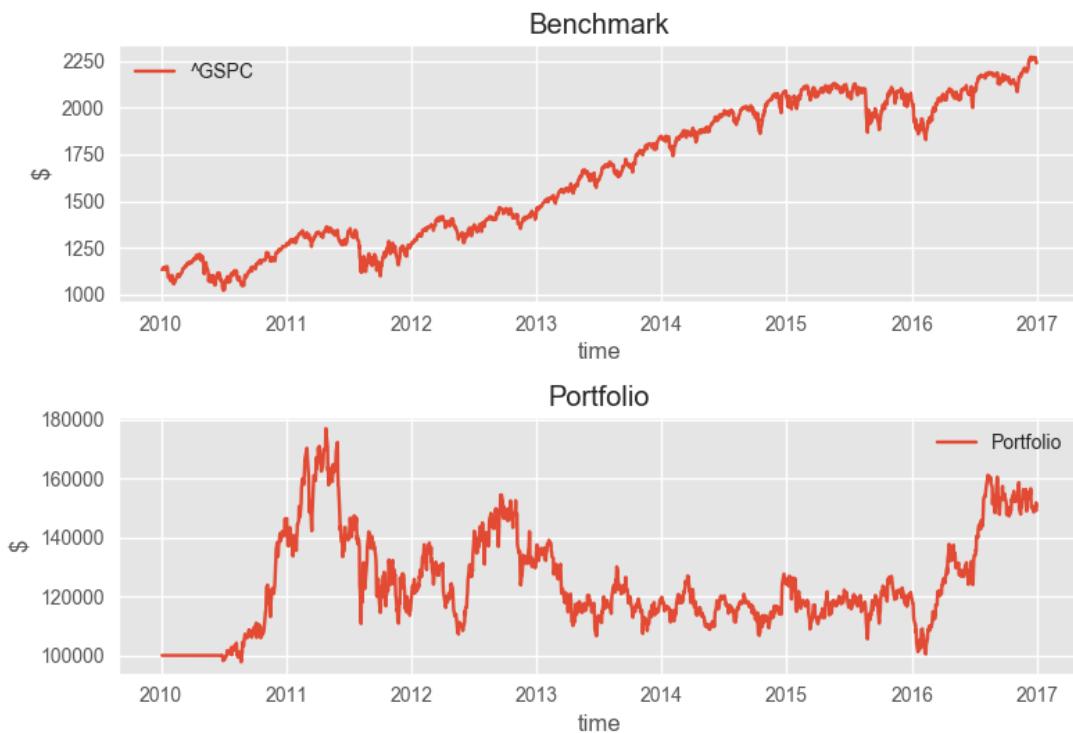


FIGURE 4.40: Machine Learning Classifier strategy with only upwards forecasts

Starting Capital	\$100,000
Total Capital Used	\$76,853.53
Sharpe Ratio	-0.548
Portfolio Value	\$20,134.519
Algorithm Period Return	-0.799
Benchmark Period Return	1.008
Algorithm Volatility	0.323
Benchmark Volatility	0.156

TABLE 4.13: Machine Learning Classifier strategy with upwards and downwards forecasts

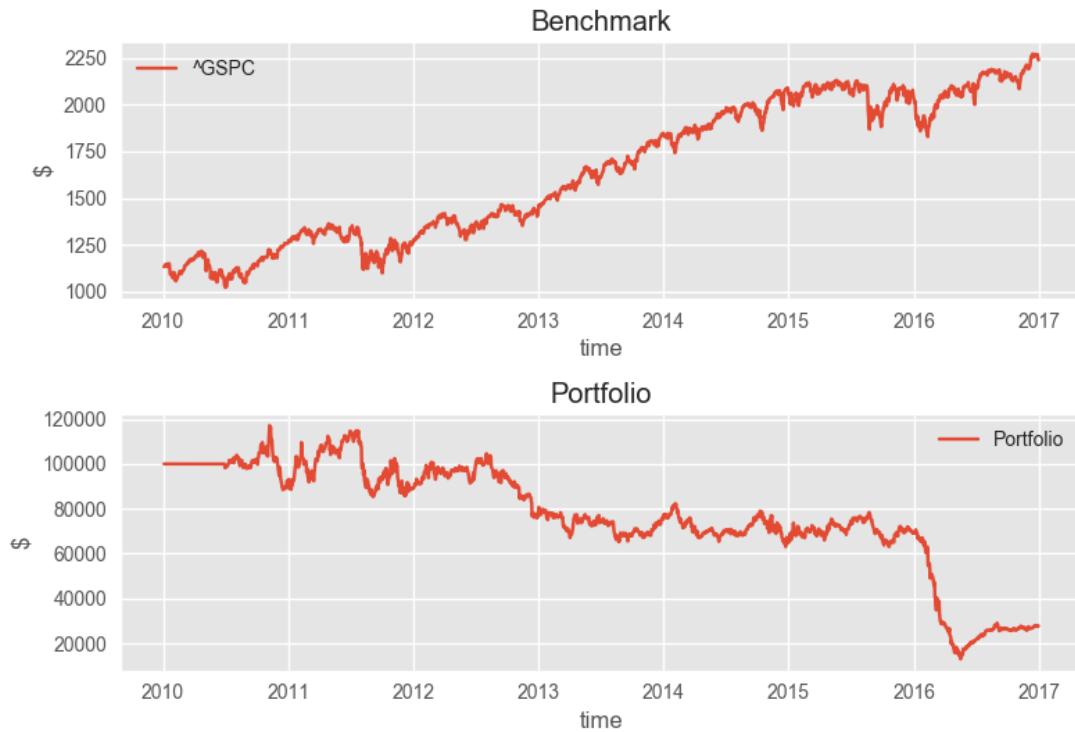


FIGURE 4.41: Machine Learning Classifier strategy with upwards and downwards forecasts

Starting Capital	\$100,000
Total Capital Used	\$295,844.66
Sharpe Ratio	0.385
Portfolio Value	\$167,474.009
Algorithm Period Return	0.675
Benchmark Period Return	1.008
Algorithm Volatility	0.393
Benchmark Volatility	0.156

TABLE 4.14: Machine Learning Classifier strategy with stop loss

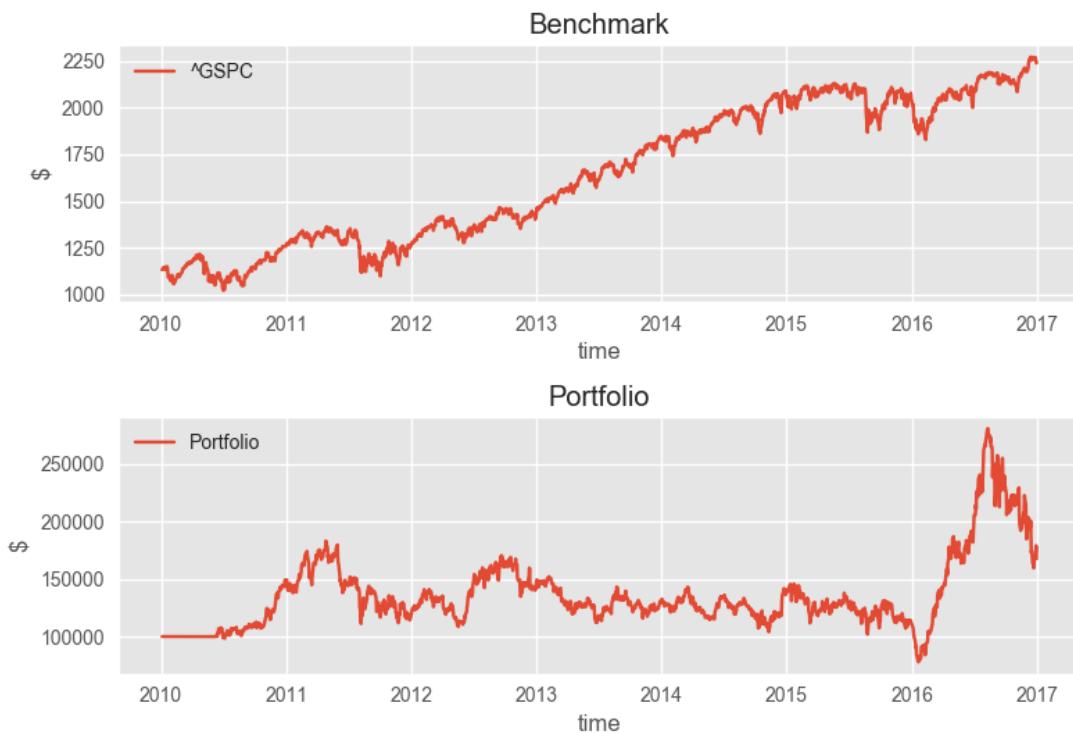


FIGURE 4.42: Machine Learning Classifier strategy with stop loss

4.4.2 Regression

Starting Capital	\$100,000
Total Capital Used	\$229,547.84
Sharpe Ratio	0.420
Portfolio Value	\$183,714.616
Algorithm Period Return	0.837
Benchmark Period Return	1.008
Algorithm Volatility	0.373
Benchmark Volatility	0.156

TABLE 4.15: Machine Learning Regression strategy with only upwards forecasts

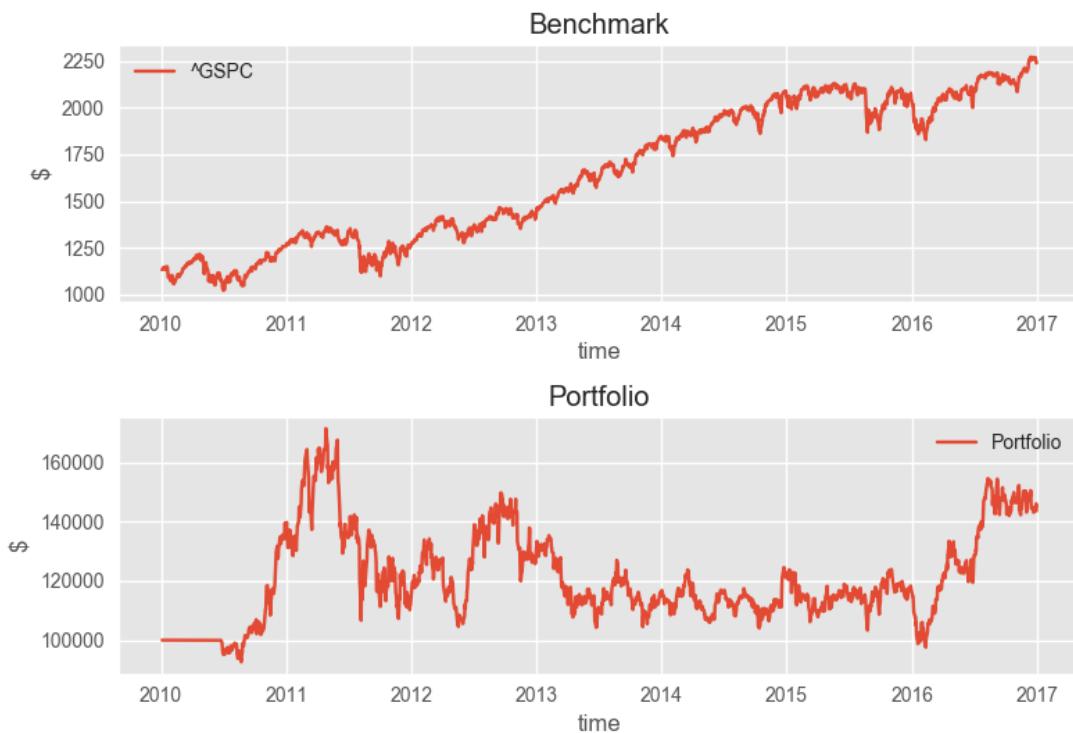


FIGURE 4.43: Machine Learning Regression strategy with only upwards forecasts

Starting Capital	\$100,000
Total Capital Used	\$80,121.51
Sharpe Ratio	-0.258
Portfolio Value	\$19,628.34
Algorithm Period Return	-0.804
Benchmark Period Return	1.008
Algorithm Volatility	0.493
Benchmark Volatility	0.156

TABLE 4.16: Machine Learning Regression strategy with upwards and downwards forecasts

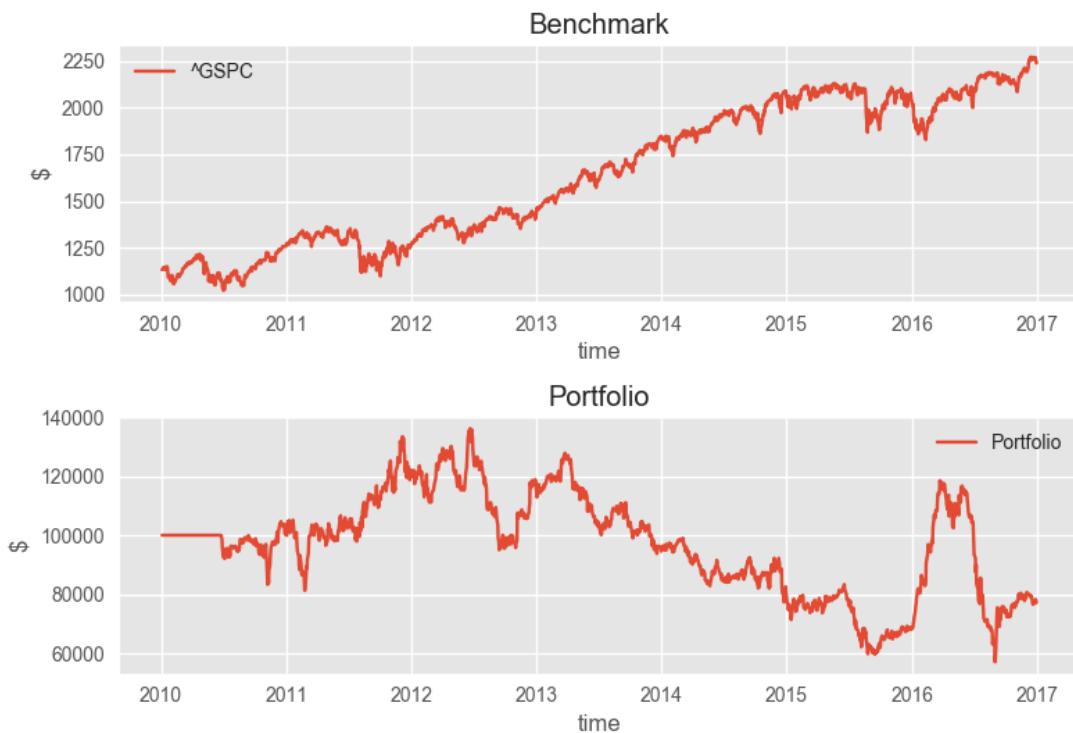


FIGURE 4.44: Machine Learning Regression strategy with upwards and downwards forecasts

Starting Capital	\$100,000
Total Capital Used	\$229,547.84
Sharpe Ratio	0.420
Portfolio Value	\$183,714.616
Algorithm Period Return	0.837
Benchmark Period Return	1.008
Algorithm Volatility	0.373
Benchmark Volatility	0.156

TABLE 4.17: Machine Learning Regression strategy with stop loss

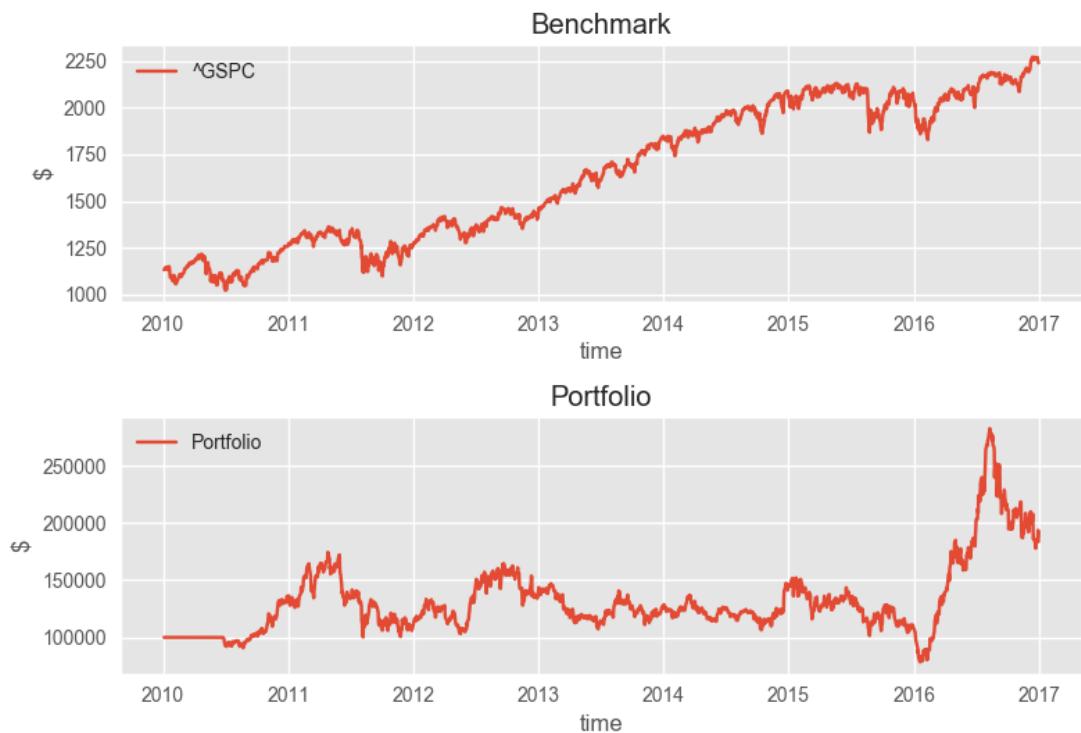


FIGURE 4.45: Machine Learning Regression strategy with stop loss

Chapter 5

Discussion and Suggestions for Future Research

5.1 Time Series Analysis

5.1.1 Random Walk

The random walk theory suggests that stock price changes have the same distribution and are independent of each other, so the past movement or trend of a stock price or market cannot be used to predict its future movement. In short, this is the idea that stocks take a random and unpredictable path.

As can be seen in figure 4.2, the correlation plots show this theory to be false and that past movement is related to future movement.

The histogram of returns show that the stocks follow a normal distribution, where some stocks showed higher returns than others as can be seen from a wider x-axis.

5.1.2 Ordinary Least Squares (OLS)

This algorithm achieved a good fit in the in-sample tests, having a prediction rate above 96% for all the stocks tested. It is evident that OLS may be used for stock market forecasting.

5.1.3 Auto Regressive (AR)

This algorithm failed to achieve a good fit in the in-sample tests, having a huge difference in sharpe ratios based on the original price returns and in-sample predicted price returns. This is also evident in the QQ and probability plots, both of which show heavy tails indicating a bad fit to the data. The algorithm failed completely in forecasting price returns in out-of-sample testing.

The histogram of returns further showed that the algorithm did not achieve a good fit, as the histogram itself was not completely symmetrical.

5.1.4 Moving Average (MA)

This algorithm faired better than AR having a lower difference in sharpe ratios based on the original price returns and in-sample predicted price returns, however still failed to achieve a good fit in the in-sample tests. The QQ and probability plots, although containing lighter tails than the previous algorithm, still show that the algorithm in question failed to achieve a good fit.

Although achieving a better fit, the histogram of returns still showed flaws as they were not perfectly symmetrical.

5.1.5 Auto Regressive Moving Average (ARMA)

This algorithm once again faired better than the previous algorithm, MA, having an even lower difference in sharpe ratios based on the original price returns and in-sample predicted price returns, however still failed to achieve a good fit in the in-sample tests. As can be seen in the time series analysis plots, ARMA showed to have very heavy tails in the QQ and probability plots. The algorithm predicted an abnormal sharp decline at the start of the forecast and showed diminishing returns over time, as was evident in figure 4.17.

As can be seen from the histogram of returns, a better fit was achieved as the histogram was more symmetrical than those of the previous algorithms.

5.1.6 Auto Regressive Integrated Moving Average (ARIMA)

ARIMA faired worse than the previous algorithm, ARMA, having a higher difference in sharpe ratios based on the original price returns and in-sample predicted price returns.

This was also evident in the QQ and probability plots, both of which showed to have heavy tails.

This can also be seen from the histogram of returns, in which its distribution was less normal than that of the previous algorithm.

5.2 Machine Learning

5.2.1 Classification

5.2.1.1 Decision Tree

This algorithm faired well in predicting stock price movements in the in-sample tests, having an accuracy score above 75% for all the stocks prices that were forecast.

5.2.1.2 Boosted Decision Tree

This algorithm yielded slightly better results when compared to the previous algorithm in predicting stock price movements in the in-sample tests.

5.2.1.3 Support Vector Machine (SVM)

This algorithm once again yielded slightly better results when compared to the previous algorithm in predicting stock price movements in the in-sample tests.

5.2.1.4 Random Forest

This algorithm scored the same accuracy as the previous algorithm.

5.2.1.5 K-Nearest Neighbour

This algorithm's accuracy was not consistent as it scored mixed results when tested against all the stocks, where the accuracy is some was much lower than others.

5.2.1.6 Logistic Regression

This algorithm faired worse than the others having not achieved an accuracy score of above 80% when tested against all the stocks.

5.2.1.7 Bernoulli Naive Bayes

This algorithm faired the worst from the algorithms tested so far, scoring accuracy scores in the lower 70th percentile.

5.2.1.8 Gaussian Naive Bayes

This algorithm also scored a lower accuracy score, showing that Naive Bayes isn't the best form of forecasting stock prices.

5.2.1.9 Neural Network

This algorithm, although did not score the best accuracy score from the other algorithms, still faired well in correctly predicting the price movements of the stocks.

5.2.1.10 Stochastic Gradient Descent

This algorithm faired the worst from the lot, varying heavily in accuracy scores, in which one stock had an accuracy score of 59%. Even though it faired much better in correctly predicting the price movements of other stocks, this algorithm proved to be unreliable.

5.2.2 Regression

5.2.2.1 Decision Tree

As can be seen in figure 4.22, the algorithm not only failed to achieve a good fit in the in-sample test, but also failed completely in predicting any values at all. Figure 4.23 also shows that the algorithm predicted a sharp fall at the end of the out-of-sample test, which could be related to the incorrect fit achieved in the in-sample test. This was reflected in the metric score, having a high number of data losses, and low accuracy scores.

5.2.2.2 Boosted Decision Tree

This algorithm also did not do very well, as can be seen in figure 4.24, the algorithm not only failed to achieve a good fit in the in-sample test, but also failed completely in predicting any values at all. The algorithm also predicted some sharp declines at the

end of the test, as can be seen in figure 4.25. The bad fit achieved is also reflected in the metric scores, in which there was a high number of data losses, and low accuracy scores.

5.2.2.3 K-Nearest Neighbour

This algorithm also didn't fair too well, showing a bad fit in figure 4.26. Similarly to the previous algorithm, the alforithm failed completely in predicting any values at all in some parts of the in-sample test.

5.2.2.4 Random Forest

Similarly to previous algorithms, this algorithm failed to achieve a good fit as can be seen in figure 4.28. The algorithm also failed to predict any values at all in the of the in-sample tests, and also predicted sharp inclines and declines at the end some of the out-of-sample tests.

5.2.2.5 Linear Regression

This algorithm performed very well in the in-sample tests, achieving a good fit with high accuracy scores above 95% with a low number of data losses.

5.2.2.6 Neural Network

This algorithm, although scored very well in some of the in-sample tests, was not very reliable as accuracy scores varied between each test, performing extremely well in some, however not too well in others.

5.2.2.7 Stochastic Gradient Descent

This algorithm performed fairly well in most in-sample tests, having high accuracy scores and low data losses escept for CDE as can be seen in figure 4.34.

5.3 Bayesian Statistics

5.3.0.1 Metropolis-Hastings

This algorithm failed to achieve a good fit in the in-sample tests, having a large difference in sharpe ratios based on the original price returns and in-sample predicted price returns.

5.3.0.2 No-U-Turn-Sampler (NUTS)

The algorithm achieved a good fit in the in-sample tests, having very similar sharpe ratios based on the original price returns and in-sample predicted price returns.

5.4 Strategy

5.4.1 Classification

When tasked with predicting rises in stock prices, the algorithm did fairly well, marking a total return of 49.1%, this however did not beat the benchmark's return of 100.8%. The algorithm underperformed immensely when tasked at also predicting stock price falls, marking a negative total return of -80%. To compensate for this, a stop loss was added to sell all positions if the stock in question falls below -20%, this resulted in a profit of 67.5%.

5.4.2 Regression

The performance of this algorithm was extremely similar to that of the previous one, this goes to show that both classifiers and regressors can be used for stock price prediction. When tasked with predicting rises in stock prices, the algorithm did fairly well, marking a slightly lower total return of 43%. The algorithm also underperformed immensely when tasked at also predicting stock price falls, marking a negative total return of -80.4%. To compensate for this, a stop loss was added to sell all positions if the stock in question falls below -20%, this resulted in a profit of 83.7%.

5.5 Future Research

Although this study provides important insights into the relationship between machine learning models and stock market forecasting, the models still failed to achieve a higher

return than the market, thus discrediting them as a viable option for investors to incorporate in their strategy. It is however the firm belief of the author that machine learning models can indeed be used to exploit the market, using methods which were not applied to this study due to a lack of data and resources.

Since only technical analysis was applied due to restrictions in data, a link may indeed be found between fundamental analysis and machine learning models, a relationship that may be well worth looking into. Many claim fundamental analysis to be extremely critical to financial analysis, providing a vast amount of strategies to implement. Since this type of analysis involves financial statements, such an increase in data could easily boost the predictive power of machine learning models.

Deep learning is also currently a growing field, making strides in many forms of applications. Being a subset of machine learning, it works similarly to other models except that its inner workings are designed to mimic the decision making of the human brain. However due to its requirements of a vast amount of data and powerful hardware, it was not possible to include it in this study. Just as deep learning has proven to perform extremely well in other applications, it may very well prove to do so in this field too.

Sentiment analysis is also another field which is currently proving to be a powerful tool in predicting outcomes. This form of analysis involves computationally identifying and categorising opinions expressed in a piece of text, especially in order to determine whether the writer's attitude towards a particular topic, is positive, negative, or neutral. This could be easily used to identify consumer sentiment which is a powerful economic indicator of the overall health of the economy as determined by consumer opinion. This could make it very easy to determine which stocks are being overvalued or undervalued, and quickly act on that information.

Chapter 6

Conclusion

Three financial forecasting methods were presented in this report, two of which showed little to no potential of ever producing any statistically significant result when the correct methodology was applied. The third method, machine learning, showed some potential in the tests carried out, which is why this method was built into an automated algorithmic strategy to trade with. The algorithm proved to be successful in forecasting future prices, using both classification and regression methods. However, the backtesting proved this method to fail in forecasting price falls. Once this factor was removed from the equation, the algorithms were very successful and reported a profit by the end of the test. This is however not always ideal as stocks which could fall in price could be catastrophic to the strategy. A stop loss would be ideal in insuring that no positions are held in downward falling stocks. It was also evident that regression methods were more successful in forecasting future price movements when compared to classification methods.

If there is anything that this report shows, is that profitable stock market prediction is an extremely tough problem. Even though the strategies reported a profit by the end of the backtest, they still did not beat the market. Whether it is at all possible to use such methods to outperform the market's returns, ultimately remains an open question. These findings support the Efficient Market Hypothesis, proving that casual investors are better off investing in passive buy and hold strategies consisting of index funds and ETFs. However, there was some evidence found showing that the Random Walk Hypothesis does not hold true for all cases, as some stocks did show signs of repeating trends.

Appendix A

An Appendix

A.1 Time Series Analysis

A.1.1 Random Walk

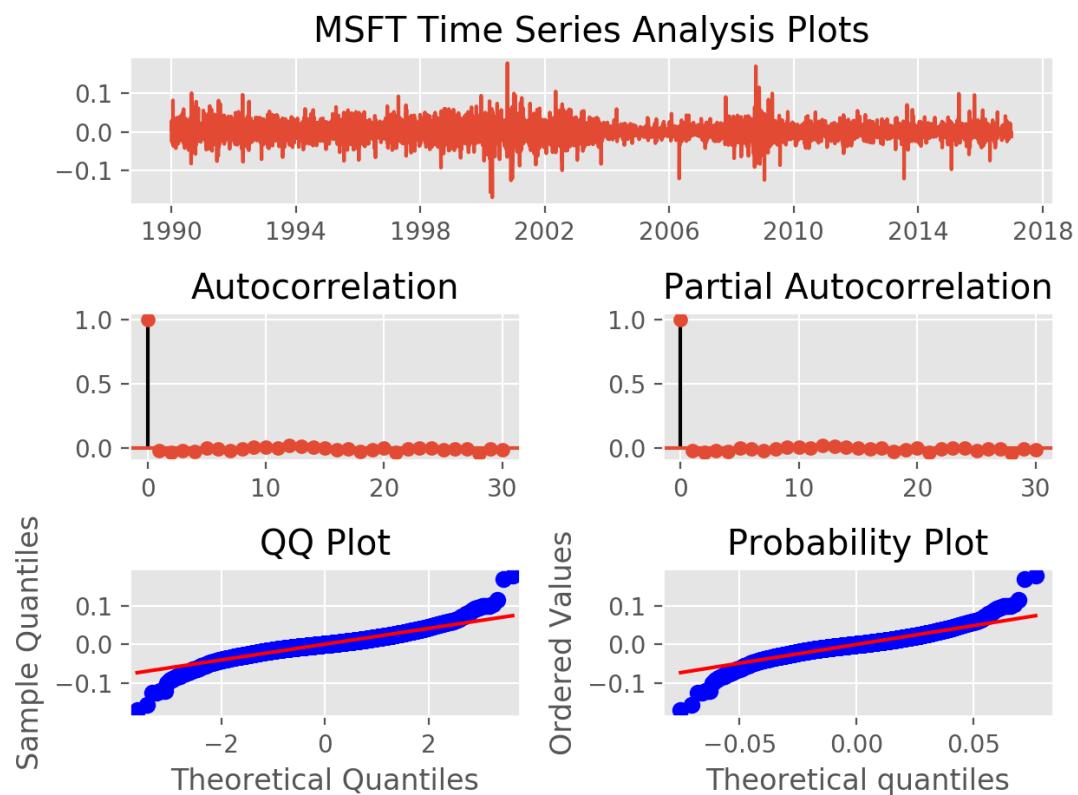


FIGURE A.1: MSFT time series analysis

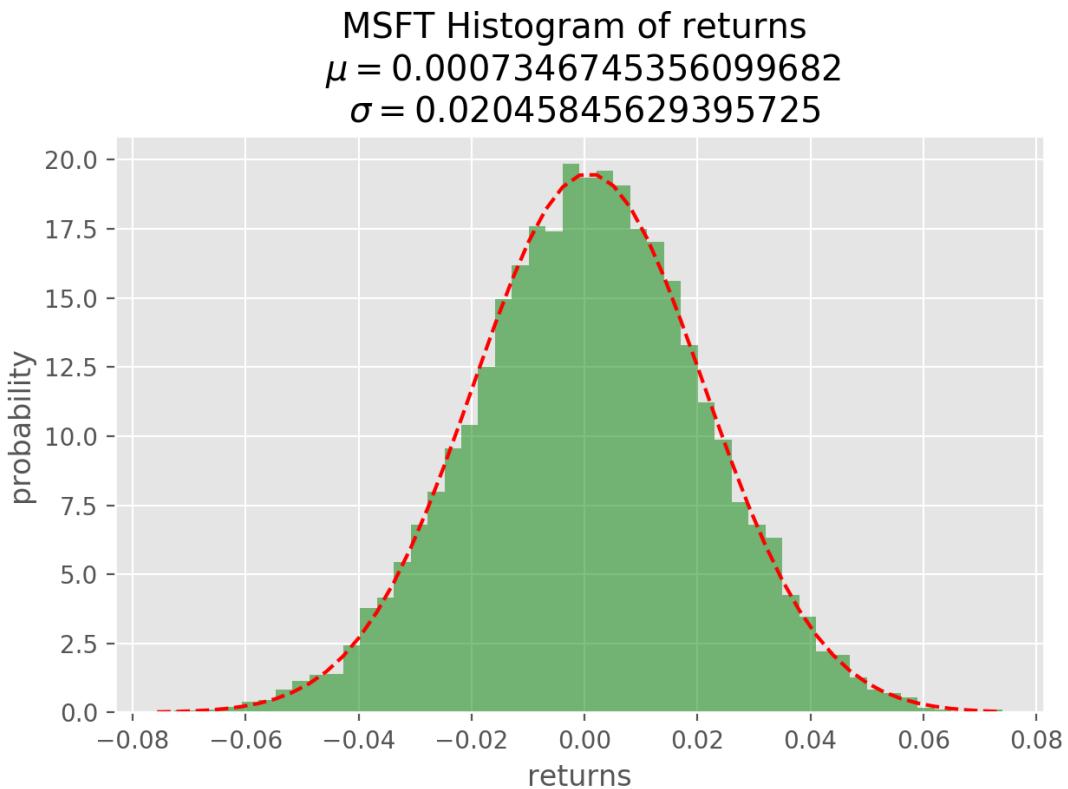


FIGURE A.2: MSFT histogram of returns

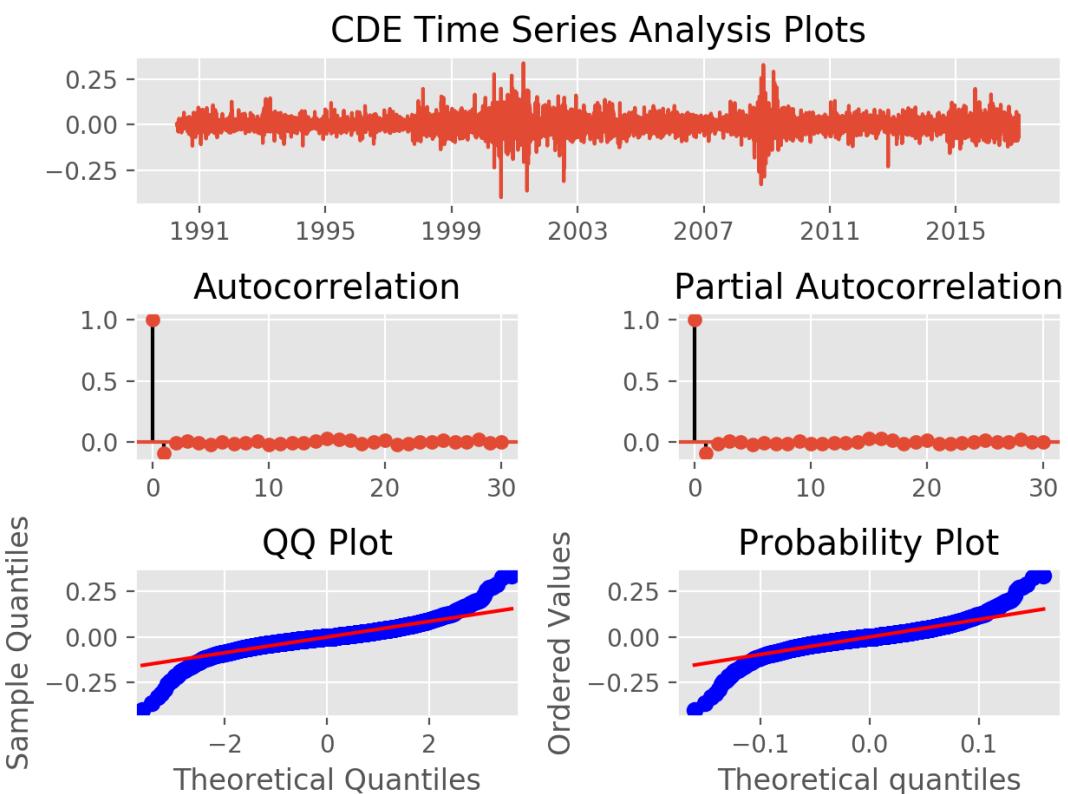


FIGURE A.3: CDE time series analysis

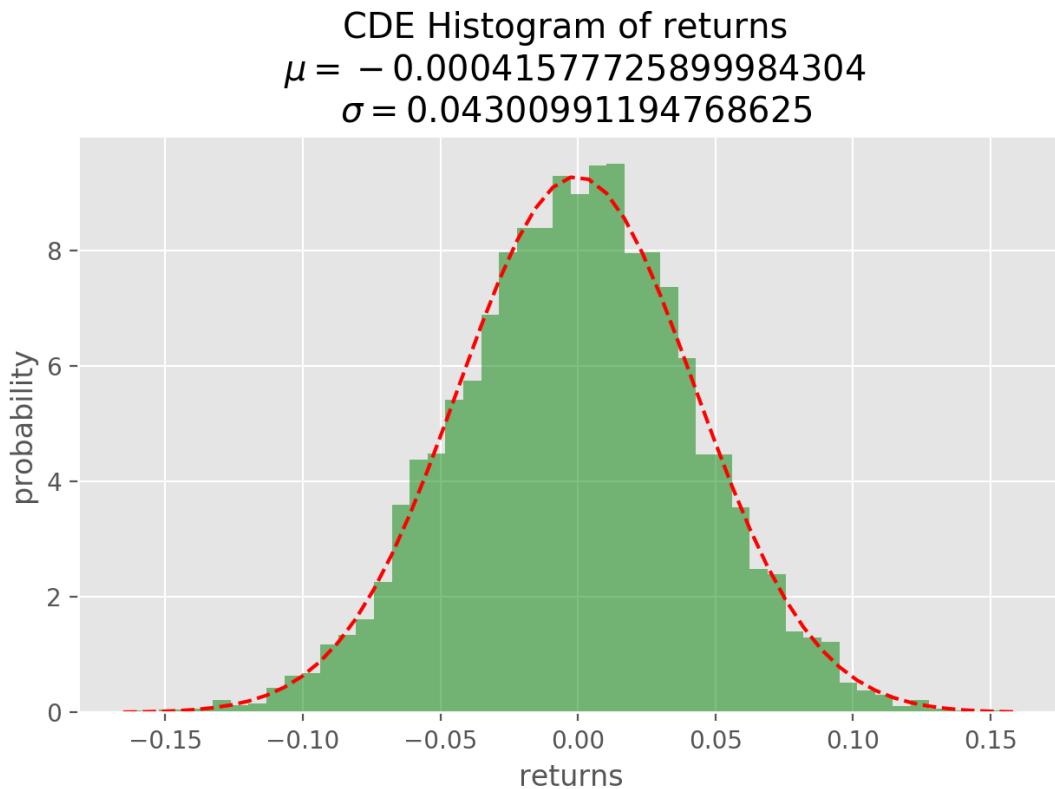


FIGURE A.4: CDE histogram of returns

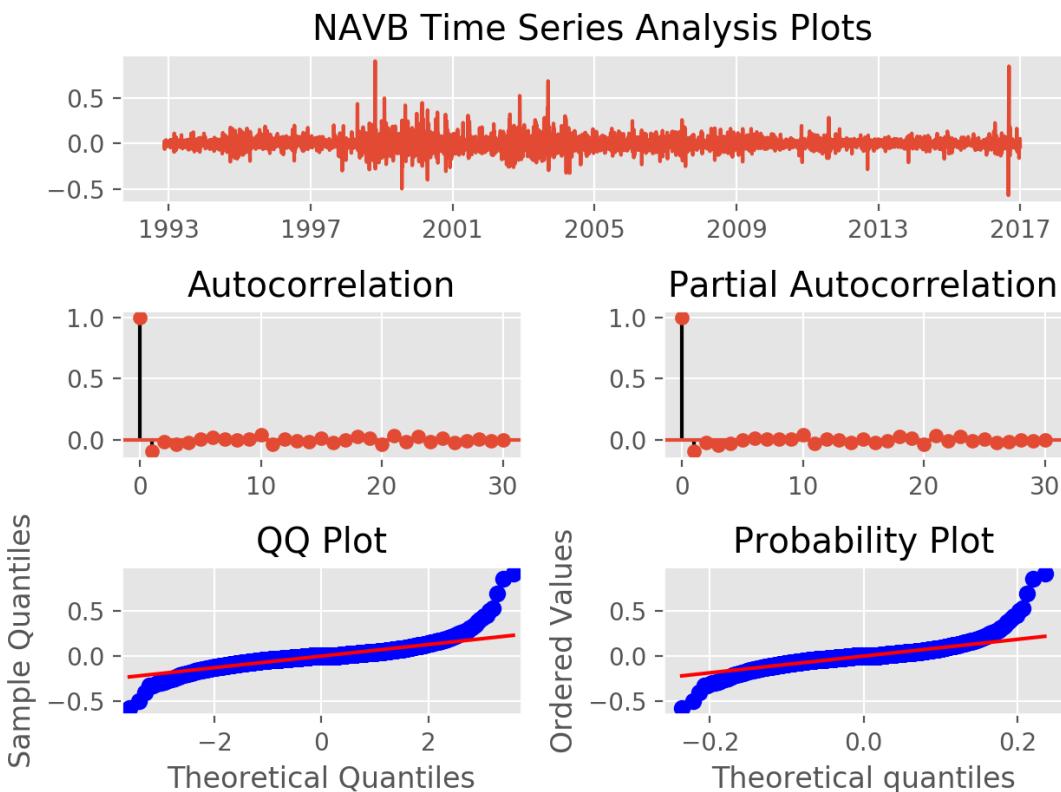


FIGURE A.5: NAVB time series analysis

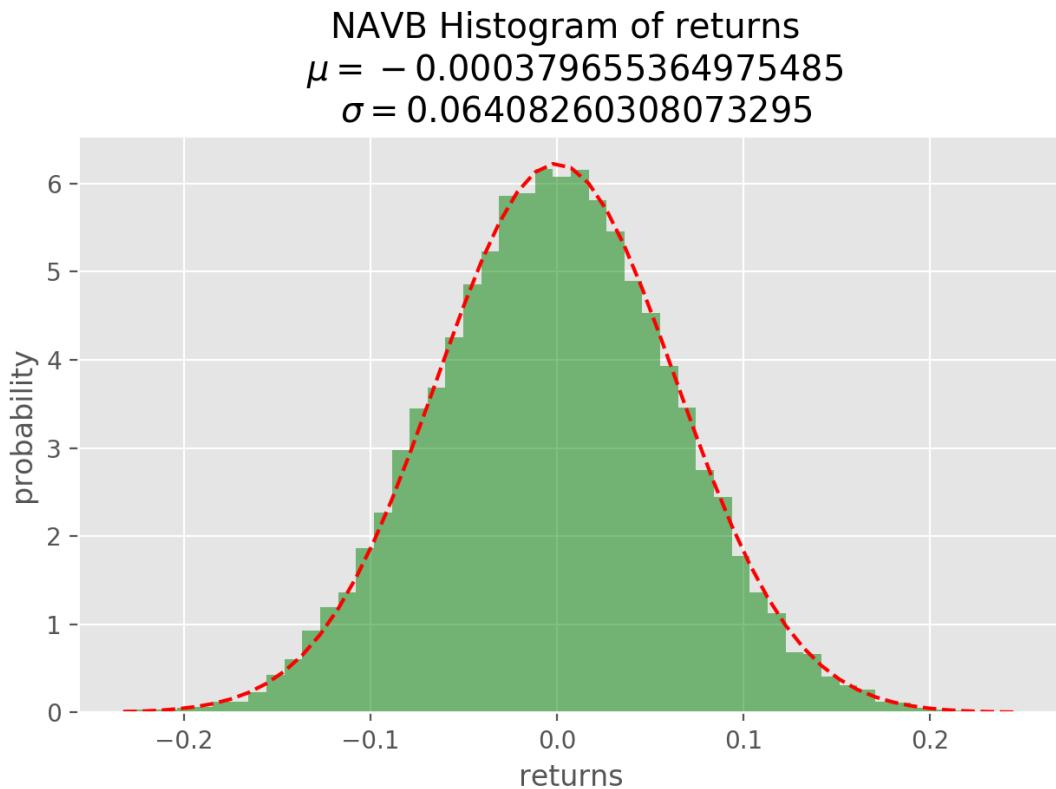


FIGURE A.6: NAVB histogram of returns

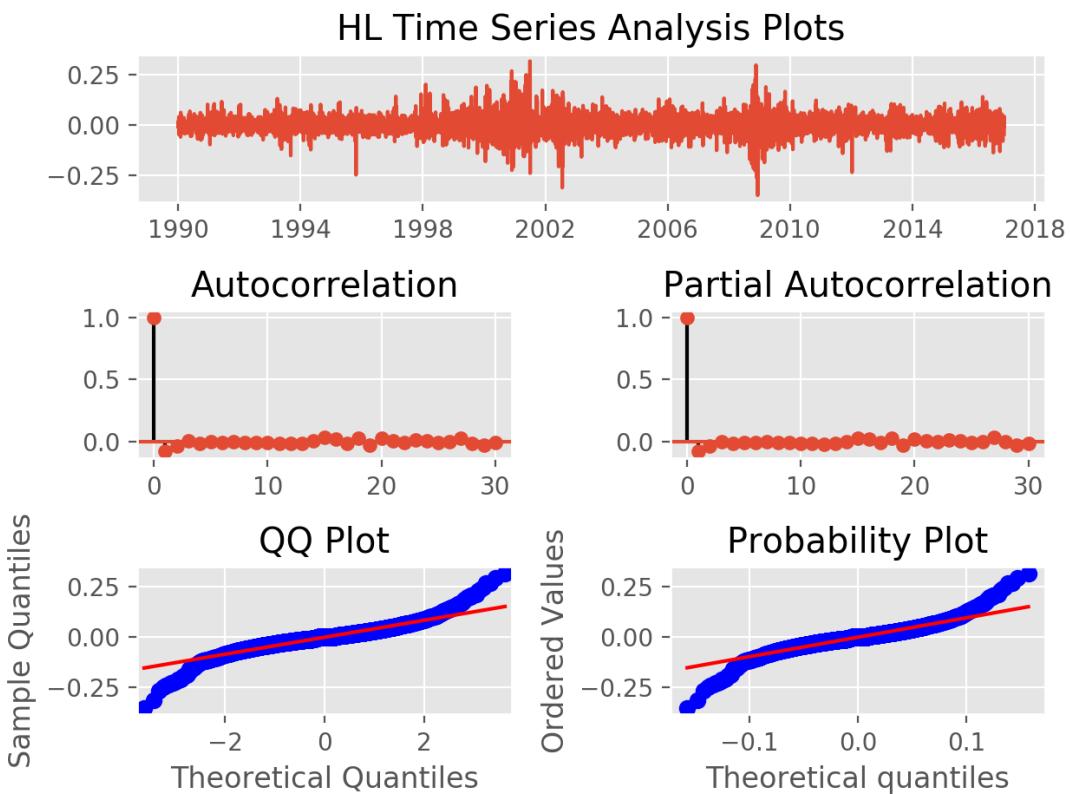


FIGURE A.7: HL time series analysis

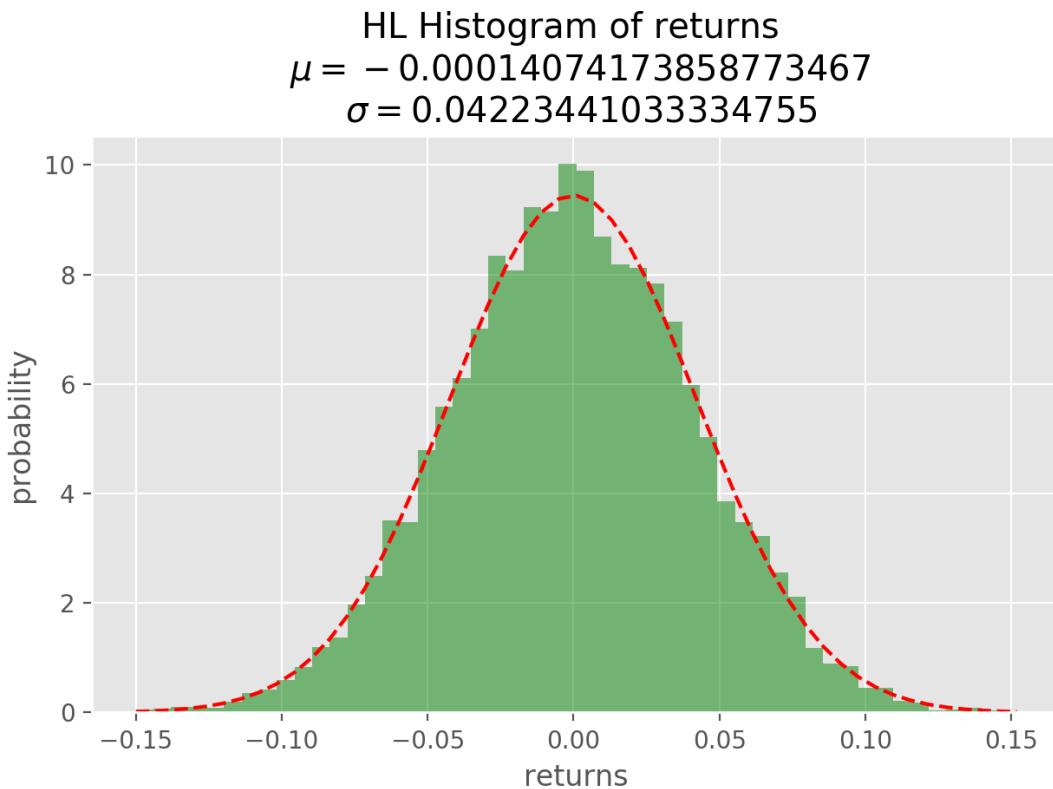


FIGURE A.8: HL histogram of returns

A.1.2 Ordinary Least Squares (OLS)

CDE scored a mean absolute error regression loss of 0.985, and a coefficient of determination of 0.973.

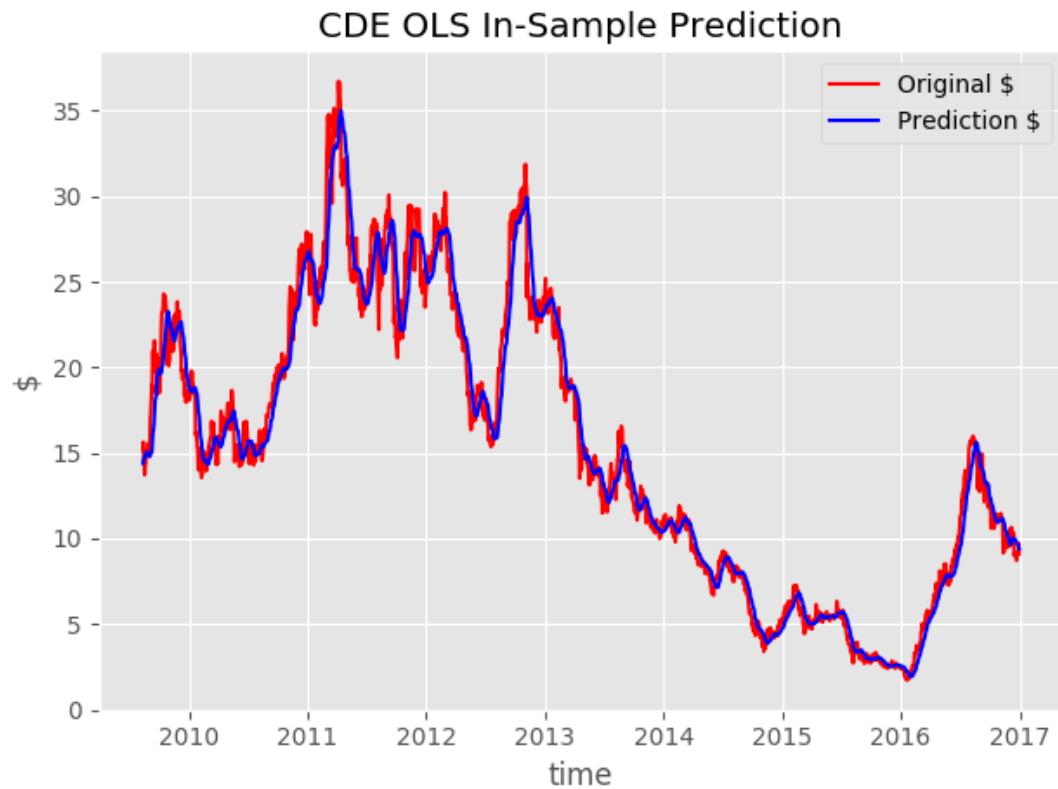


FIGURE A.9: CDE OLS in-sample prediction

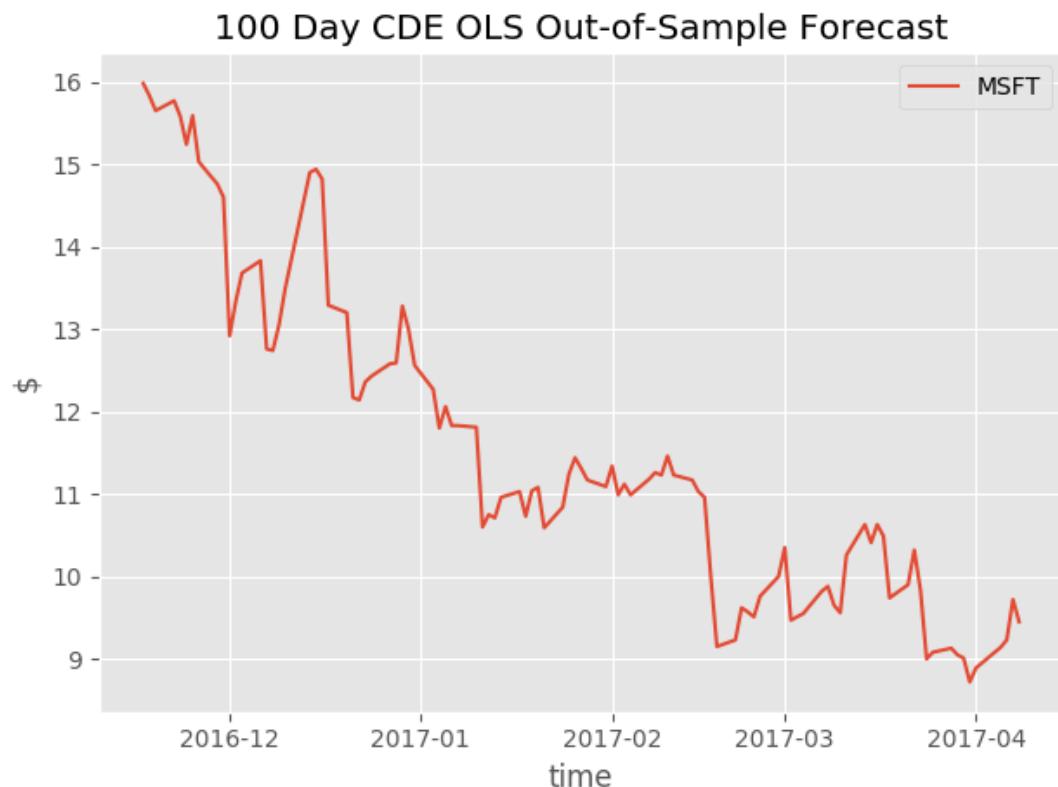


FIGURE A.10: 100 Day CDE OLS out of sample forecast

NAVB scored a mean absolute error regression loss of 0.124, and a coefficient of determination of 0.966.

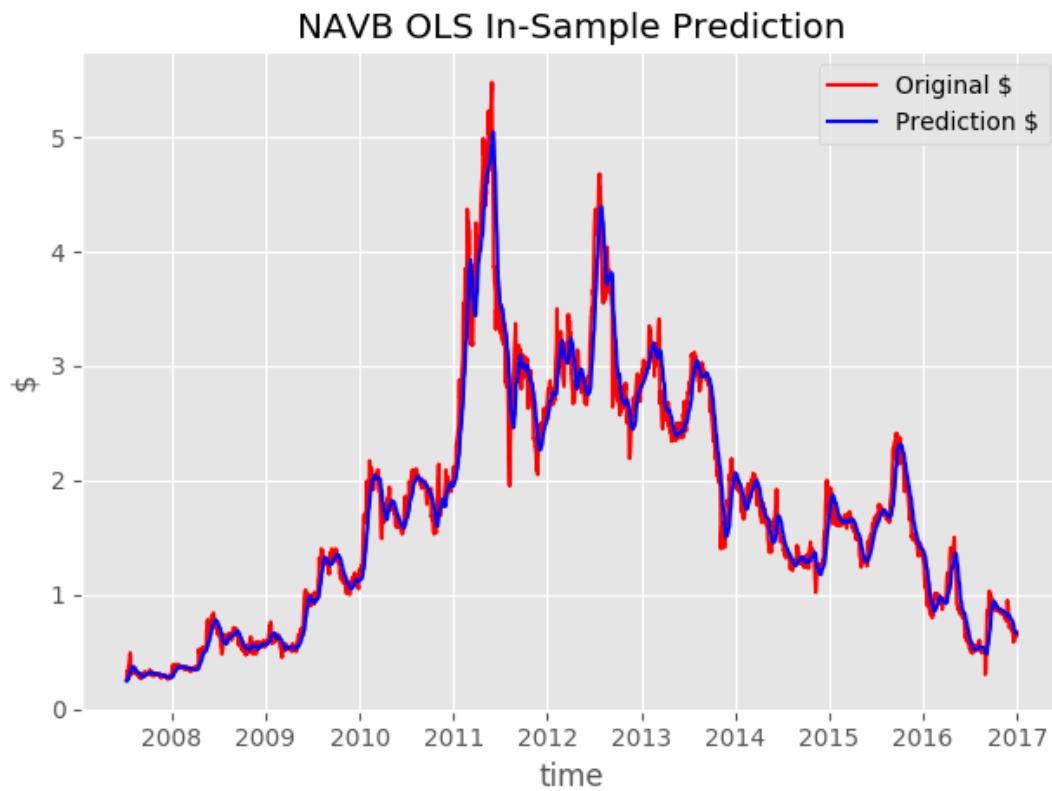


FIGURE A.11: NAVB OLS in-sample prediction

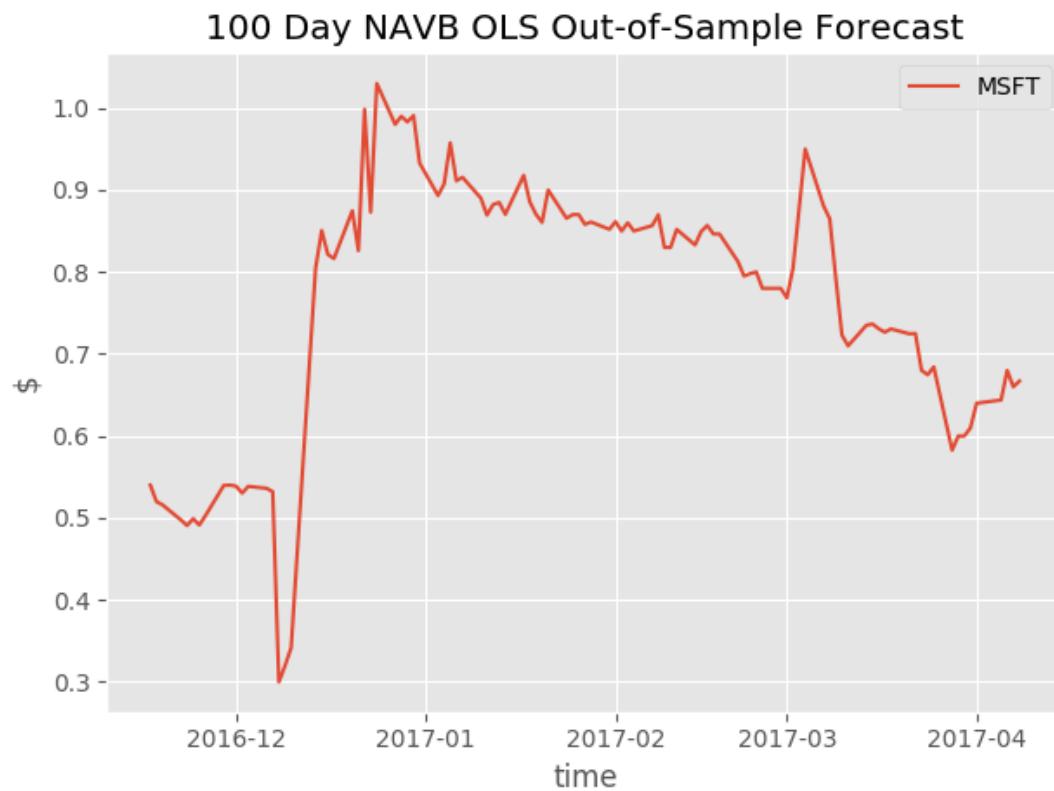


FIGURE A.12: 100 day NAVB OLS in-sample forecast

HRG scored a mean absolute error regression loss of 0.291, and a coefficient of determination of 0.986.

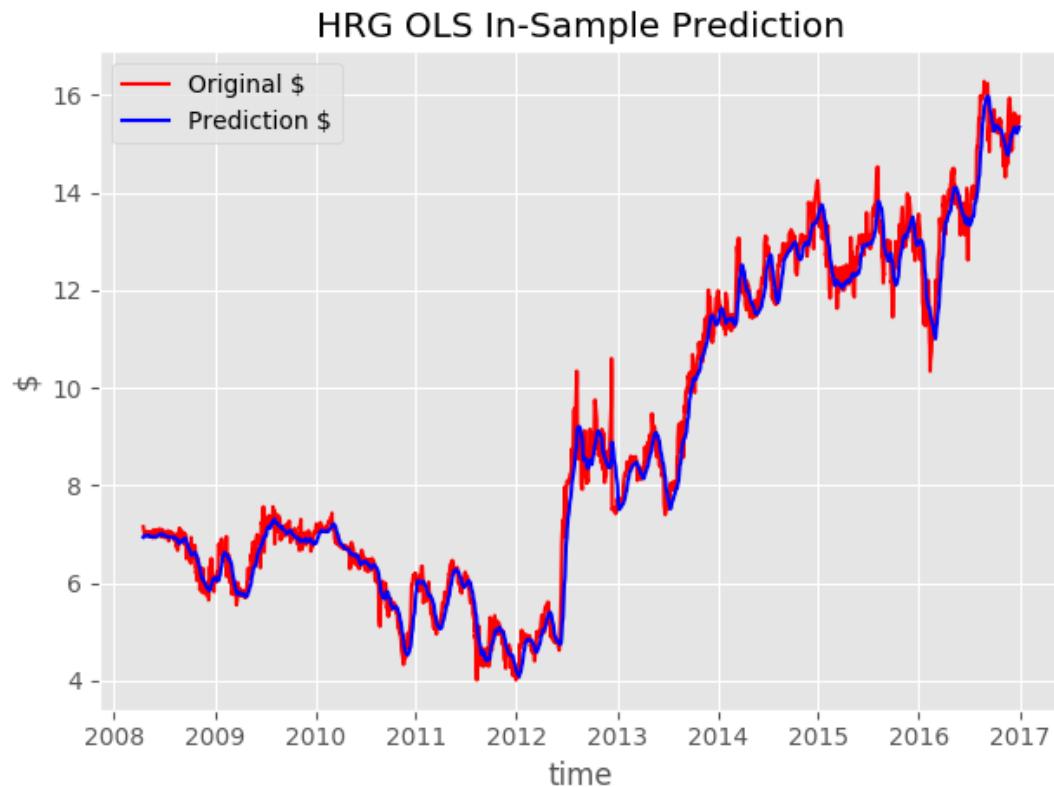


FIGURE A.13: HRG OLS in-sample prediction

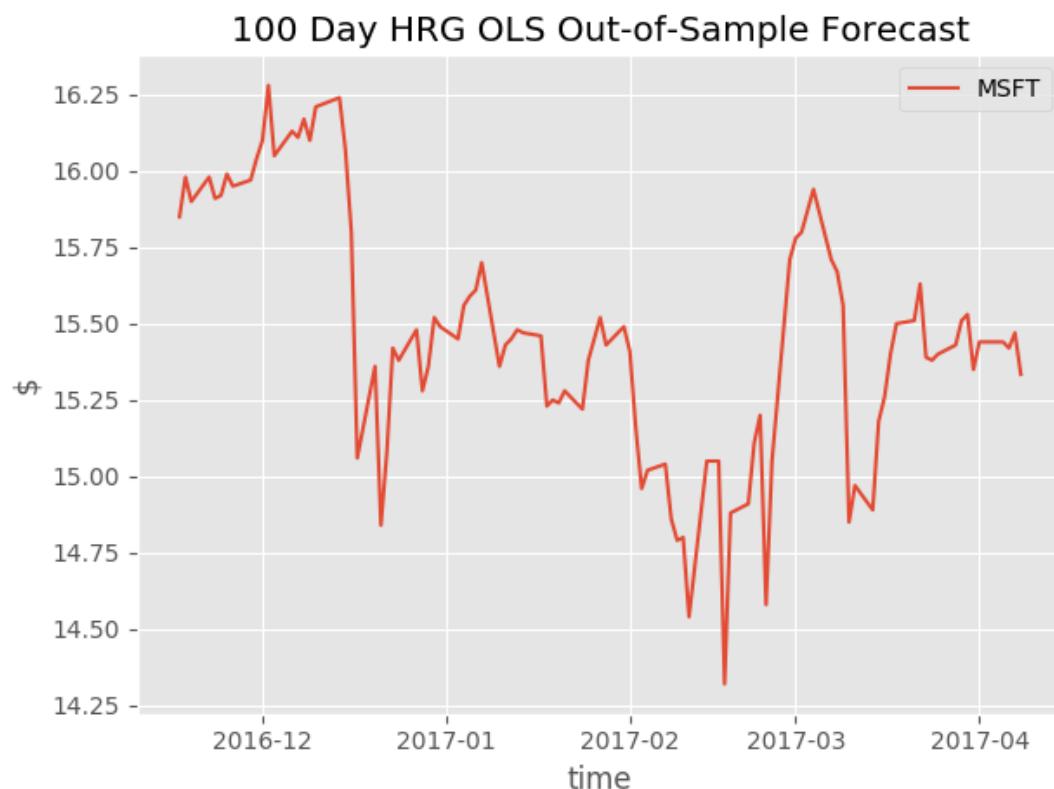


FIGURE A.14: 100 day HRG OLS in-sample forecast

HL scored a mean absolute error regression loss of 0.336, and a coefficient of determination of 0.963.

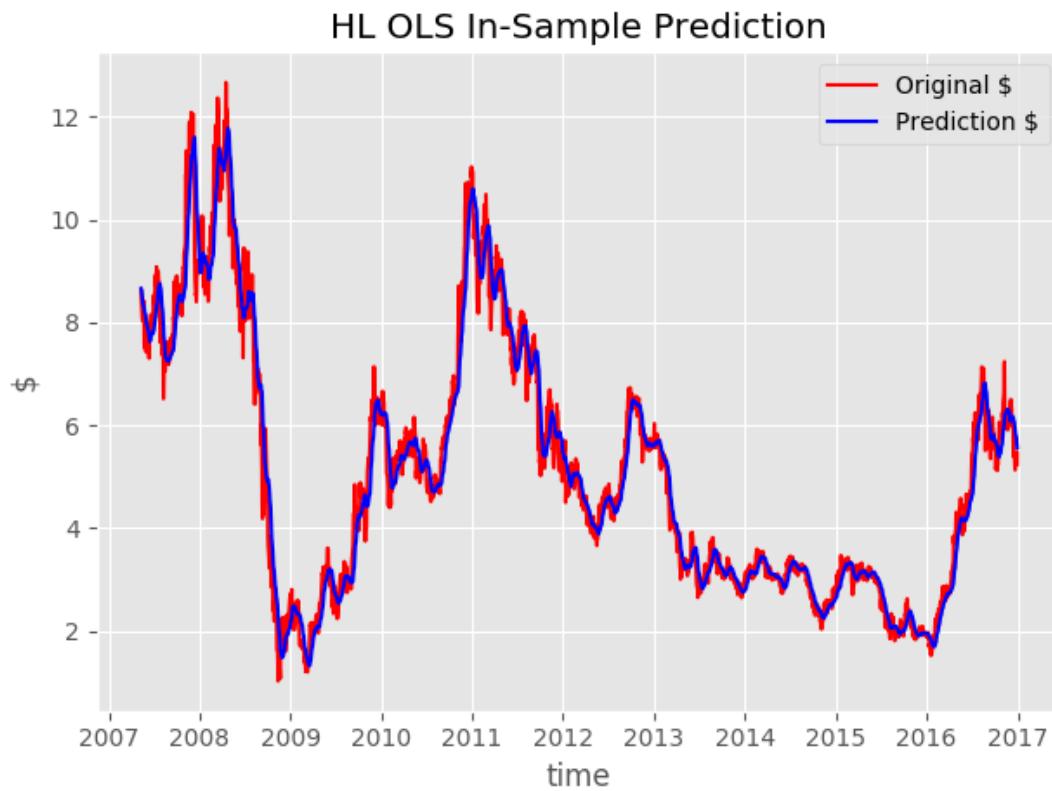


FIGURE A.15: HL OLS in-sample prediction

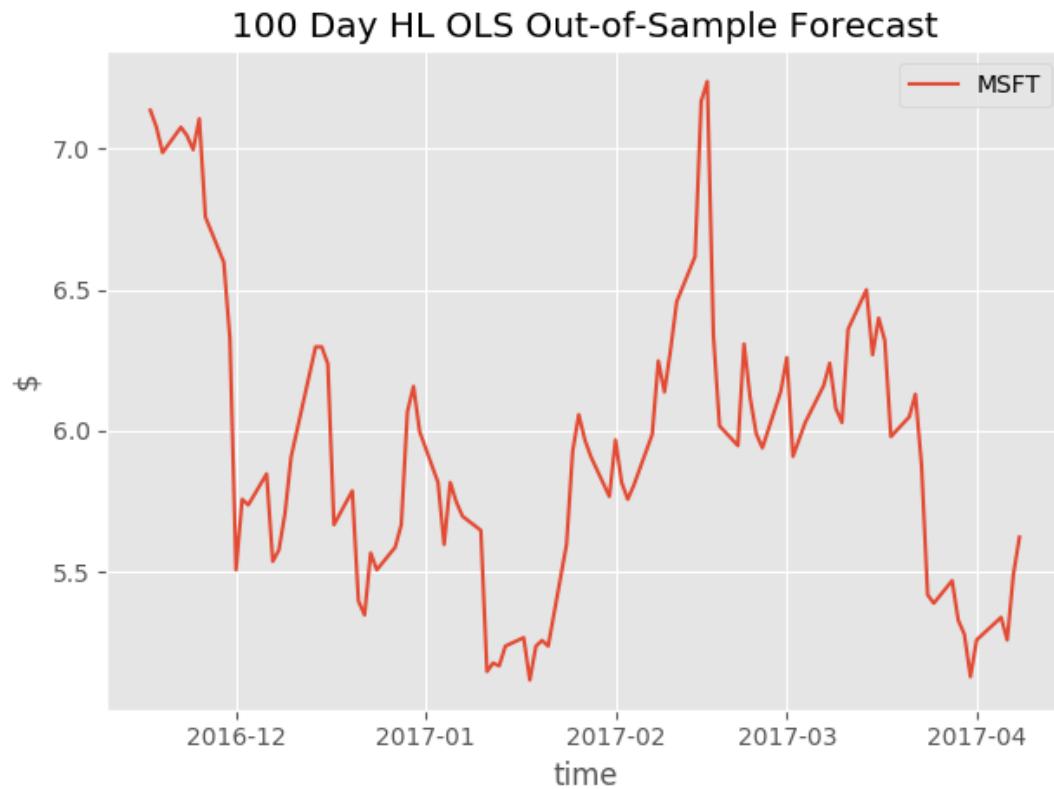


FIGURE A.16: 100 day HL OLS in-sample forecast

A.1.3 Ordinary Least Squares (OLS)

CDE scored a mean absolute error regression loss of 0.985, and a coefficient of determination of 0.973.

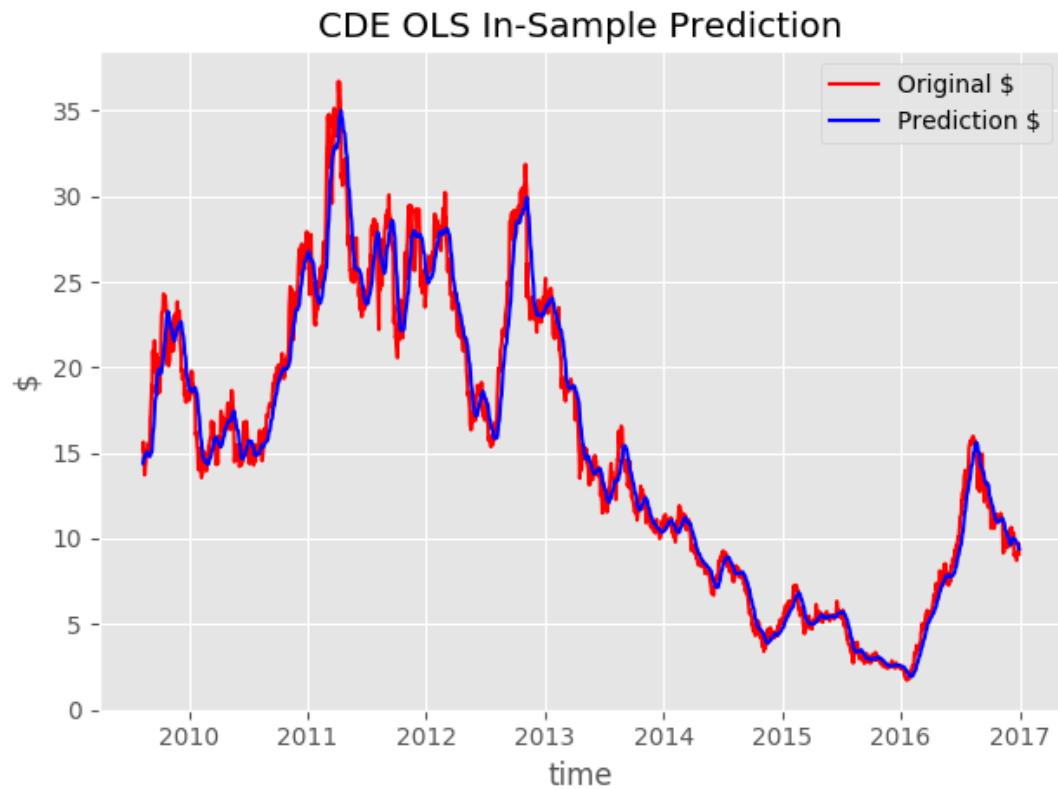


FIGURE A.17: CDE OLS in-sample prediction

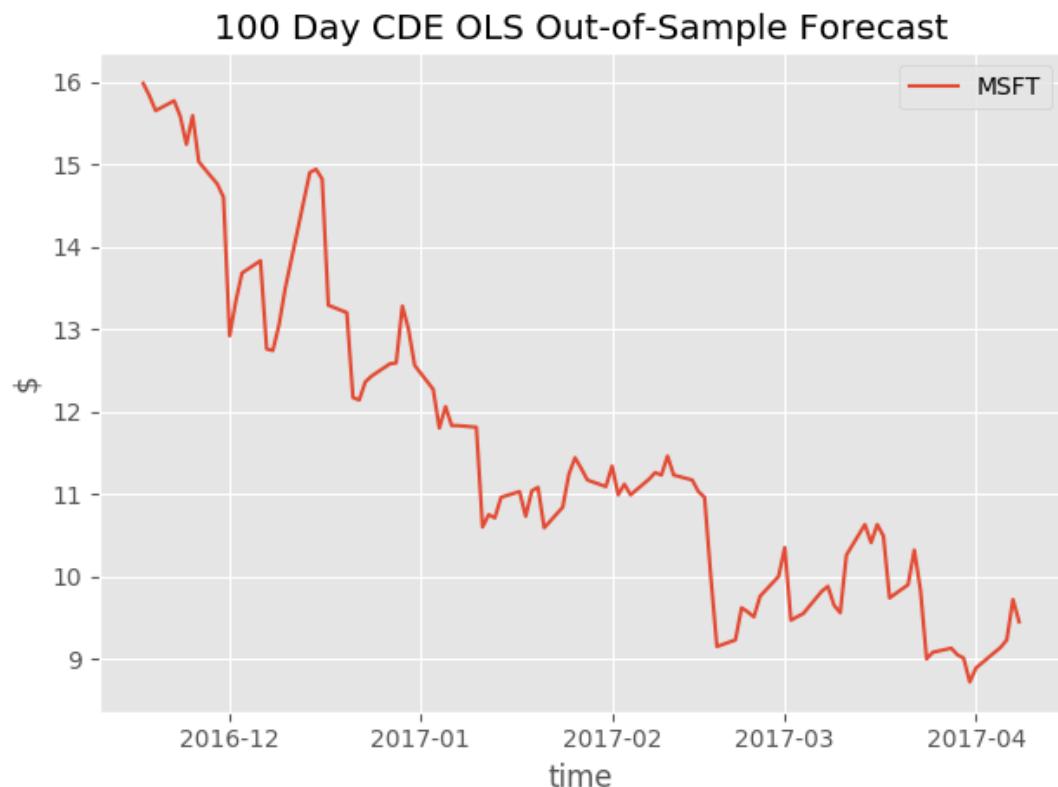


FIGURE A.18: 100 Day CDE OLS out of sample forecast

NAVB scored a mean absolute error regression loss of 0.124, and a coefficient of determination of 0.966.

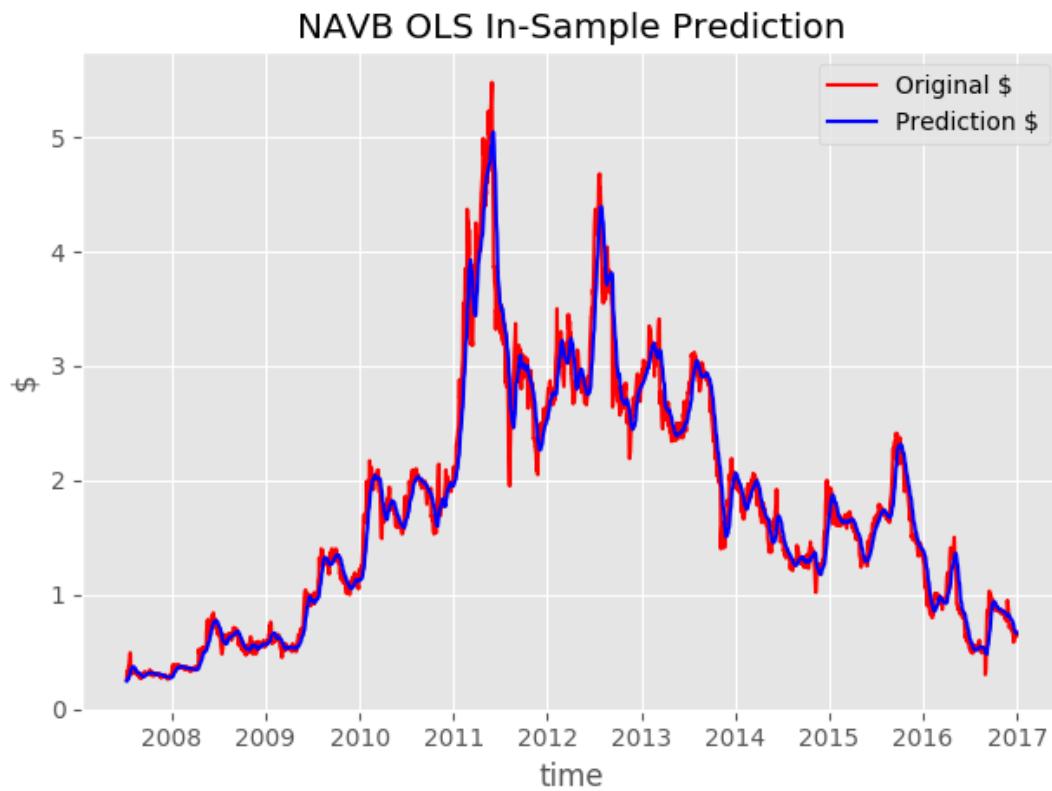


FIGURE A.19: NAVB OLS in-sample prediction

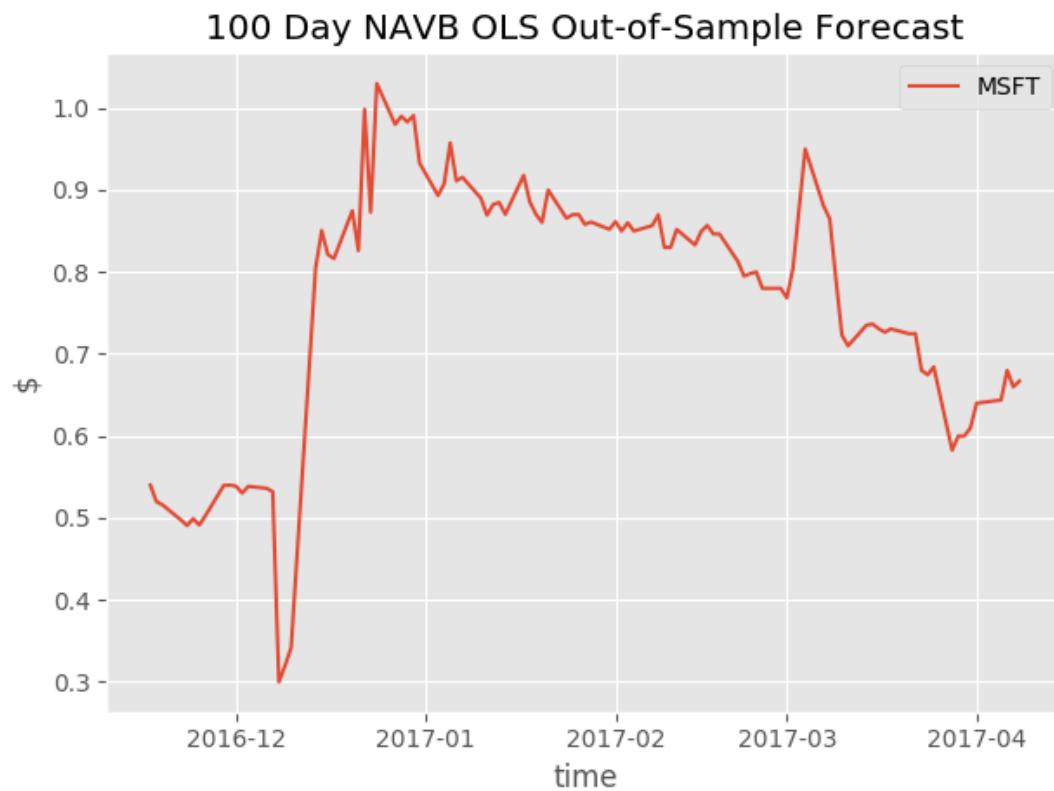


FIGURE A.20: 100 day NAVB OLS in-sample forecast

HRG scored a mean absolute error regression loss of 0.291, and a coefficient of determination of 0.986.

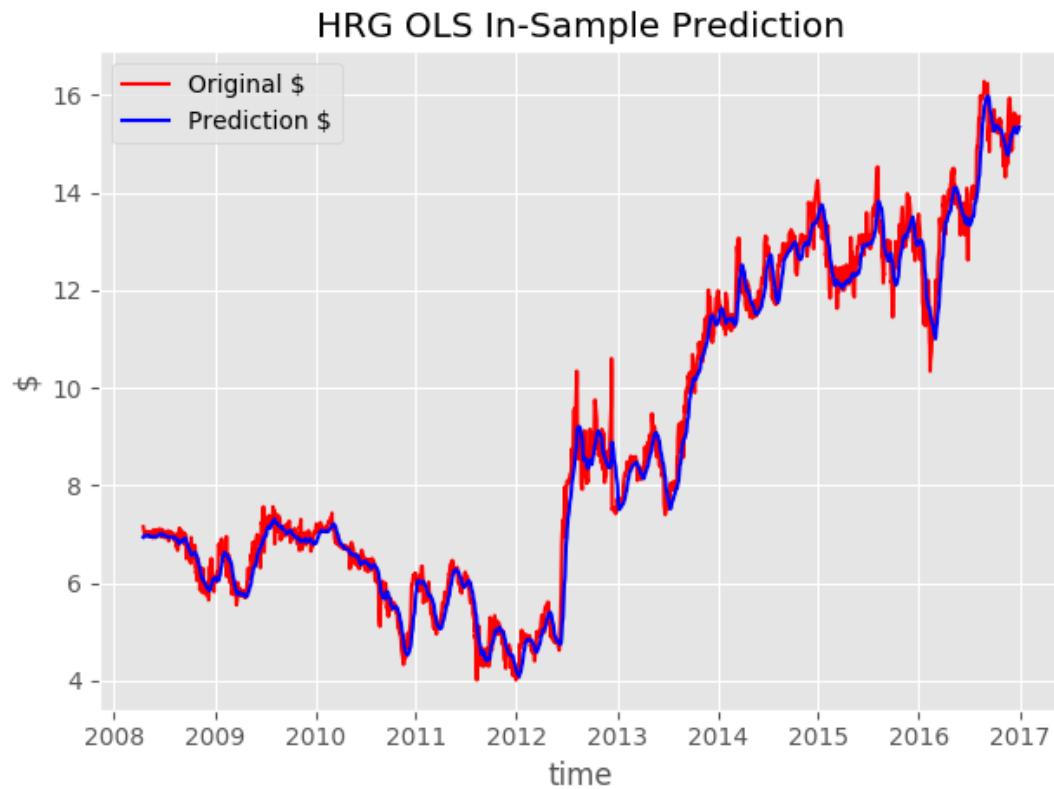


FIGURE A.21: HRG OLS in-sample prediction

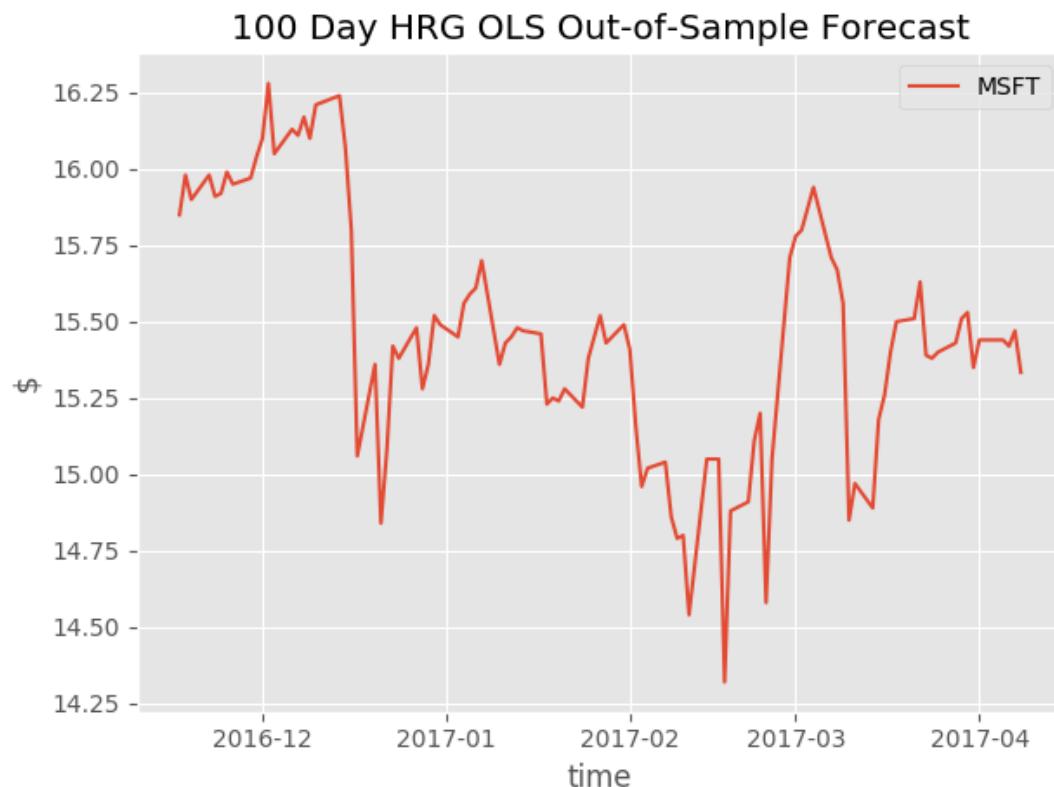


FIGURE A.22: 100 day HRG OLS in-sample forecast

HL scored a mean absolute error regression loss of 0.336, and a coefficient of determination of 0.963.

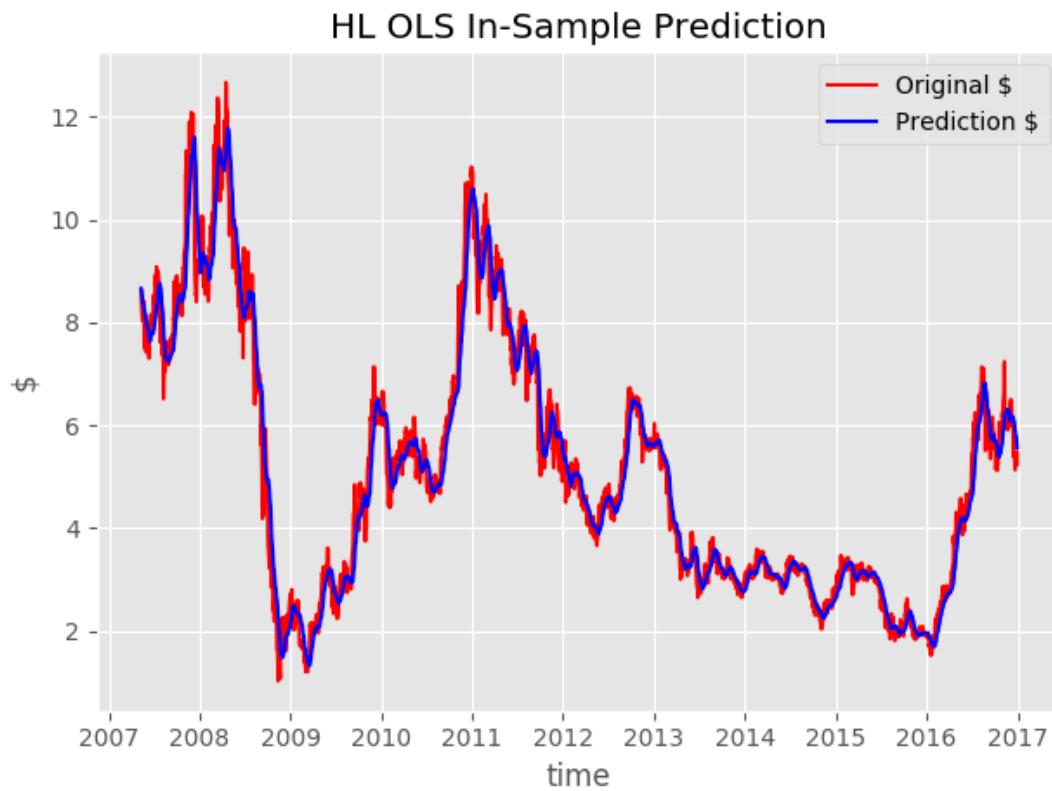


FIGURE A.23: HL OLS in-sample prediction

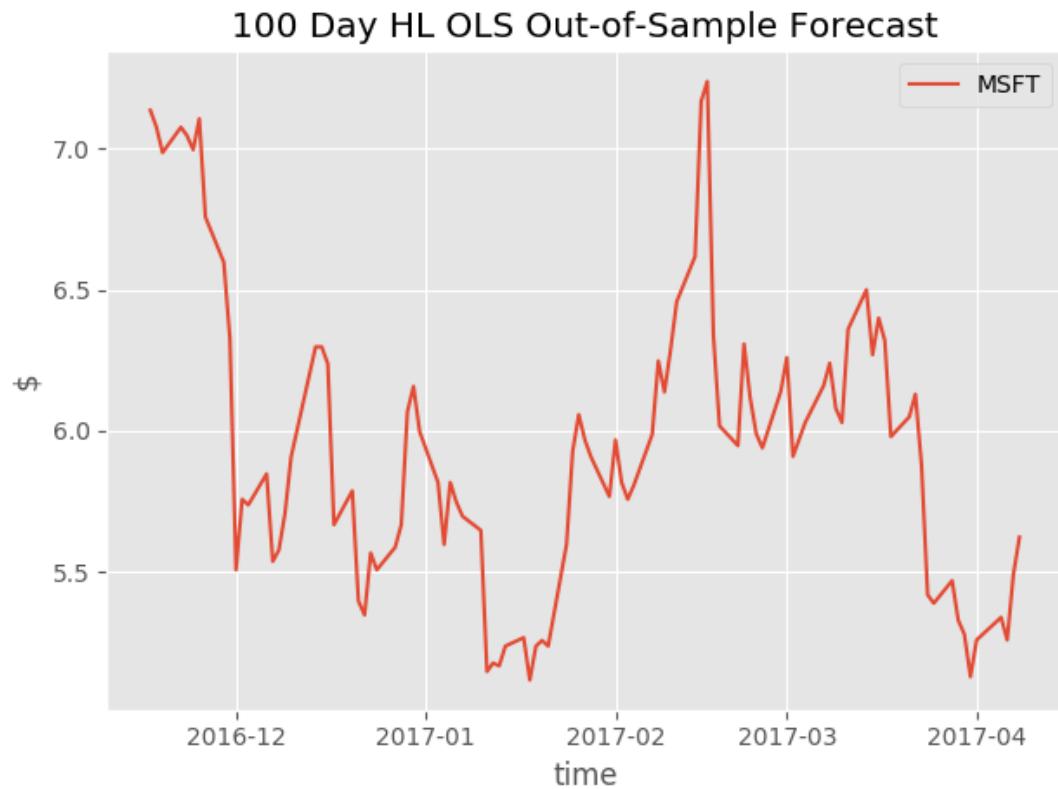


FIGURE A.24: 100 day HL OLS in-sample forecast

A.1.4 Auto Regressive (AR)

CDE scored sharpe ratios of -1.050 for the original returns, and -0.239 for the predicted returns in the in-sample test.

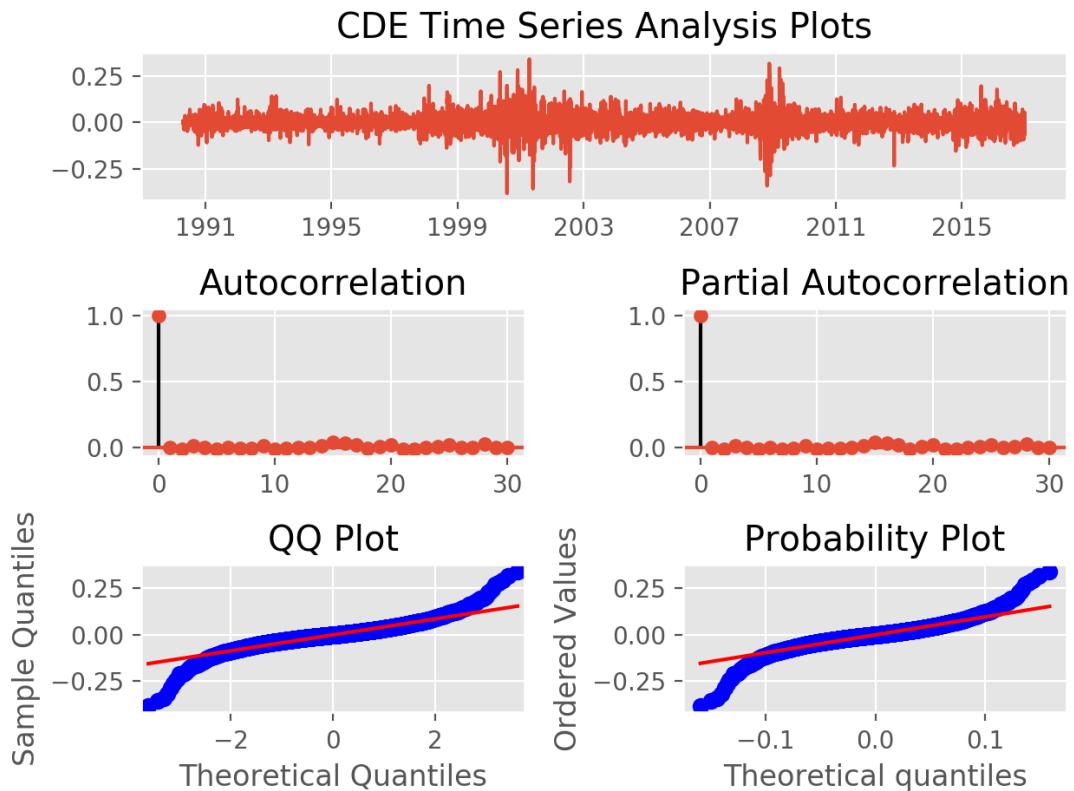


FIGURE A.25: CDE AR time series analysis

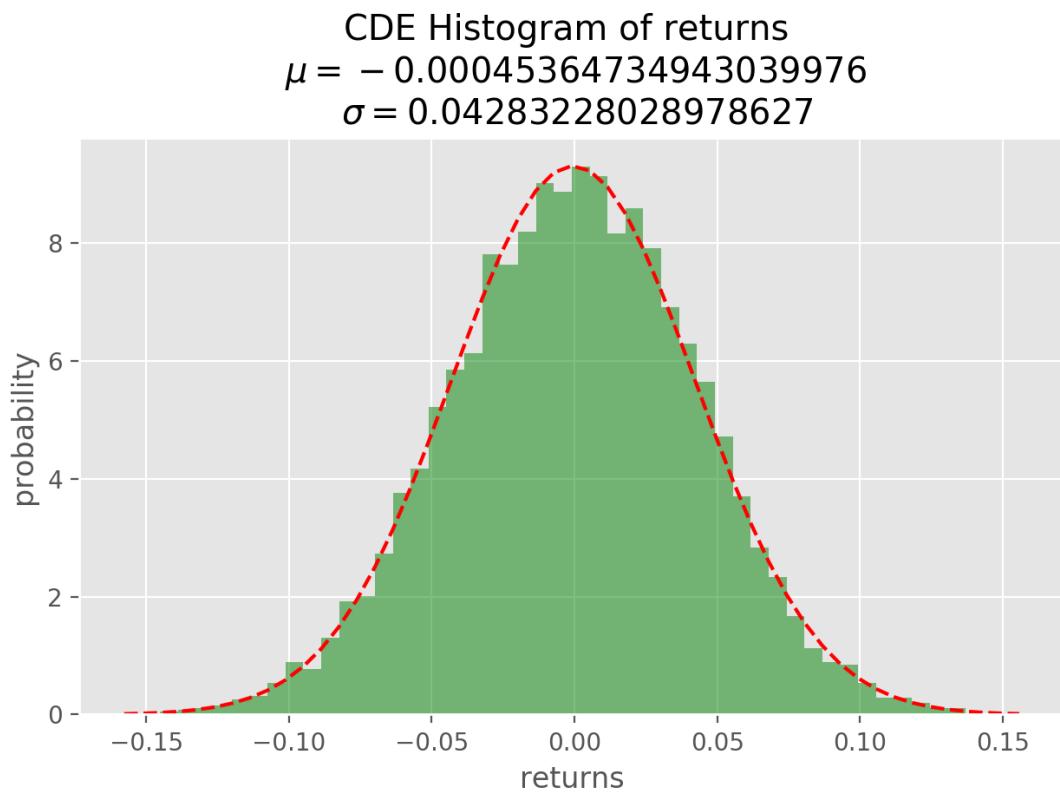


FIGURE A.26: CDE AR histogram of returns

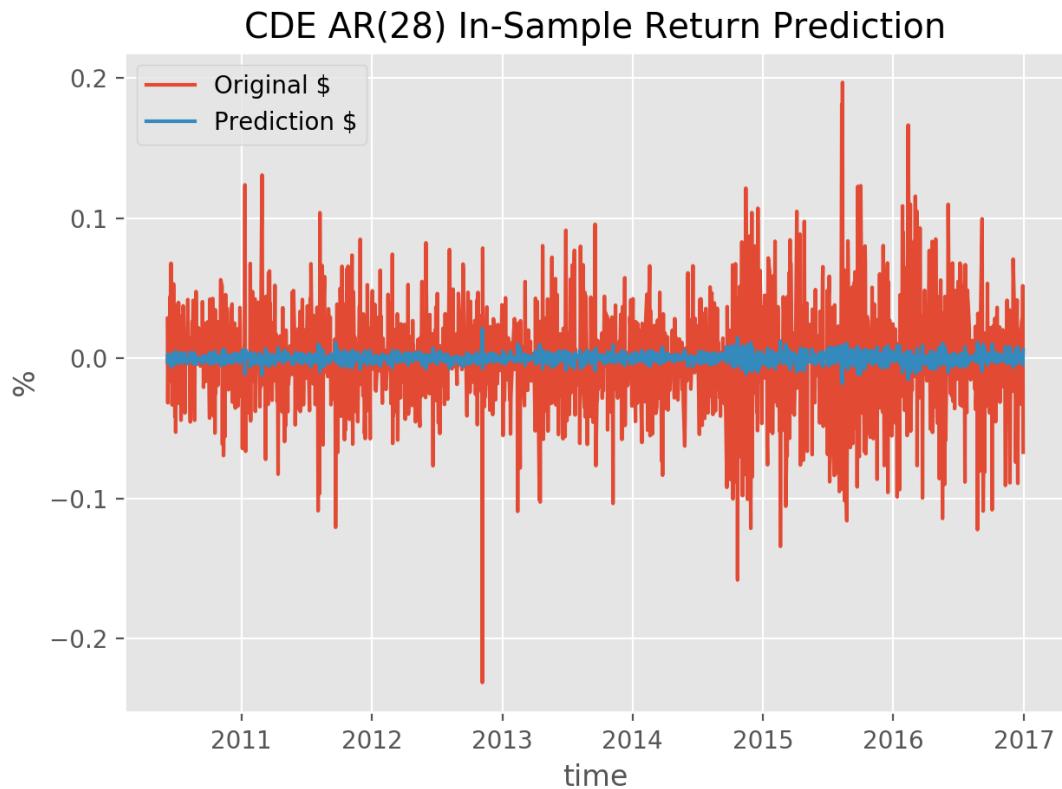


FIGURE A.27: CDE AR in-sample returns prediction

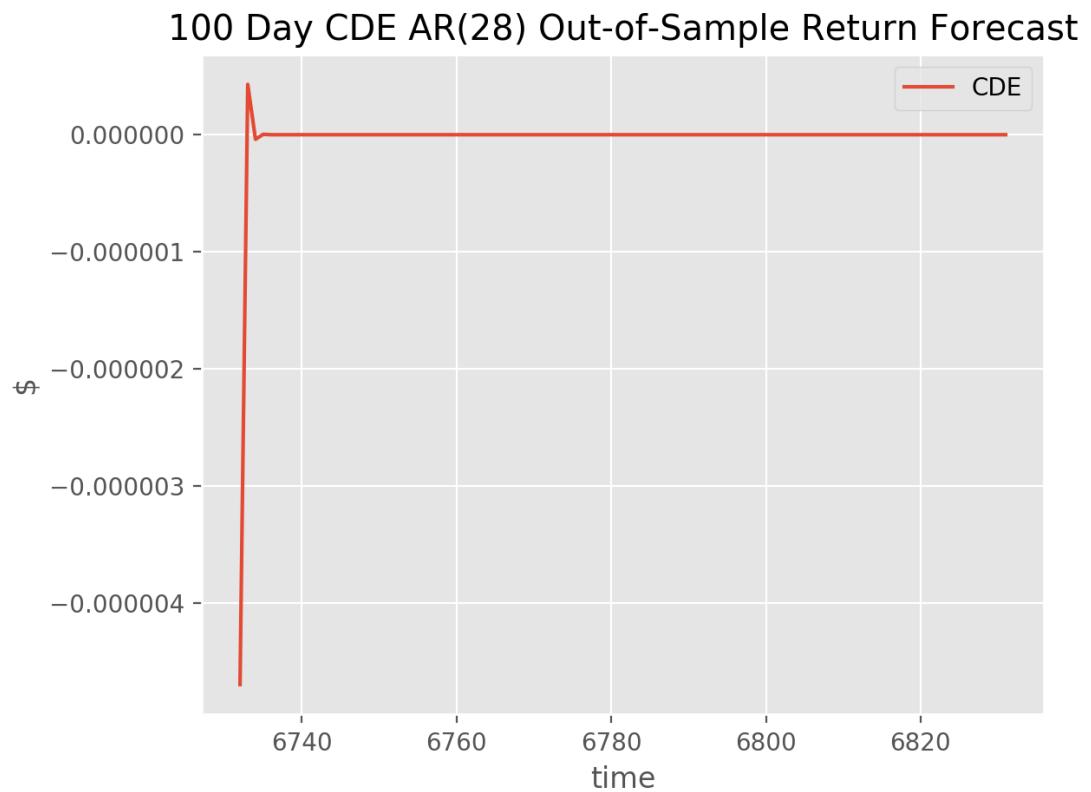


FIGURE A.28: 100 day CDE AR in-sample returns forecast

NAV B scored sharpe ratios of -0.410 for the original returns, and 0.124 for the predicted returns in the in-sample test.

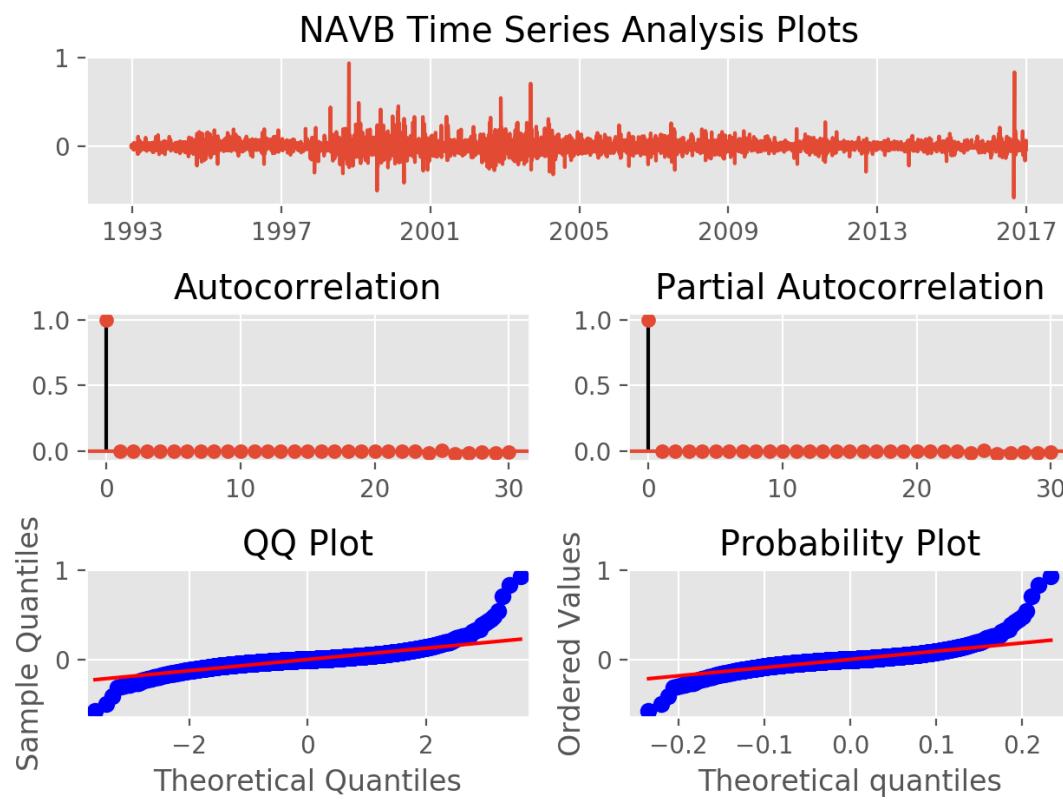


FIGURE A.29: NAVB AR time series analysis

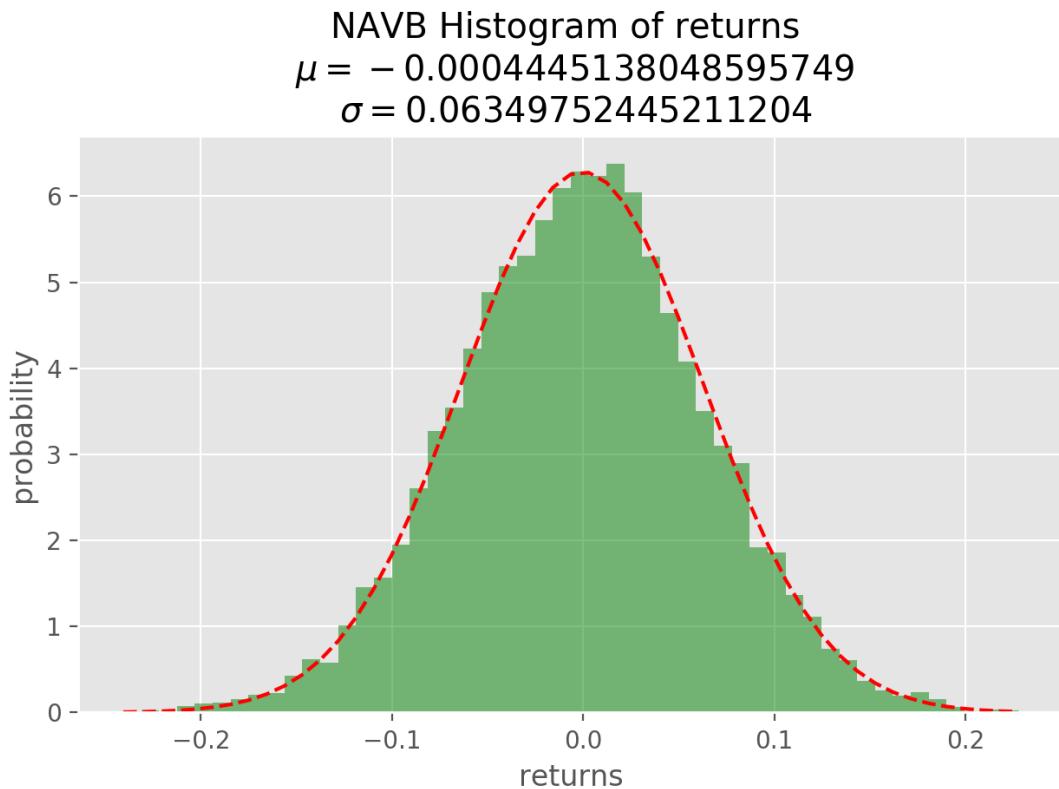


FIGURE A.30: NAVB AR histogram of returns

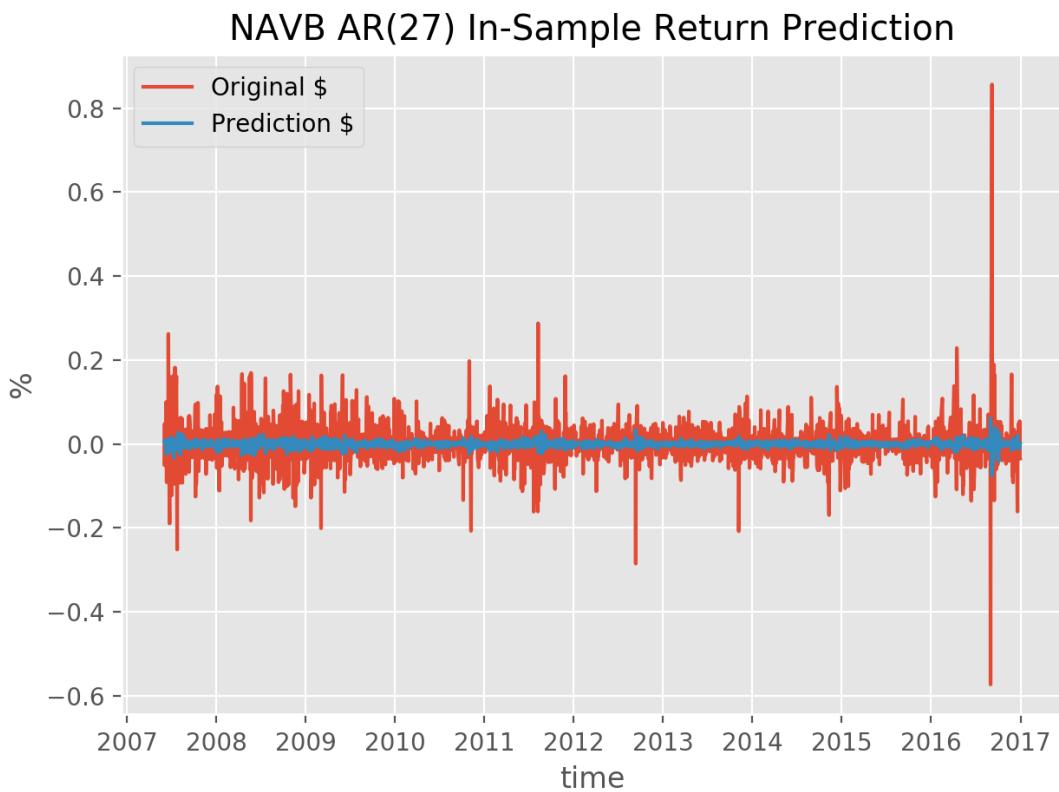


FIGURE A.31: NAVB AR in-sample returns prediction

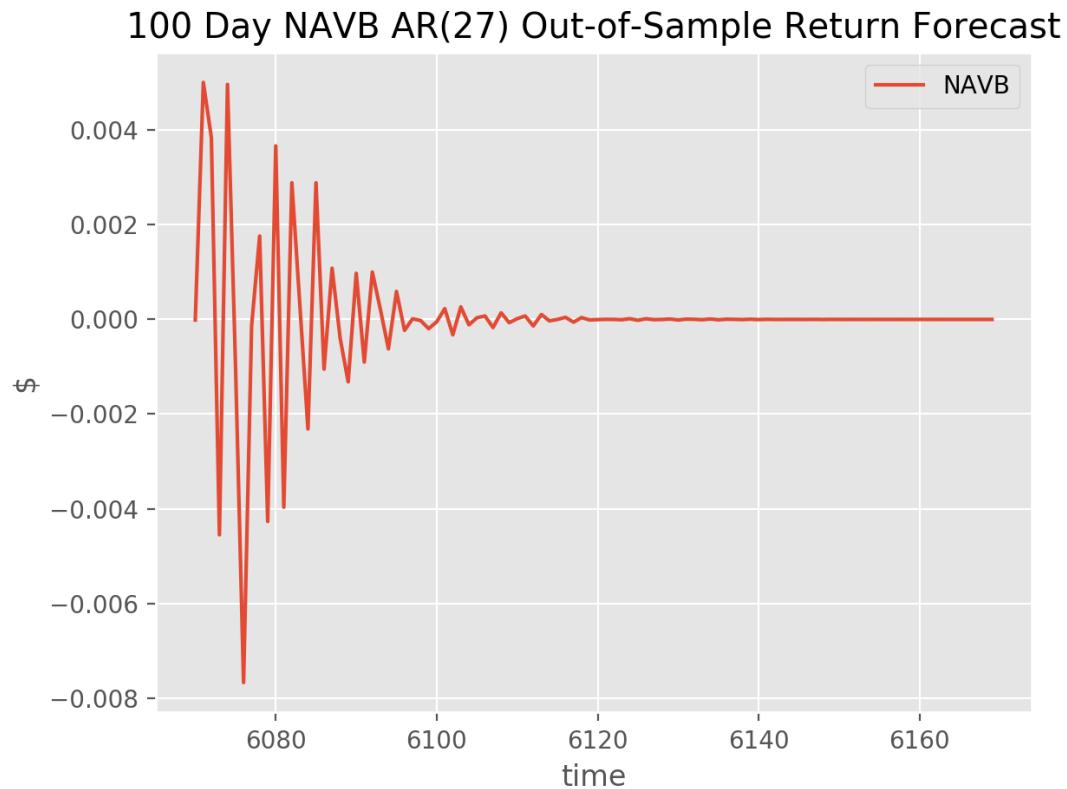


FIGURE A.32: 100 day NAVB AR in-sample returns forecast

HRG scored sharpe ratios of 0.627 for the original returns, and -2.603 for the predicted returns in the in-sample test.

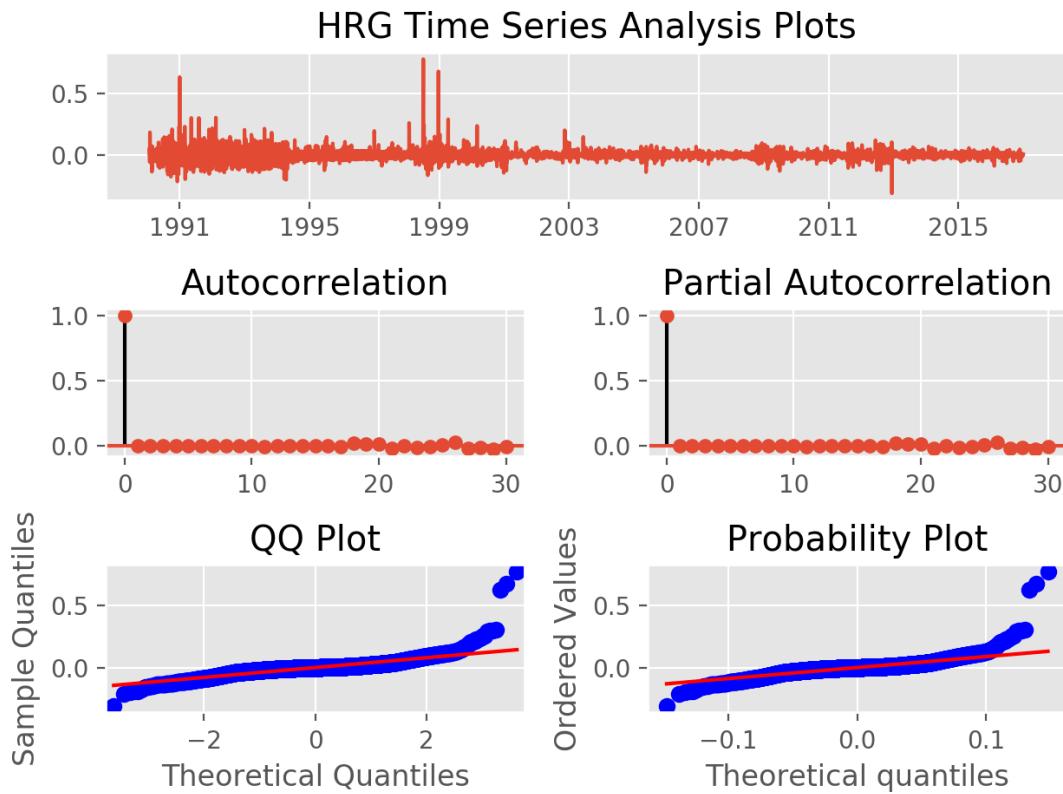


FIGURE A.33: HRG AR time series analysis

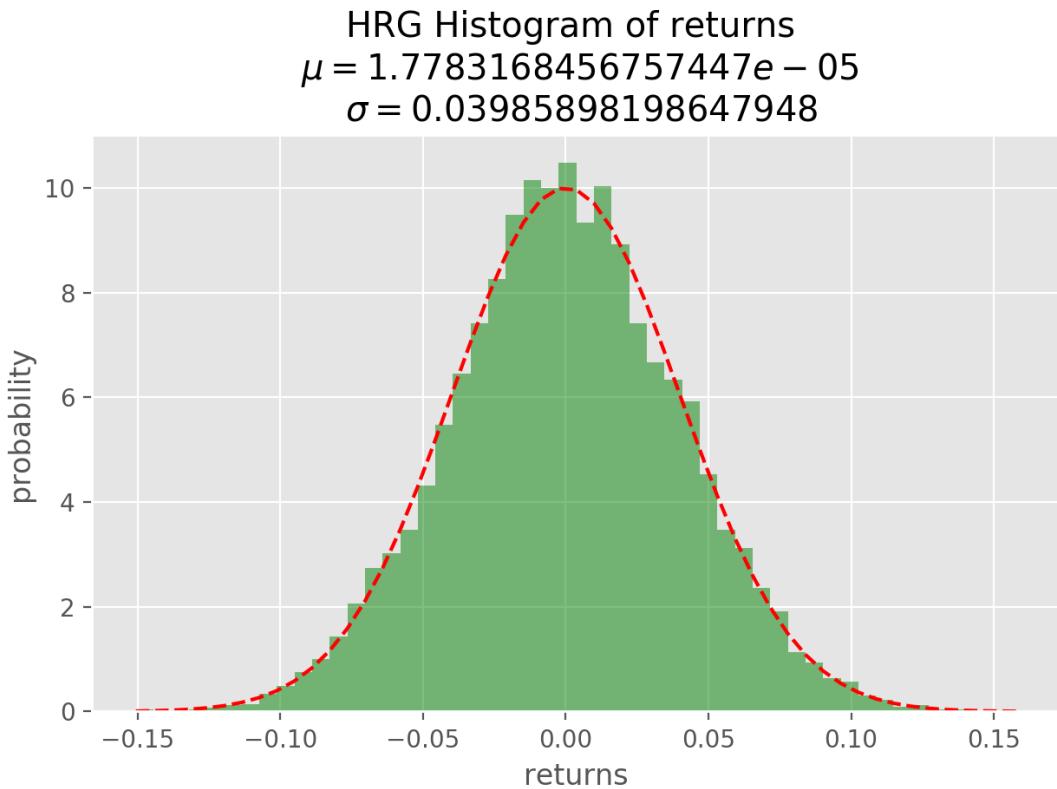


FIGURE A.34: HRG AR histogram of returns

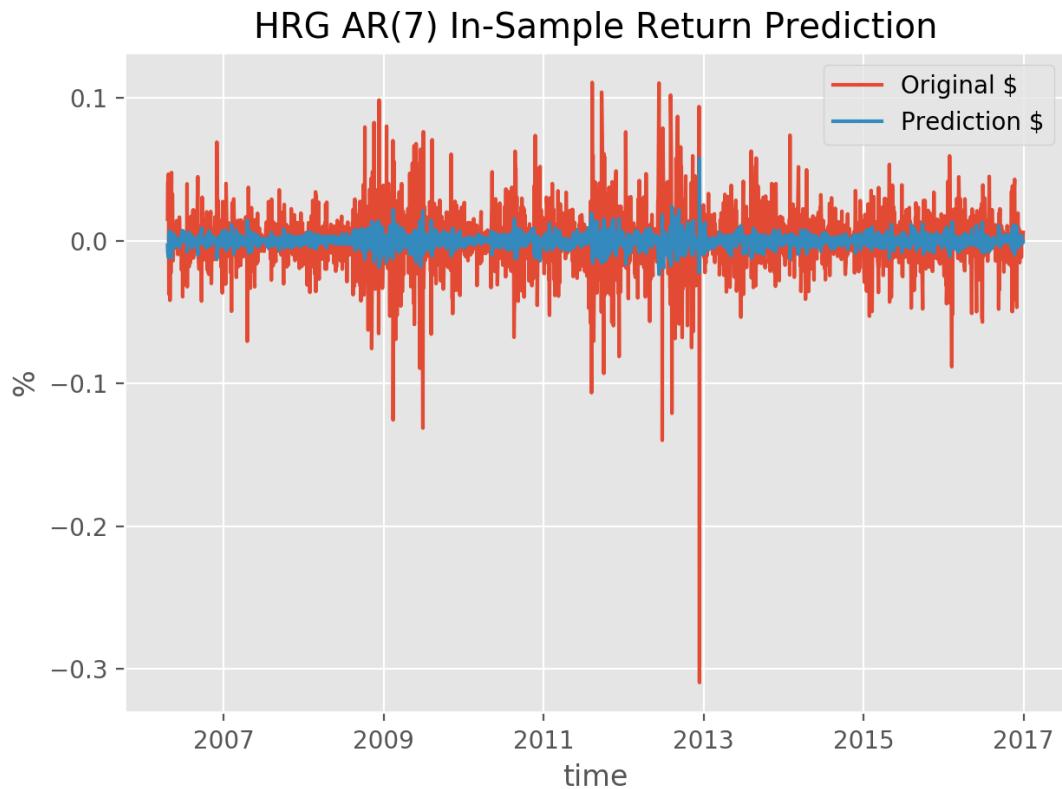


FIGURE A.35: HL AR in-sample returns prediction

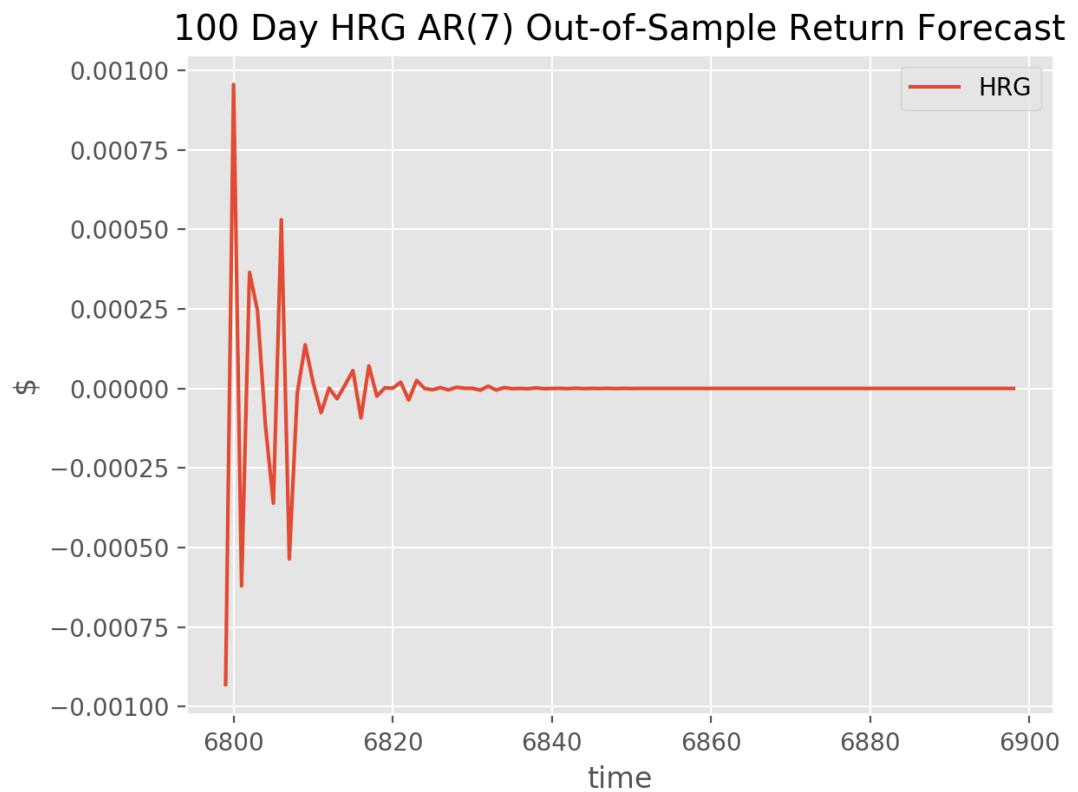


FIGURE A.36: 100 day HRG AR in-sample returns forecast

HL scored sharpe ratios of 0.695 for the original returns, and -1.880 for the predicted returns in the in-sample test.

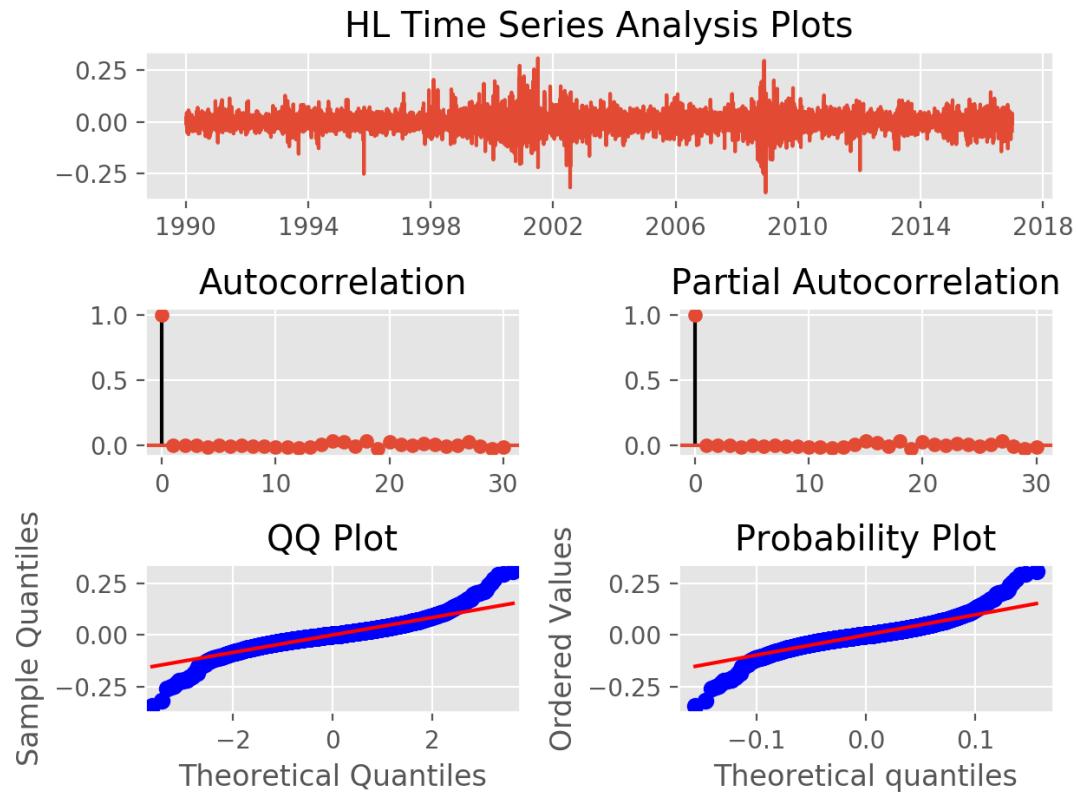


FIGURE A.37: HL AR time series analysis

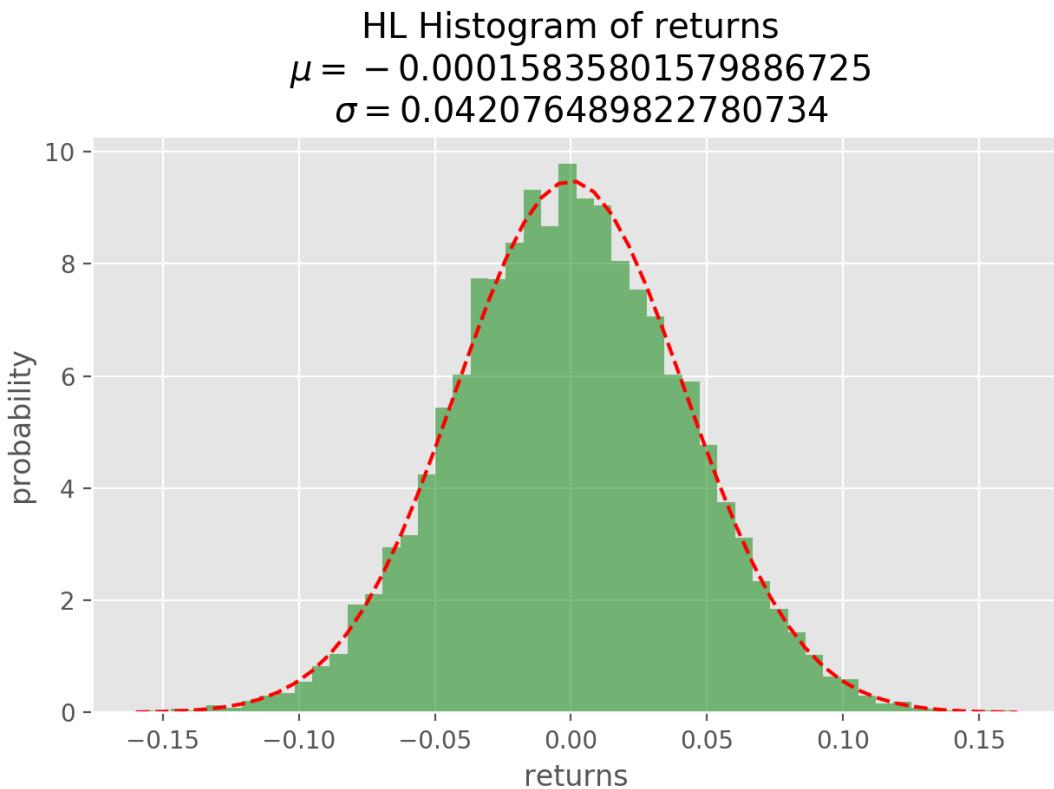


FIGURE A.38: HL AR histogram of returns

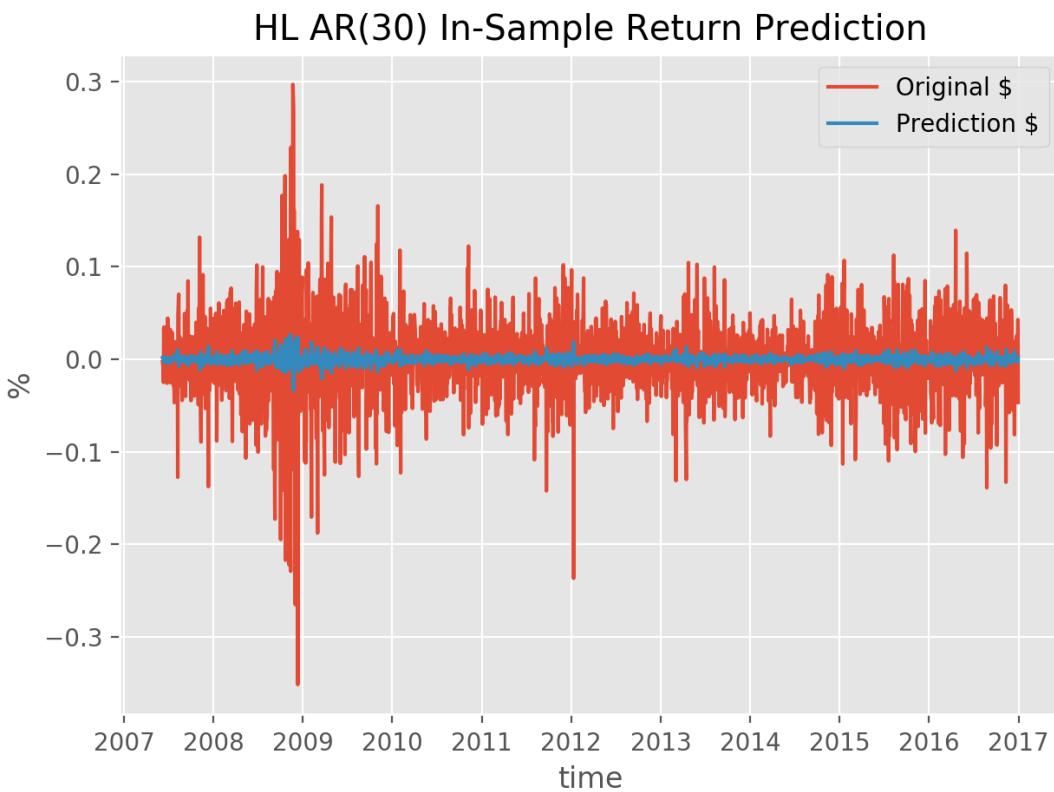


FIGURE A.39: HL AR in-sample returns prediction

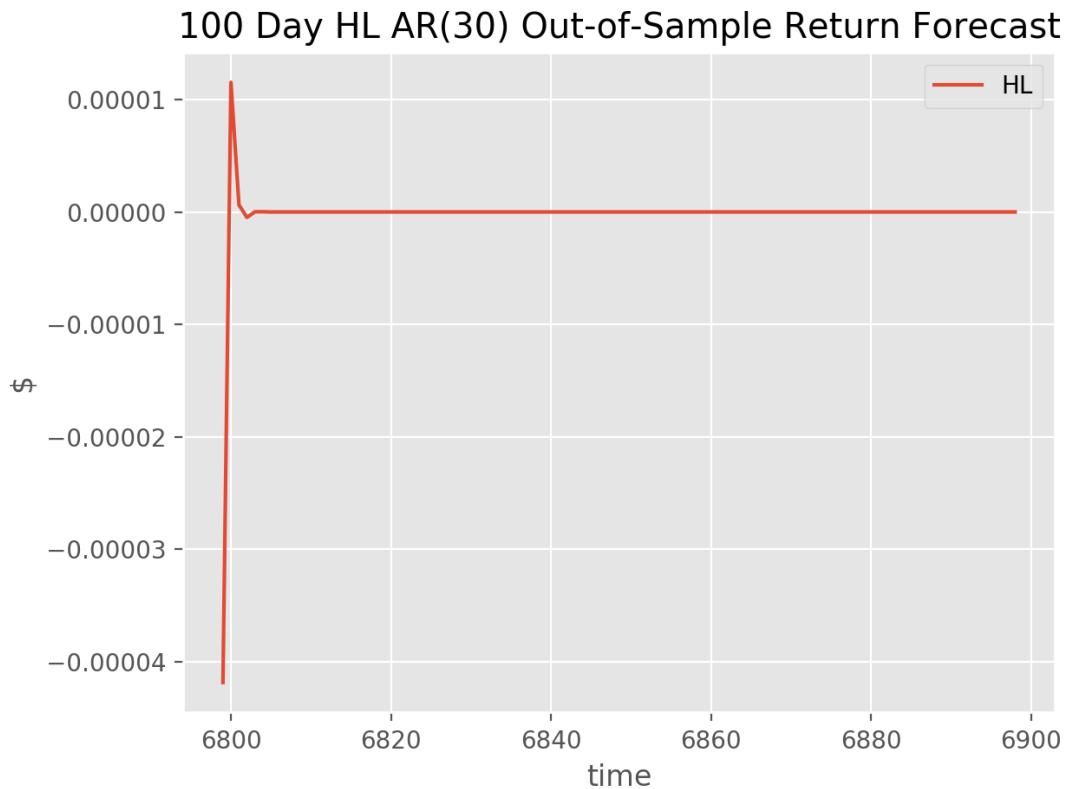


FIGURE A.40: 100 day HL AR in-sample returns forecast

A.1.5 Moving Average (MA)

CDE scored sharpe ratios of -0.199 for the original returns, and -0.760 for the predicted returns in the in-sample test.

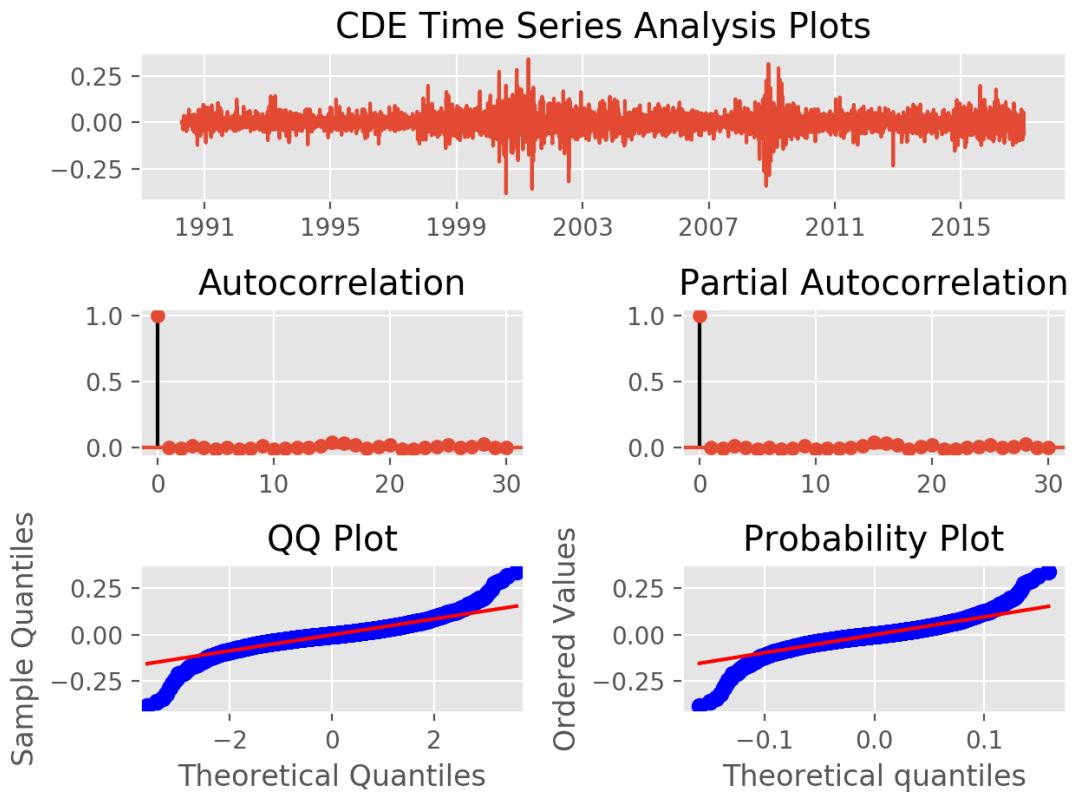


FIGURE A.41: CDE MA time series analysis

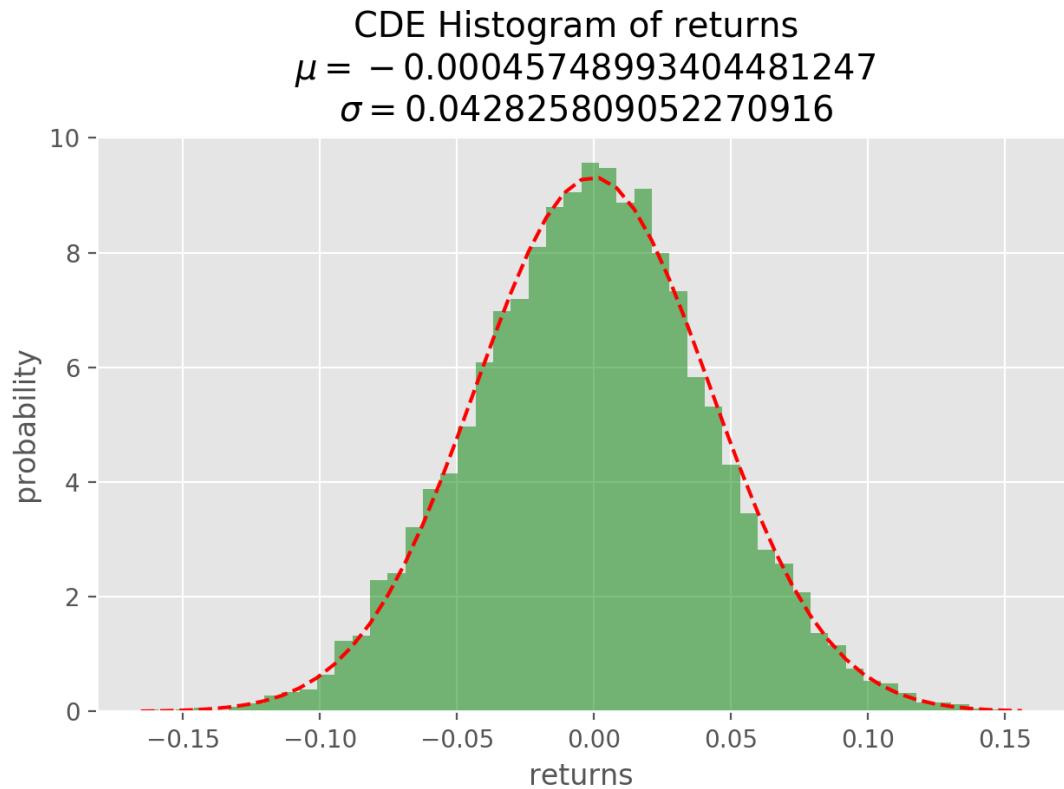


FIGURE A.42: CDE MA histogram of returns

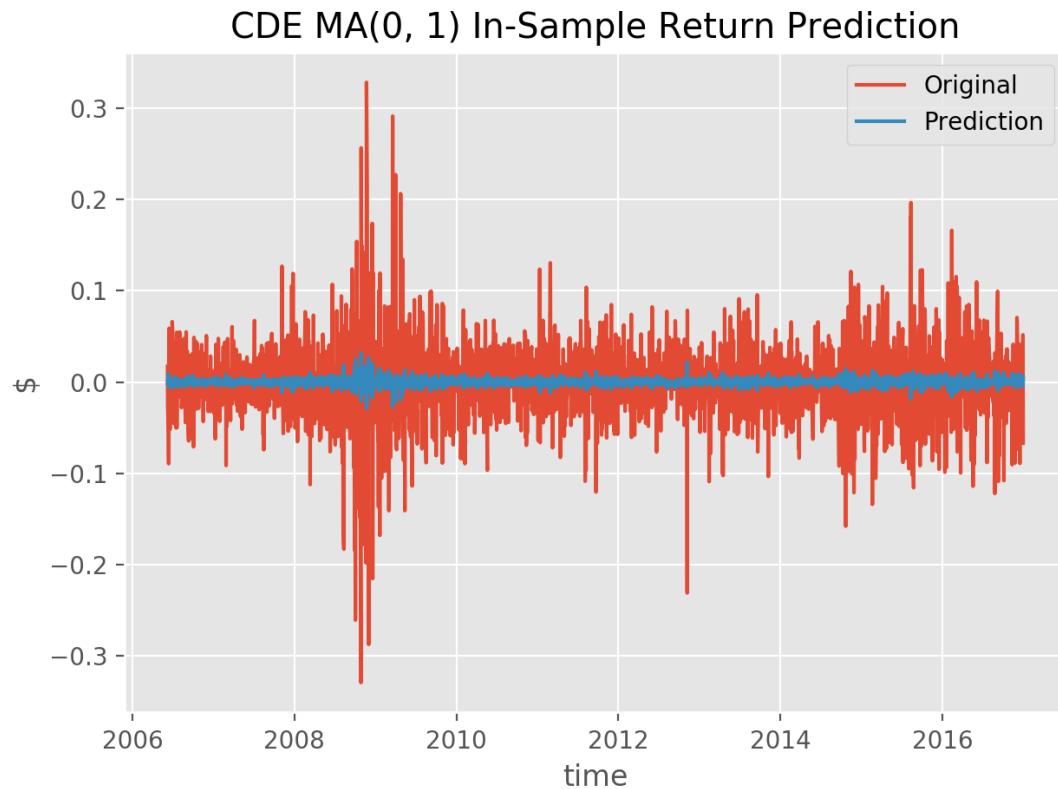


FIGURE A.43: CDE MA in-sample returns prediction

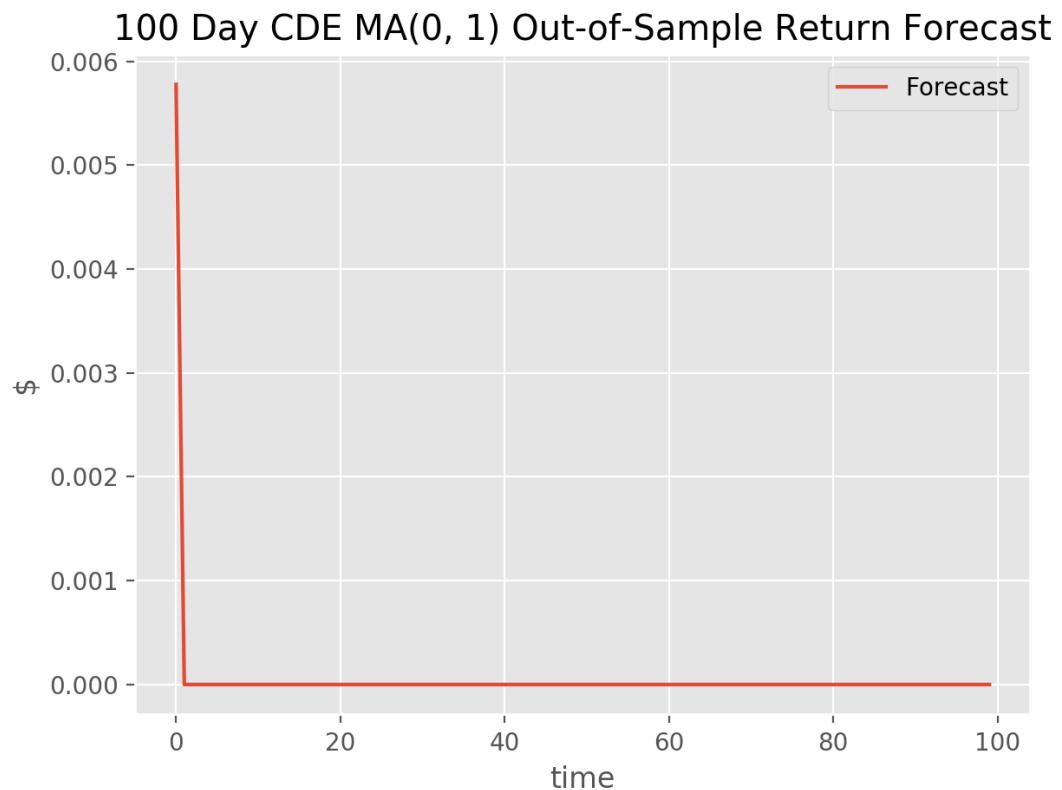


FIGURE A.44: 100 day CDE MA in-sample returns forecast

NAV B scored sharpe ratios of -0.067 for the original returns, and 0.590 for the predicted returns in the in-sample test.

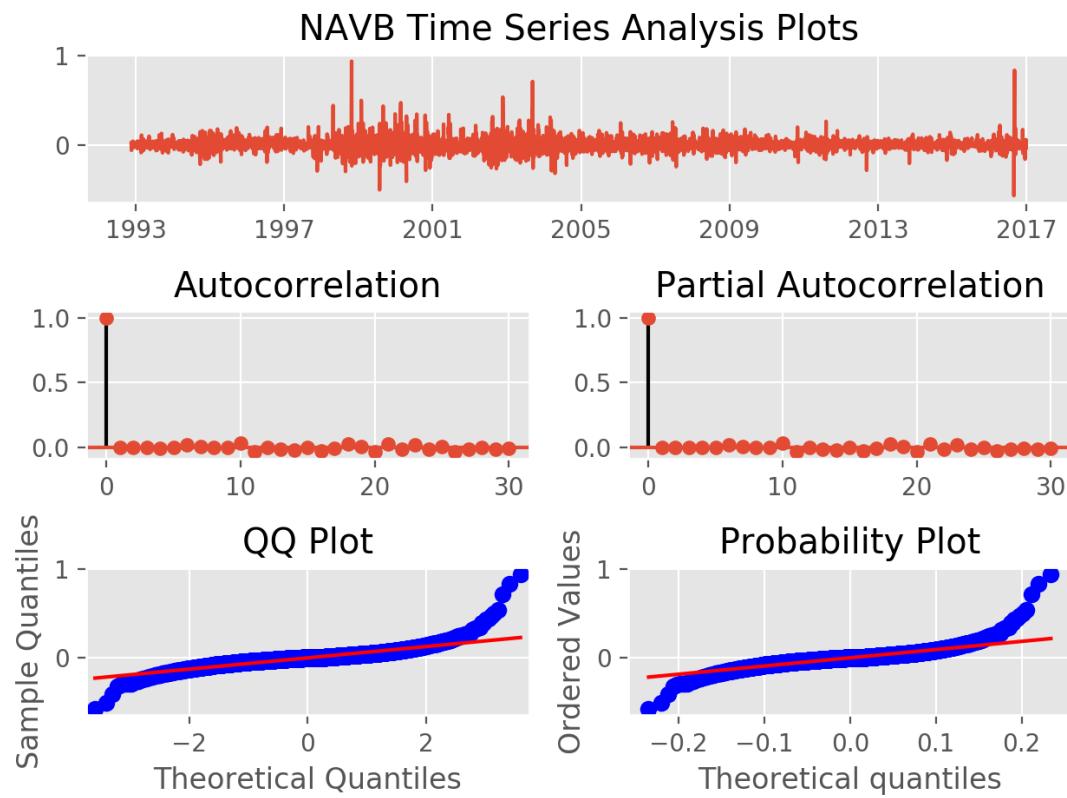


FIGURE A.45: NAVB MA time series analysis

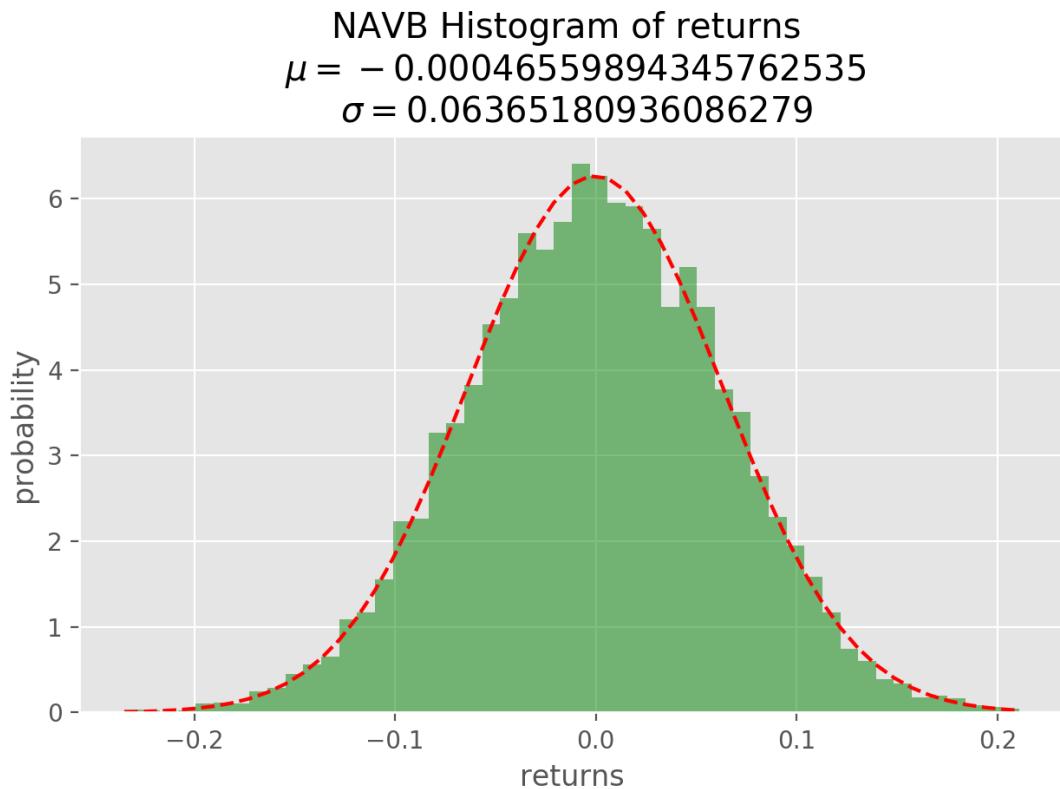


FIGURE A.46: NAVB MA histogram of returns

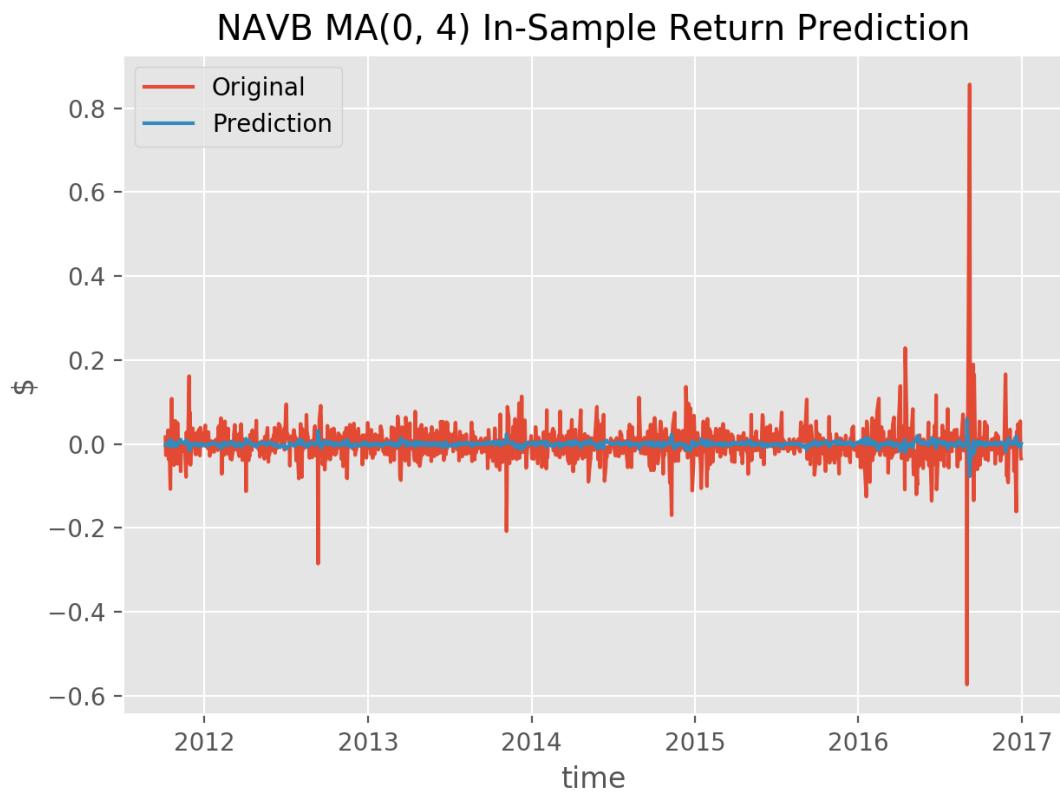


FIGURE A.47: NAVB MA in-sample returns prediction

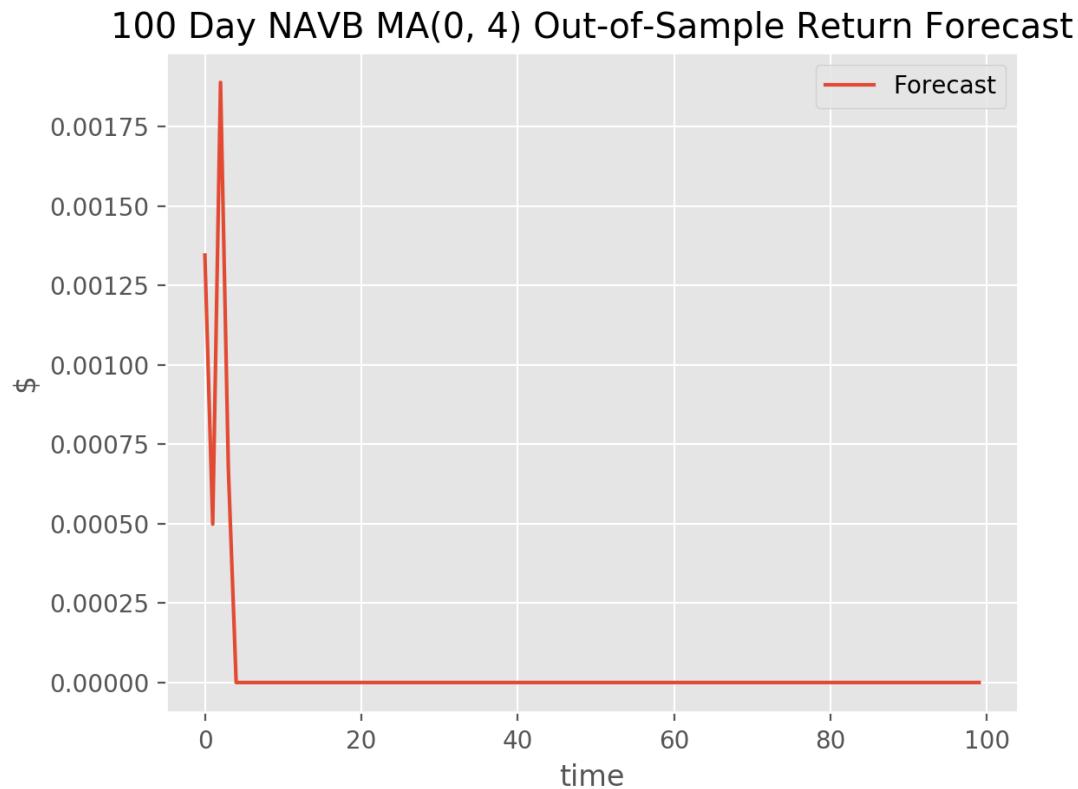


FIGURE A.48: 100 day NAVB MA in-sample returns forecast

HRG scored sharpe ratios of 0.084 for the original returns, and -1.020 for the predicted returns in the in-sample test.

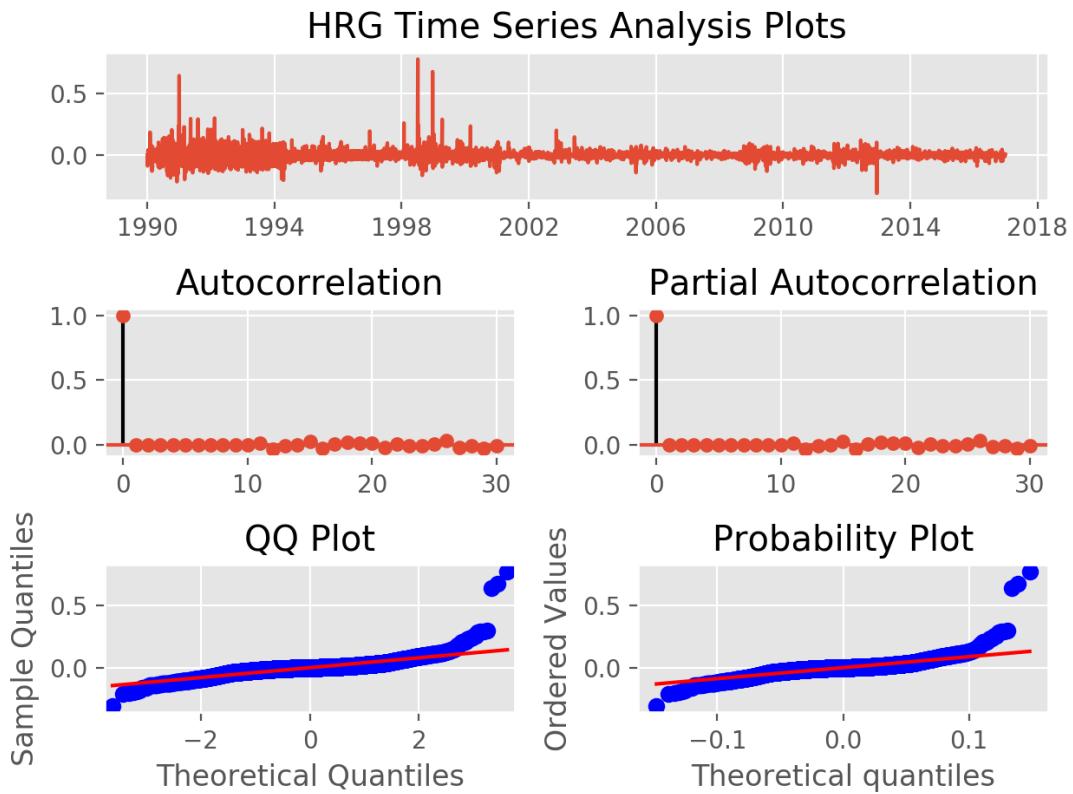


FIGURE A.49: HRG MA time series analysis

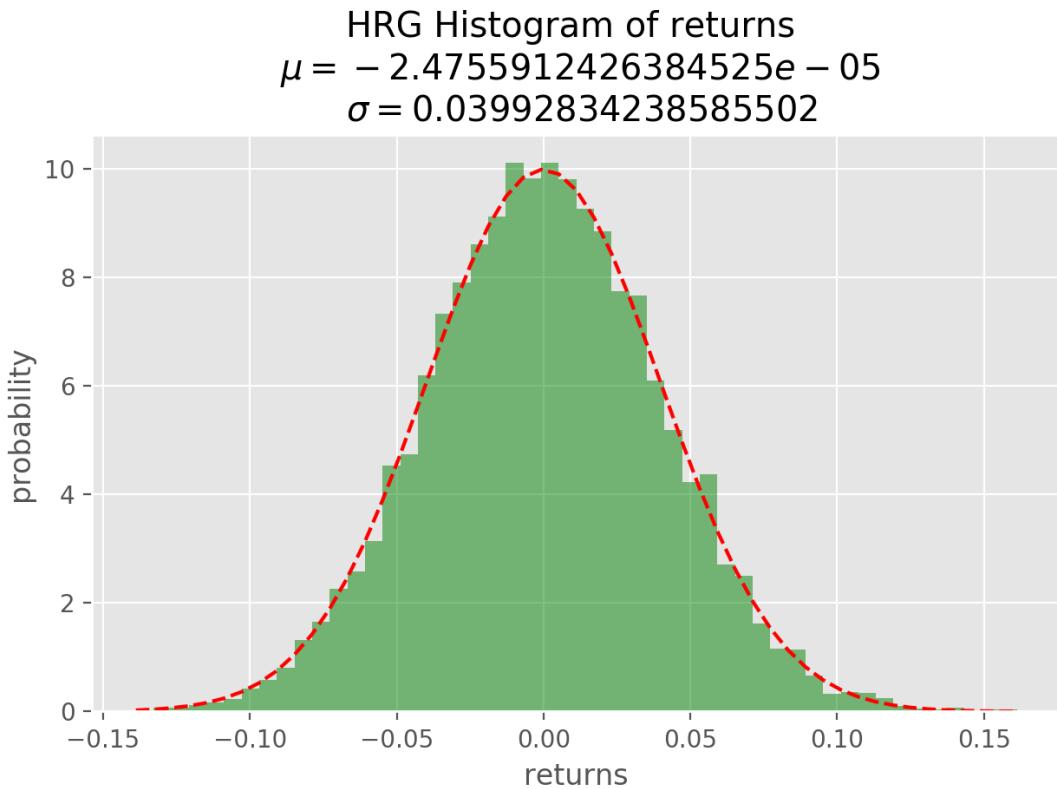


FIGURE A.50: HRG MA histogram of returns

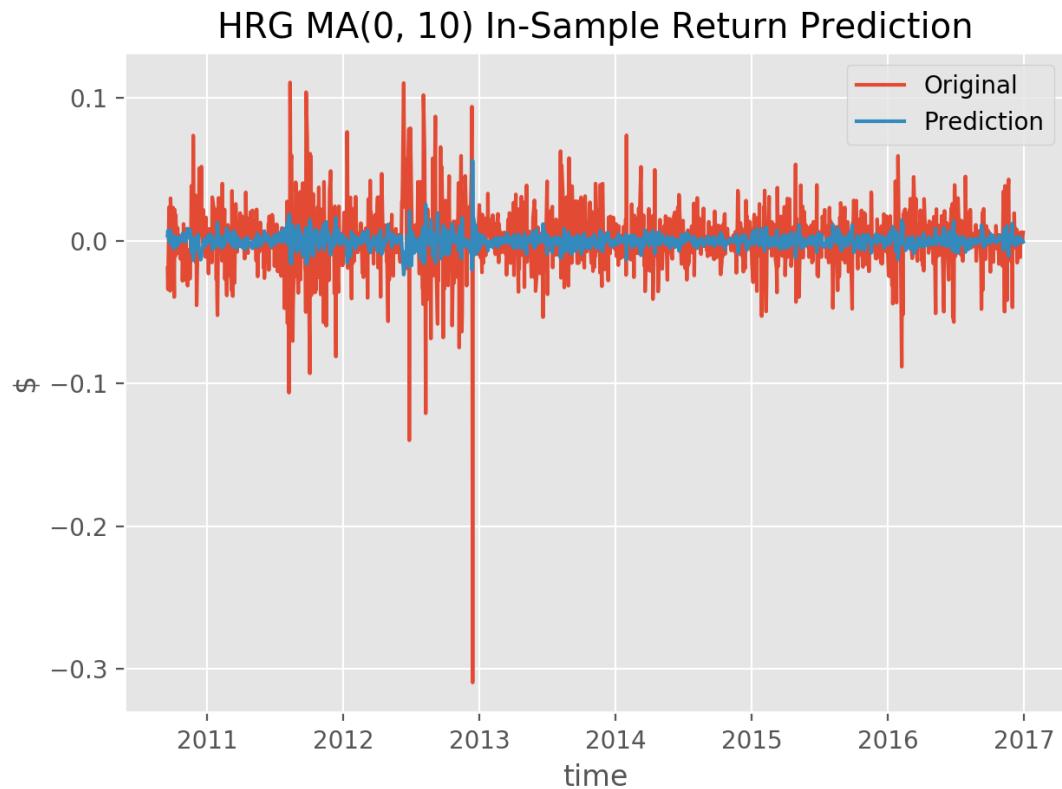


FIGURE A.51: HRG MA in-sample returns prediction

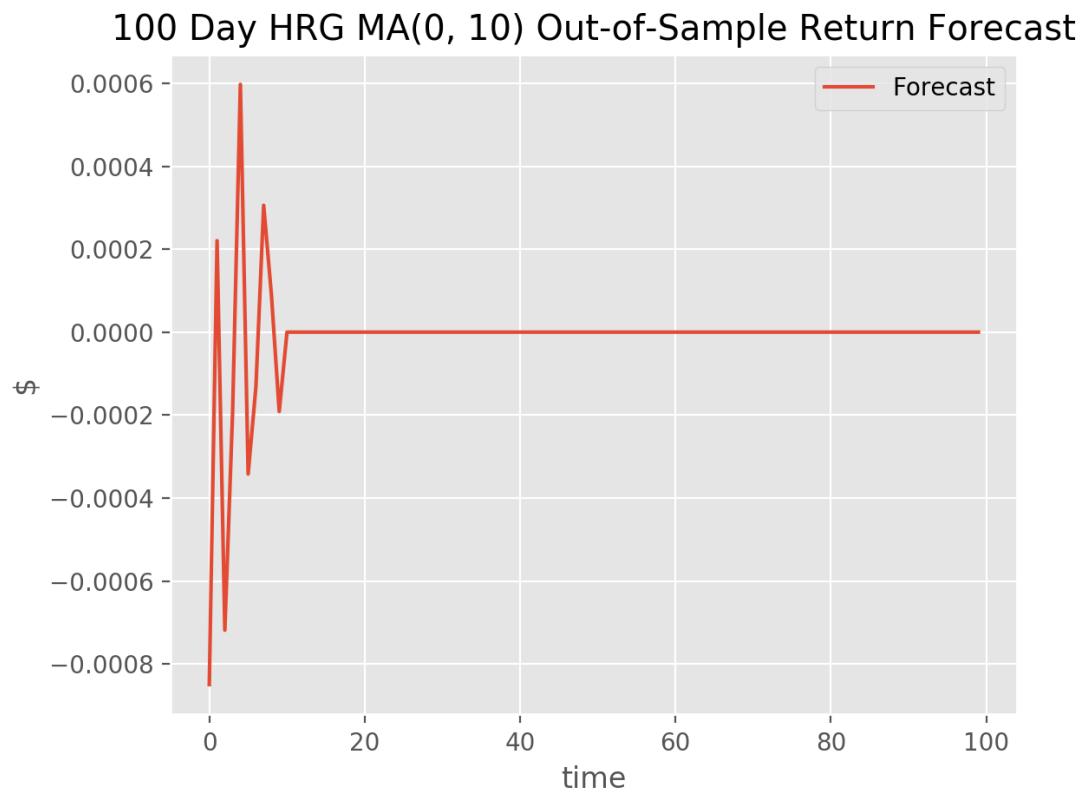


FIGURE A.52: 100 day HRG MA in-sample returns forecast

HL scored sharpe ratios of -0.128 for the original returns, and -0.977 for the predicted returns in the in-sample test.

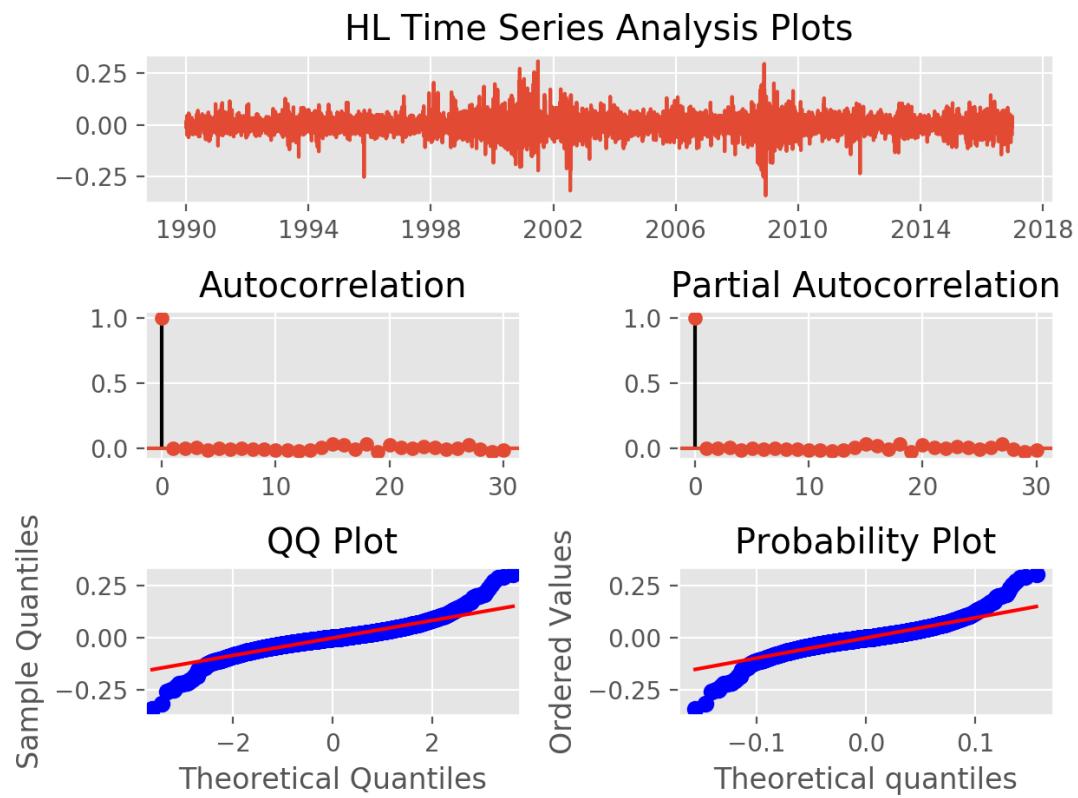


FIGURE A.53: HL MA time series analysis

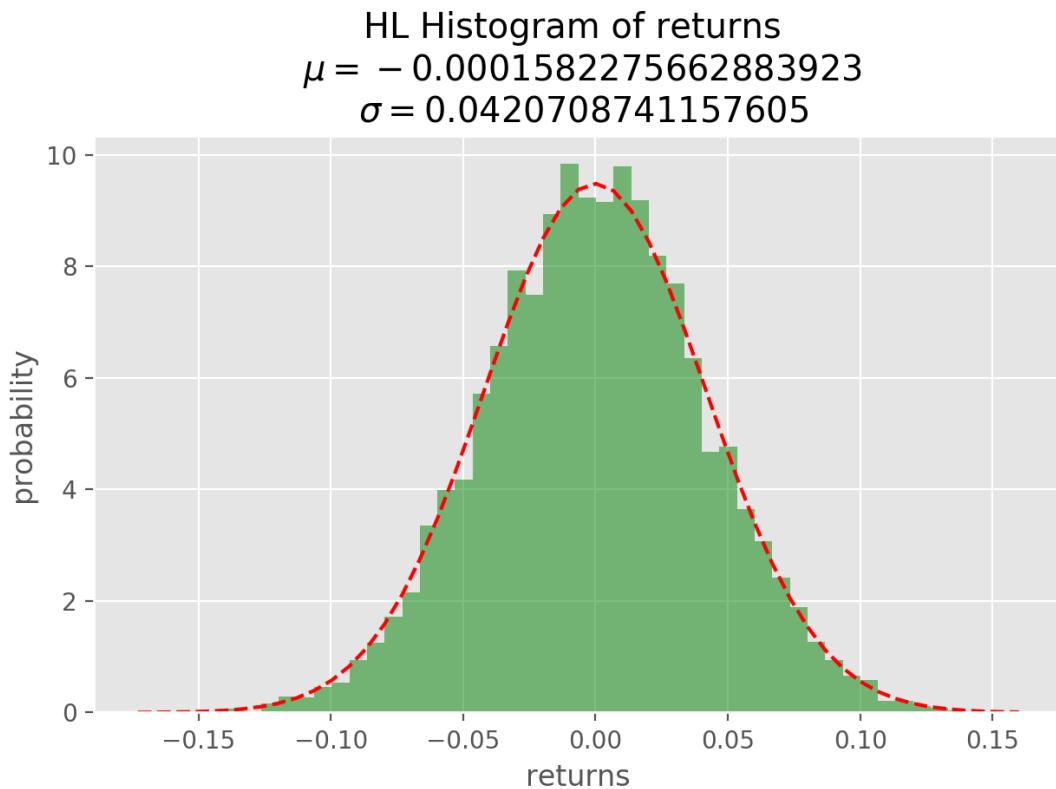


FIGURE A.54: HL MA histogram of returns

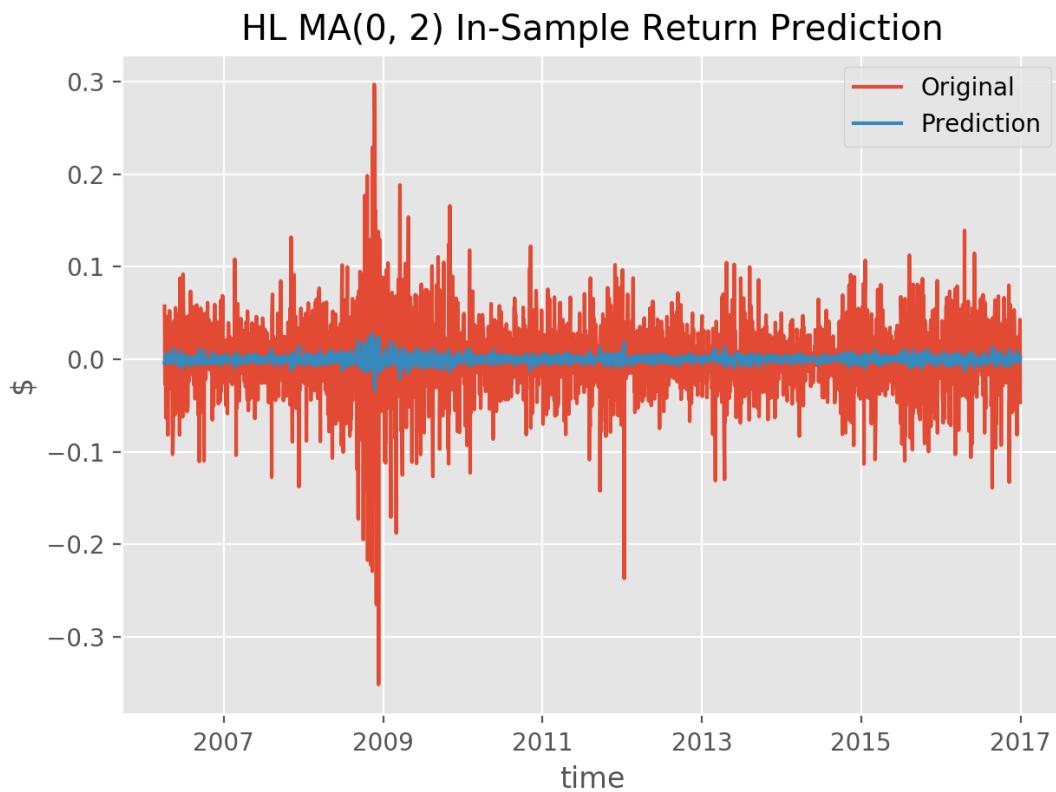


FIGURE A.55: HL MA in-sample returns prediction

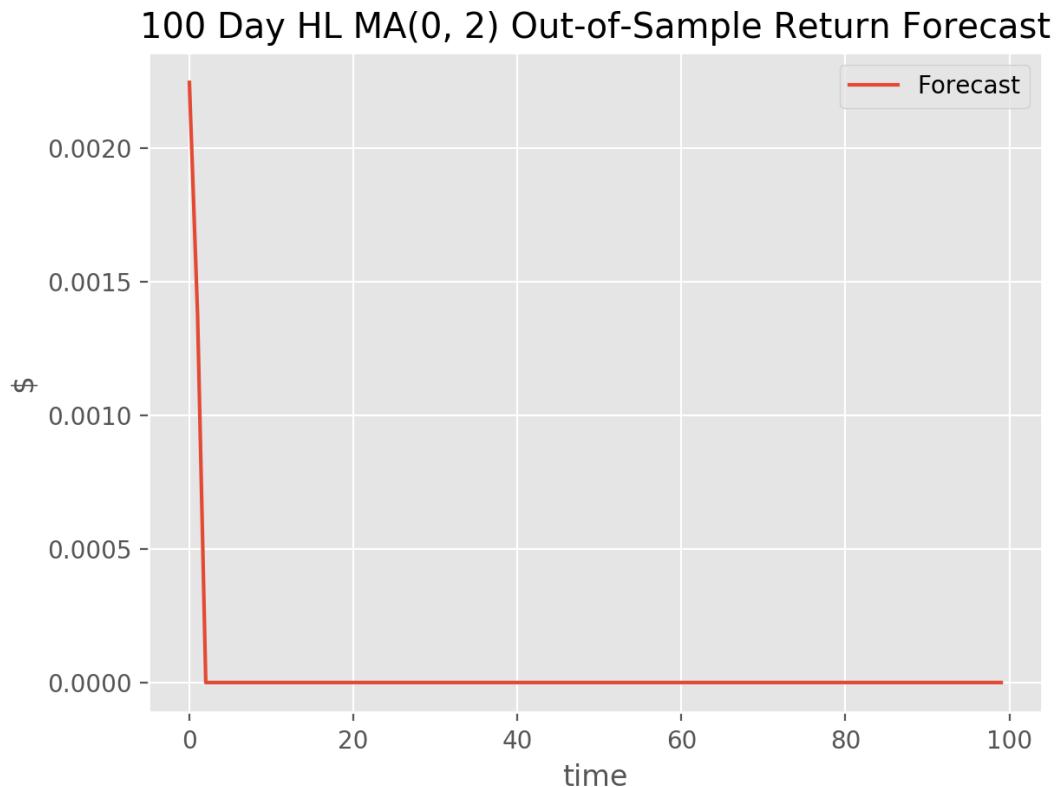


FIGURE A.56: 100 day HL MA in-sample returns forecast

A.1.6 Auto Regressive Moving Average (ARMA)

MSFT scored sharpe ratios of 0.537 for the original returns, and -6.485 for the predicted returns in the in-sample test.

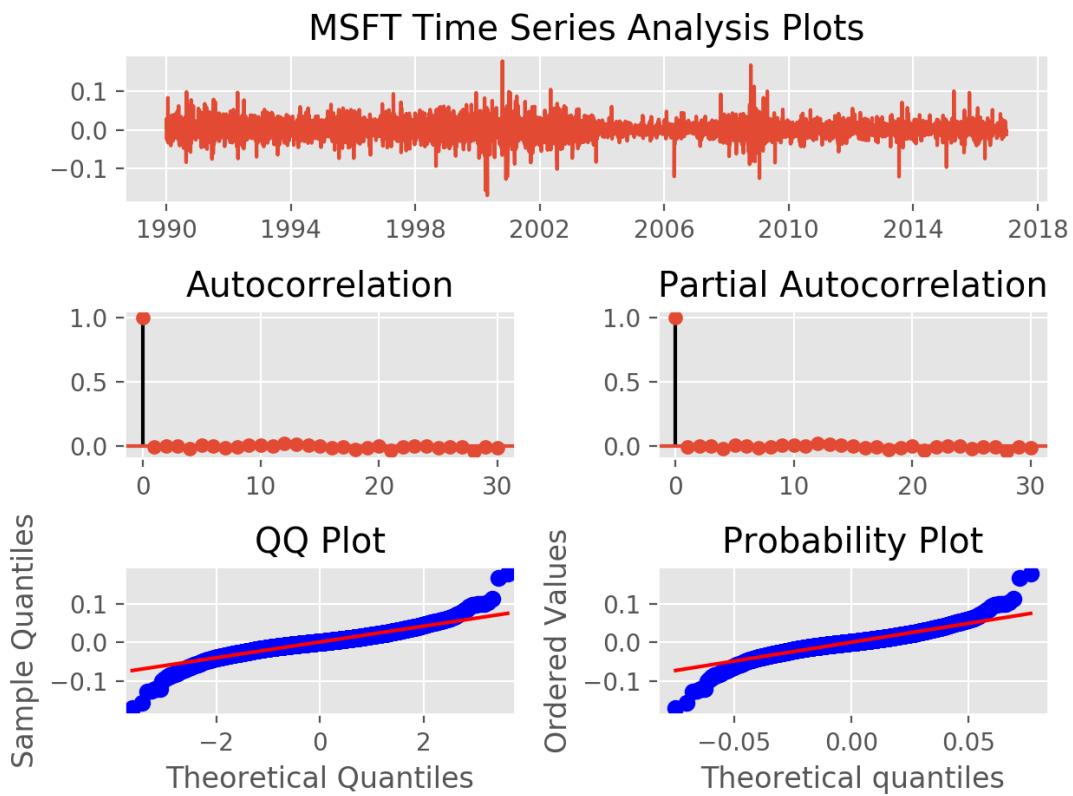


FIGURE A.57: MSFT ARMA time series analysis

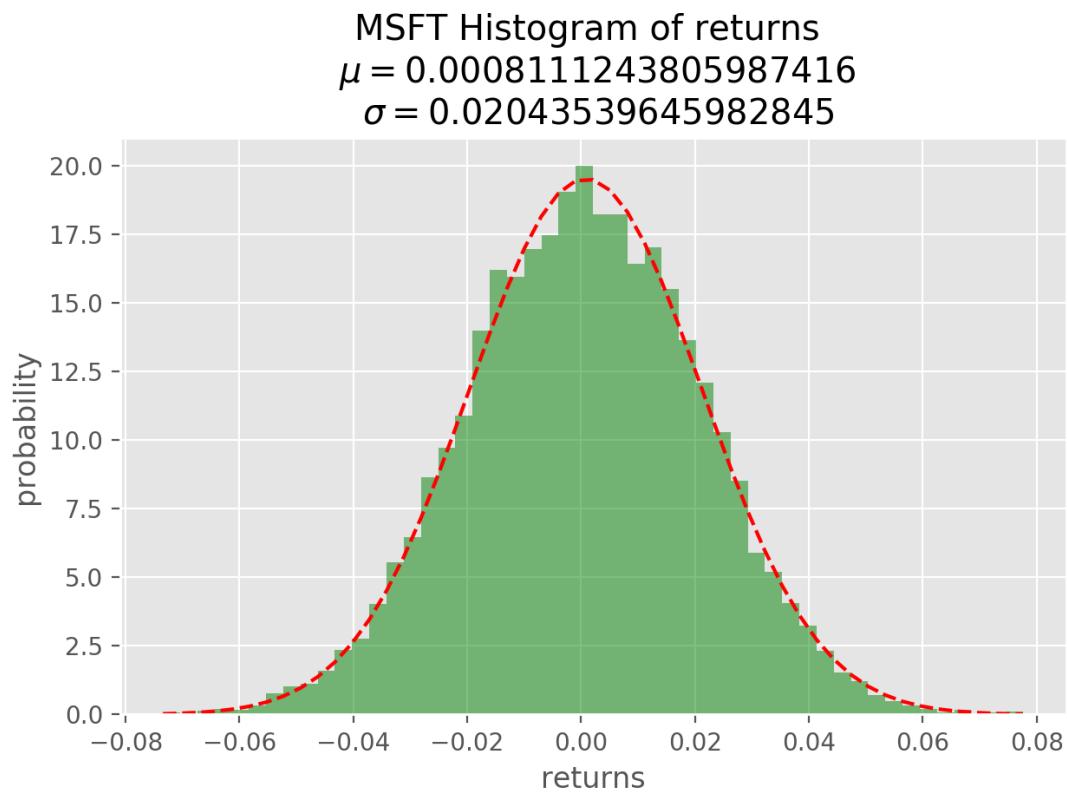


FIGURE A.58: MSFT ARMA histogram of returns

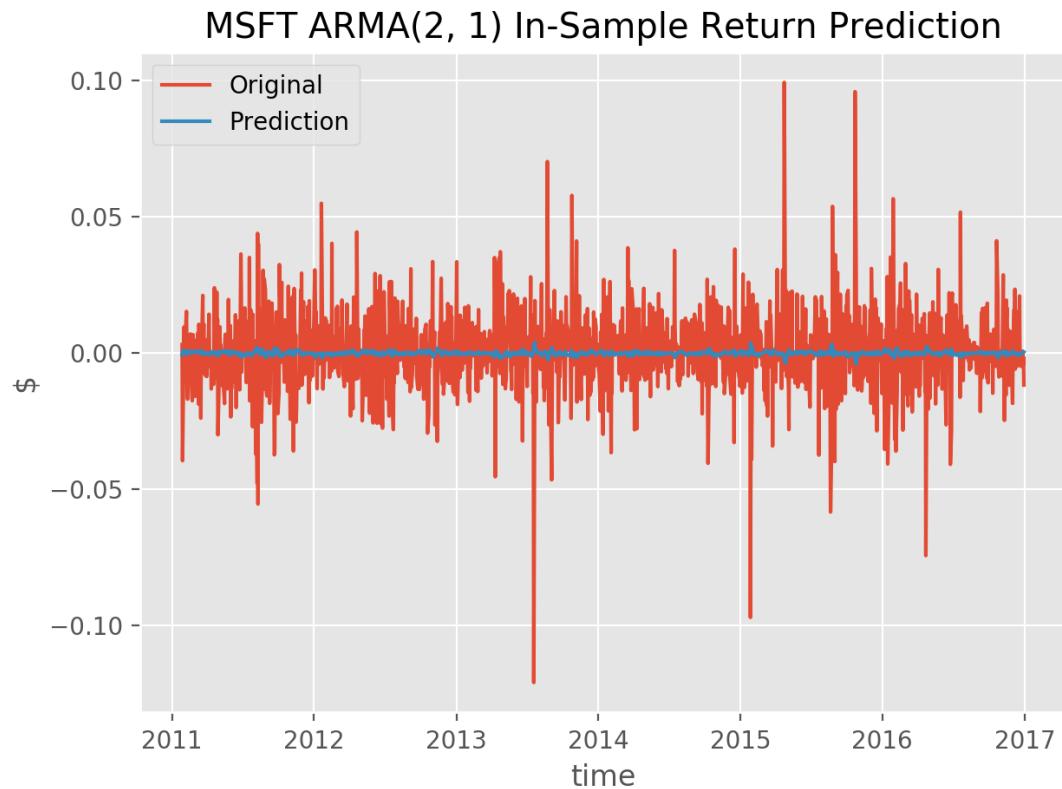


FIGURE A.59: MSFT ARMA in-sample returns prediction

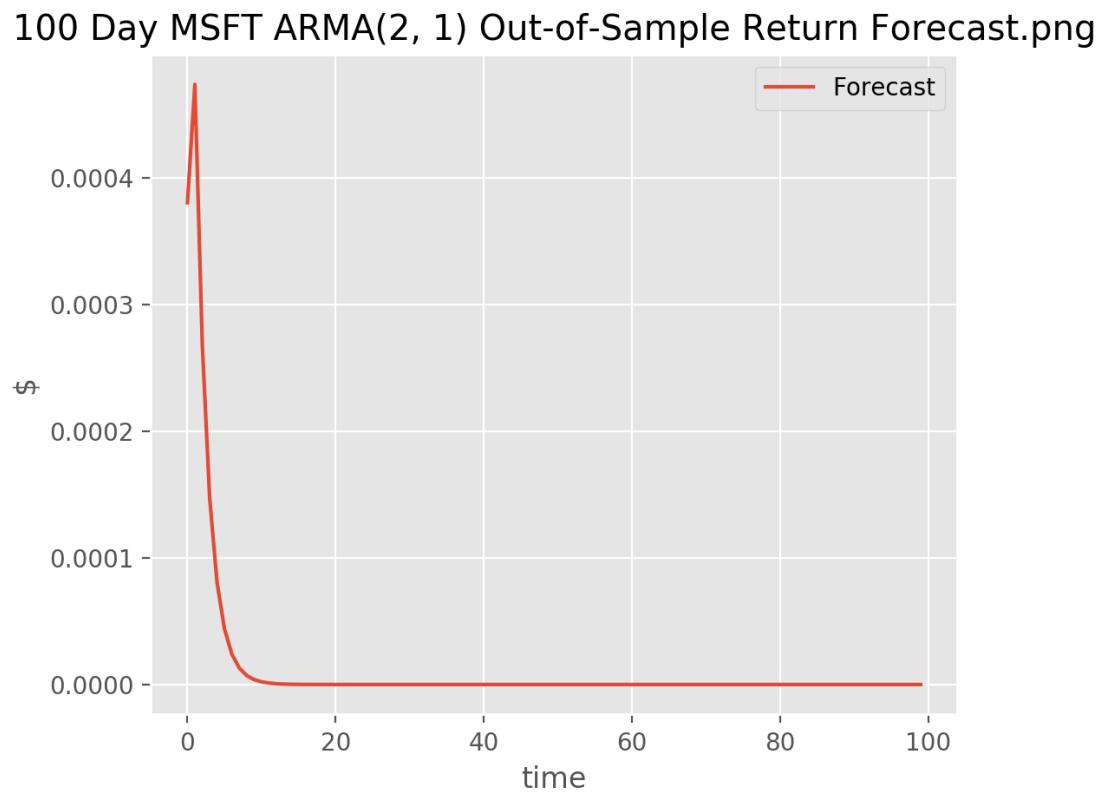


FIGURE A.60: 100 day MSFT ARMA in-sample returns forecast

NAVB scored sharpe ratios of -0.245 for the original returns, and -0.614 for the predicted returns in the in-sample test.

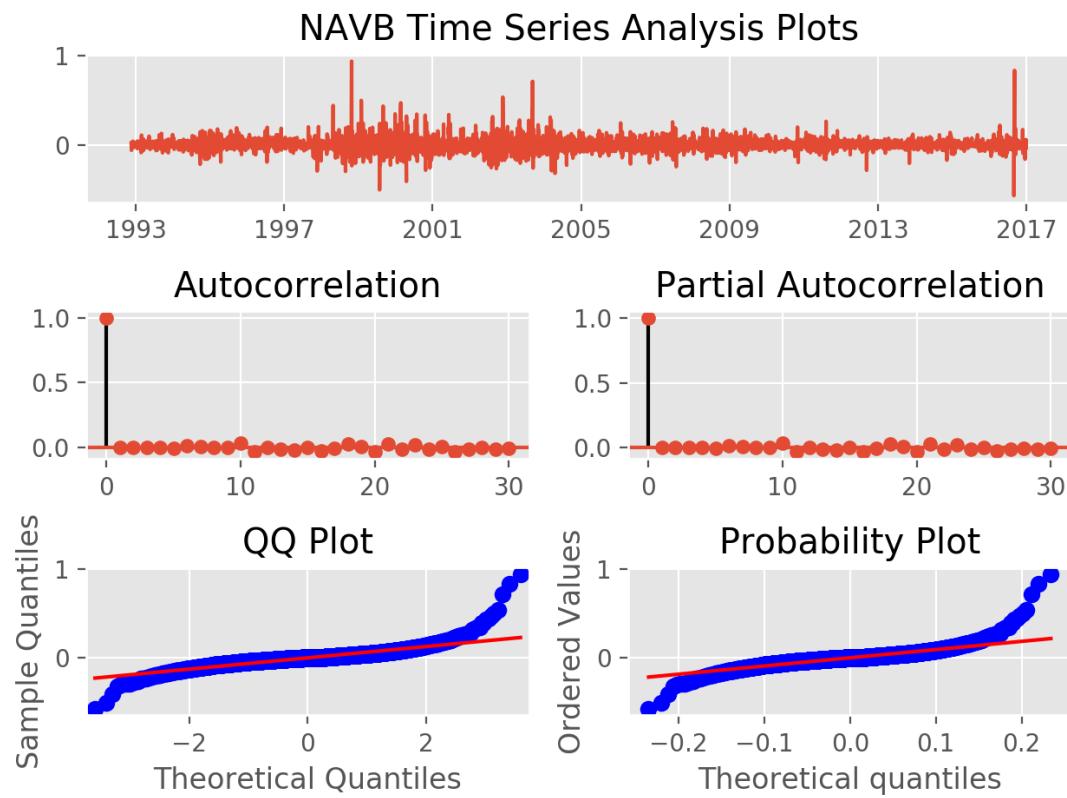


FIGURE A.61: NAVB ARMA time series analysis

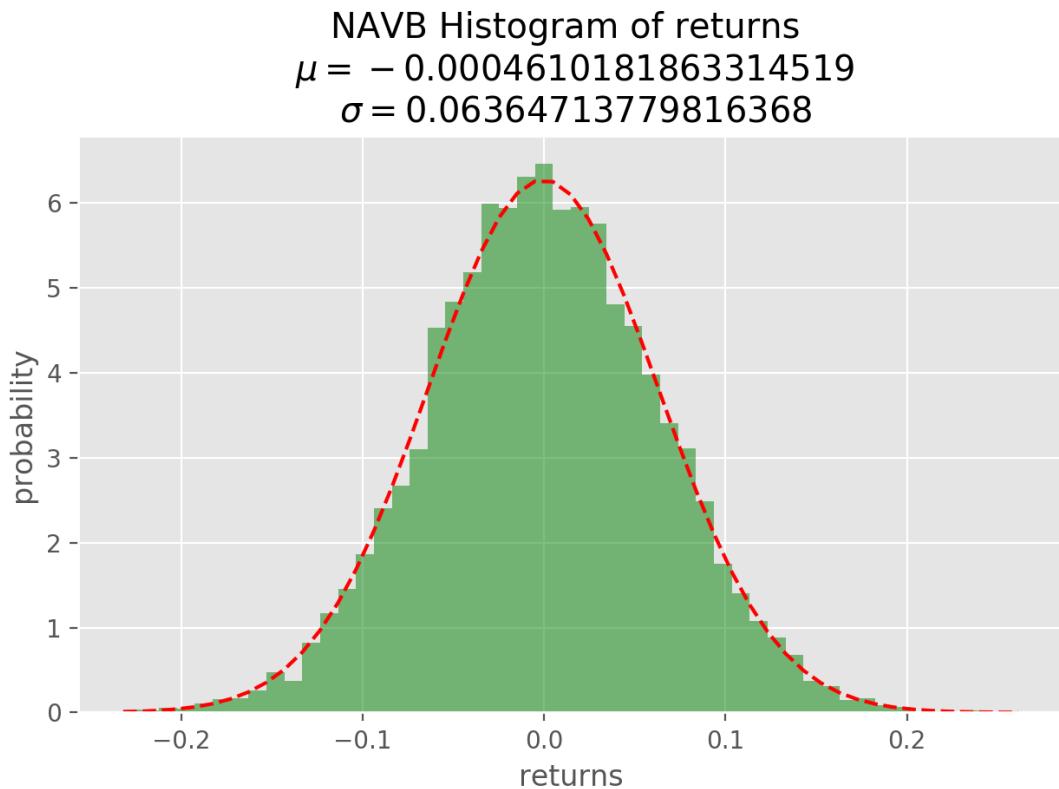


FIGURE A.62: NAVB ARMA histogram of returns

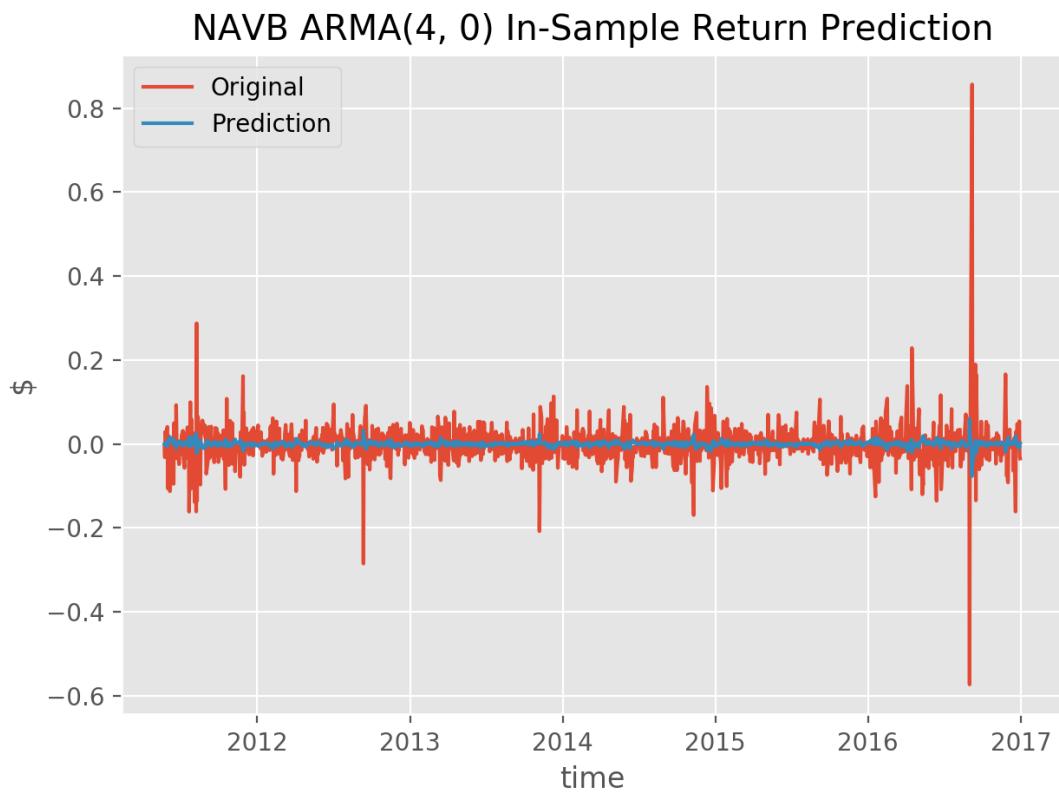


FIGURE A.63: NAVB ARMA in-sample returns prediction

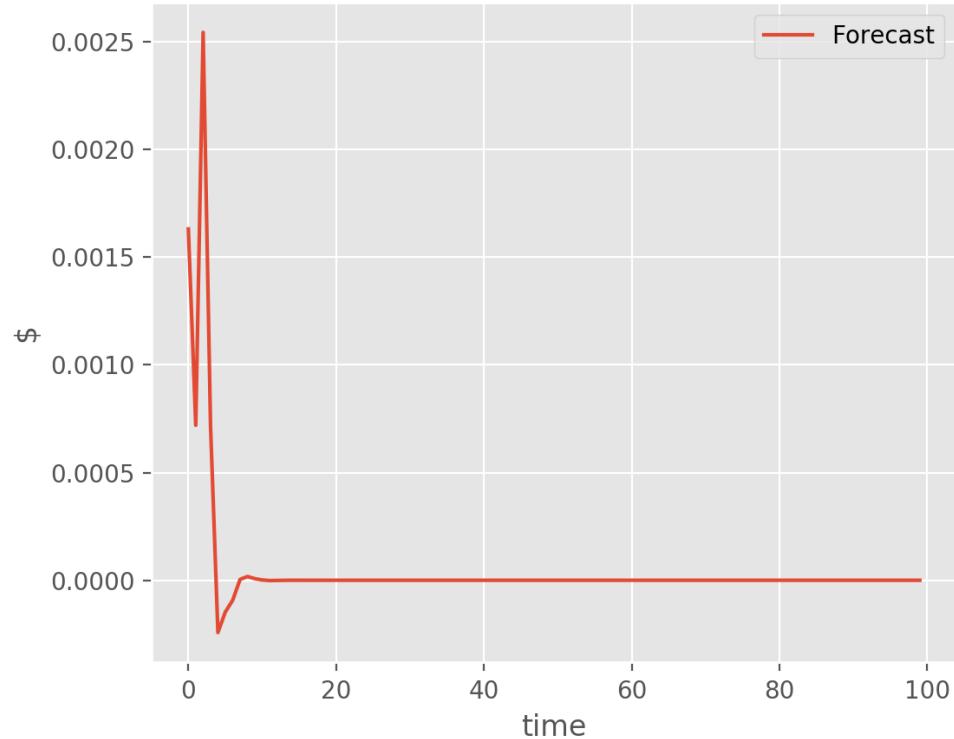
100 Day NAVB ARMA(4, 0) Out-of-Sample Return Forecast.png

FIGURE A.64: 100 day NAVB ARMA in-sample returns forecast

HRG scored sharpe ratios of -0.032 for the original returns, and -0.641 for the predicted returns in the in-sample test.

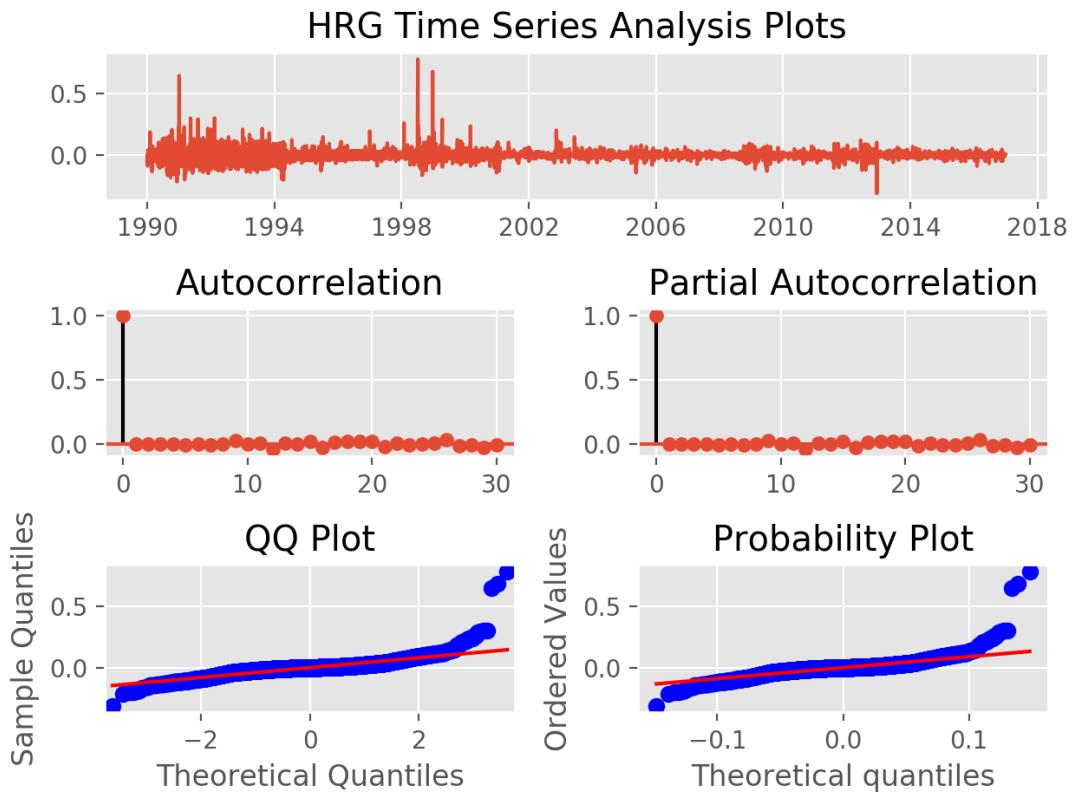


FIGURE A.65: HRG ARMA time series analysis

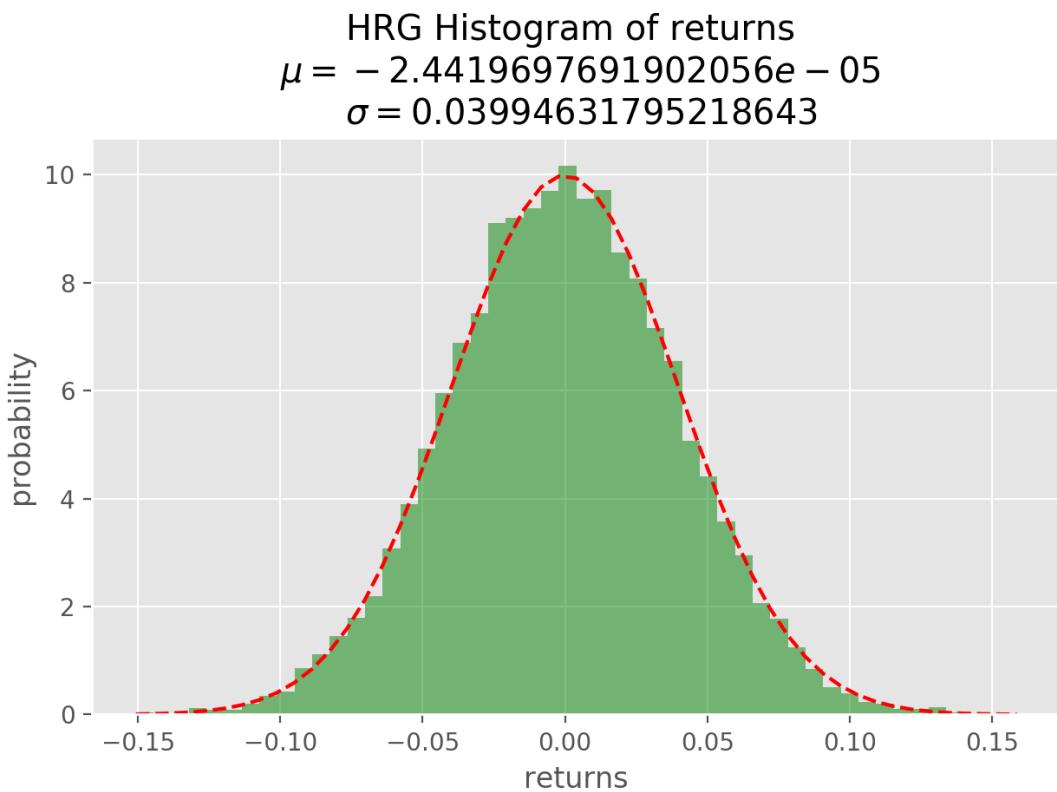


FIGURE A.66: HRG ARMA histogram of returns

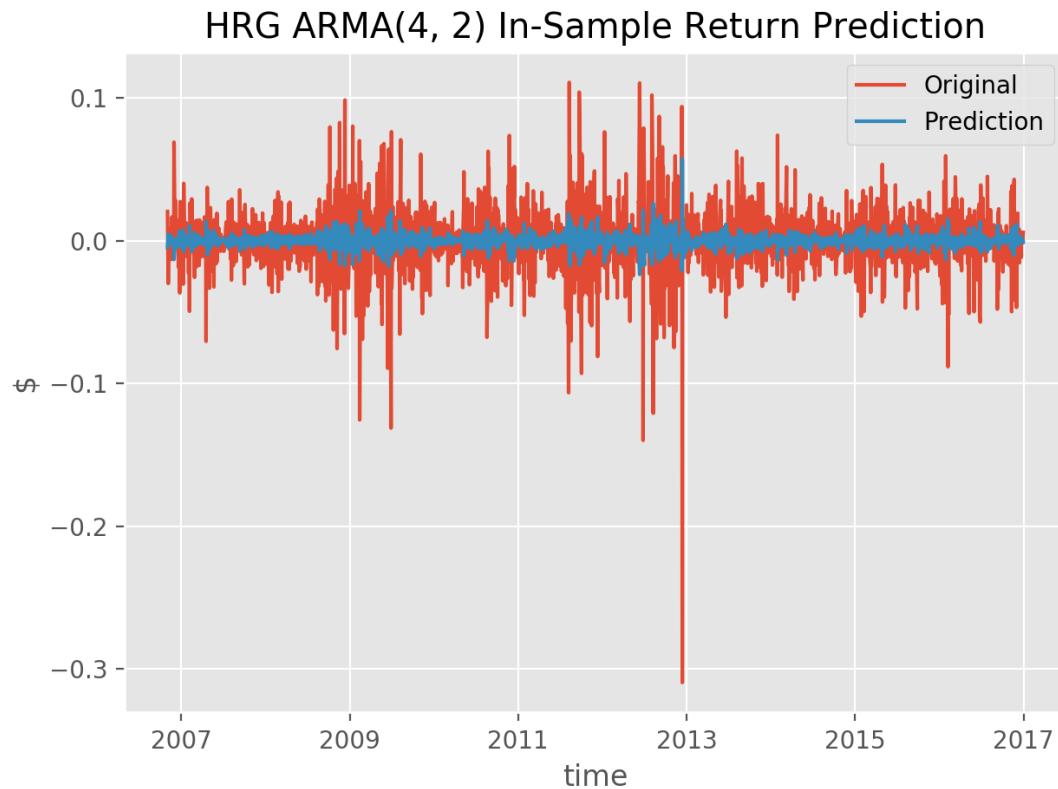


FIGURE A.67: HRG ARMA in-sample returns prediction

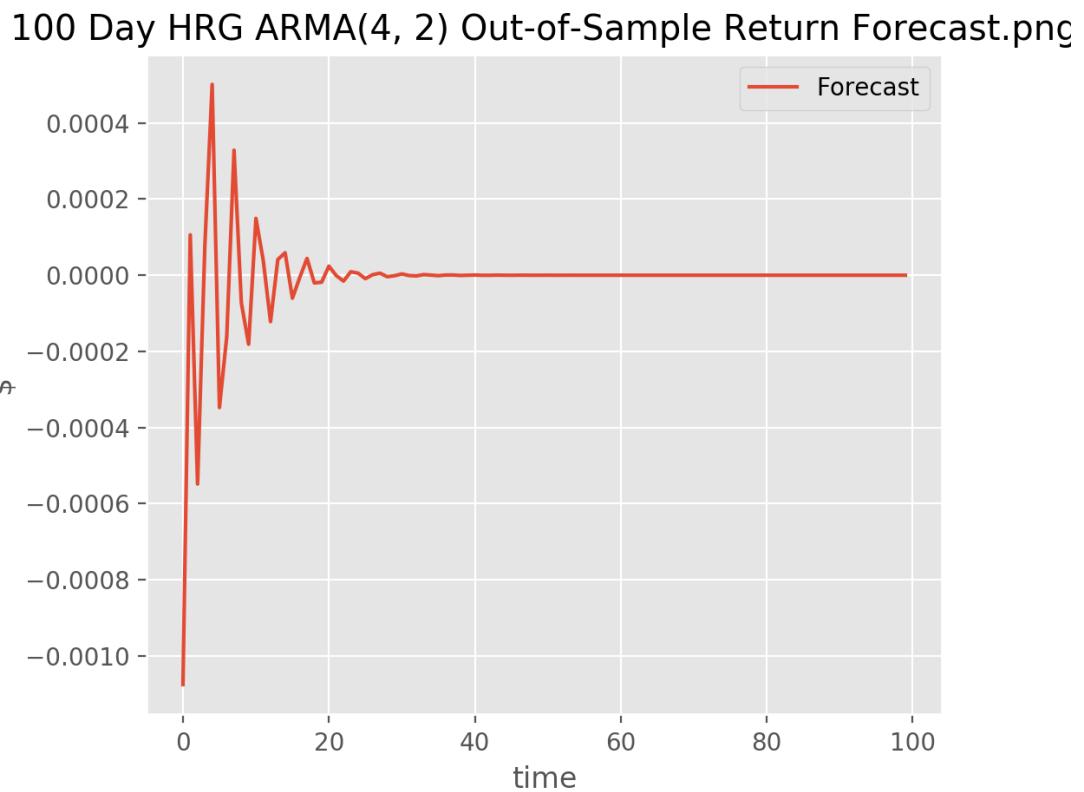


FIGURE A.68: 100 day HRG ARMA in-sample returns forecast

HL scored sharpe ratios of 0.111 for the original returns, and -1.023 for the predicted returns in the in-sample test.

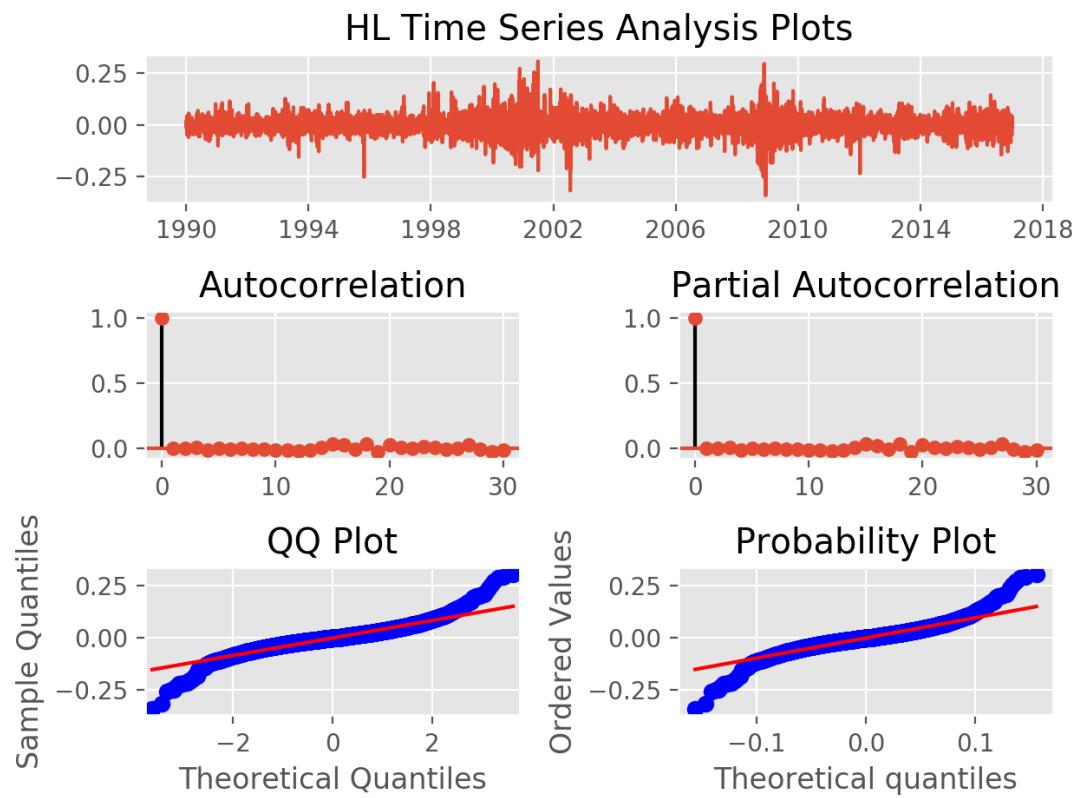


FIGURE A.69: HL ARMA time series analysis

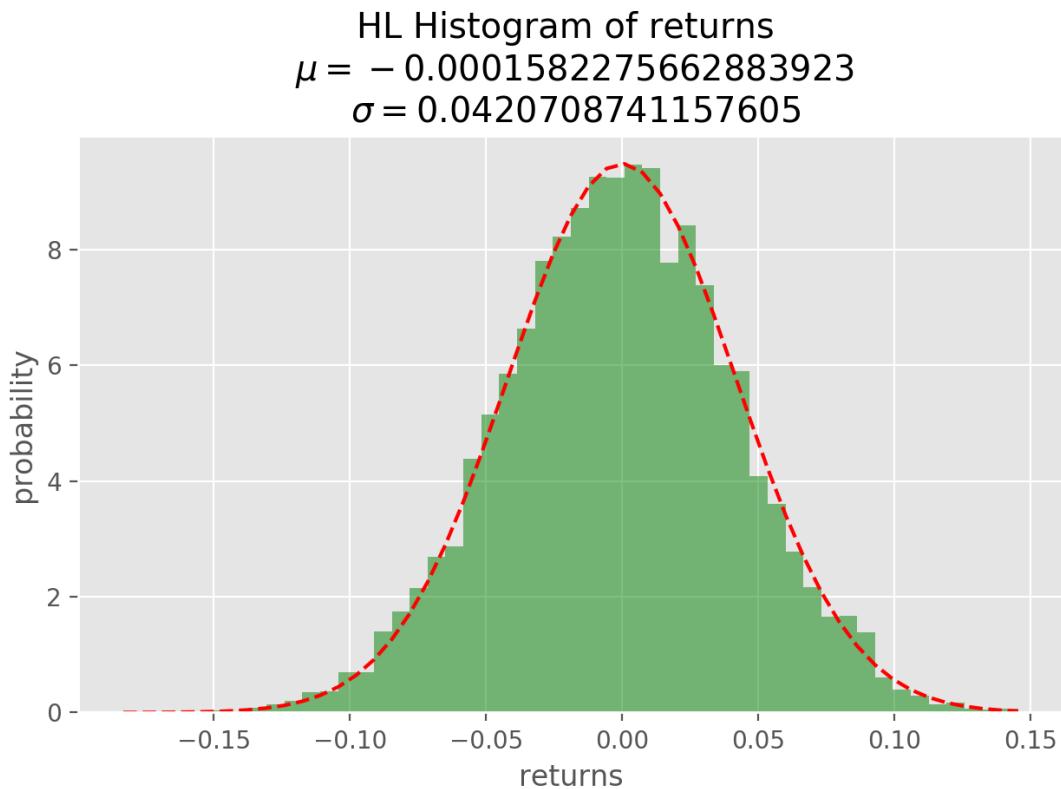


FIGURE A.70: HL ARMA histogram of returns

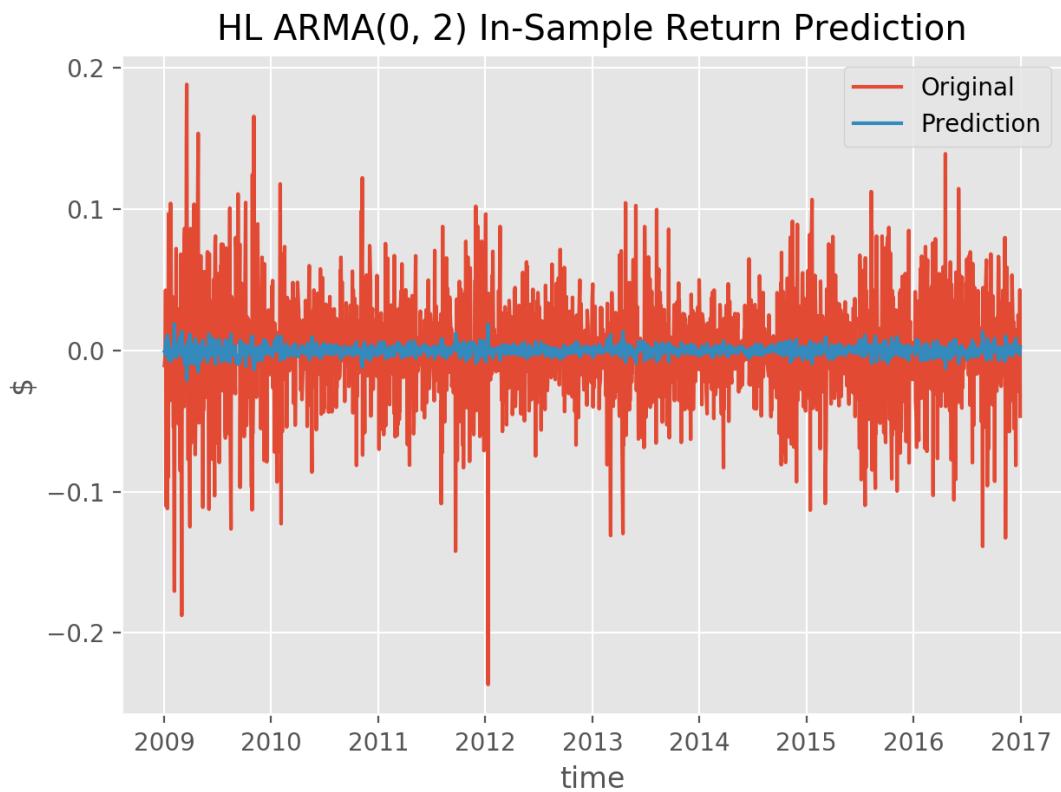


FIGURE A.71: HL ARMA in-sample returns prediction

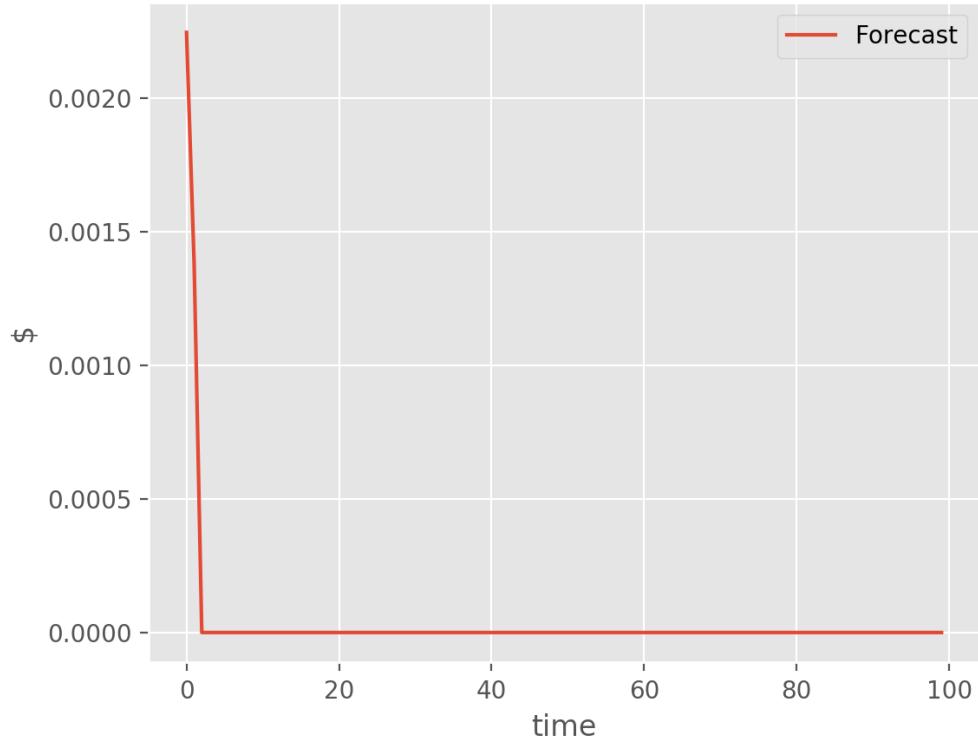
100 Day HL ARMA(0, 2) Out-of-Sample Return Forecast.png

FIGURE A.72: 100 day HL ARMA in-sample returns forecast

A.1.7 Auto Regressive Integrated Moving Average (ARIMA)

CDE scored sharpe ratios of -0.133 for the original returns, and -0.698 for the predicted returns in the in-sample test.

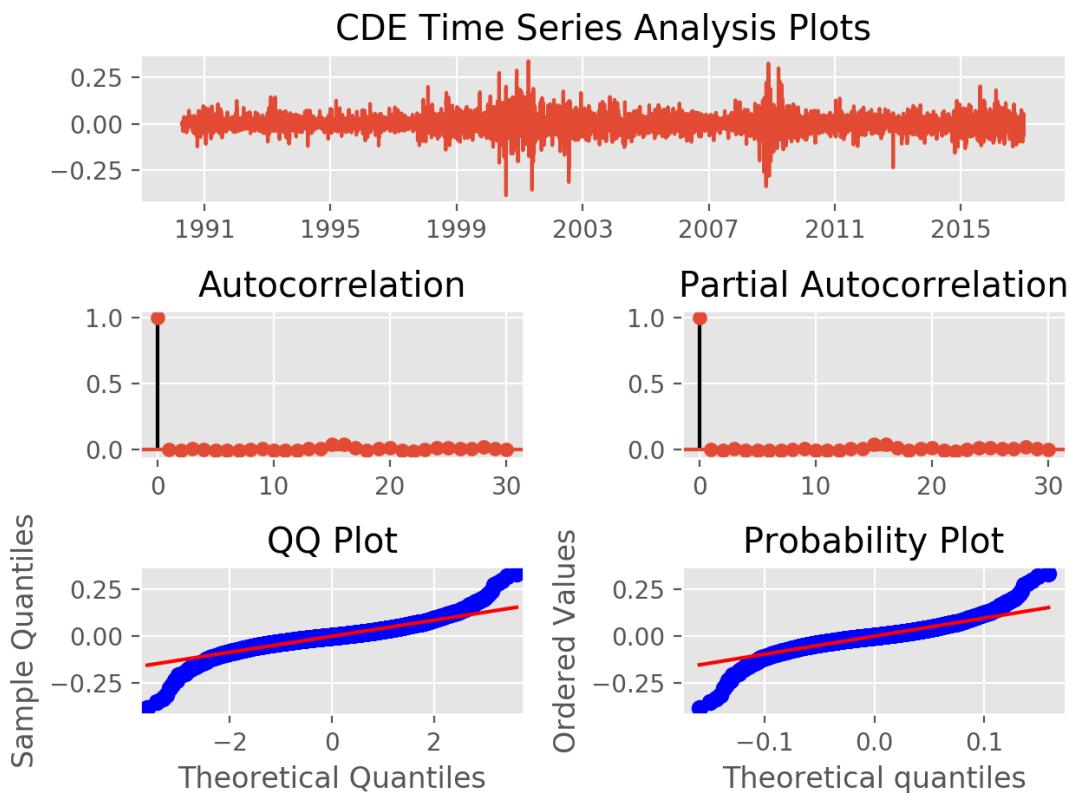


FIGURE A.73: CDE ARIMA time series analysis

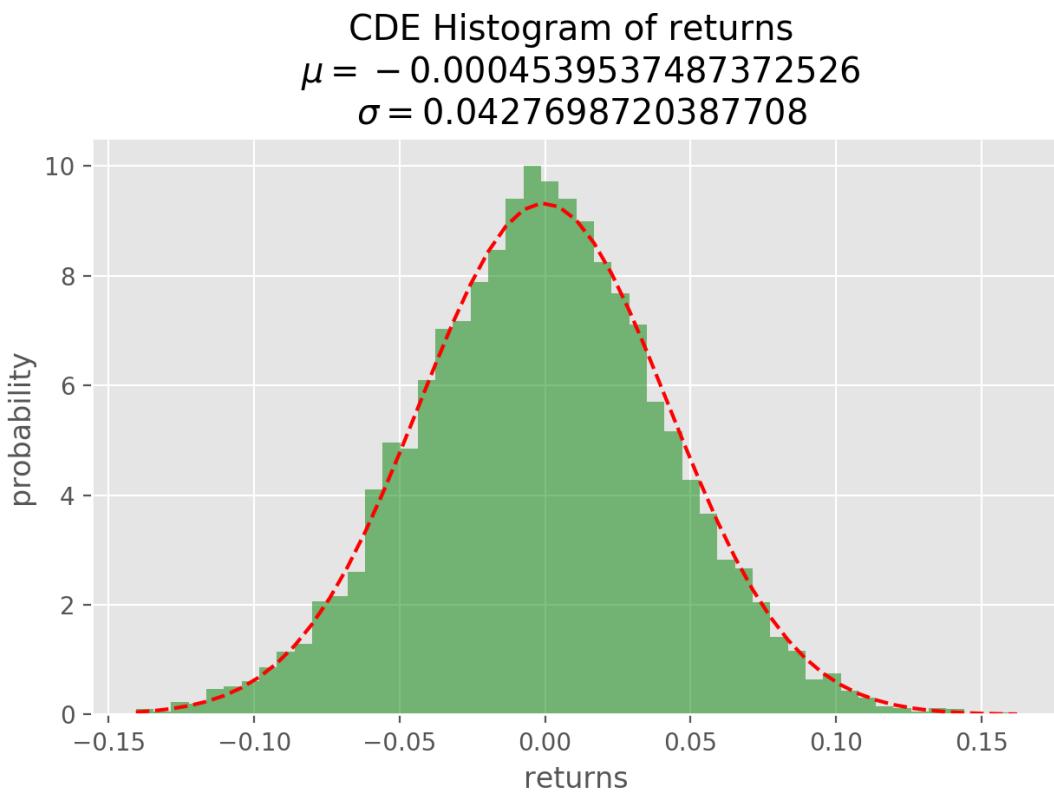


FIGURE A.74: CDE ARIMA histogram of returns

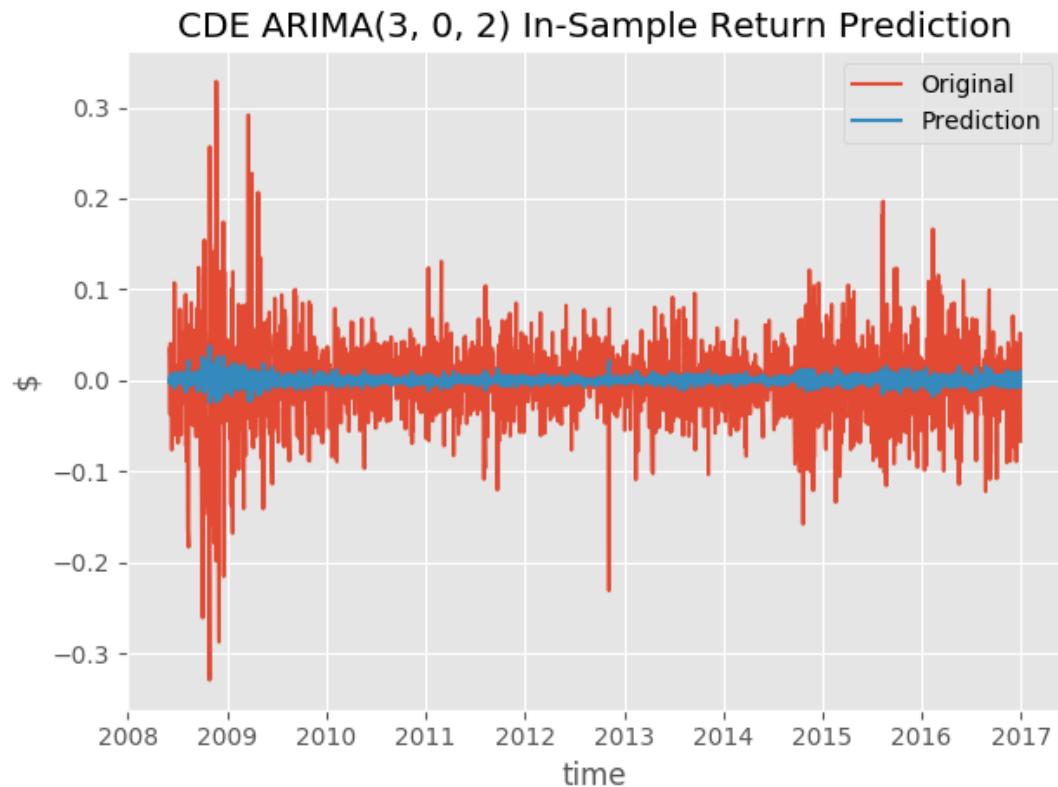


FIGURE A.75: CDE ARIMA in-sample returns prediction

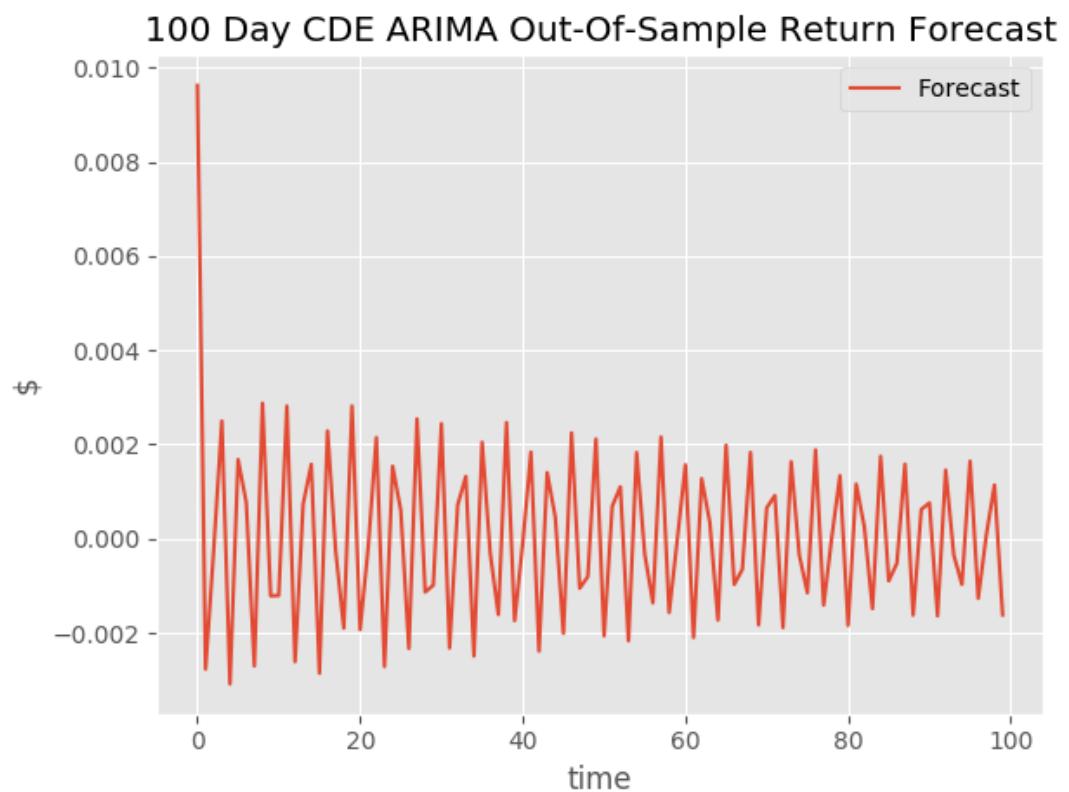


FIGURE A.76: 100 day CDE ARIMA in-sample returns forecast

NAVB scored sharpe ratios of 0.015 for the original returns, and -0.708 for the predicted returns in the in-sample test.

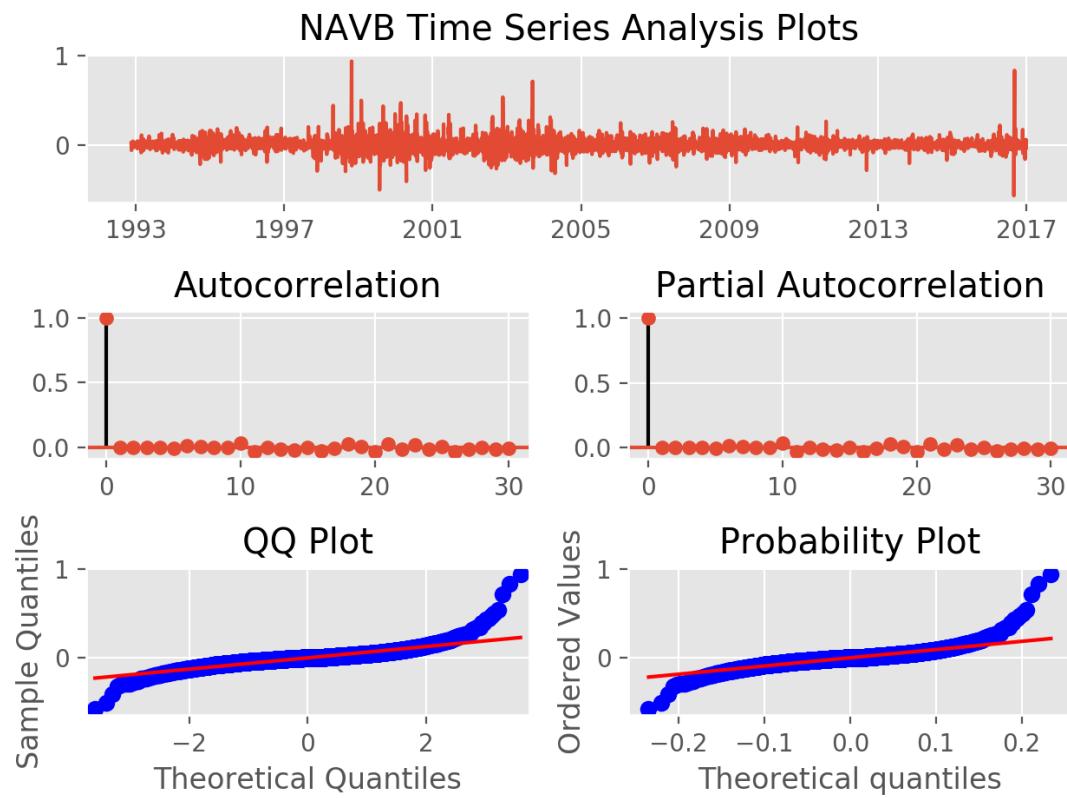


FIGURE A.77: NAVB ARIMA time series analysis

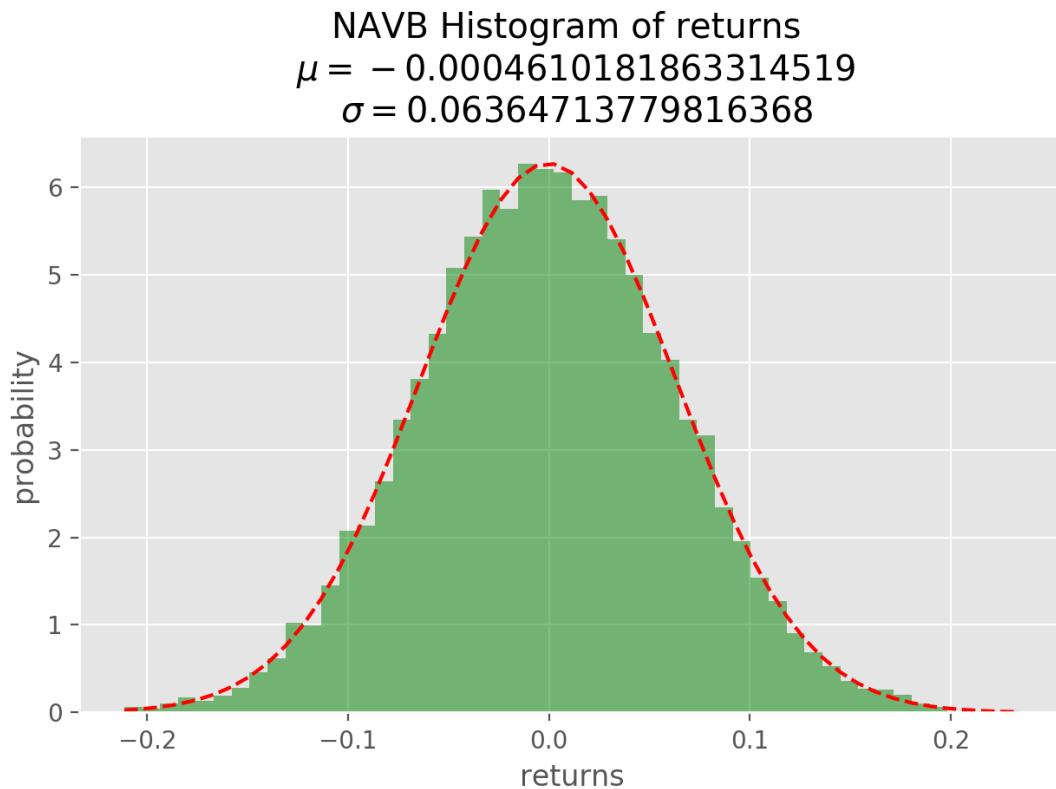


FIGURE A.78: NAVB ARIMA histogram of returns

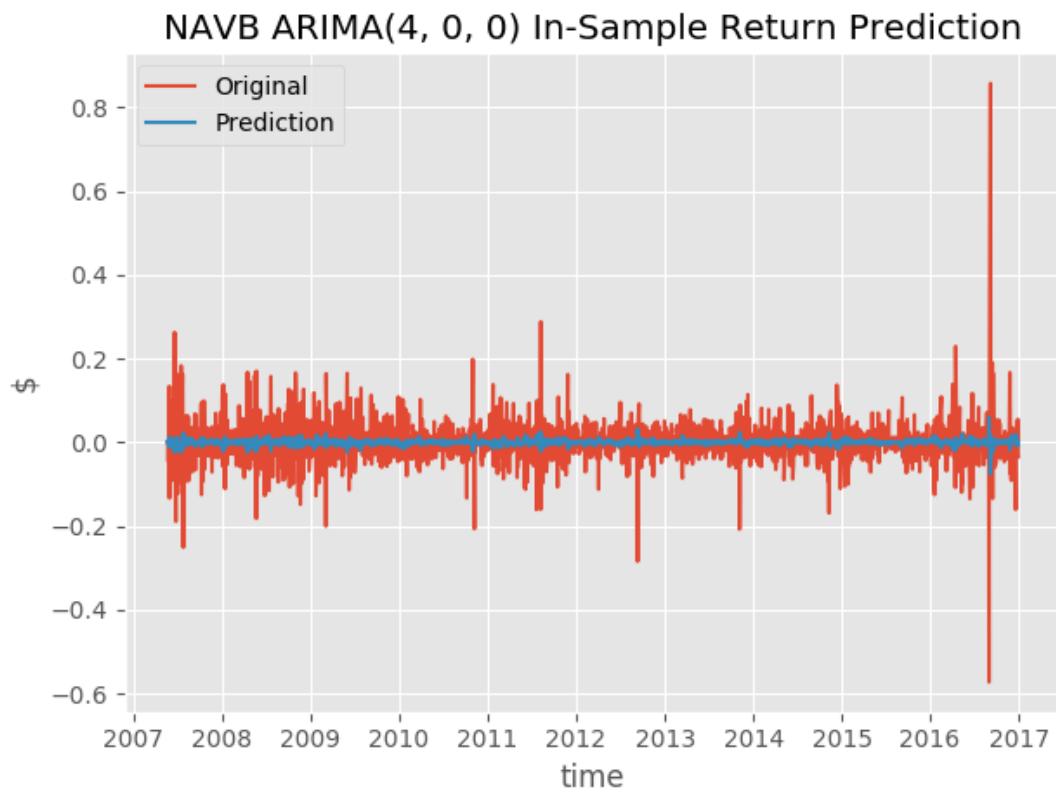


FIGURE A.79: NAVB ARIMA in-sample returns prediction

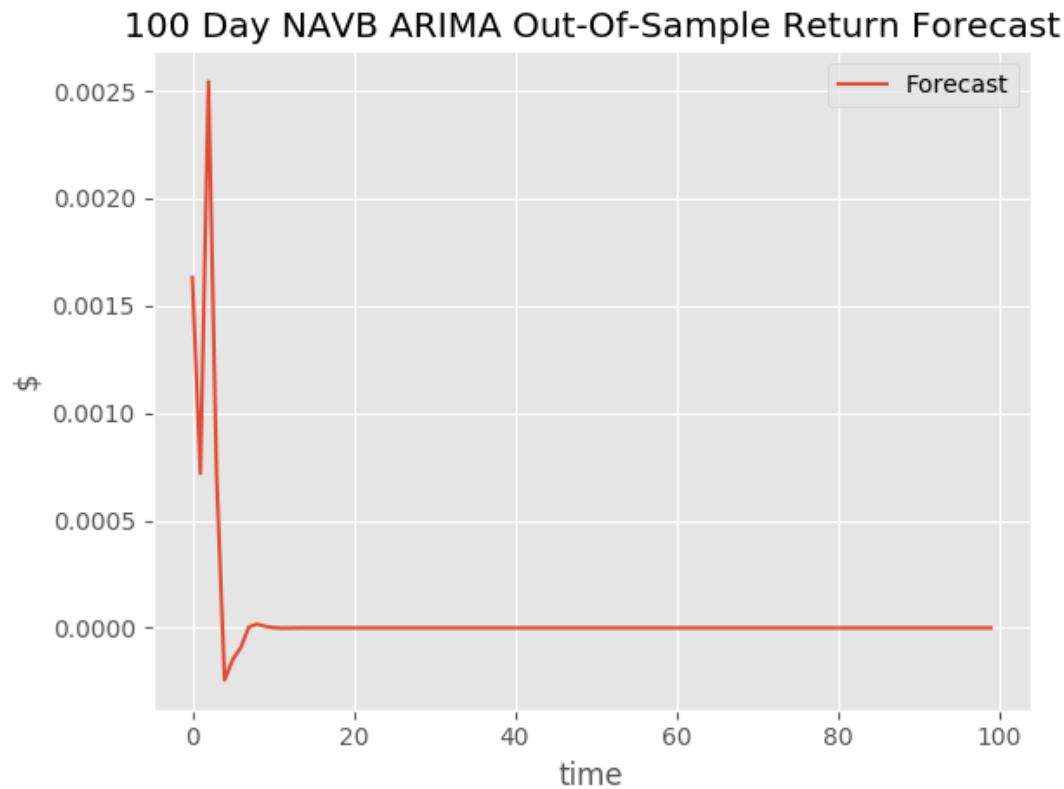


FIGURE A.80: 100 day NAVB ARIMA in-sample returns forecast

HRG scored sharpe ratios of 0.401 for the original returns, and -1.354 for the predicted returns in the in-sample test.

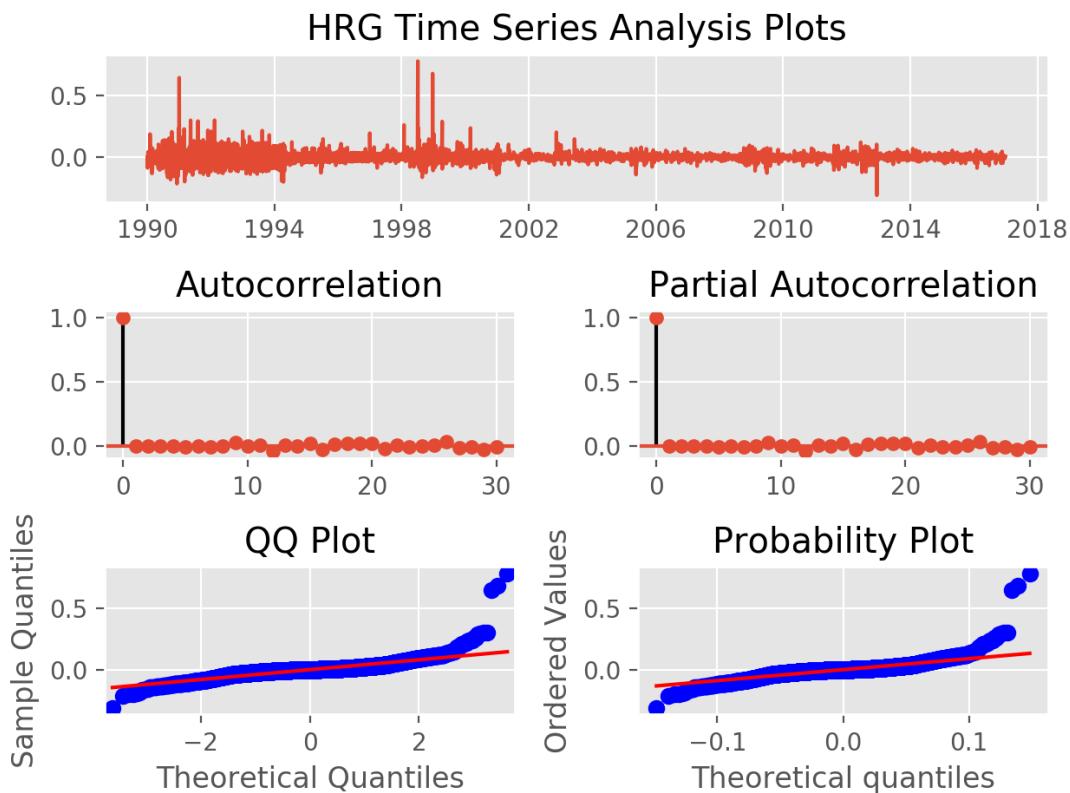


FIGURE A.81: HRG ARIMA time series analysis

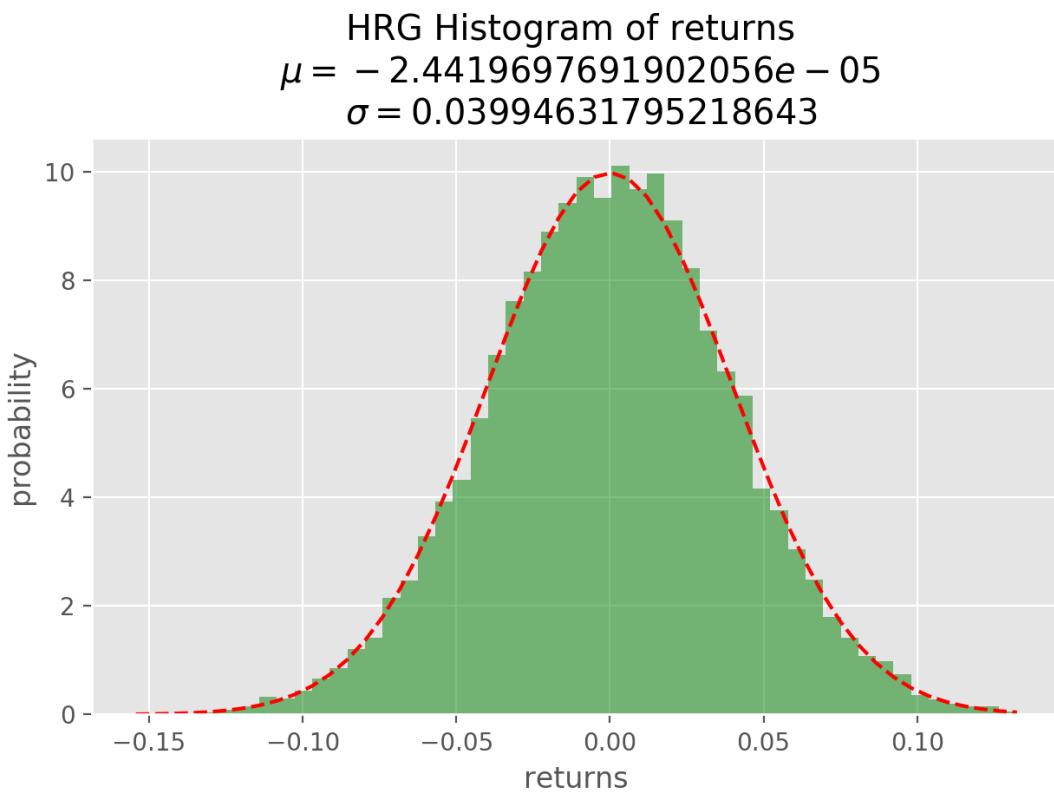


FIGURE A.82: HRG ARIMA histogram of returns

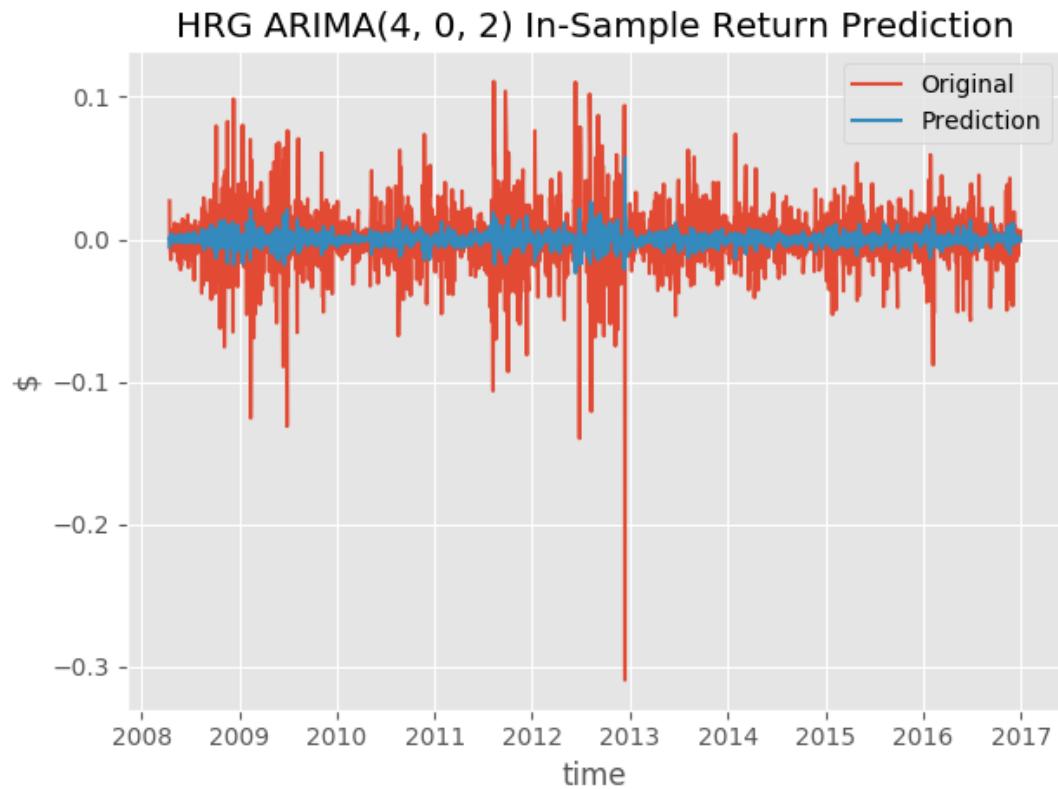


FIGURE A.83: HRG ARIMA in-sample returns prediction

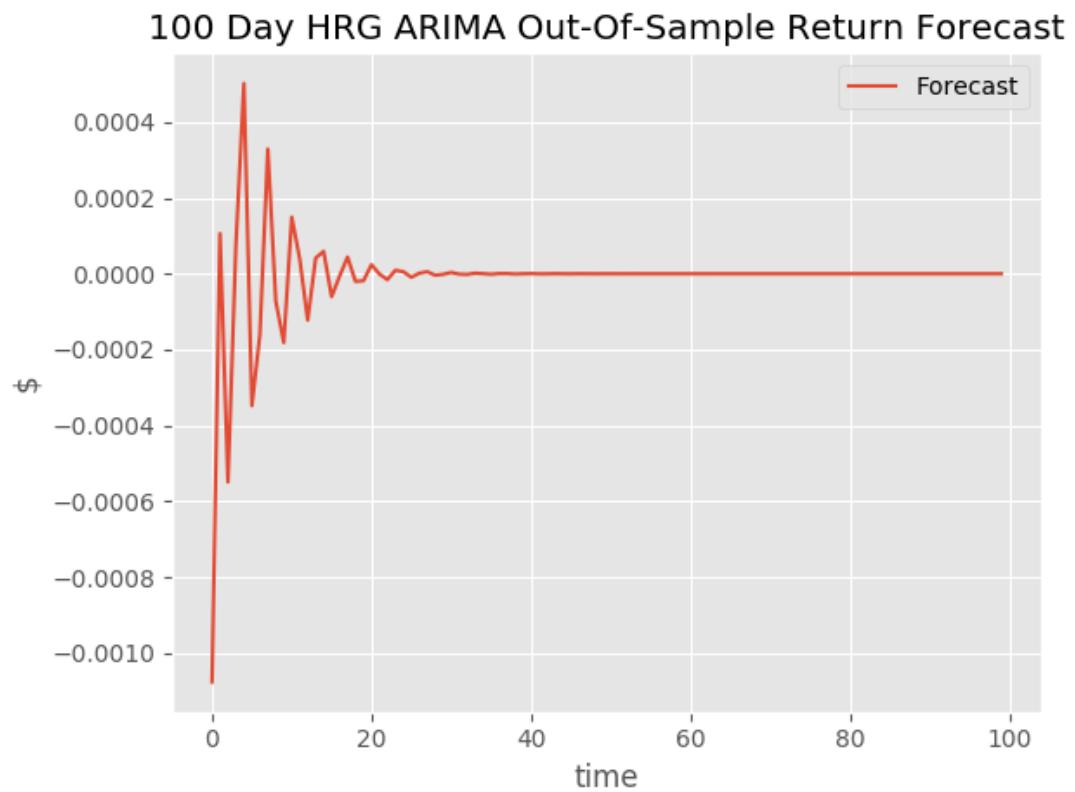


FIGURE A.84: 100 day HRG ARIMA in-sample returns forecast

HL scored sharpe ratios of -0.127 for the original returns, and -0.974 for the predicted returns in the in-sample test.

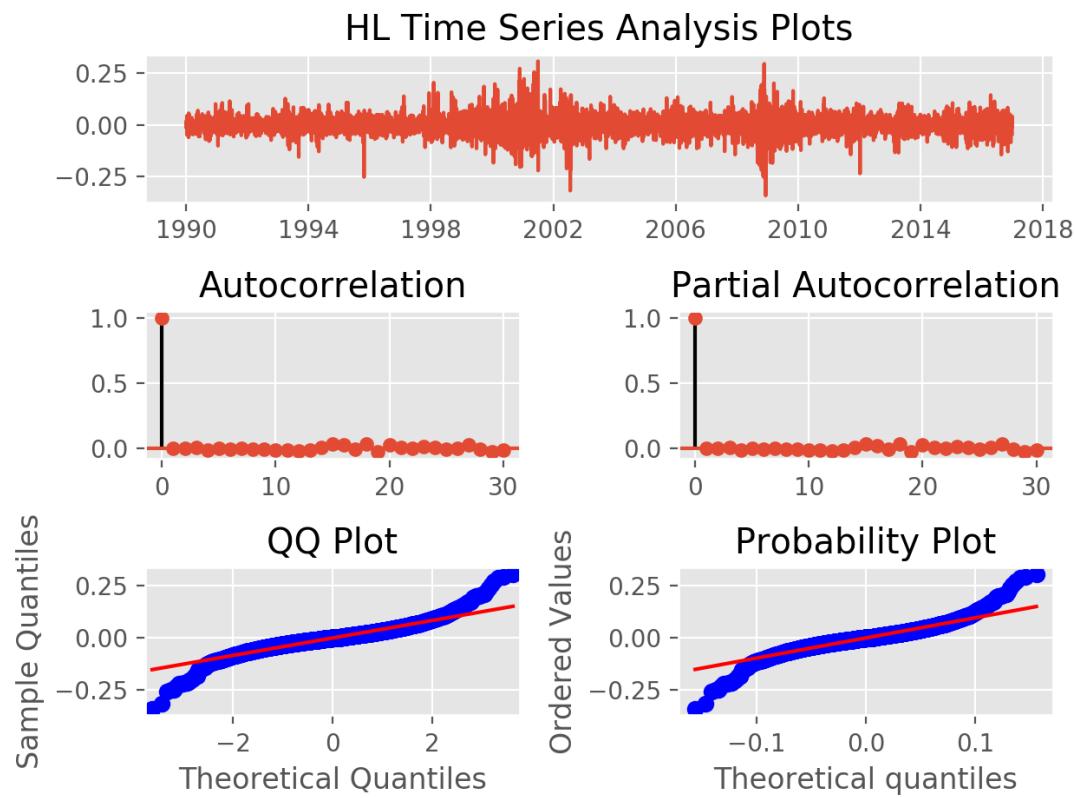


FIGURE A.85: HL ARIMA time series analysis

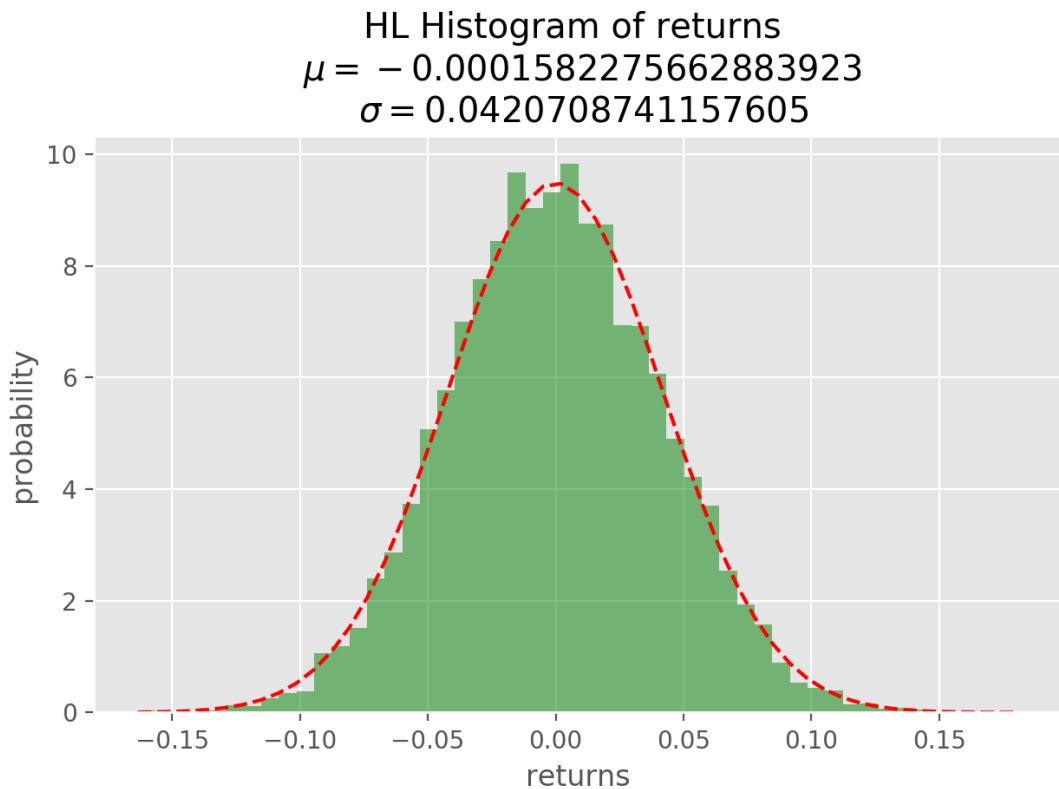


FIGURE A.86: HL ARIMA histogram of returns

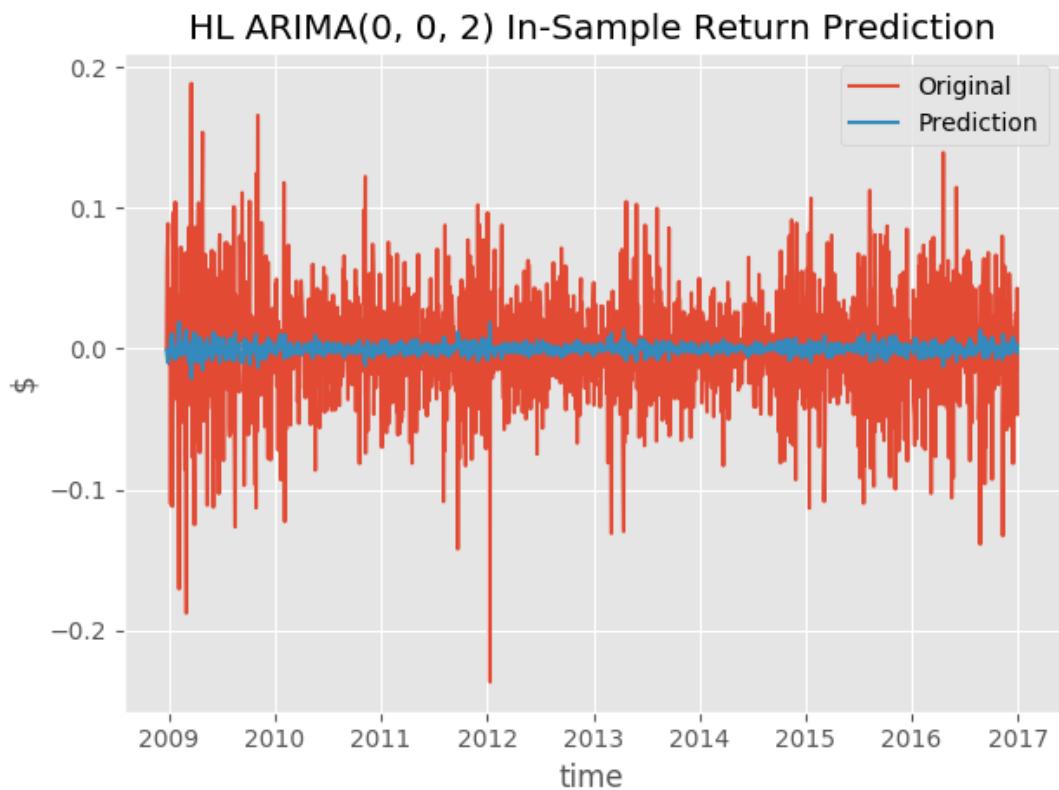


FIGURE A.87: HL ARIMA in-sample returns prediction

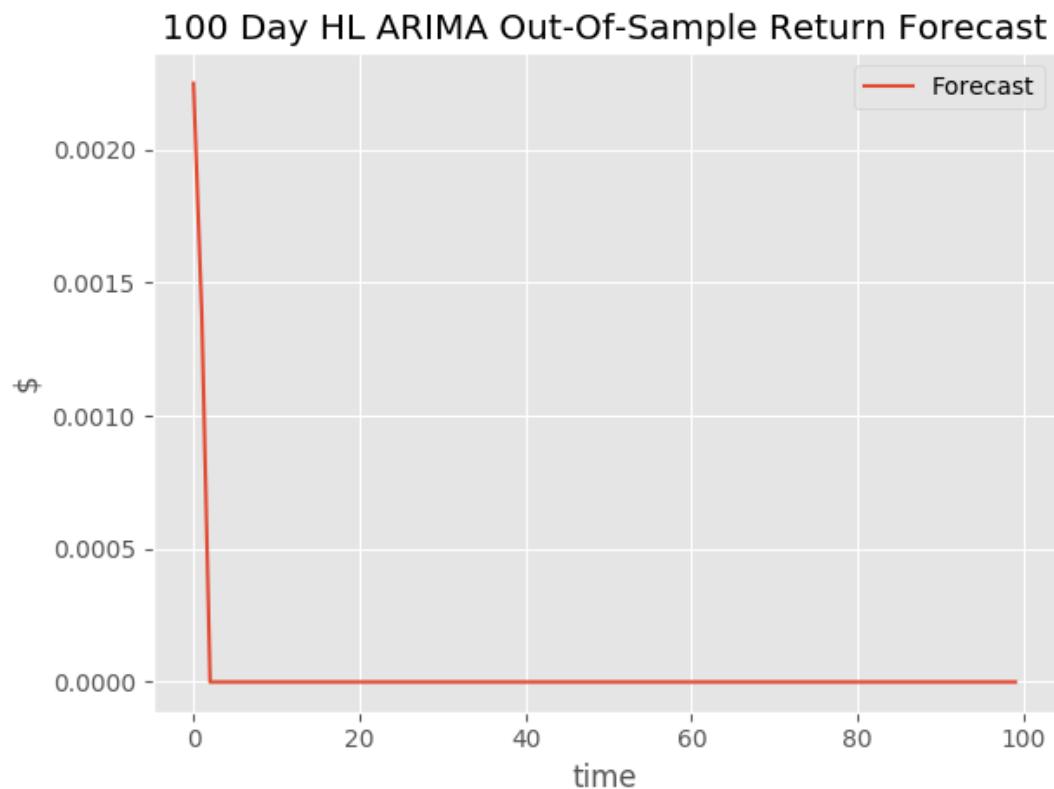


FIGURE A.88: 100 day HL ARIMA in-sample returns forecast

A.2 Machine Learning

A.2.0.1 Decision Tree

CDE scored a mean absolute error regression loss of 2.906, and a coefficient of determination of 0.816.

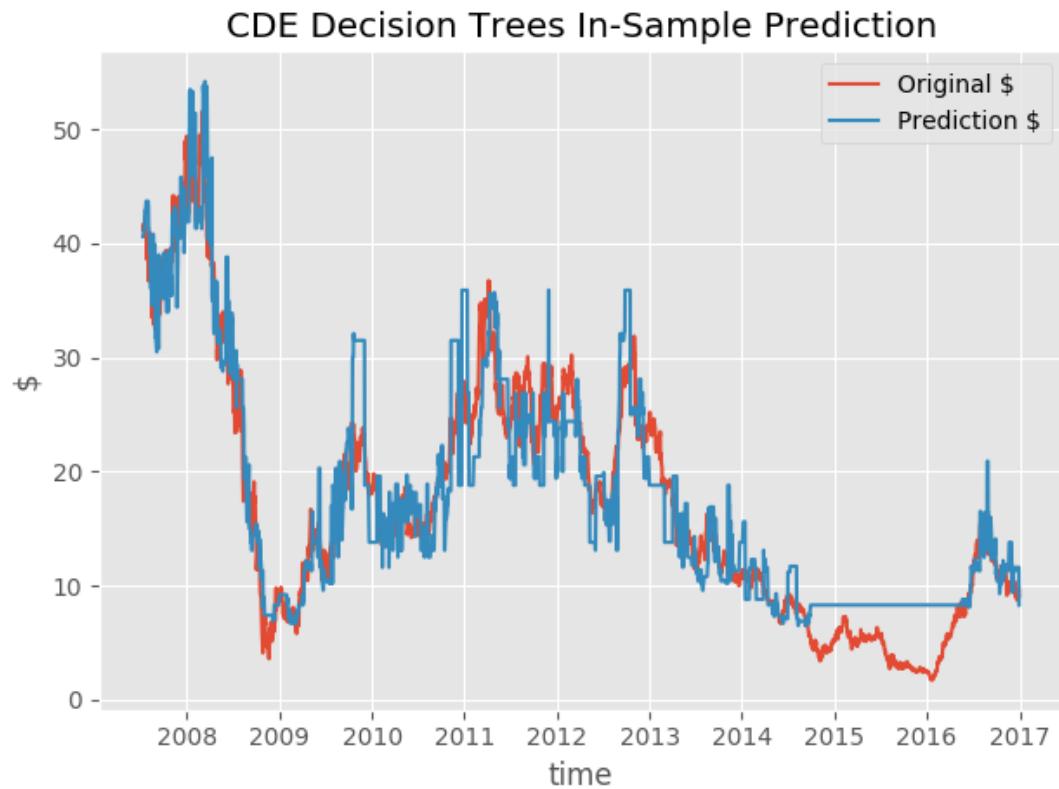


FIGURE A.89: CDE Decision Trees in-sample prediction

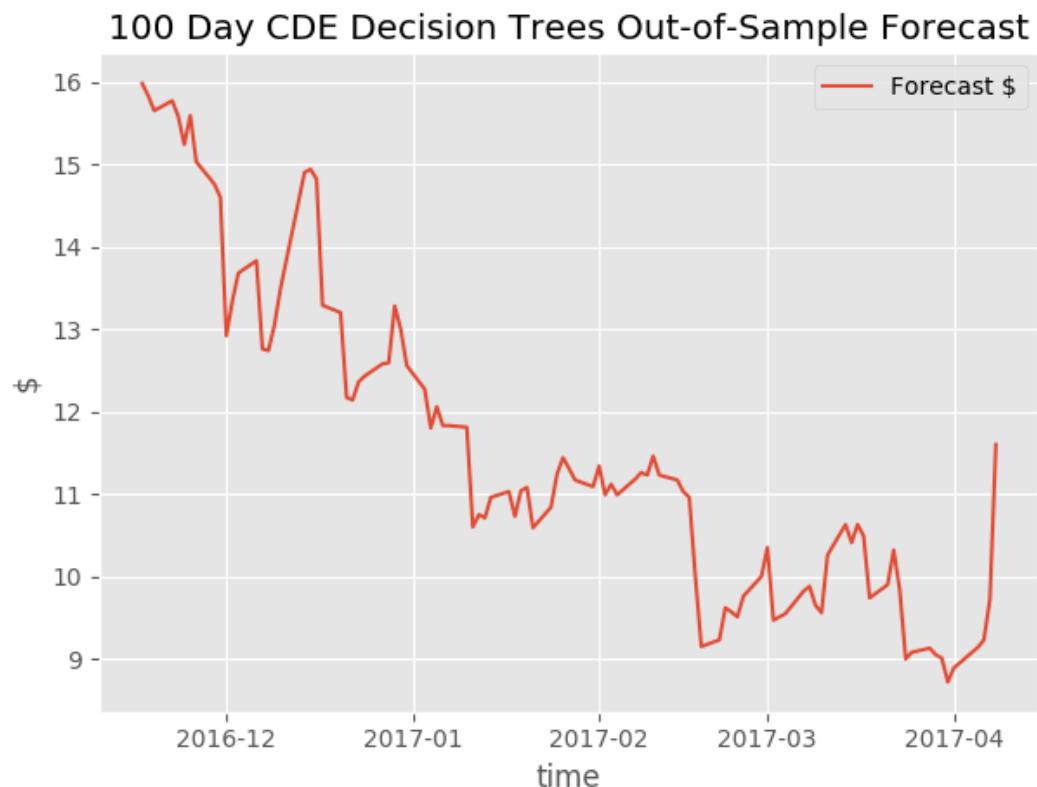


FIGURE A.90: 100 day CDE Decision Trees out-of-sample forecast

NAVB scored a mean absolute error regression loss of 0.422, and a coefficient of determination of 0.693.

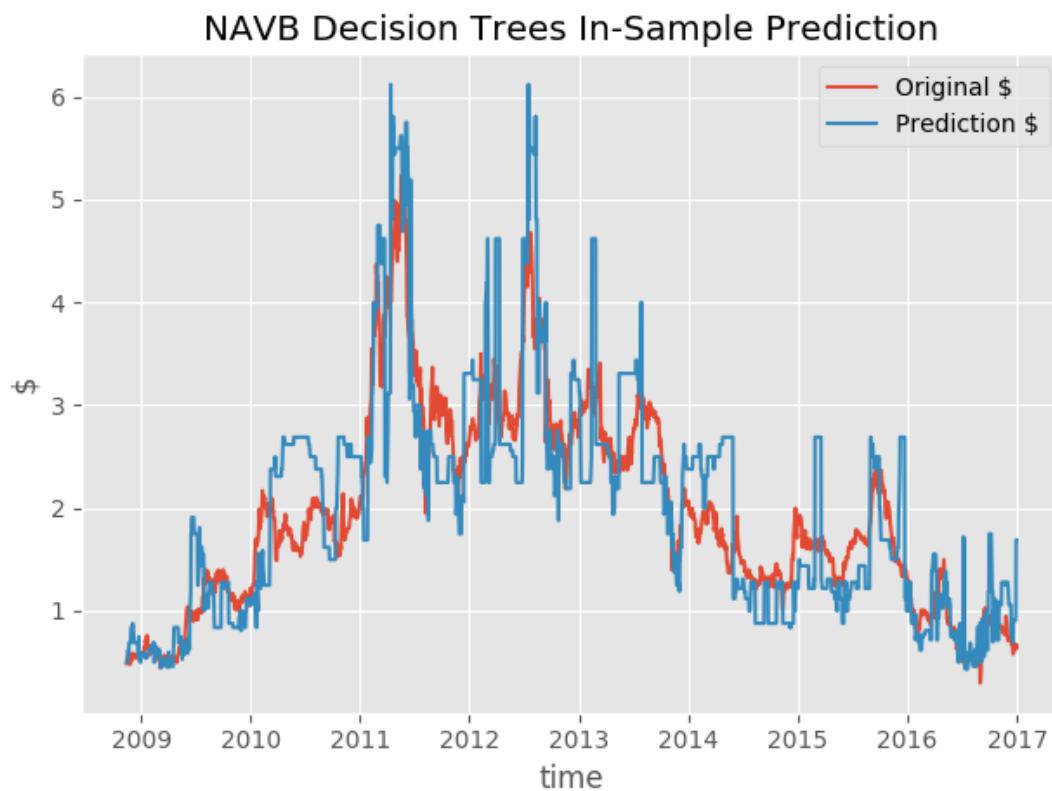


FIGURE A.91: NAVB Decision Trees in-sample prediction

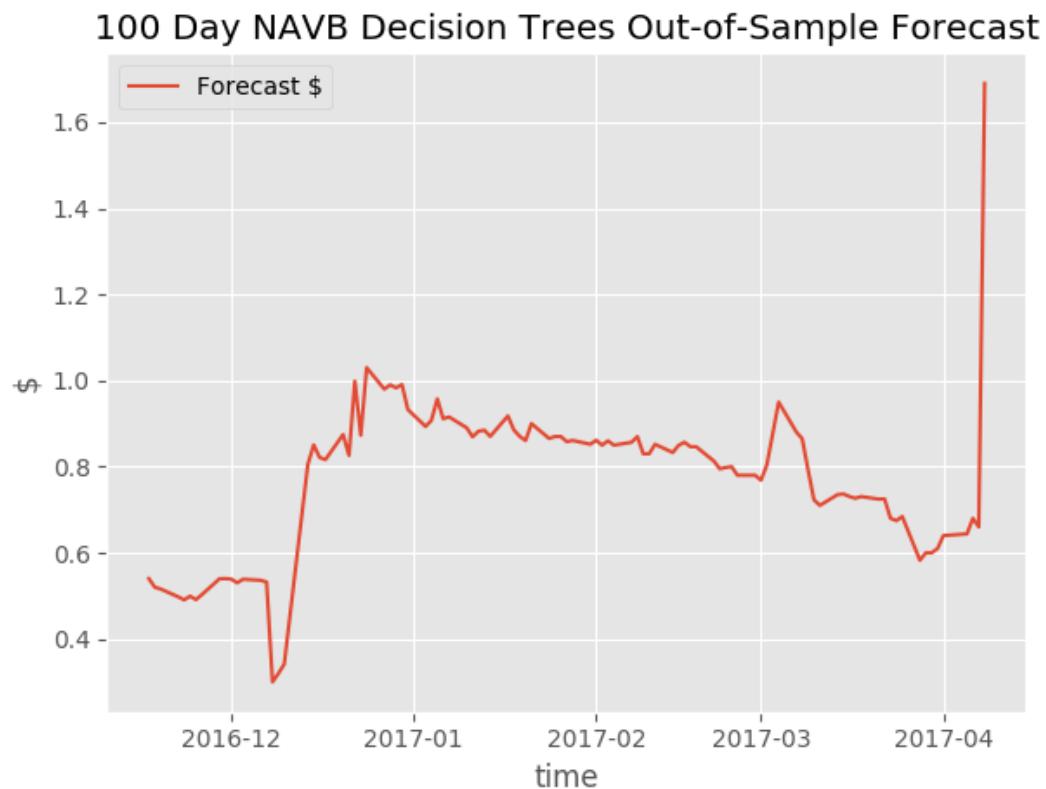


FIGURE A.92: 100 day NAVB Decision Trees out-of-sample forecast

HRG scored a mean absolute error regression loss of 1.415, and a coefficient of determination of 0.464.

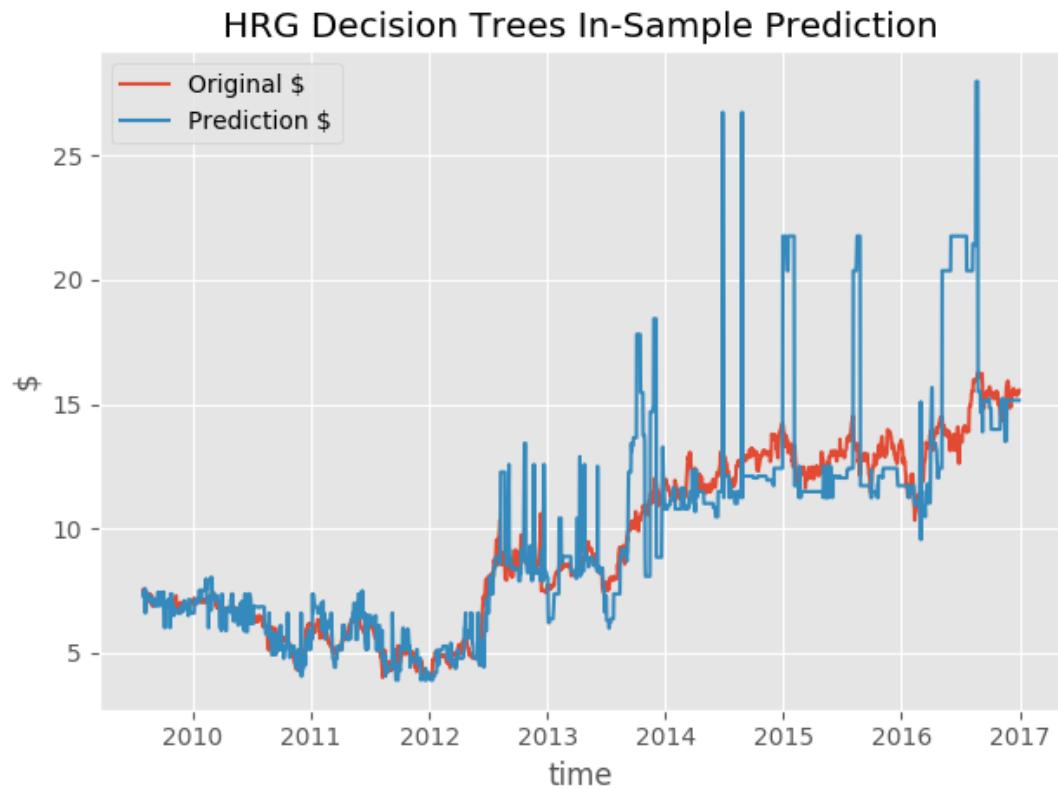


FIGURE A.93: HRG Decision Trees in-sample prediction

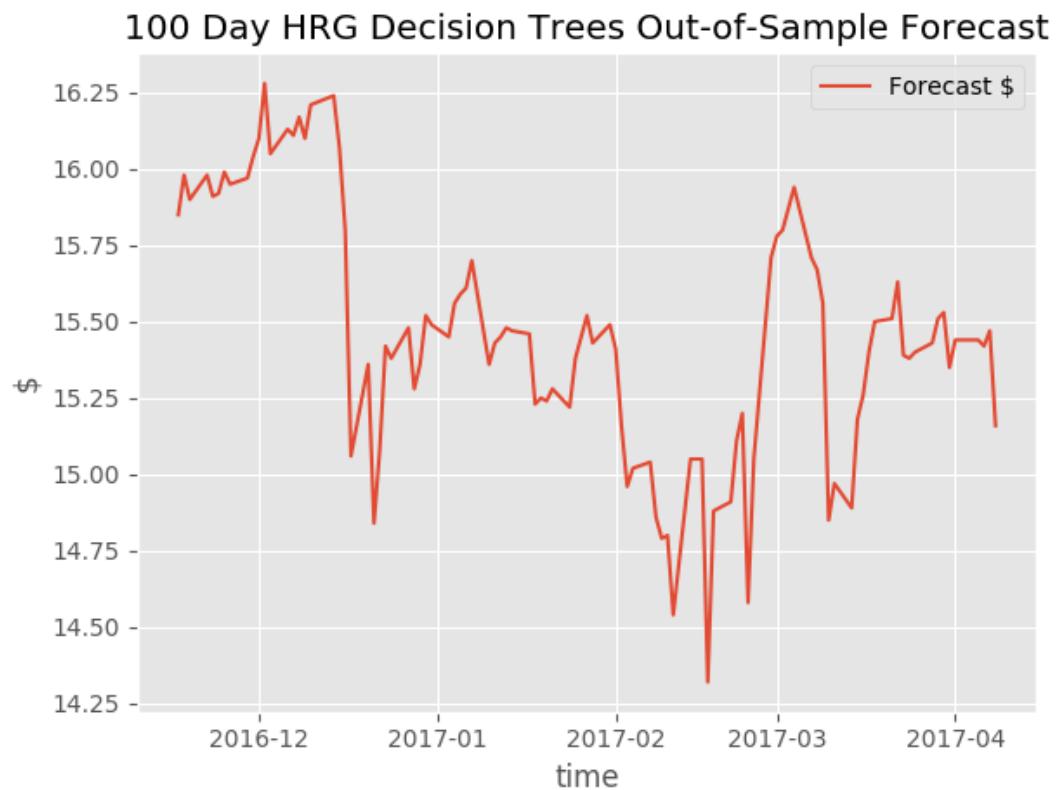


FIGURE A.94: 100 day HRG Decision Trees out-of-sample forecast

HL scored a mean absolute error regression loss of 0.506, and a coefficient of determination of 0.923.

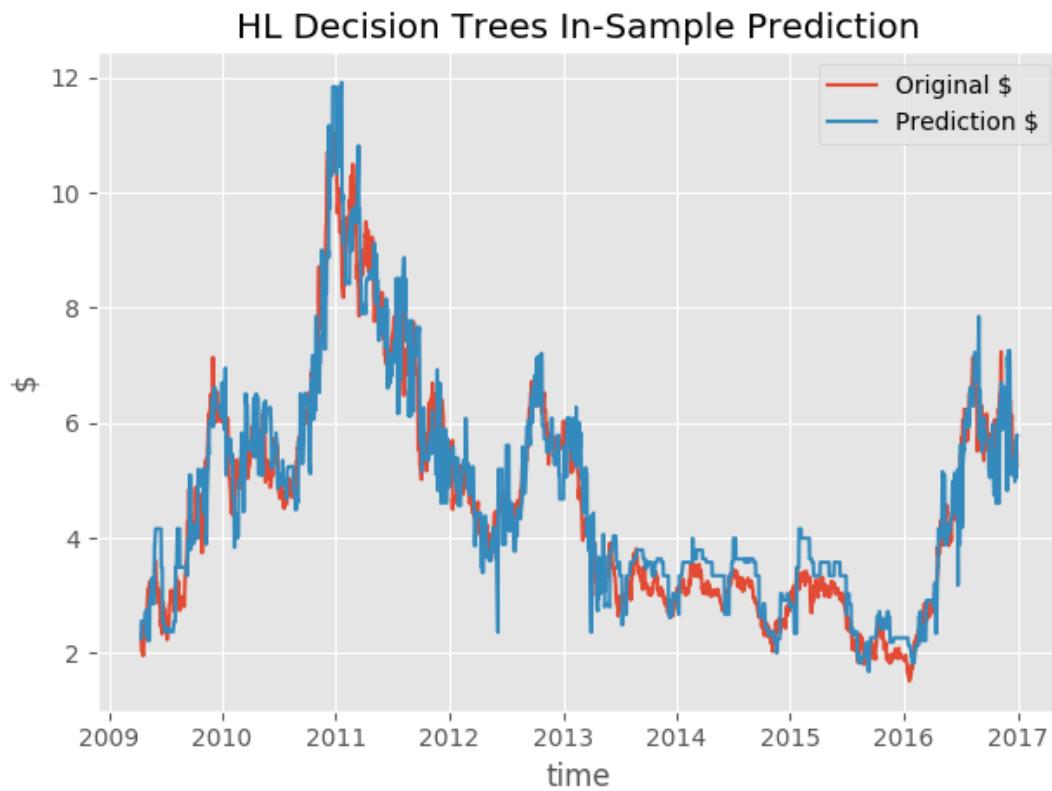


FIGURE A.95: HL Decision Trees in-sample prediction

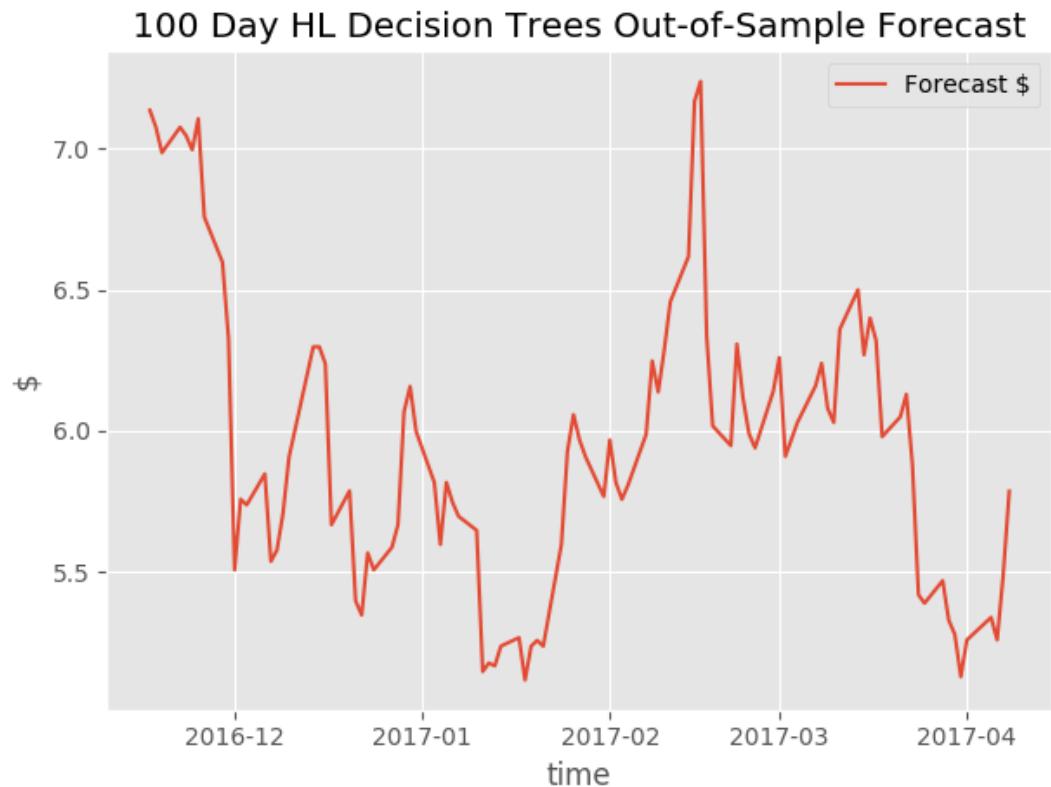


FIGURE A.96: 100 day HL Decision Trees out-of-sample forecast

A.2.0.2 Boosted Decision Tree

CDE scored a mean absolute error regression loss of 5.452, and a coefficient of determination of 0.581.

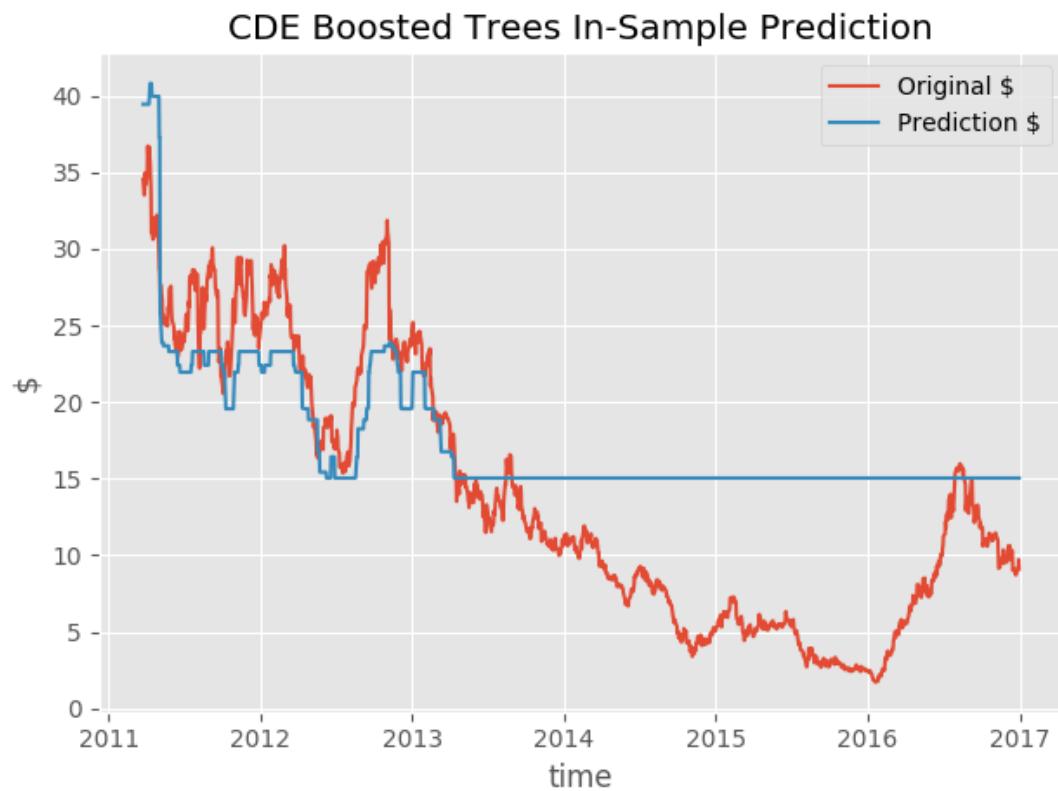


FIGURE A.97: CDE Boosted Decision Trees in-sample prediction

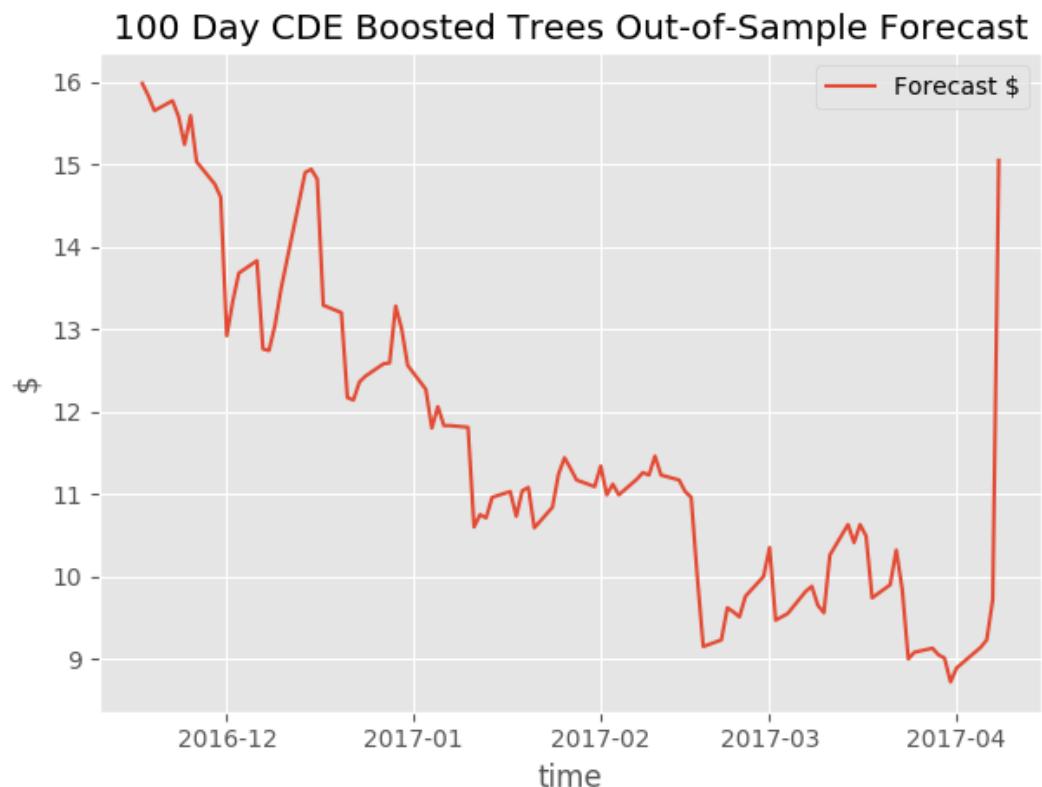


FIGURE A.98: 100 day CDE Boosted Decision Trees out-of-sample forecast

NAVB scored a mean absolute error regression loss of 0.579, and a coefficient of determination of 0.509.

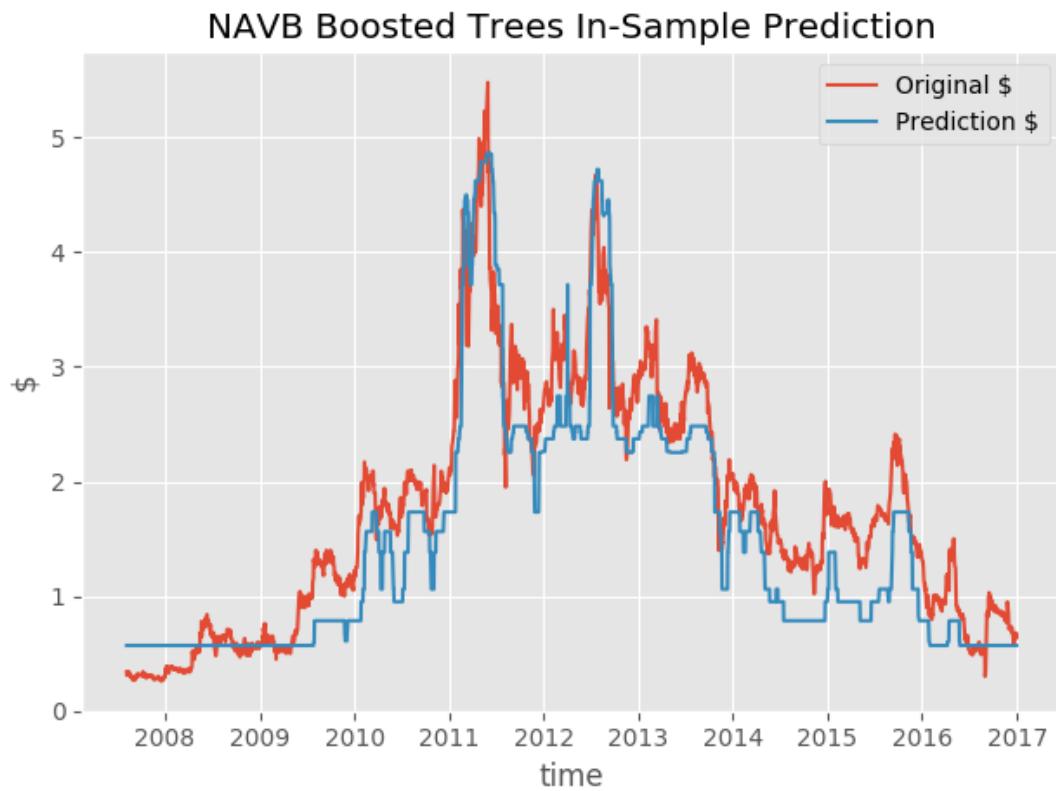


FIGURE A.99: NAVB Boosted Decision Trees in-sample prediction

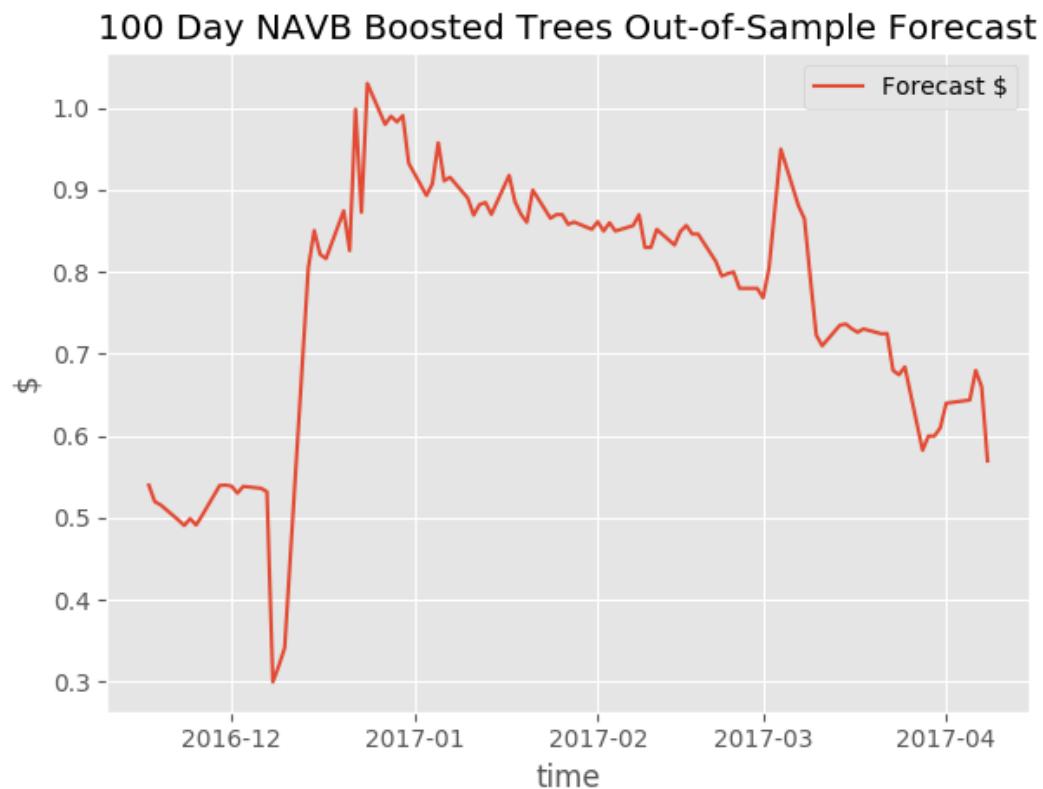


FIGURE A.100: 100 day NAVB Boosted Decision Trees out-of-sample forecast

HRG scored a mean absolute error regression loss of 0.980, and a coefficient of determination of 0.855.

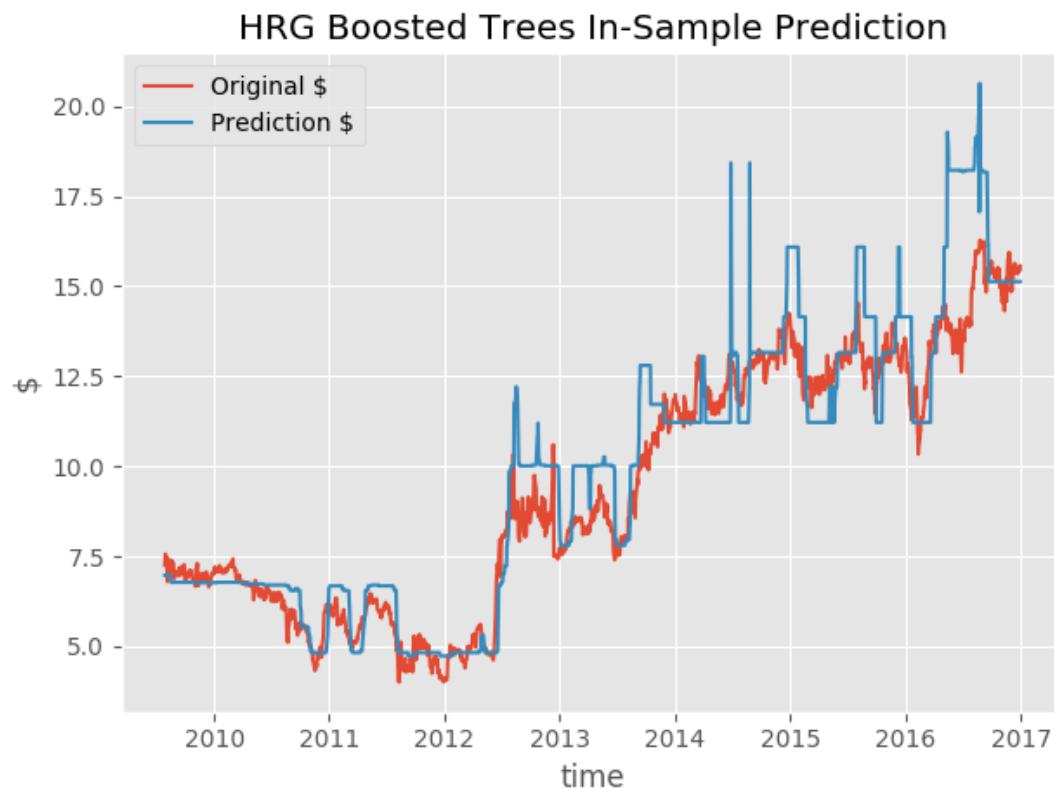


FIGURE A.101: HRG Boosted Decision Trees in-sample prediction

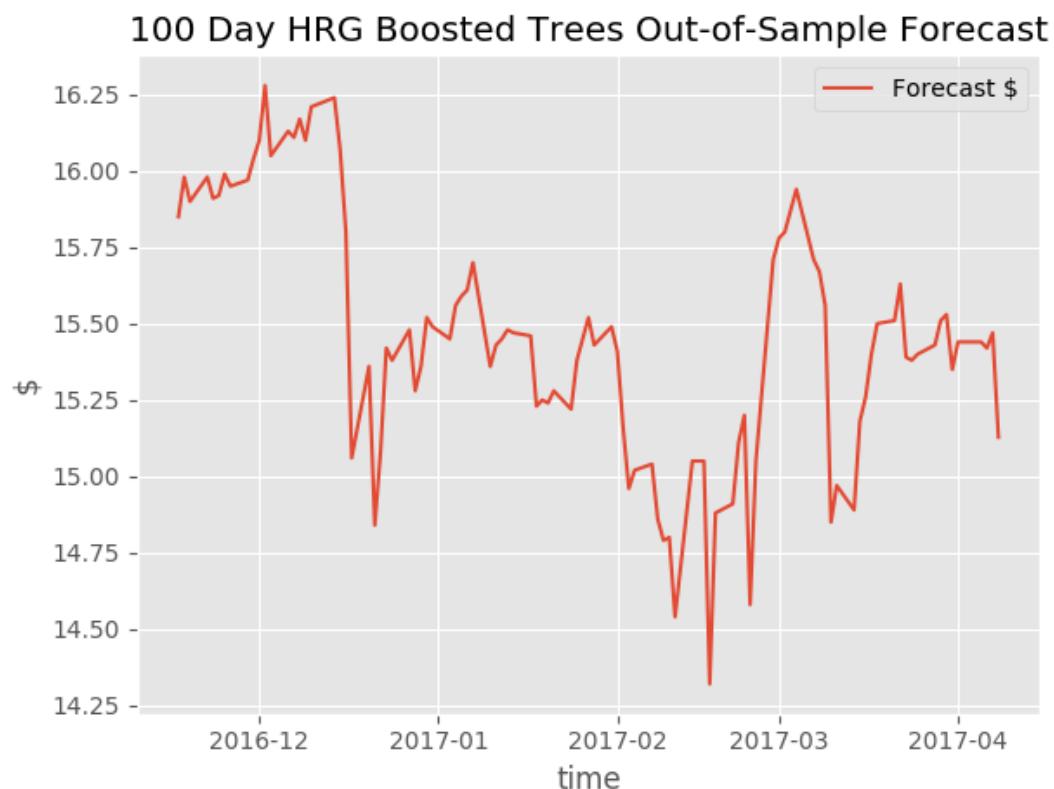


FIGURE A.102: 100 day HRG Boosted Decision Trees out-of-sample forecast

HL scored a mean absolute error regression loss of 0.410, and a coefficient of determination of 0.952.

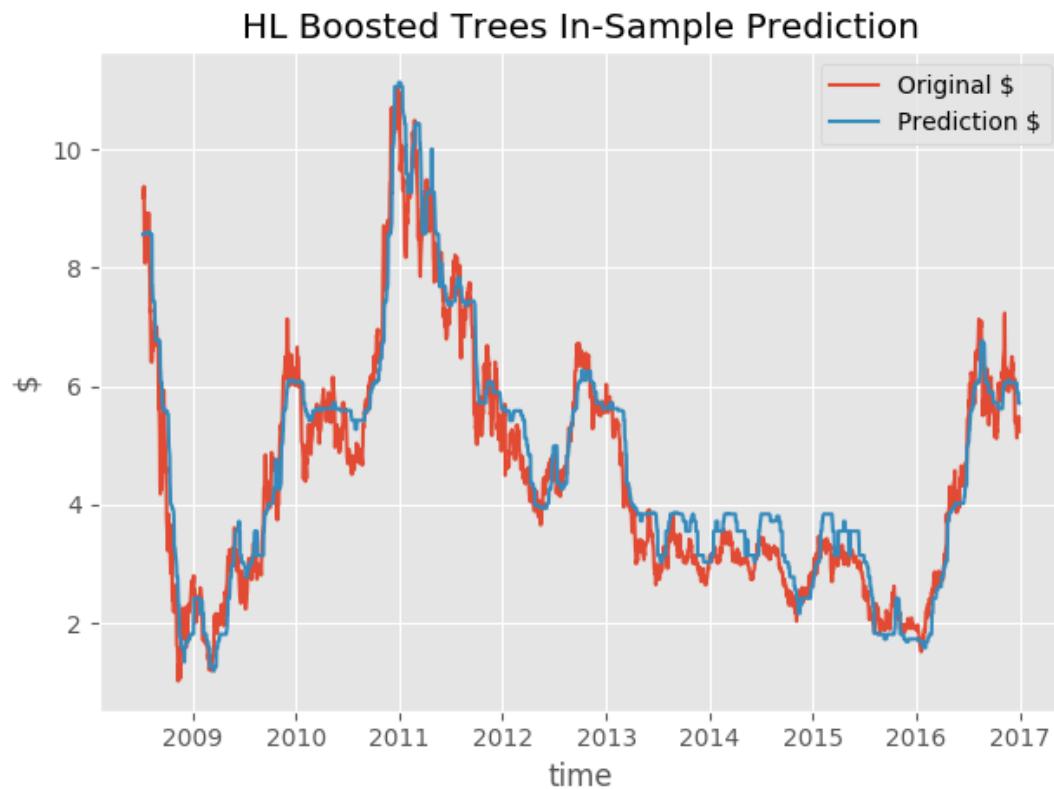


FIGURE A.103: HL Boosted Decision Trees in-sample prediction

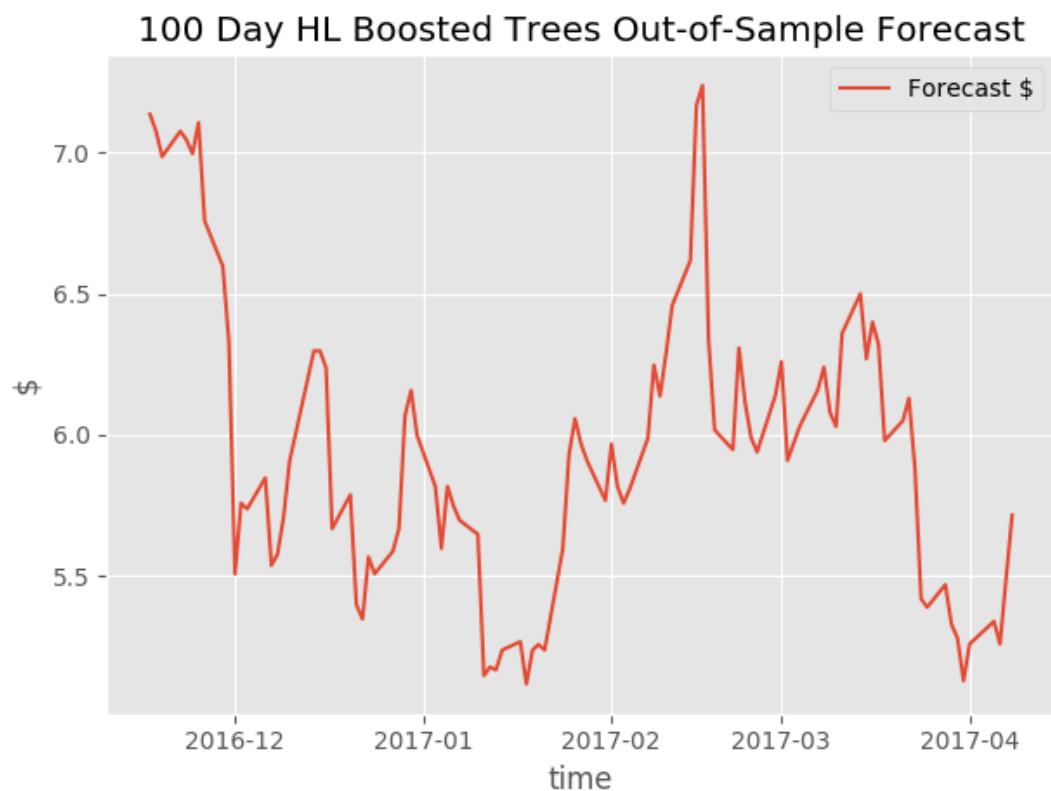


FIGURE A.104: 100 day HL Boosted Decision Trees out-of-sample forecast

A.2.0.3 K-Nearest Neighbour

CDE scored a mean absolute error regression loss of 5.452, and a coefficient of determination of 0.581.

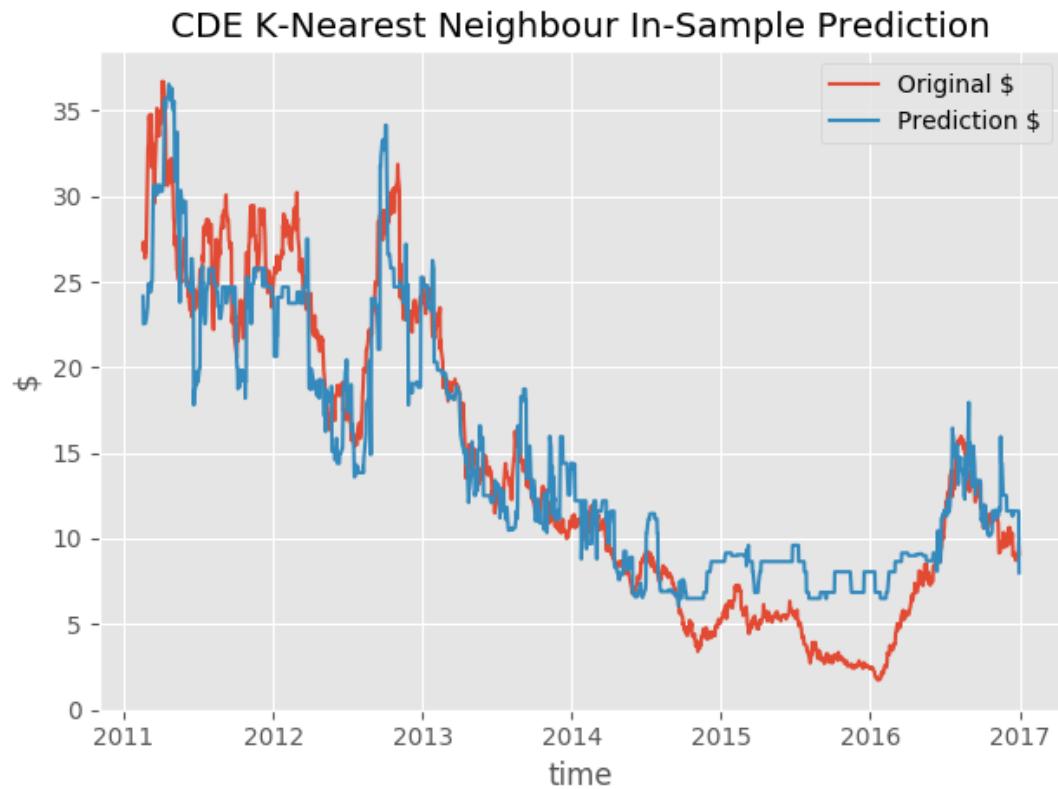


FIGURE A.105: CDE K-Nearest Neighbour in-sample prediction

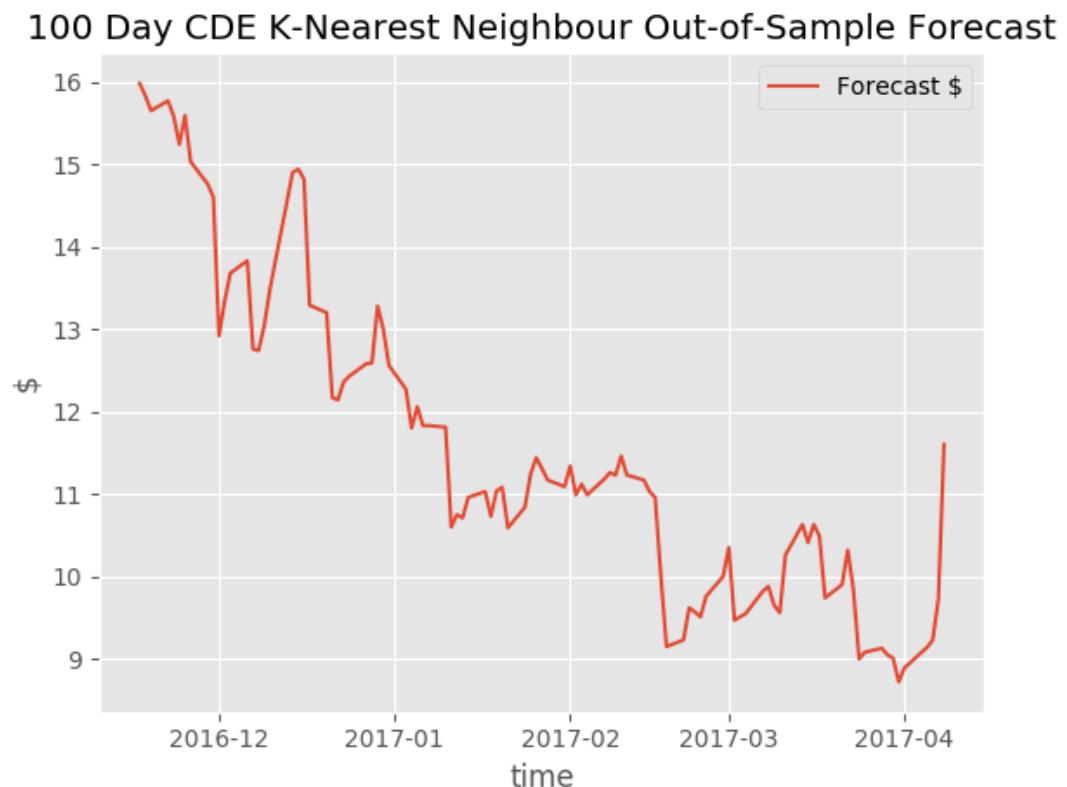


FIGURE A.106: 100 day CDE K-Nearest Neighbour out-of-sample forecast

NAVB scored a mean absolute error regression loss of 0.579, and a coefficient of determination of 0.509.

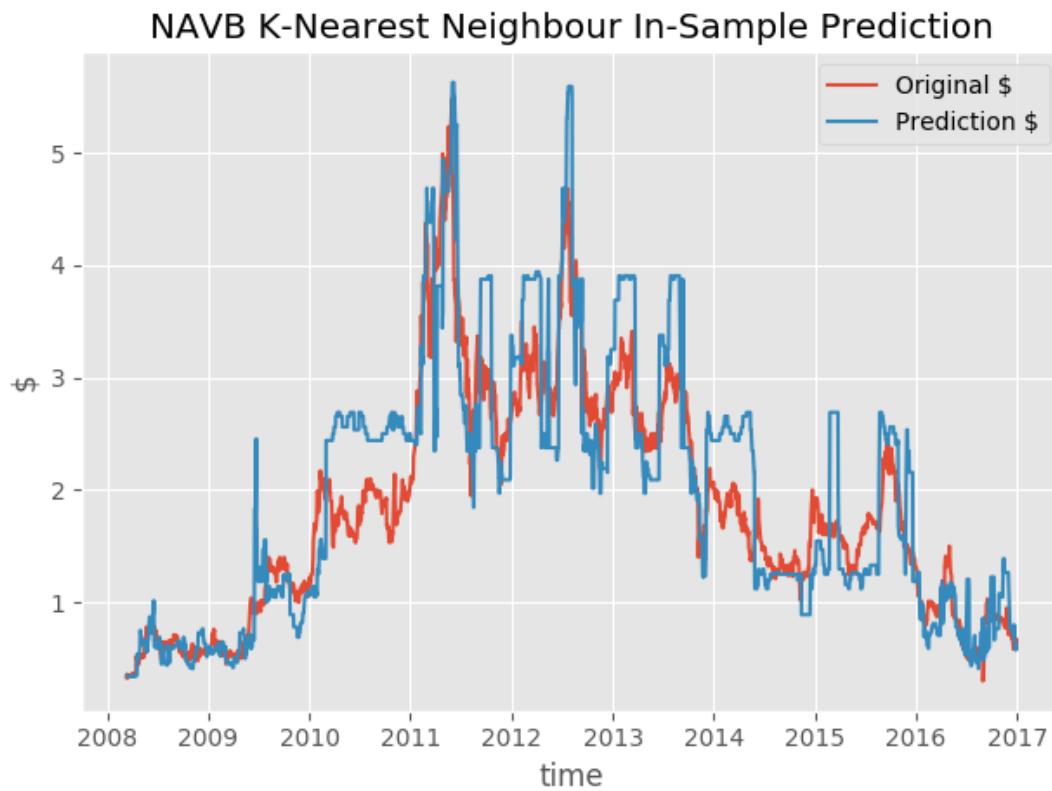


FIGURE A.107: NAVB K-Nearest Neighbour in-sample prediction

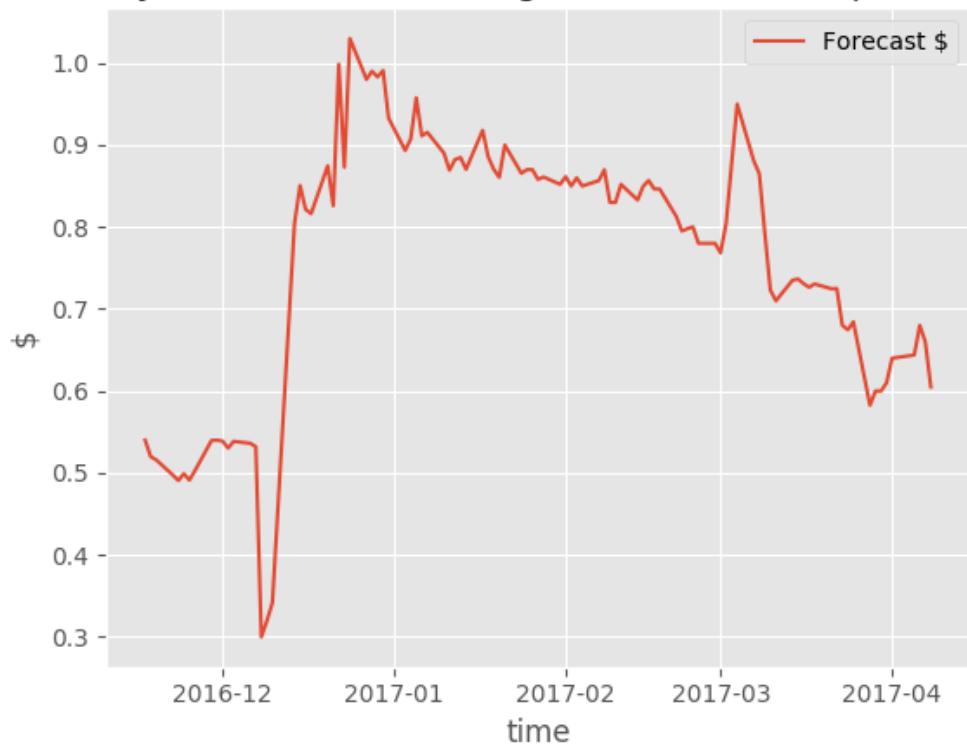
100 Day NAVB K-Nearest Neighbour Out-of-Sample Forecast

FIGURE A.108: 100 day NAVB K-Nearest Neighbour out-of-sample forecast

HRG scored a mean absolute error regression loss of 0.980, and a coefficient of determination of 0.855.

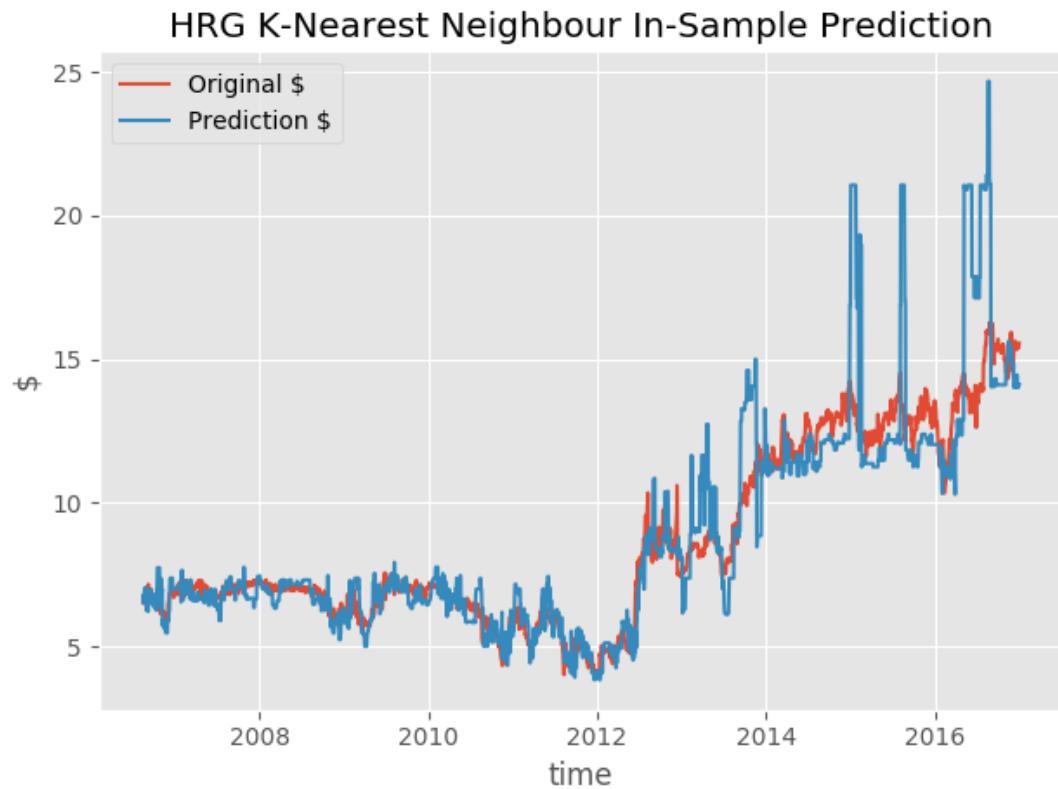


FIGURE A.109: HRG K-Nearest Neighbour in-sample prediction

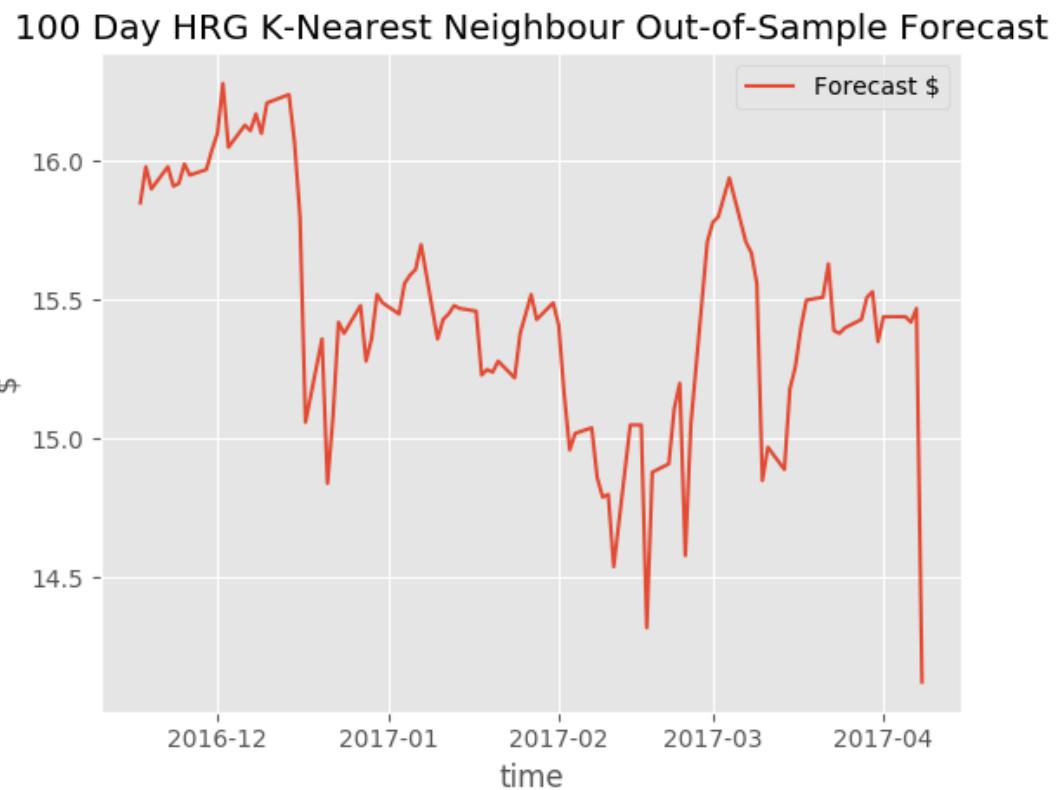


FIGURE A.110: 100 day HRG K-Nearest Neighbour out-of-sample forecast

HL scored a mean absolute error regression loss of 0.410, and a coefficient of determination of 0.952.

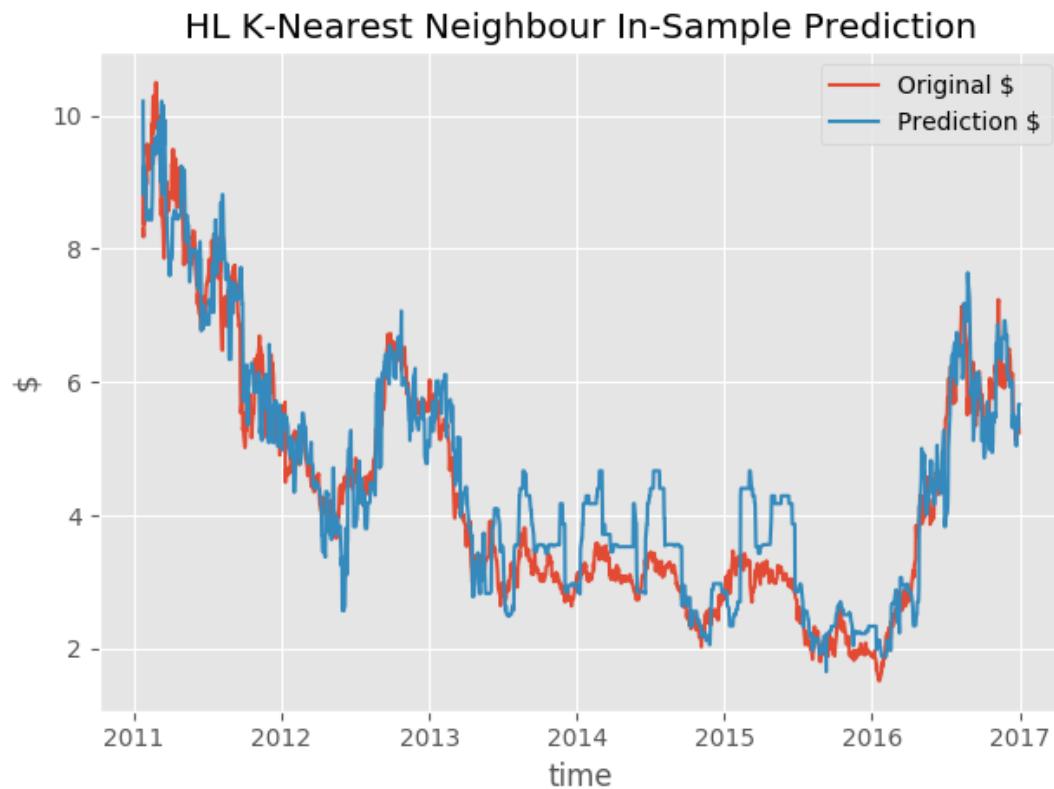


FIGURE A.111: HL K-Nearest Neighbour in-sample prediction

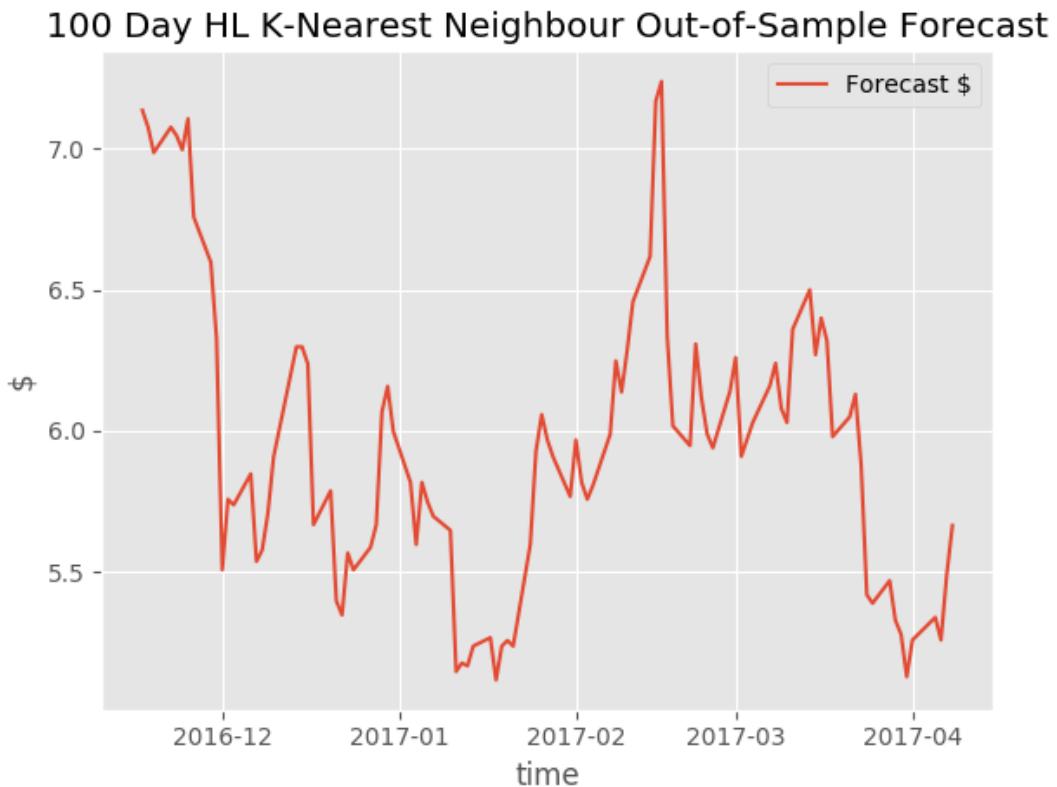


FIGURE A.112: 100 day HL K-Nearest Neighbour out-of-sample forecast

A.2.0.4 Random Forest

CDE scored a mean absolute error regression loss of 2.505, and a coefficient of determination of 0.861.

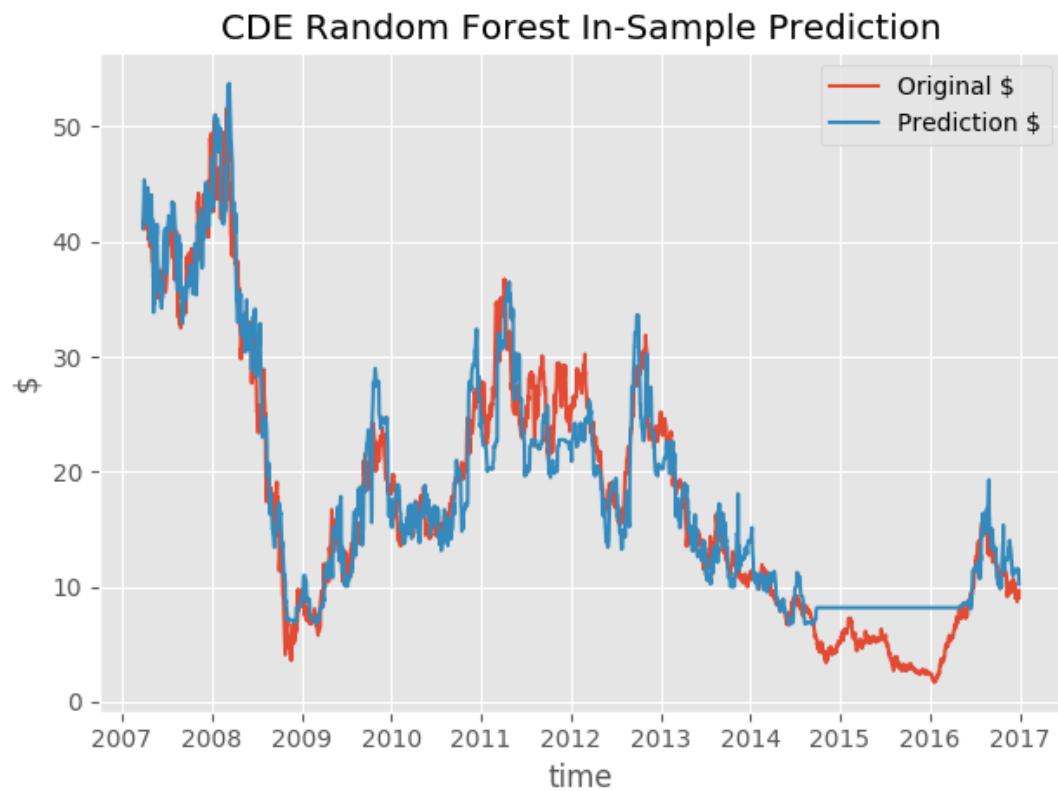


FIGURE A.113: CDE Random Forest in-sample prediction

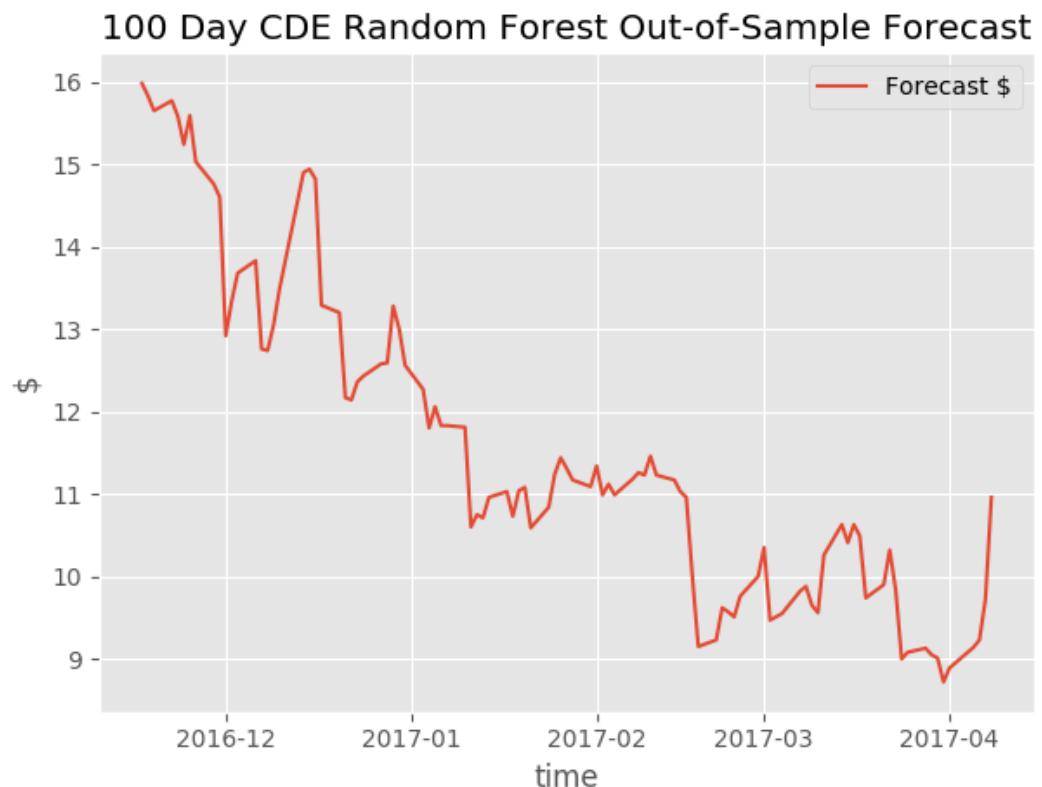


FIGURE A.114: 100 day CDE Random Forest out-of-sample forecast

NAVB scored a mean absolute error regression loss of 0.293, and a coefficient of determination of 0.856.

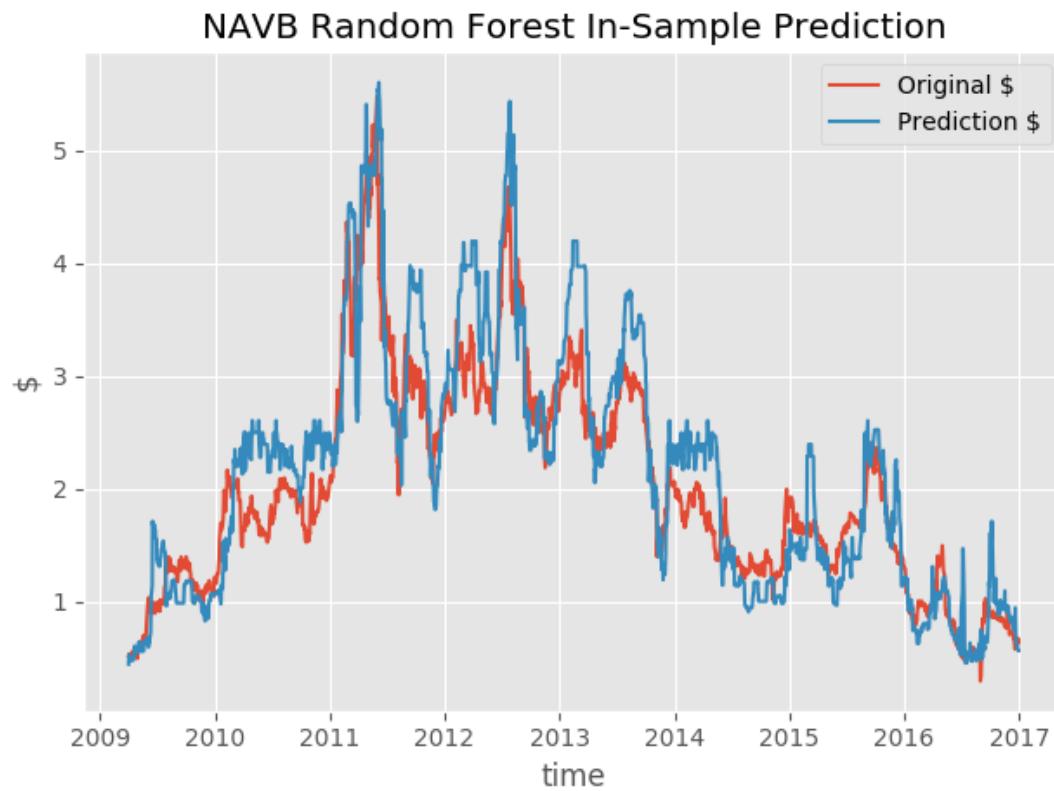


FIGURE A.115: NAVB Random Forest in-sample prediction

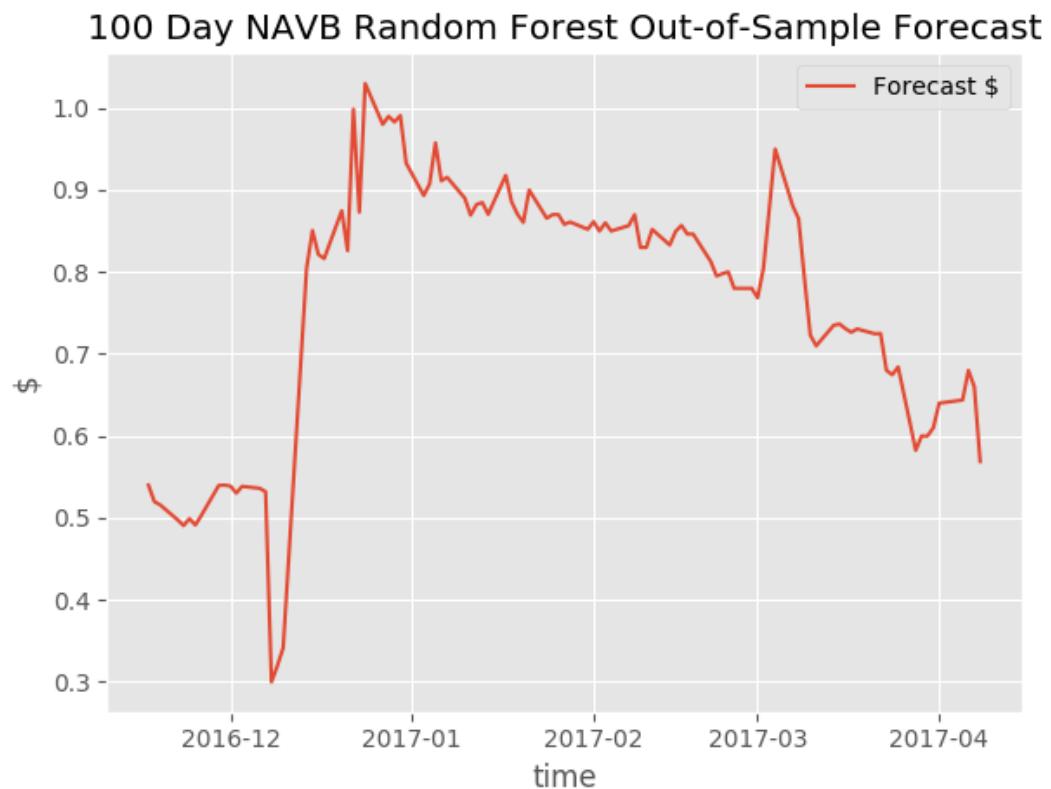


FIGURE A.116: 100 day NAVB Random Forest out-of-sample forecast

HRG scored a mean absolute error regression loss of 0.675, and a coefficient of determination of 0.907.

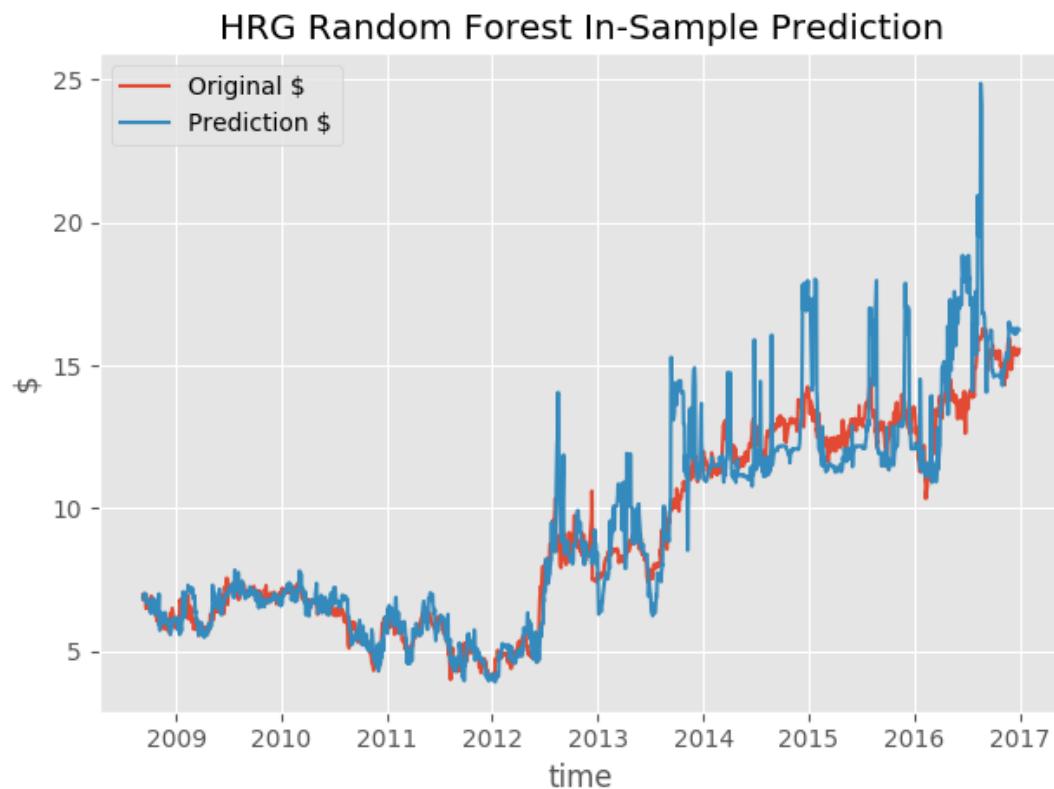


FIGURE A.117: HRG Random Forest in-sample prediction

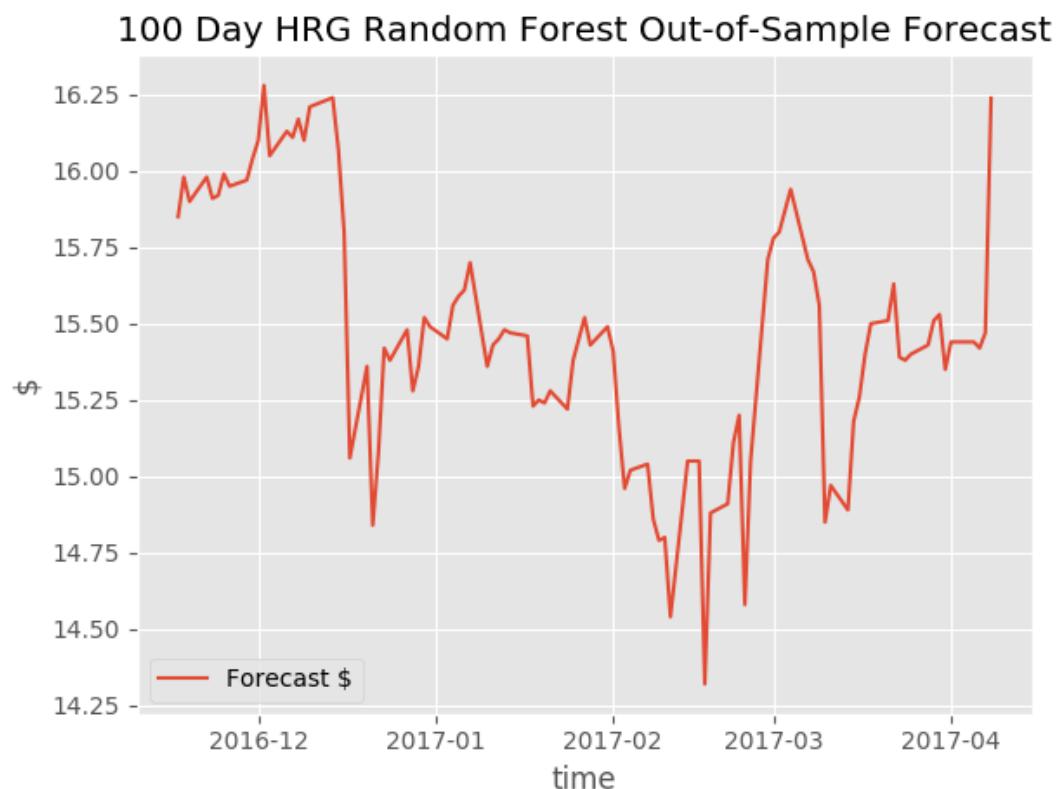


FIGURE A.118: 100 day HRG Random Forest out-of-sample forecast

HL scored a mean absolute error regression loss of 0.434, and a coefficient of determination of 0.922.

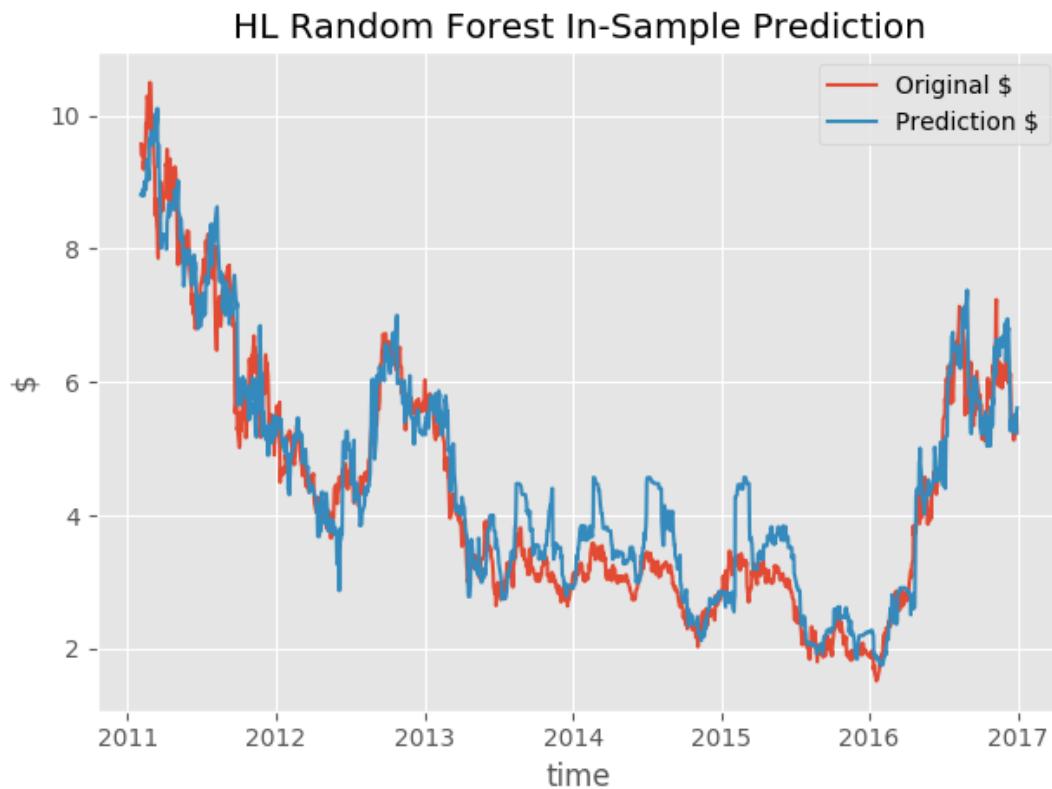


FIGURE A.119: HL Random Forest in-sample prediction

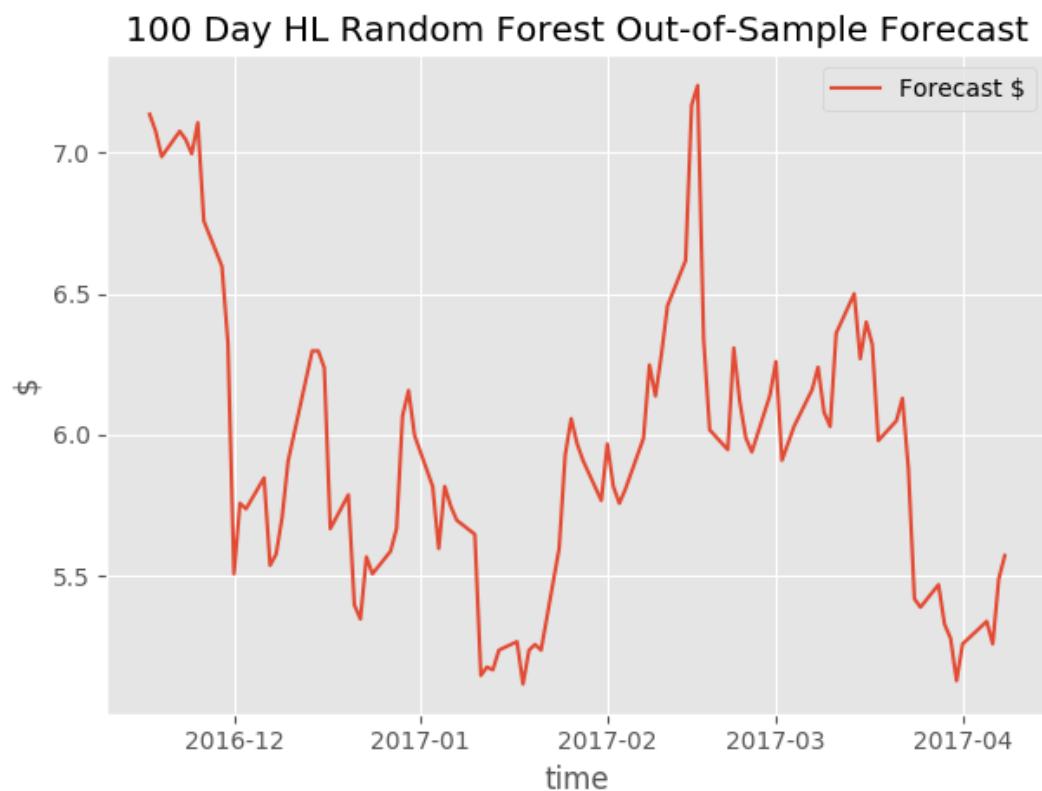


FIGURE A.120: 100 day HL Random Forest out-of-sample forecast

A.2.0.5 Linear Regression

MSFT scored a mean absolute error regression loss of 0.844, and a coefficient of determination of 0.990.

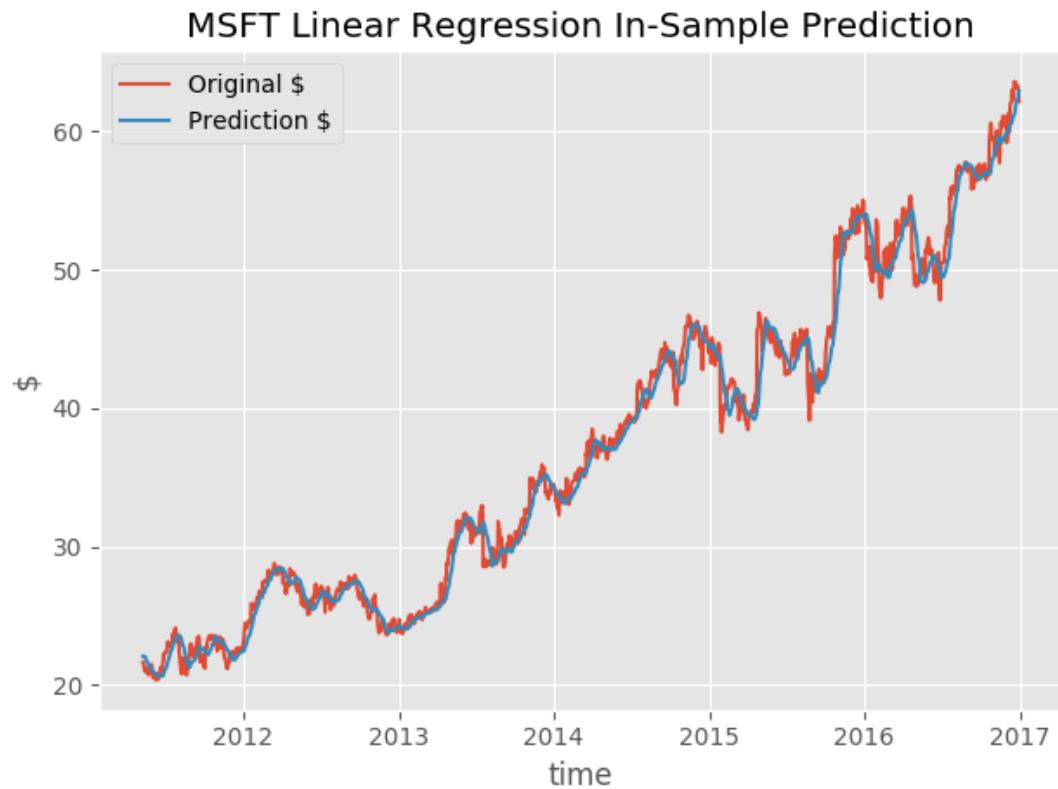


FIGURE A.121: MSFT Linear Regression in-sample prediction

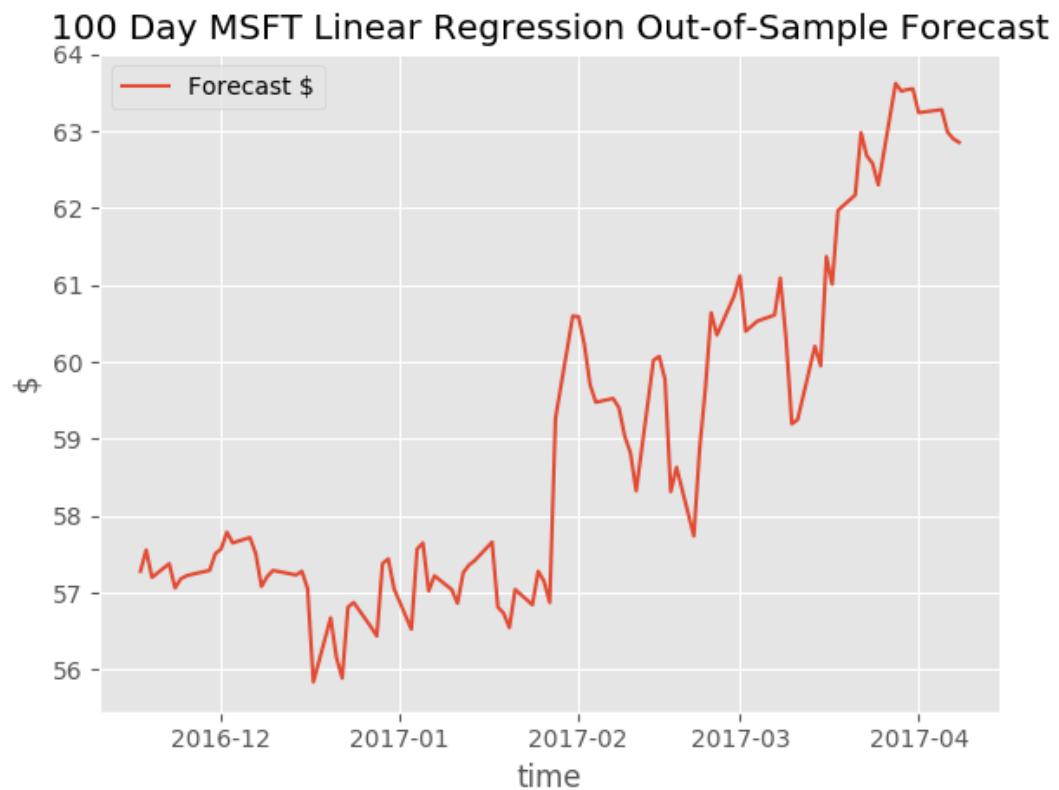


FIGURE A.122: 100 day MSFT Linear Regression out-of-sample forecast

CDE scored a mean absolute error regression loss of 0.990, and a coefficient of determination of 0.972.

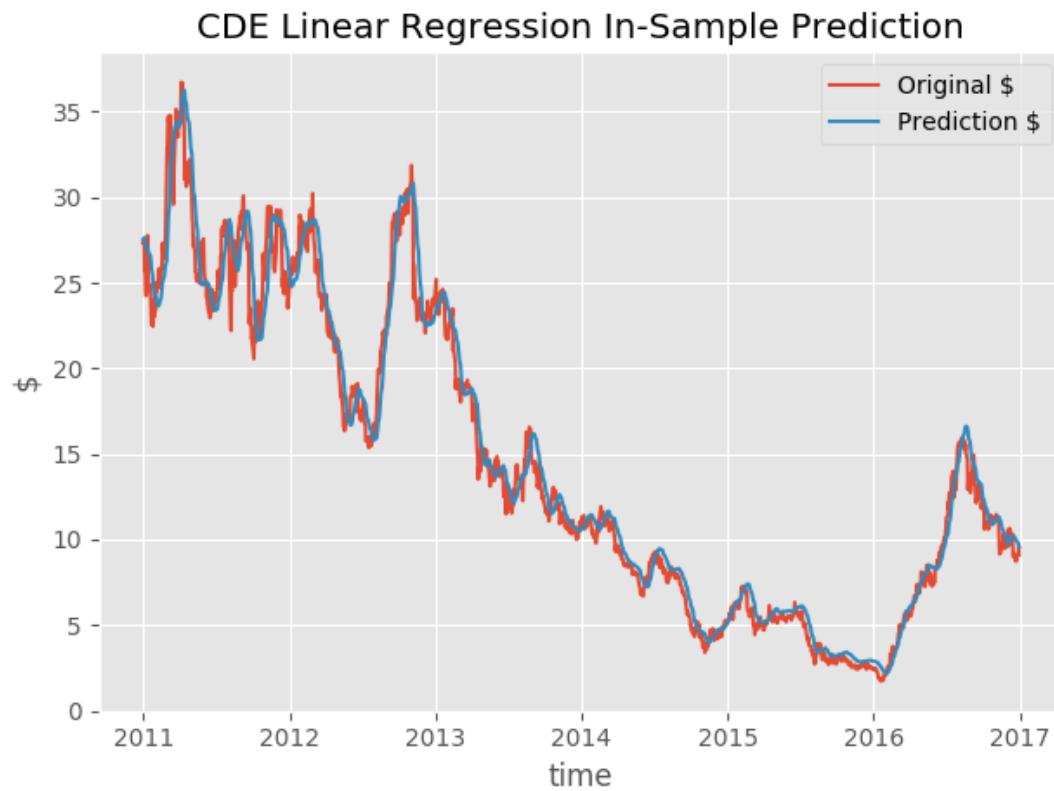


FIGURE A.123: CDE Linear Regression in-sample prediction

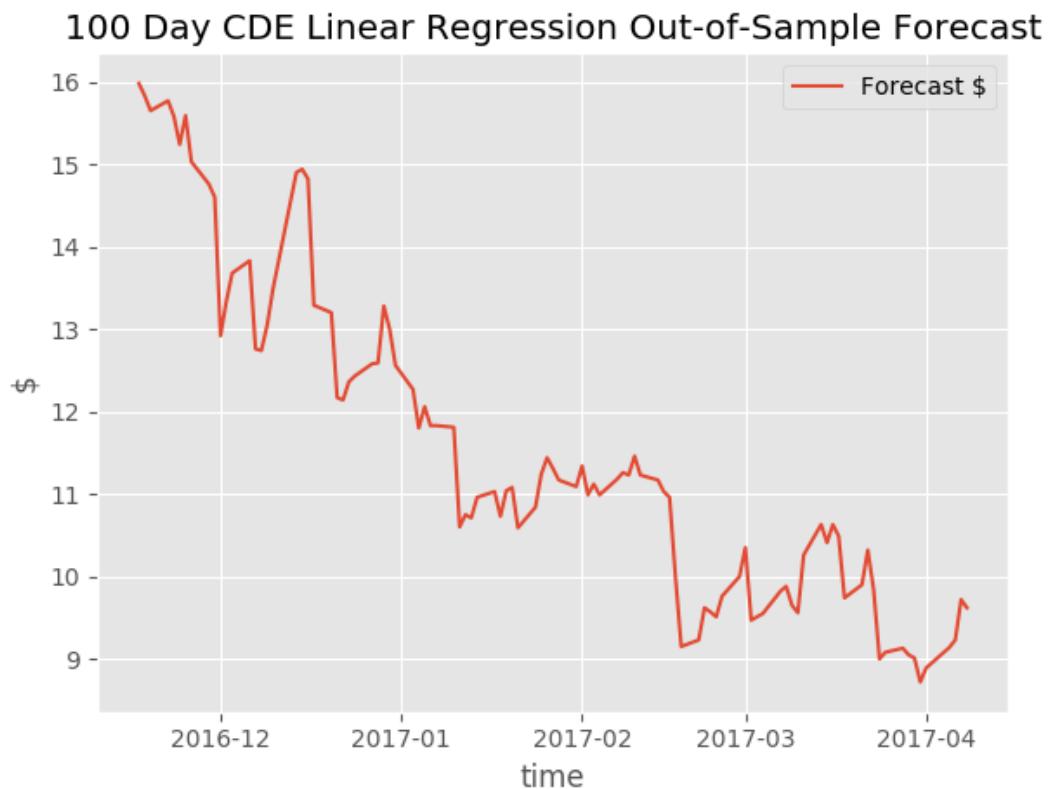


FIGURE A.124: 100 day CDE Linear Regression out-of-sample forecast

HRG scored a mean absolute error regression loss of 0.324, and a coefficient of determination of 0.984.

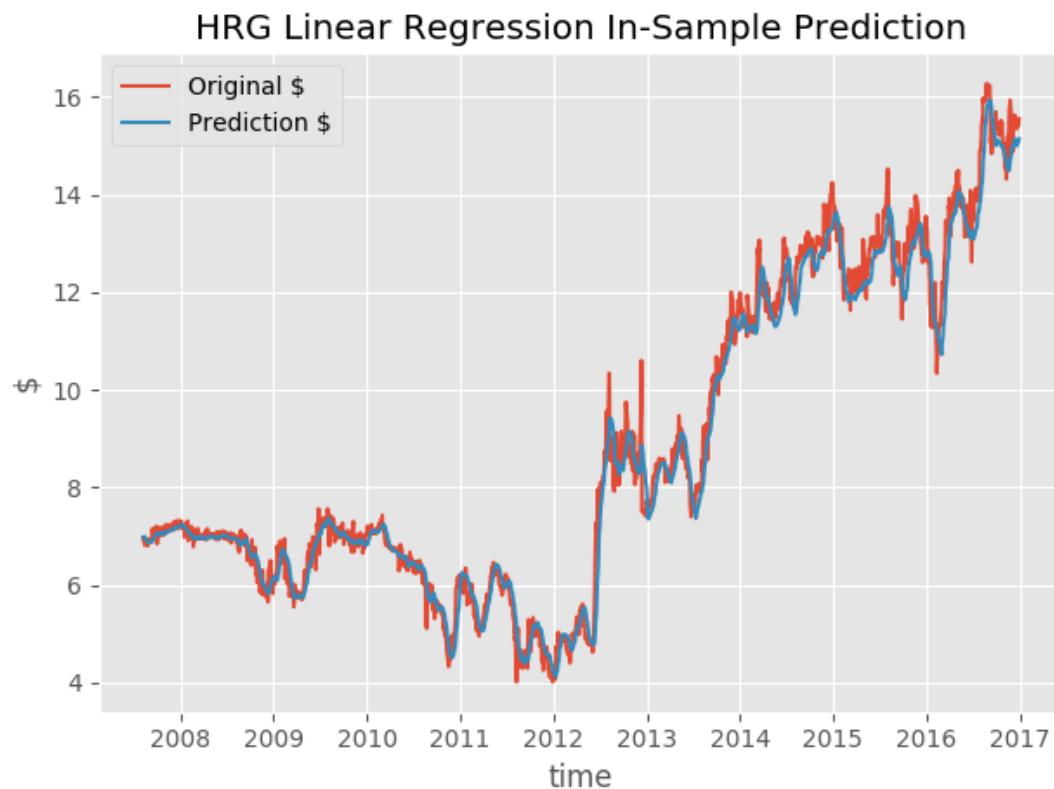


FIGURE A.125: HRG Linear Regression in-sample prediction

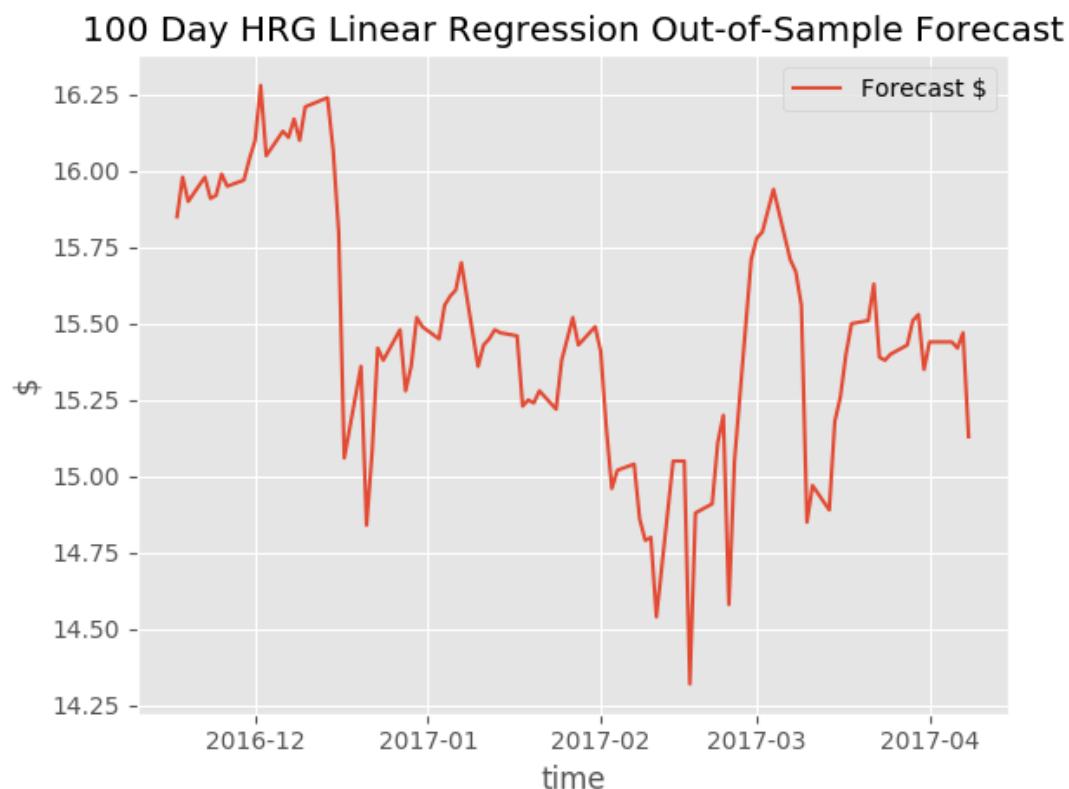


FIGURE A.126: 100 day HRG Linear Regression out-of-sample forecast

HL scored a mean absolute error regression loss of 0.243, and a coefficient of determination of 0.964.

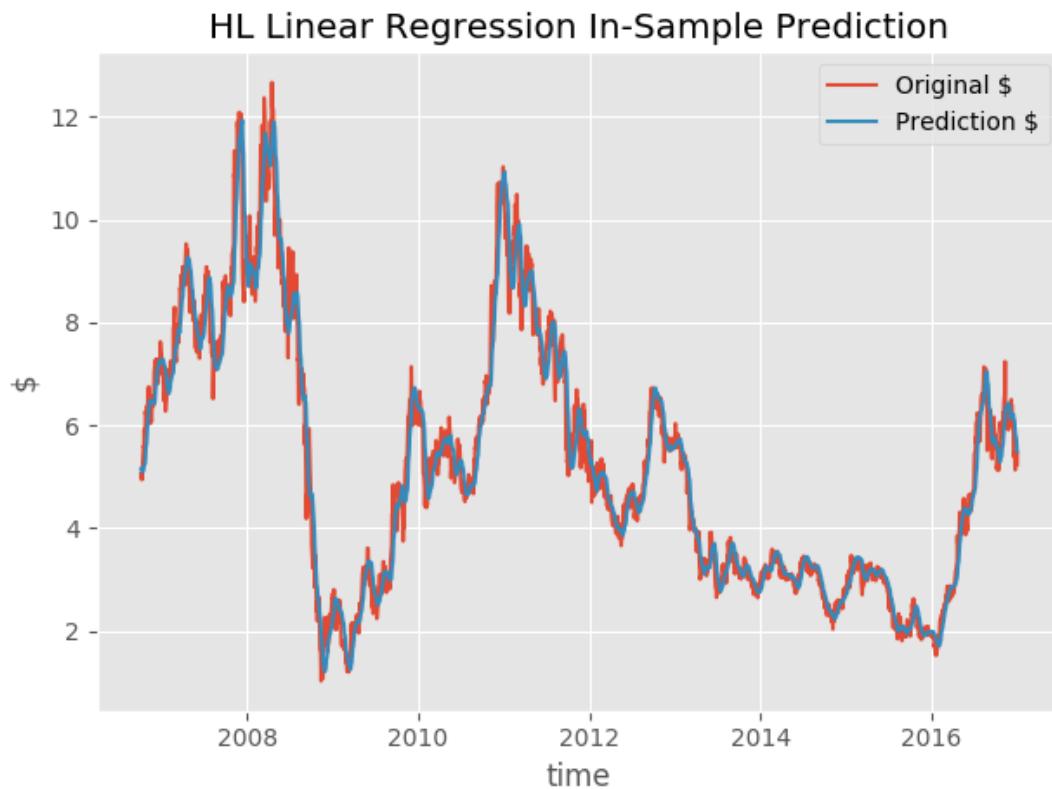


FIGURE A.127: HL Linear Regression in-sample prediction

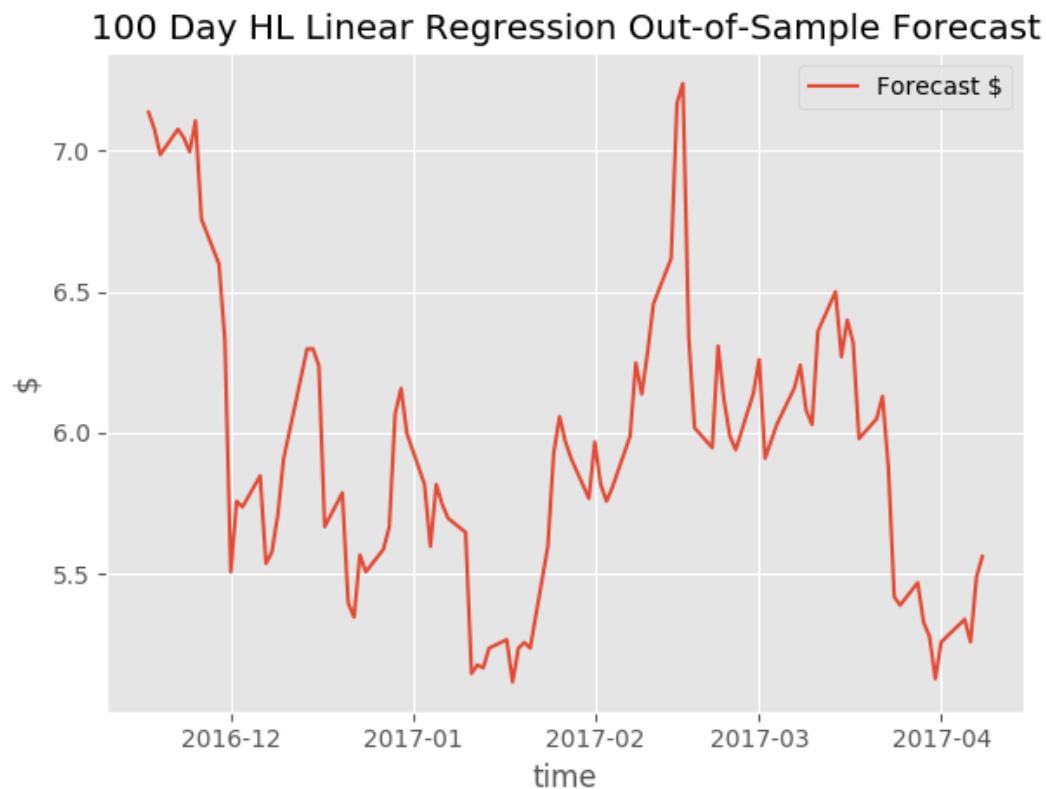


FIGURE A.128: 100 day HL Linear Regression out-of-sample forecast

A.2.0.6 Neural Network

MSFT scored a mean absolute error regression loss of 1.073, and a coefficient of determination of 0.992.

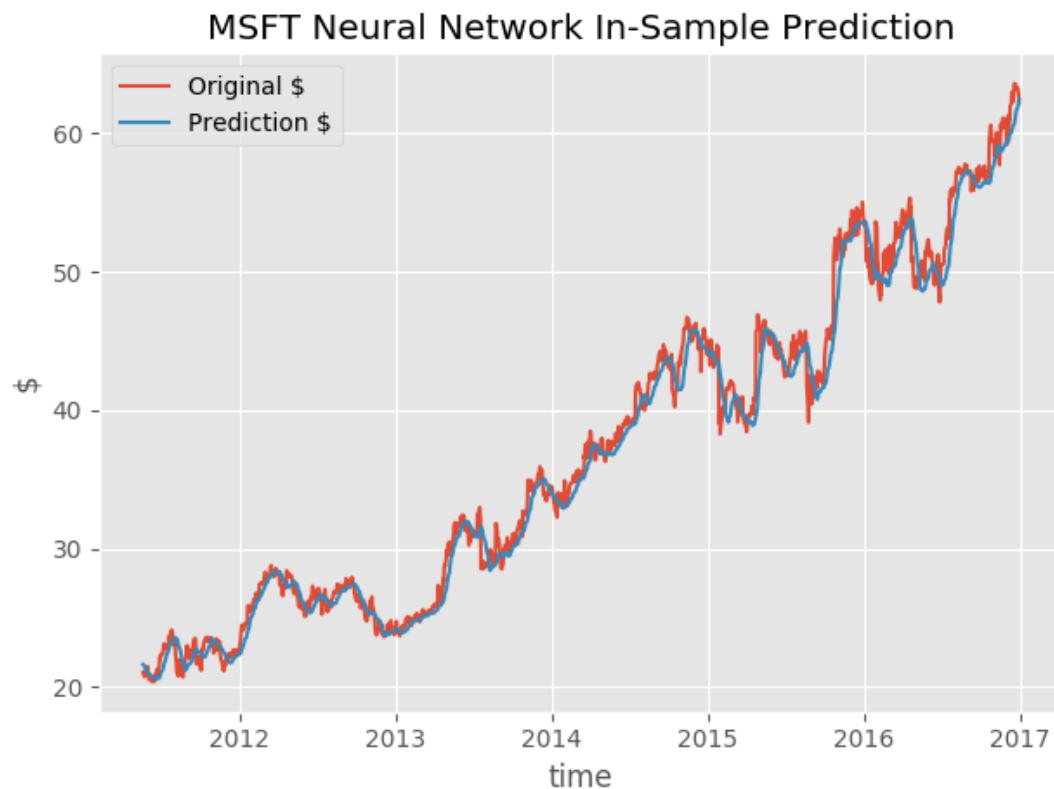


FIGURE A.129: MSFT Neural Network in-sample prediction

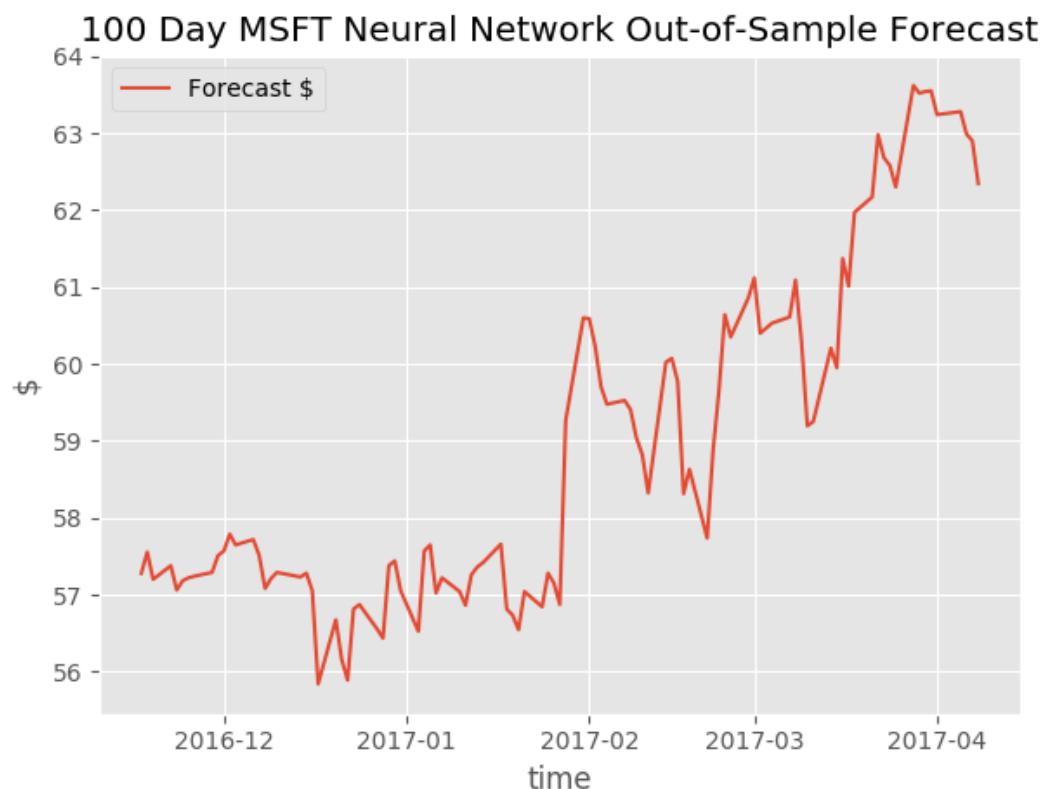


FIGURE A.130: 100 day MSFT Neural Network out-of-sample forecast

NAVB scored a mean absolute error regression loss of 0.422, and a coefficient of determination of 0.693.

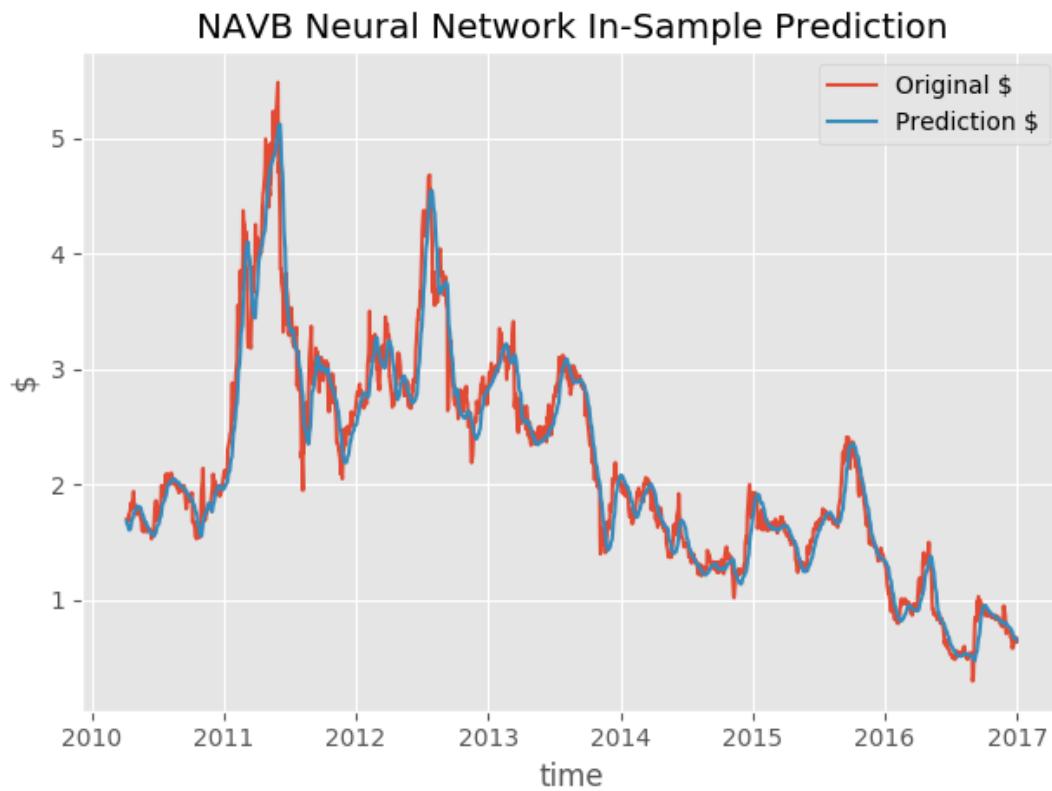


FIGURE A.131: NAVB Neural Network in-sample prediction

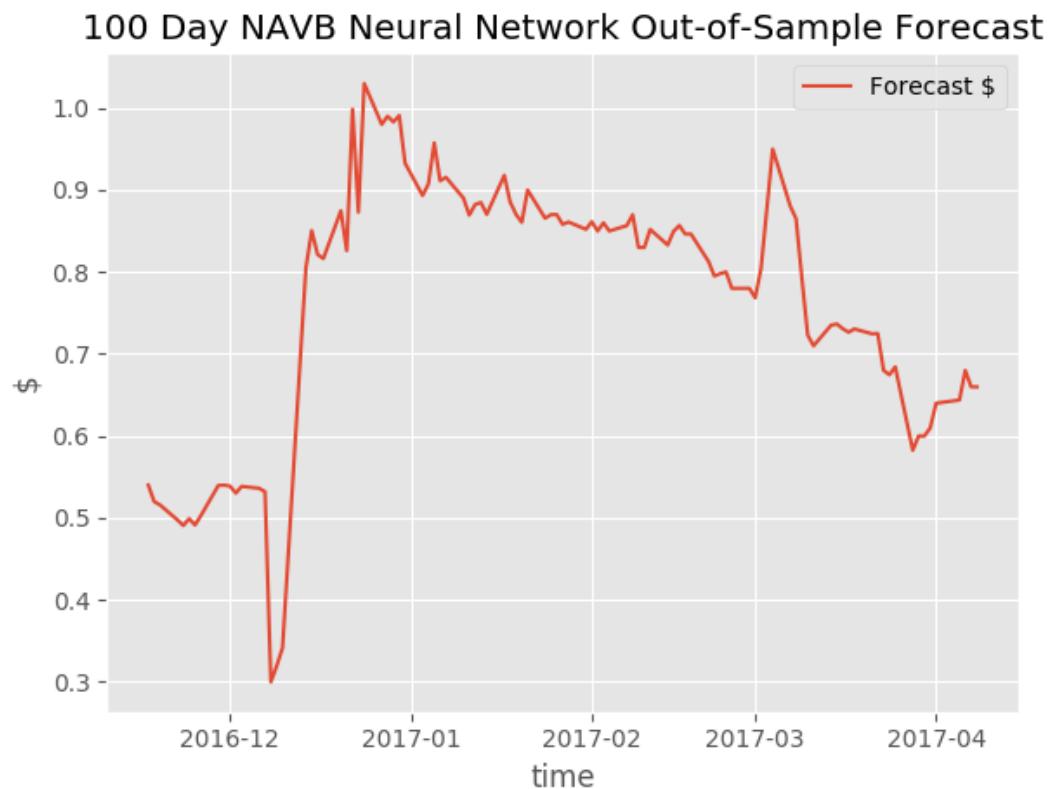


FIGURE A.132: 100 day NAVB Neural Network out-of-sample forecast

HRG scored a mean absolute error regression loss of 1.415, and a coefficient of determination of 0.464.

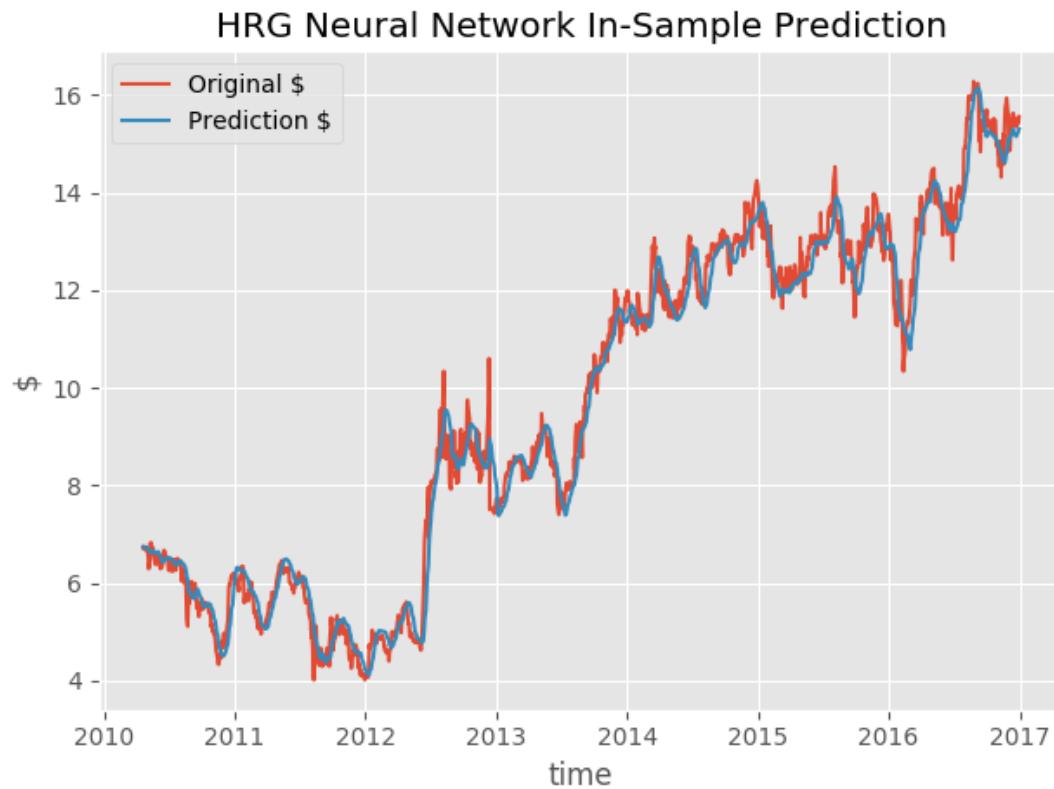


FIGURE A.133: HRG Neural Network in-sample prediction

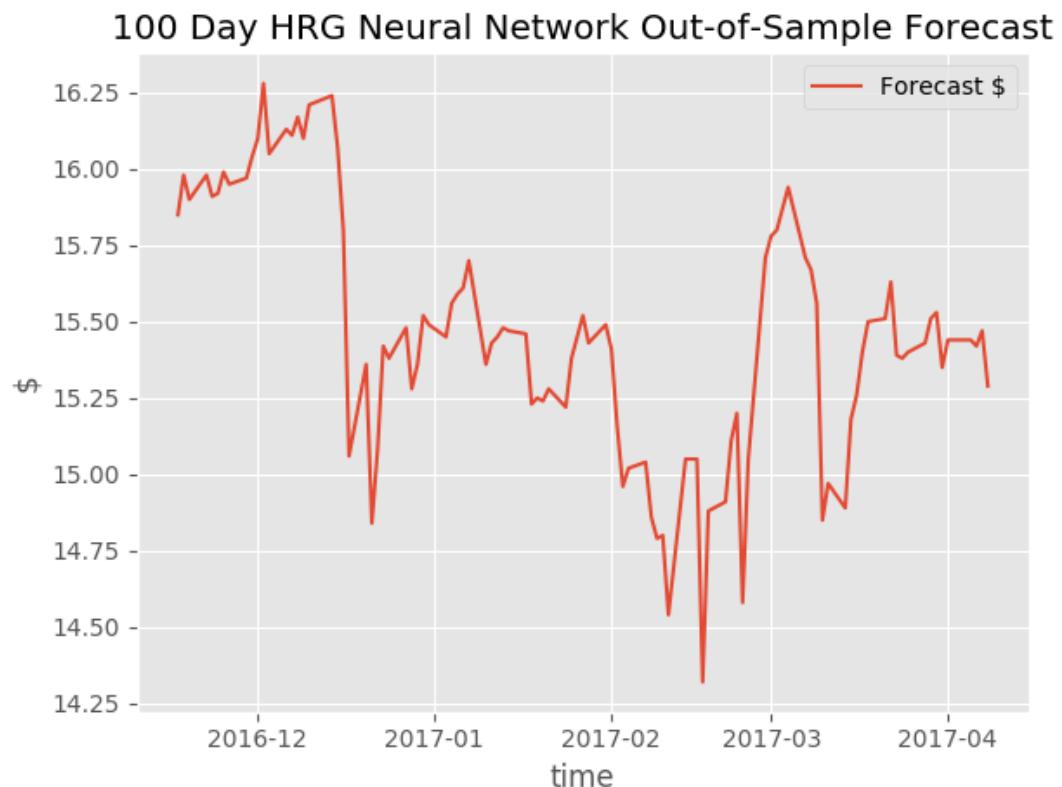


FIGURE A.134: 100 day HRG Neural Network out-of-sample forecast

HL scored a mean absolute error regression loss of 0.506, and a coefficient of determination of 0.923.

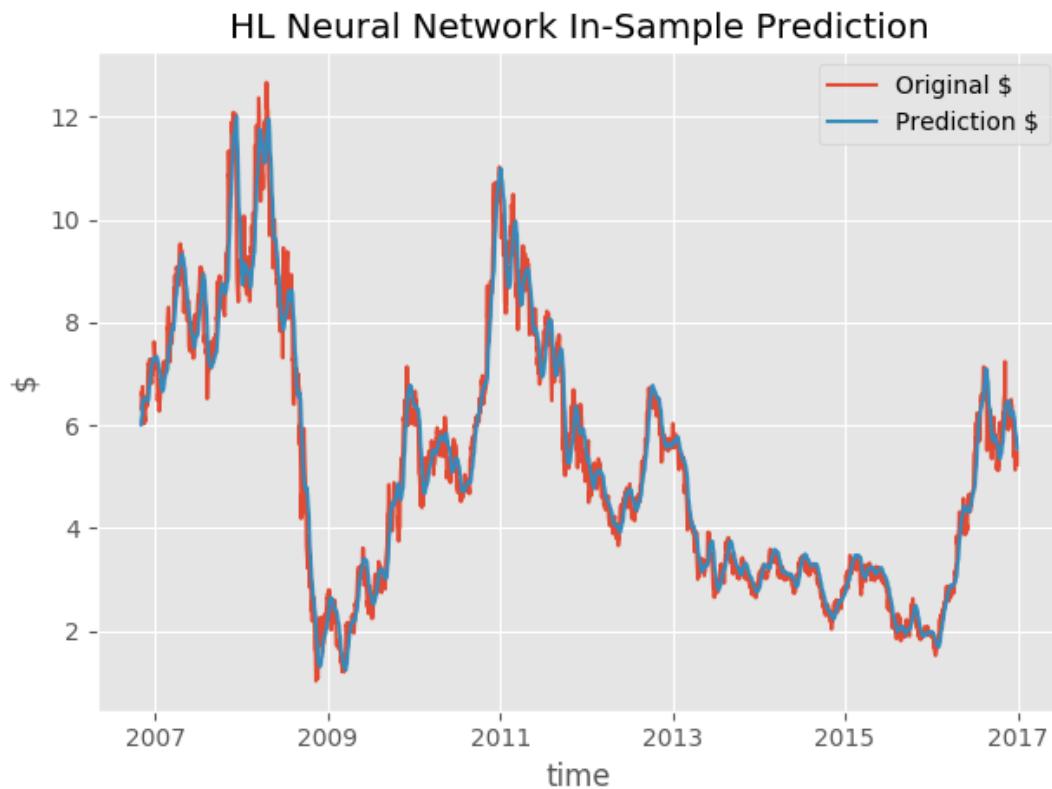


FIGURE A.135: HL Neural Network in-sample prediction

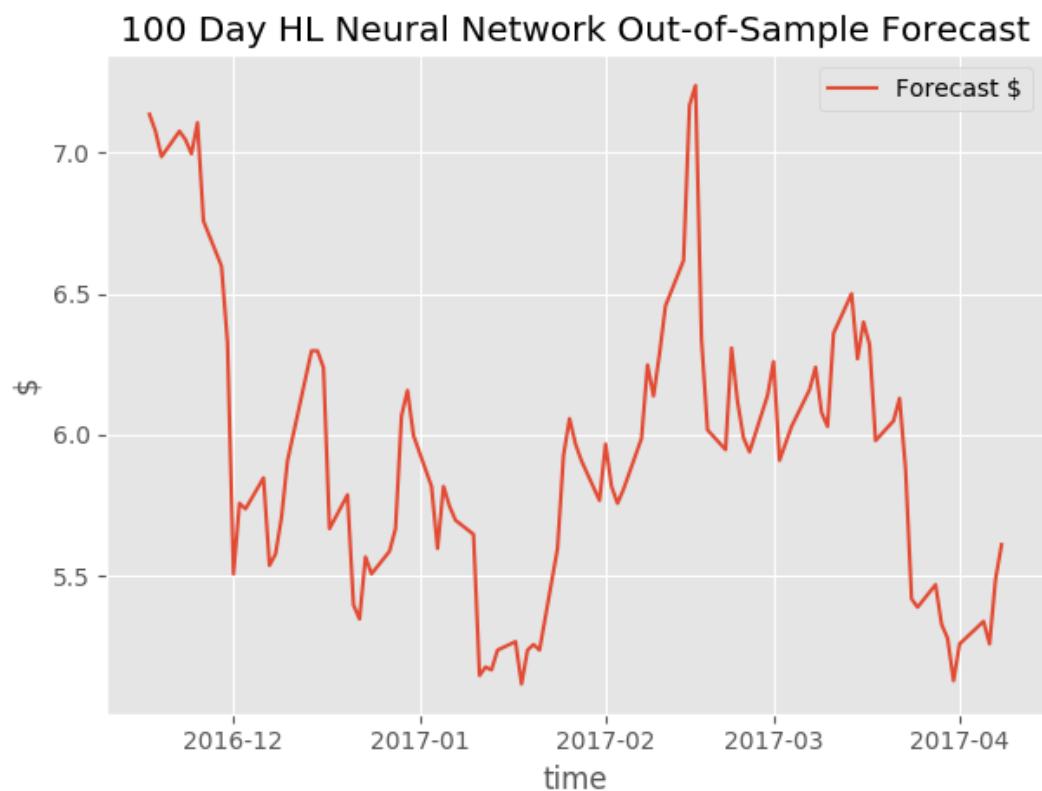


FIGURE A.136: 100 day HL Neural Network out-of-sample forecast

A.2.0.7 Stochastic Gradient Descent

MSFT scored a mean absolute error regression loss of 0.829, and a coefficient of determination of 0.990.

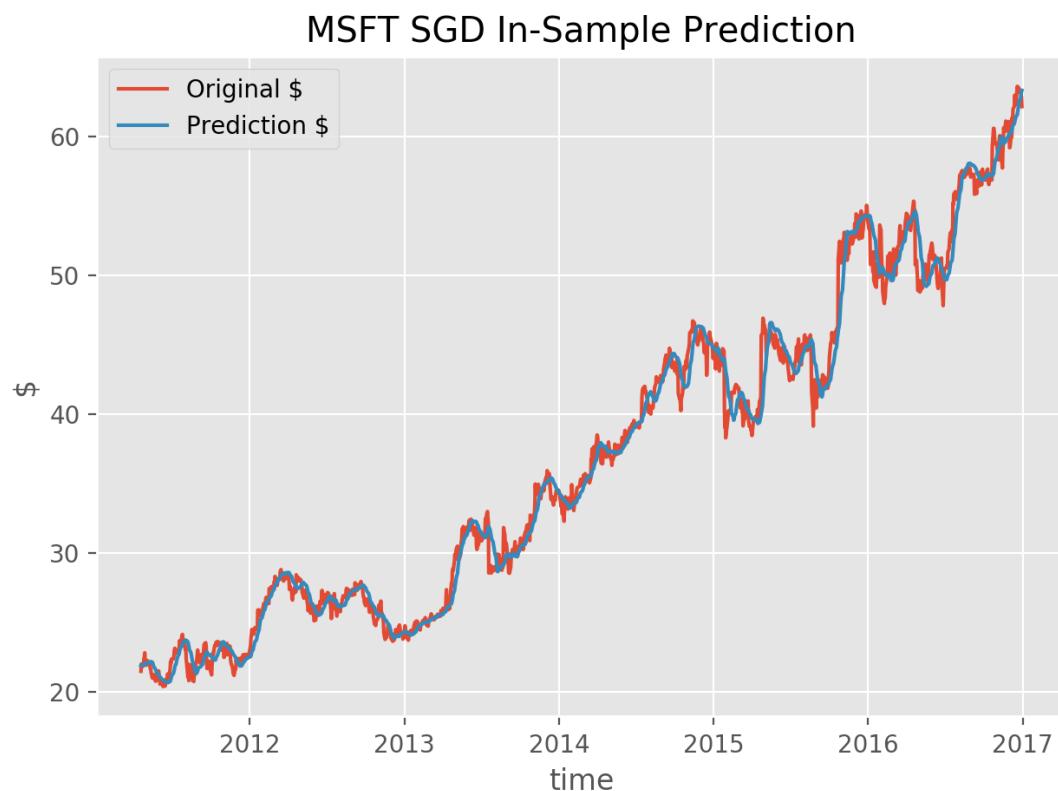


FIGURE A.137: MSFT SGD in-sample prediction

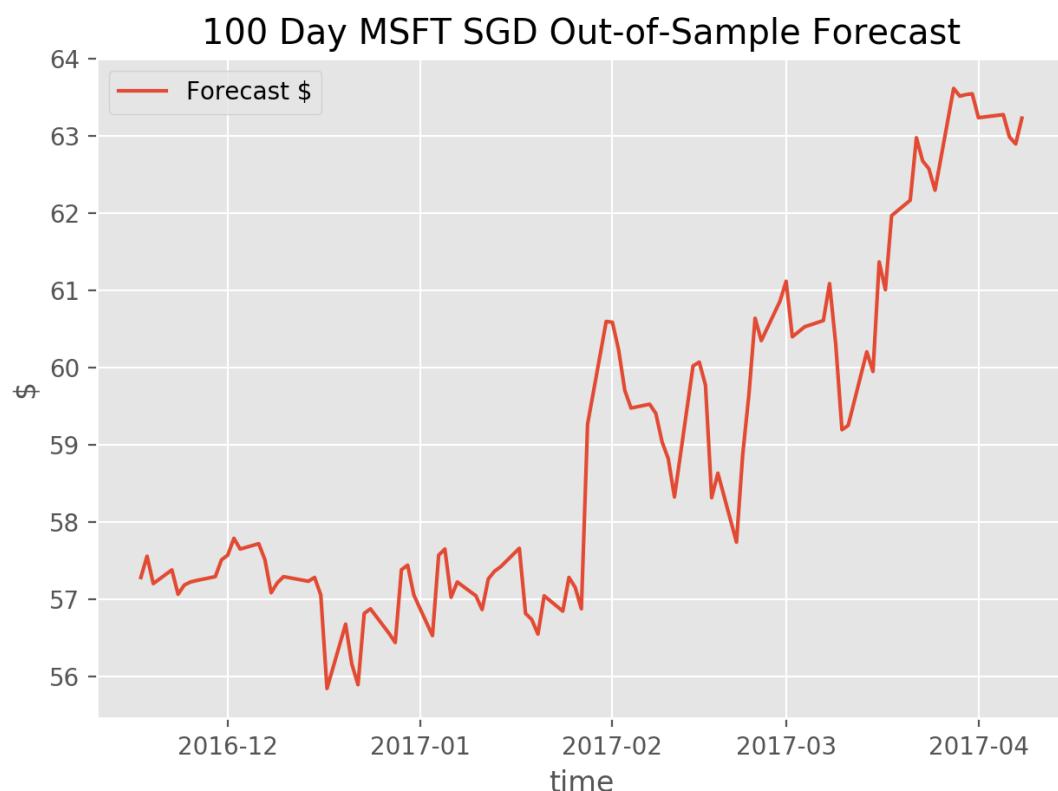


FIGURE A.138: 100 day MSFT SGD out-of-sample forecast

NAV_B scored a mean absolute error regression loss of 0.142, and a coefficient of determination of 0.955.

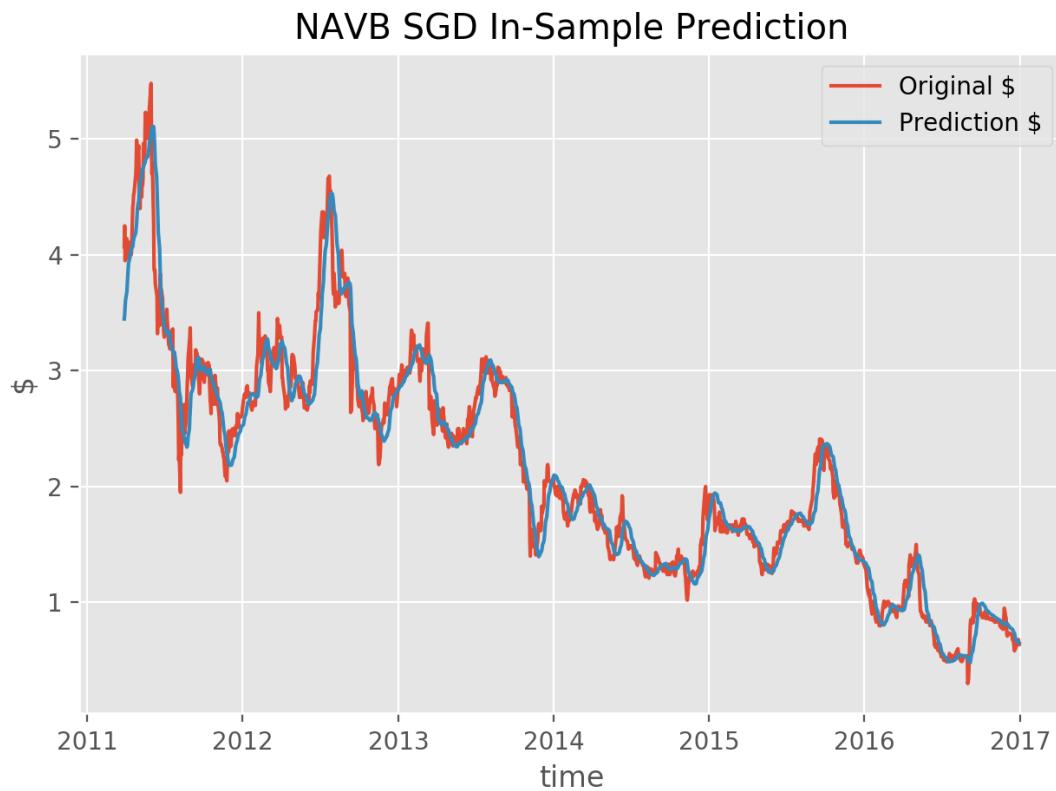


FIGURE A.139: NAV_B SGD in-sample prediction

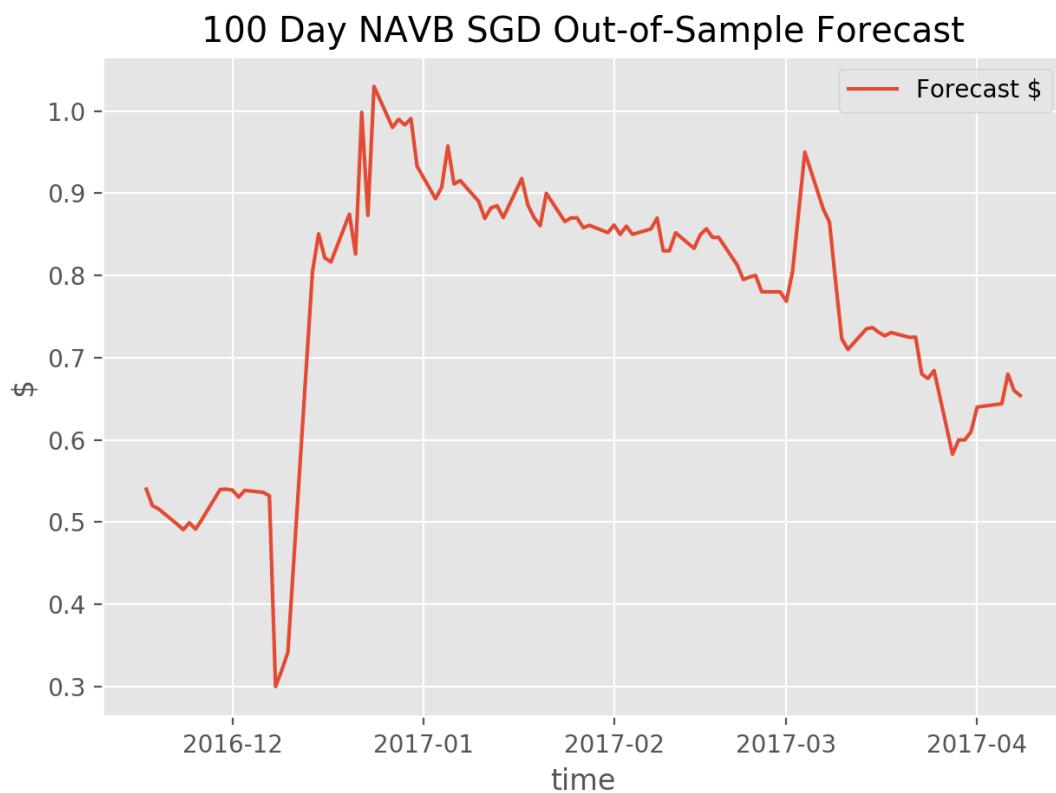


FIGURE A.140: 100 day NAVB SGD out-of-sample forecast

HRG scored a mean absolute error regression loss of 0.292, and a coefficient of determination of 0.987.

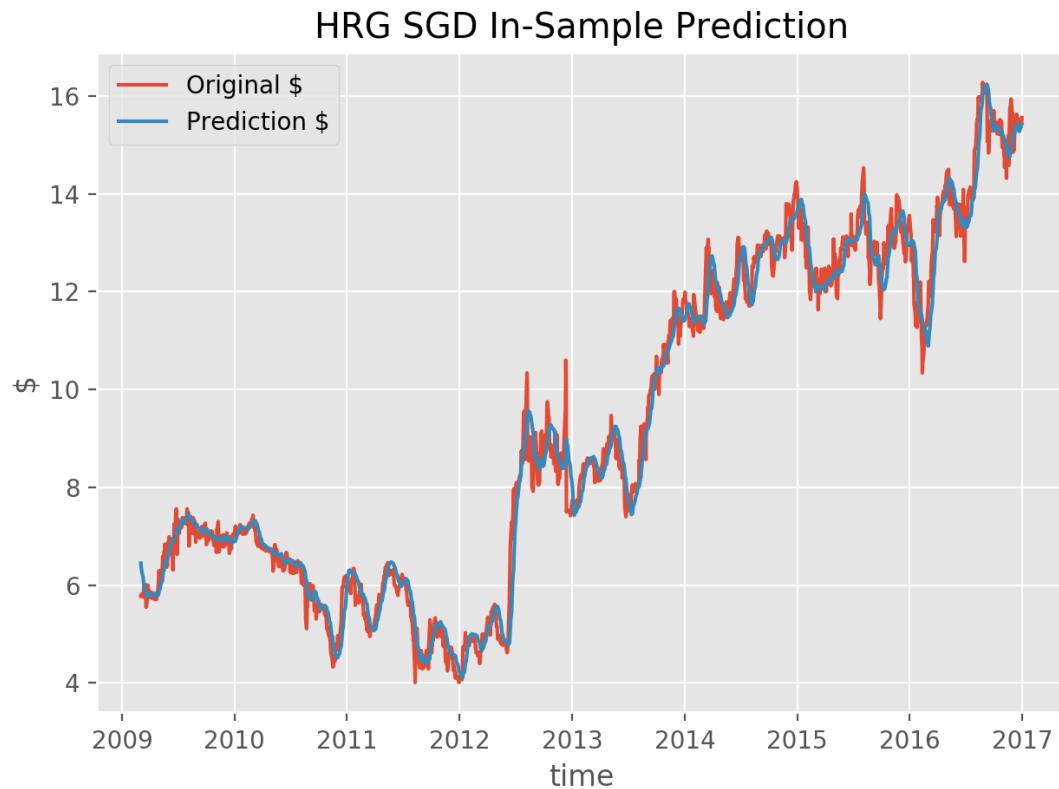


FIGURE A.141: HRG SGD in-sample prediction

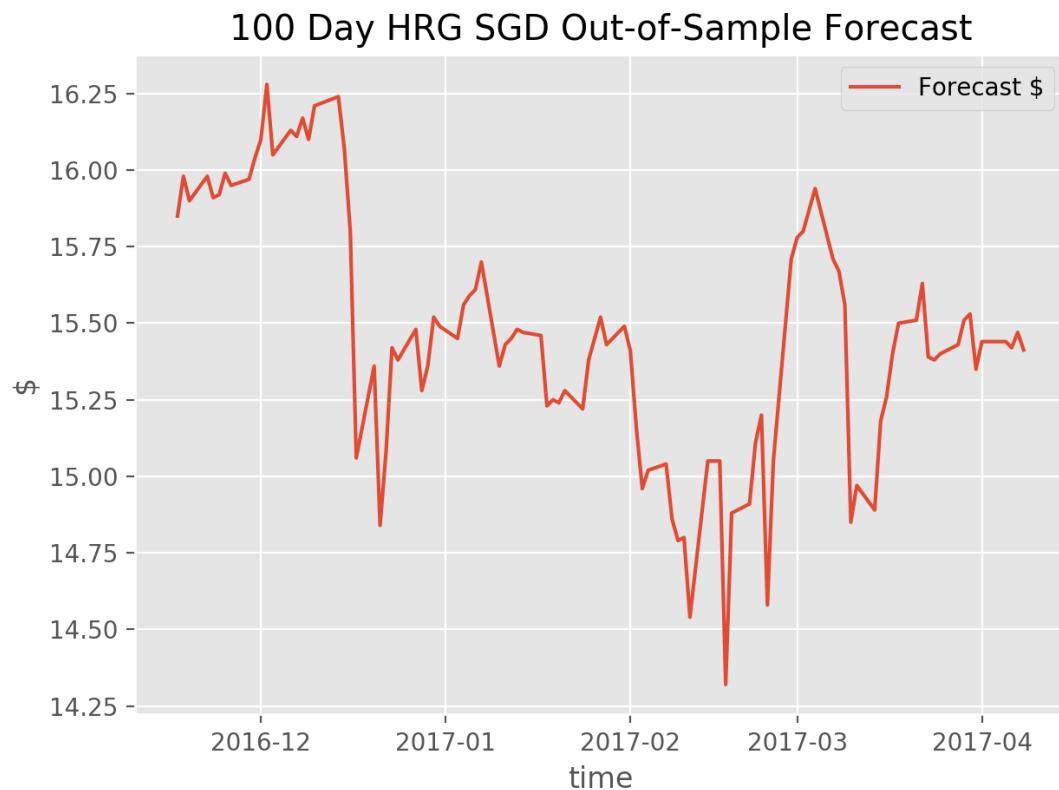


FIGURE A.142: 100 day HRG SGD out-of-sample forecast

HL scored a mean absolute error regression loss of 0.275, and a coefficient of determination of 0.962.

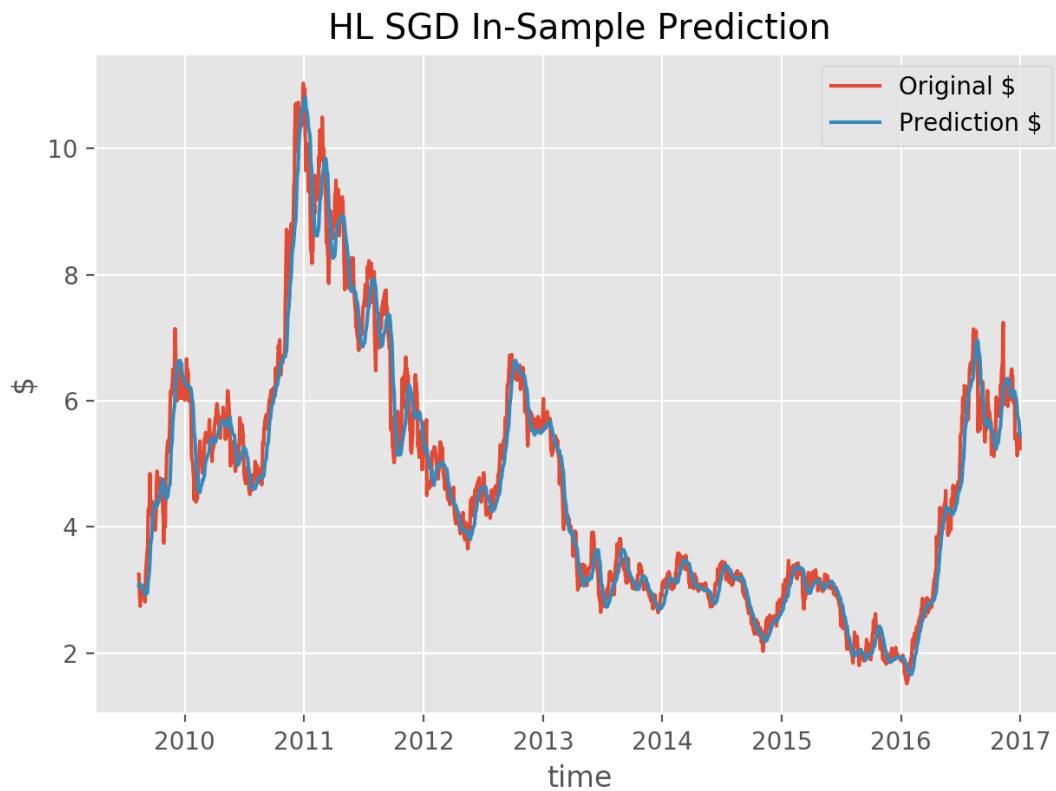


FIGURE A.143: HL SGD in-sample prediction

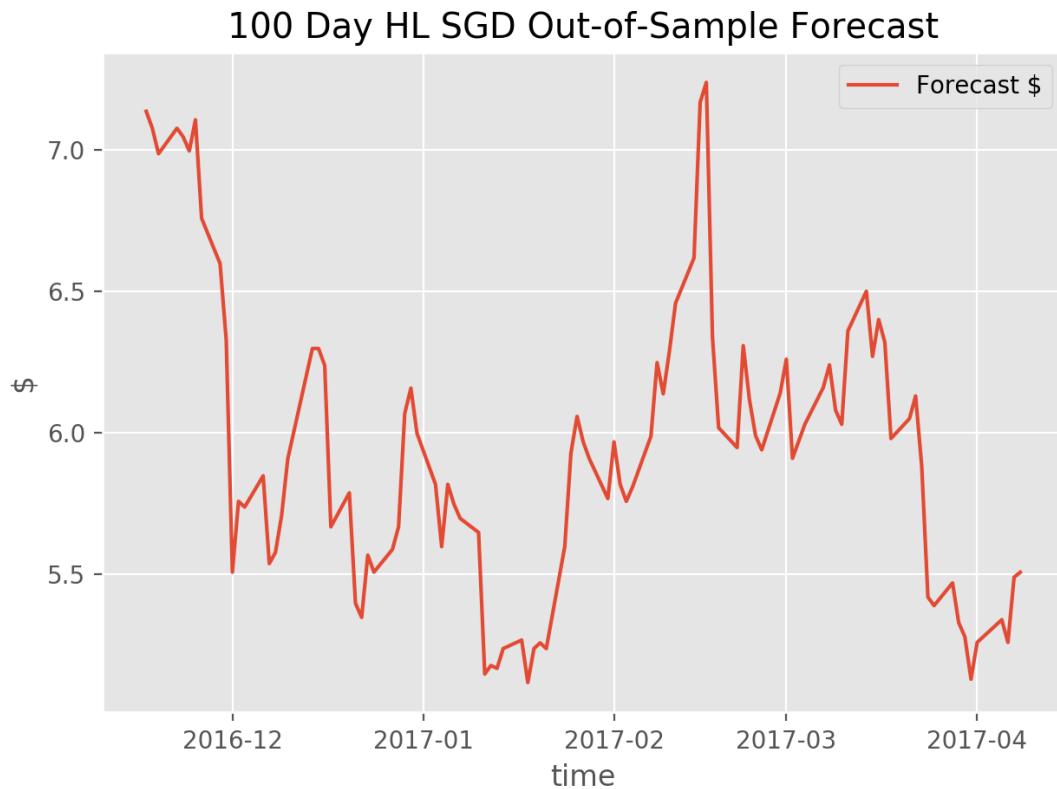


FIGURE A.144: 100 day HL SGD out-of-sample forecast

A.3 Bayesian Statistics

A.3.1 Metropolis-Hastings

CDE scored sharpe ratios of 2.028 for the original returns, and 4.075 for the predicted returns in the in-sample test.

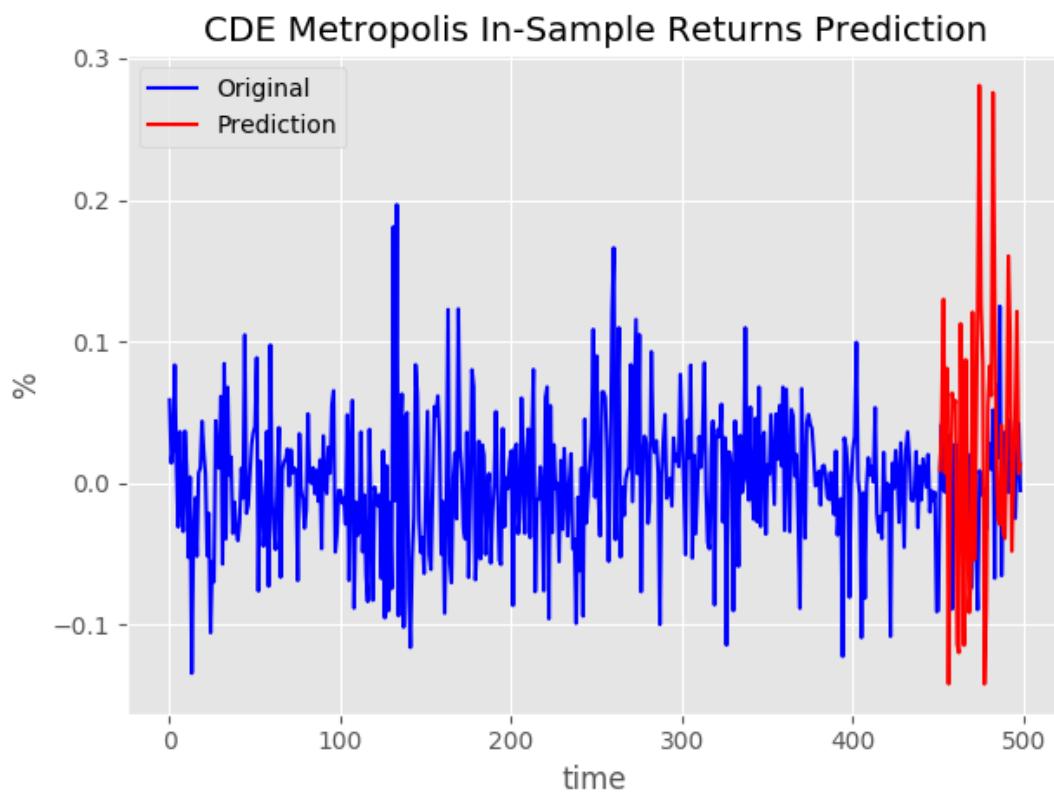


FIGURE A.145: CDE Metropolis-Hastings in-sample prediction

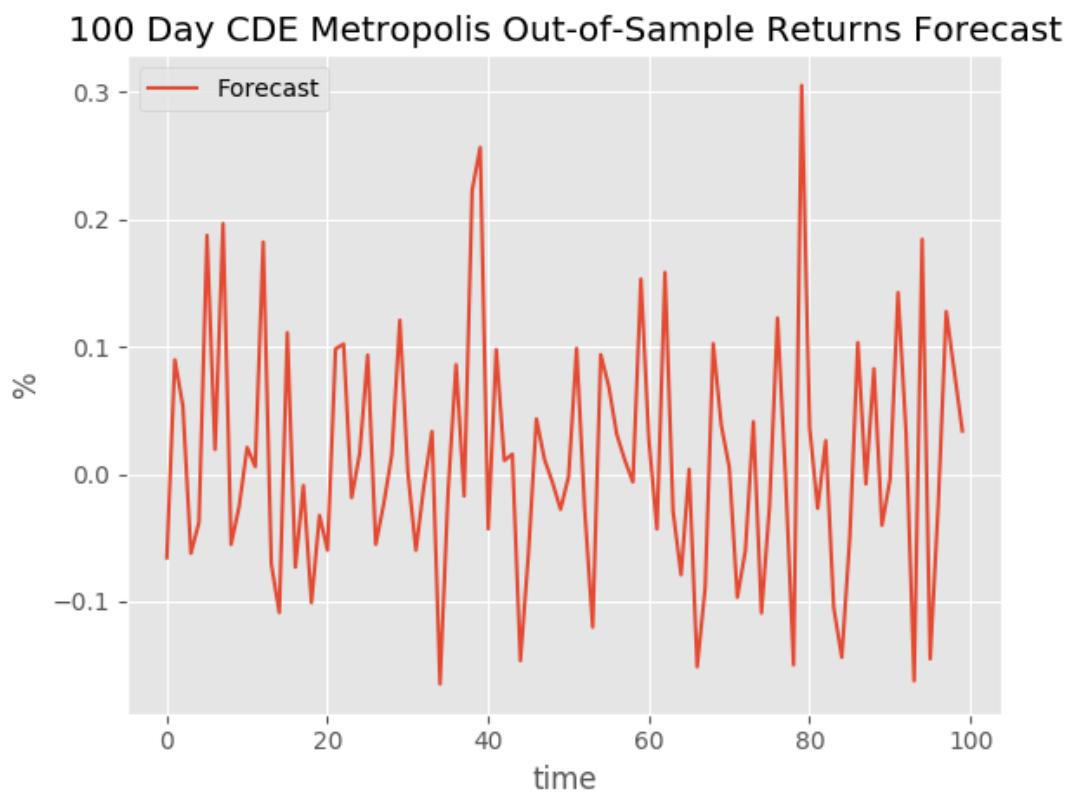


FIGURE A.146: 100 day CDE Metropolis-Hastings out-of-sample forecast

NAVB scored sharpe ratios of -2.800 for the original returns, and 2.998 for the predicted returns in the in-sample test.

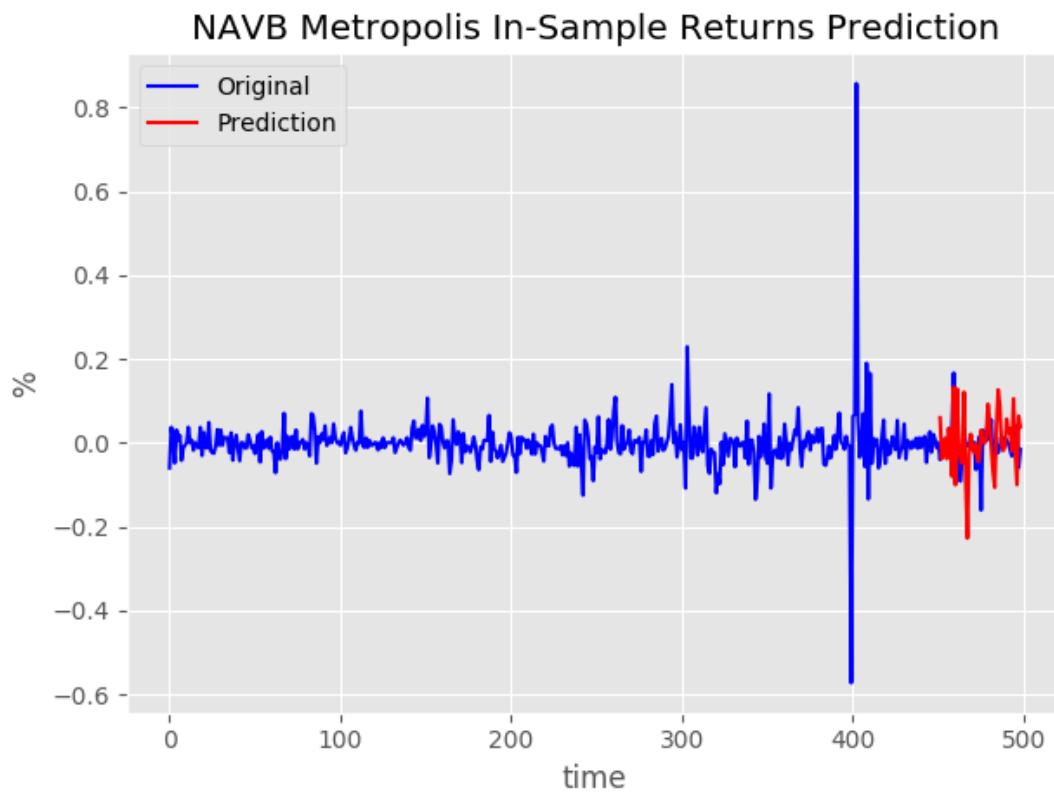


FIGURE A.147: NAVB Metropolis-Hastings in-sample prediction

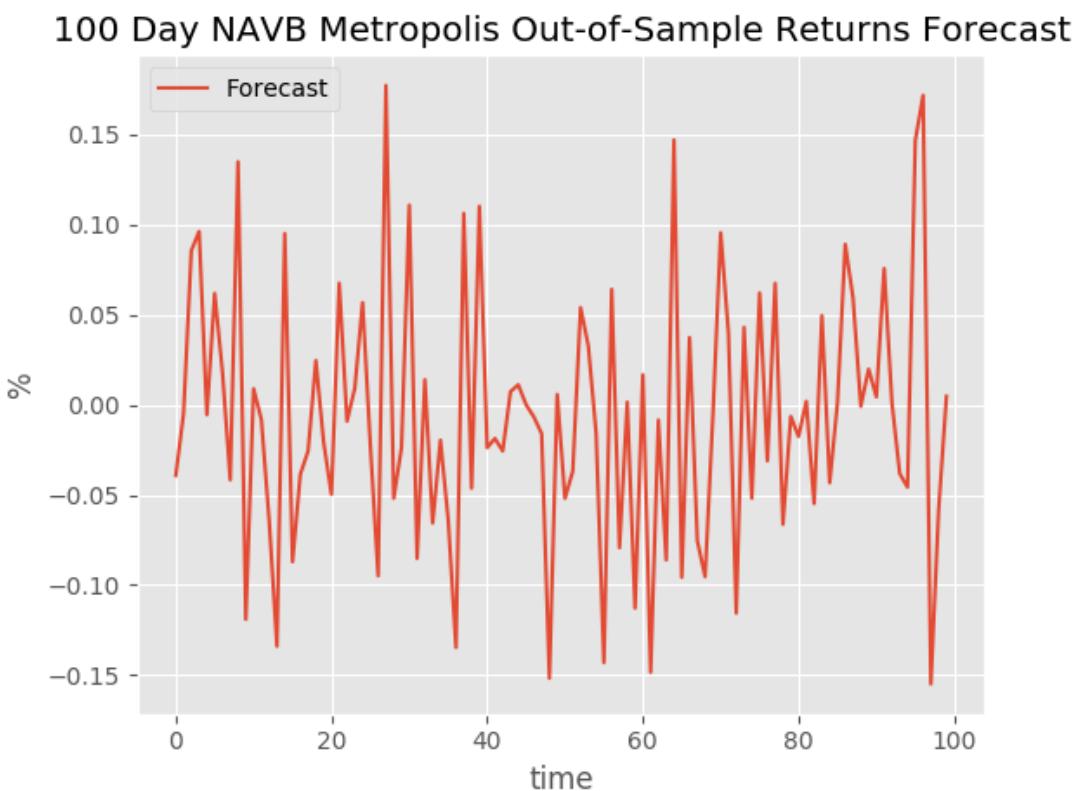


FIGURE A.148: 100 day NAVB Metropolis-Hastings out-of-sample forecast

HRG scored sharpe ratios of 1.362 for the original returns, and -2.173 for the predicted returns in the in-sample test.

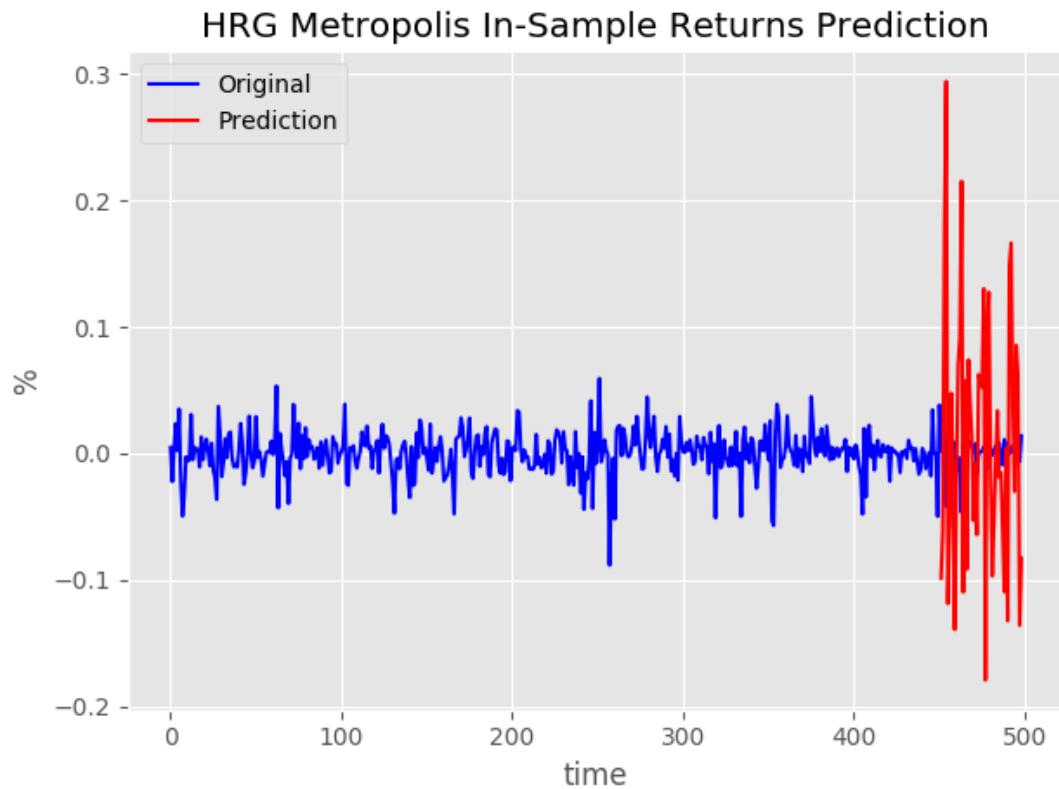


FIGURE A.149: HRG Metropolis-Hastings in-sample prediction

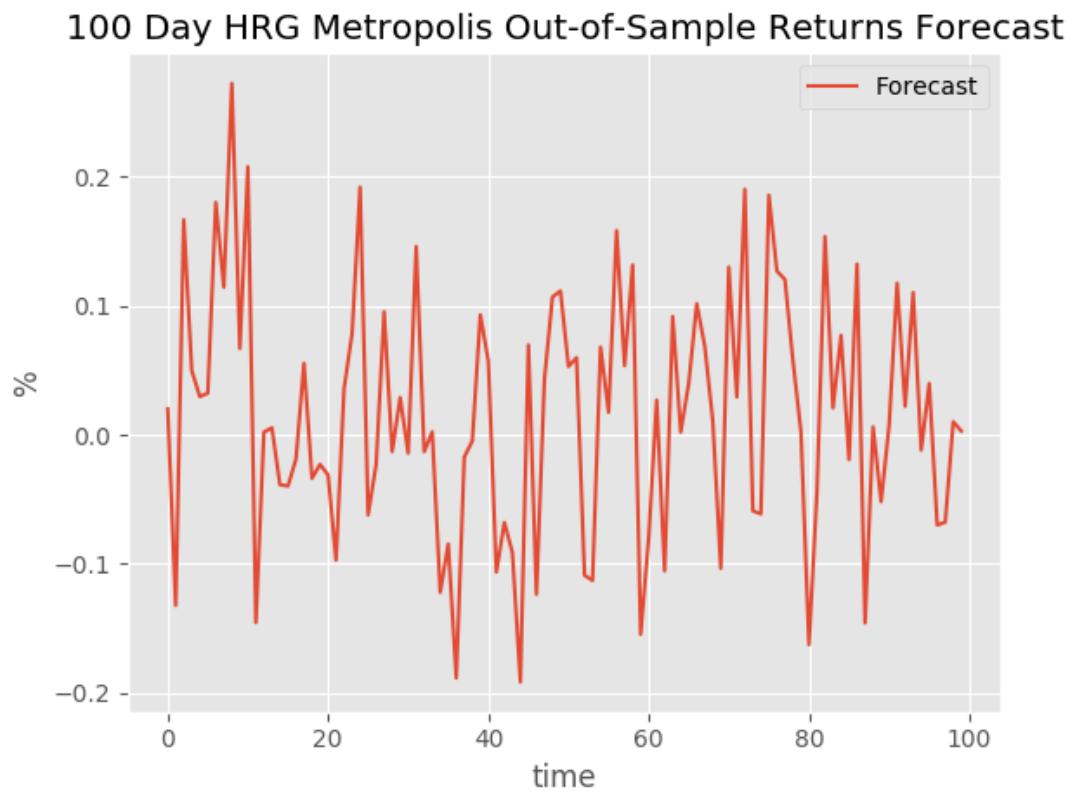


FIGURE A.150: 100 day HRG Metropolis-Hastings out-of-sample forecast

HL scored sharpe ratios of 0.439 for the original returns, and -2.283 for the predicted returns in the in-sample test.

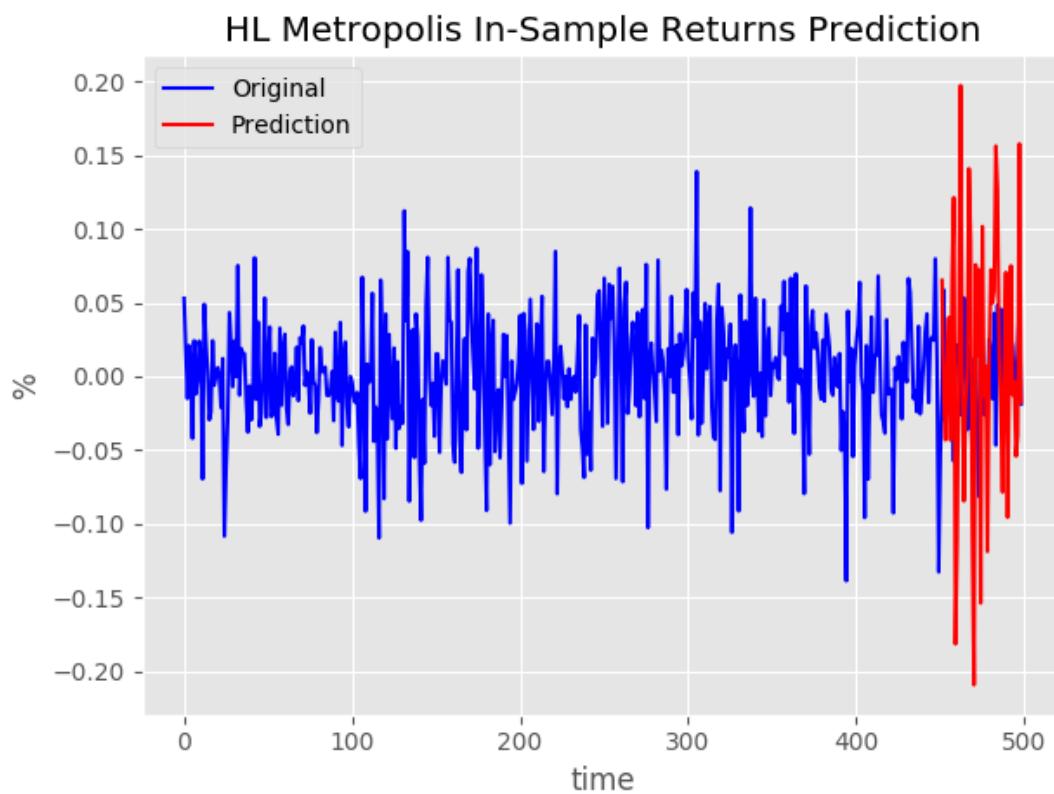


FIGURE A.151: HL Metropolis-Hastings in-sample prediction

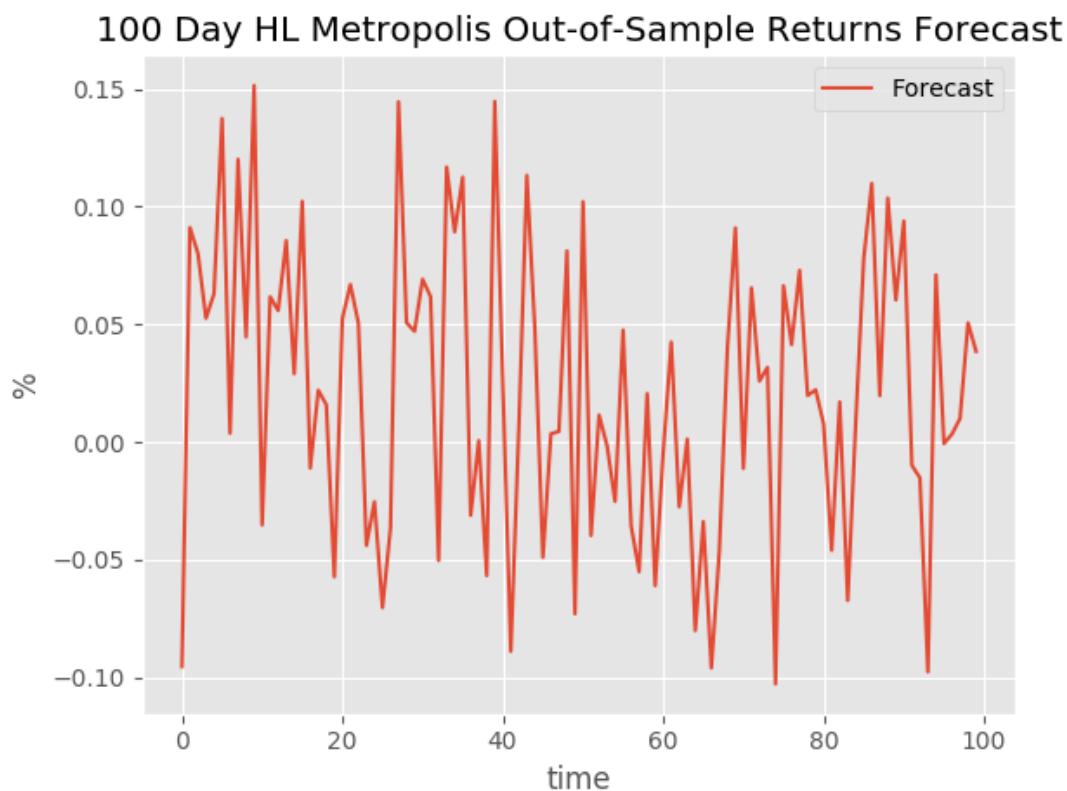


FIGURE A.152: 100 day HL Metropolis-Hastings out-of-sample forecast

A.3.2 No-U-Turn Sampler (NUTS)

MSFT scored sharpe ratios of 2.499 for the original returns, and -1.355 for the predicted returns in the in-sample test.

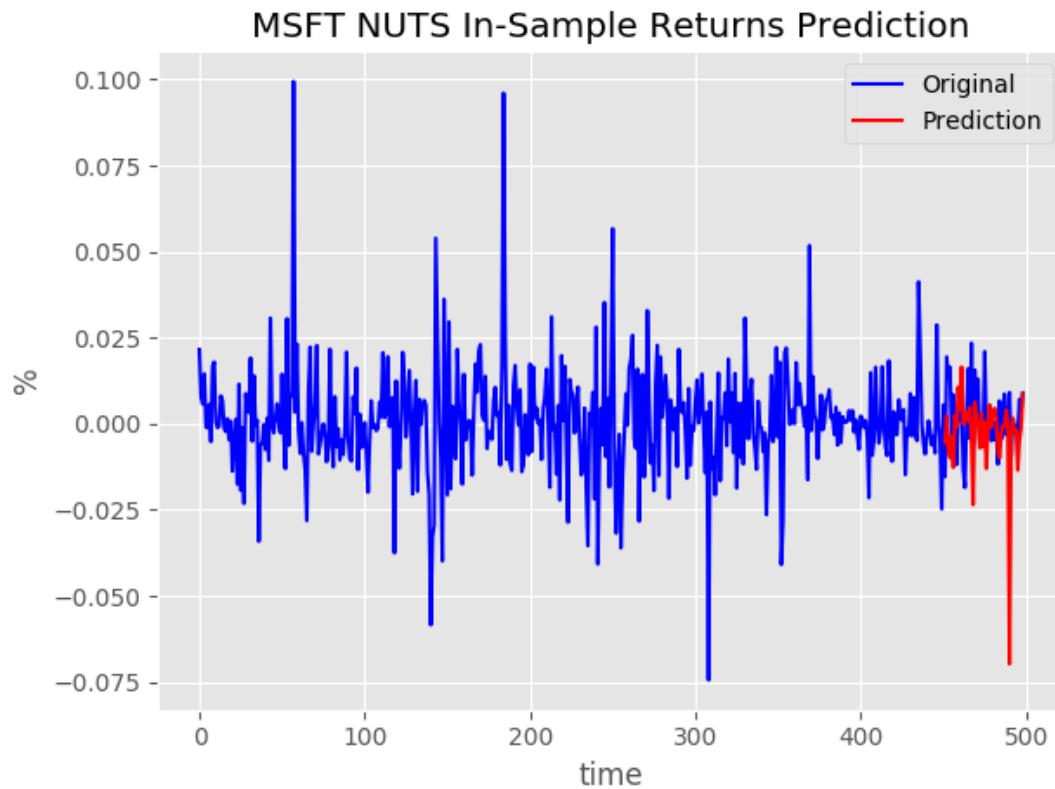


FIGURE A.153: MSFT NUTS in-sample prediction

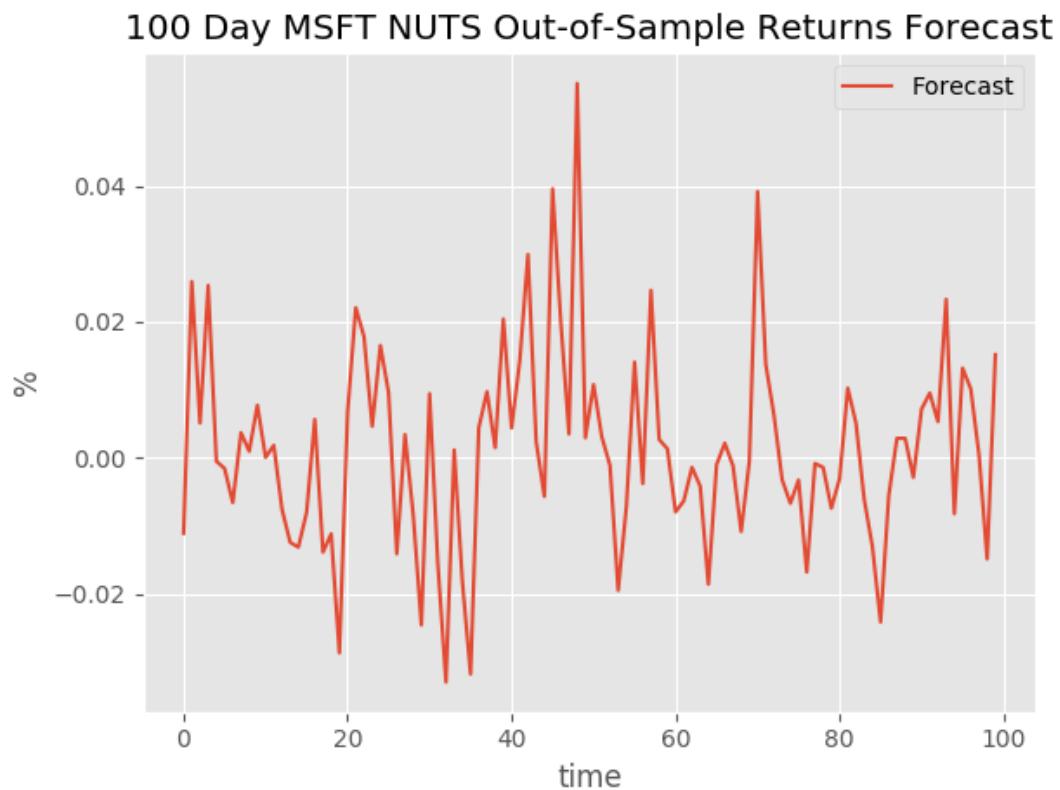


FIGURE A.154: 100 day MSFT NUTS out-of-sample forecast

CDE scored sharpe ratios of 2.028 for the original returns, and 3.066 for the predicted returns in the in-sample test.

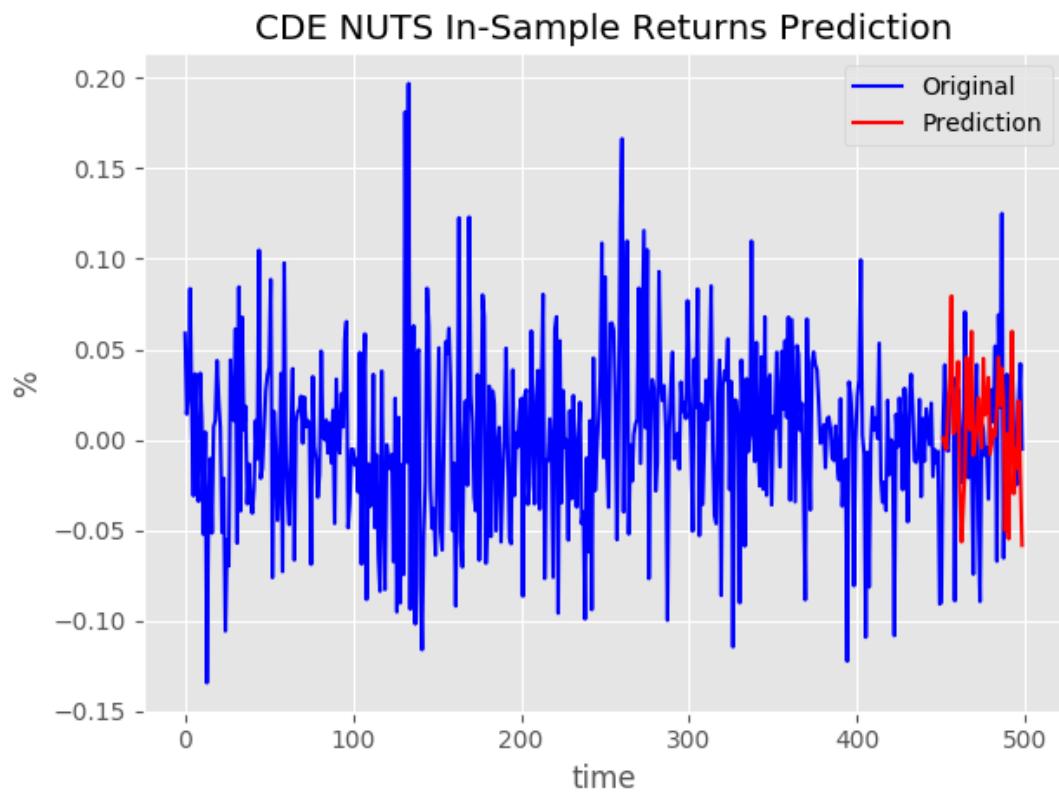


FIGURE A.155: CDE NUTS in-sample prediction

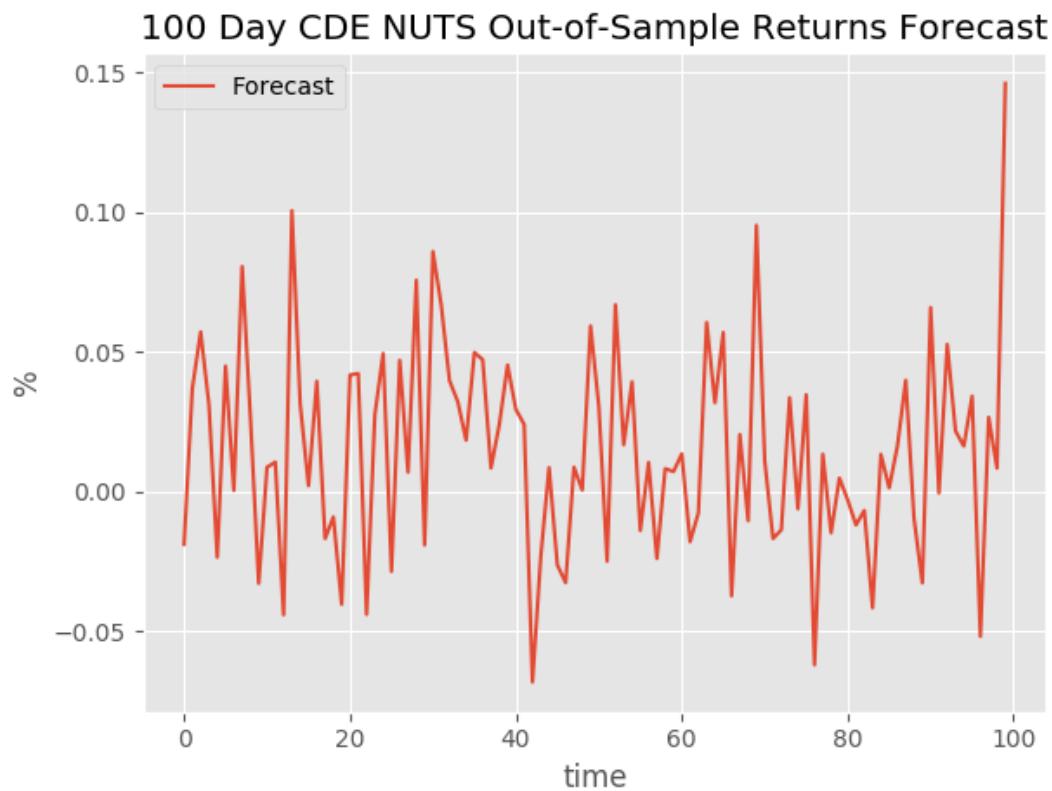


FIGURE A.156: 100 day CDE NUTS out-of-sample forecast

NAV B scored sharpe ratios of -2.800 for the original returns, and -2.914 for the predicted returns in the in-sample test.

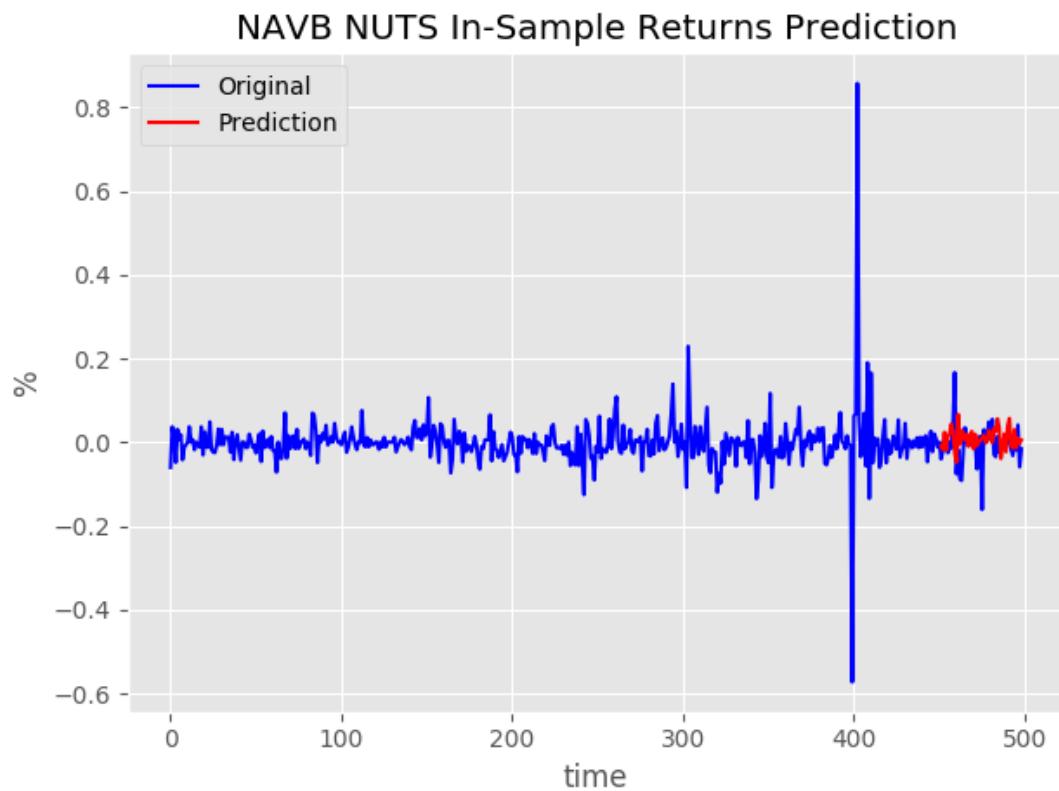


FIGURE A.157: NAVB NUTS in-sample prediction

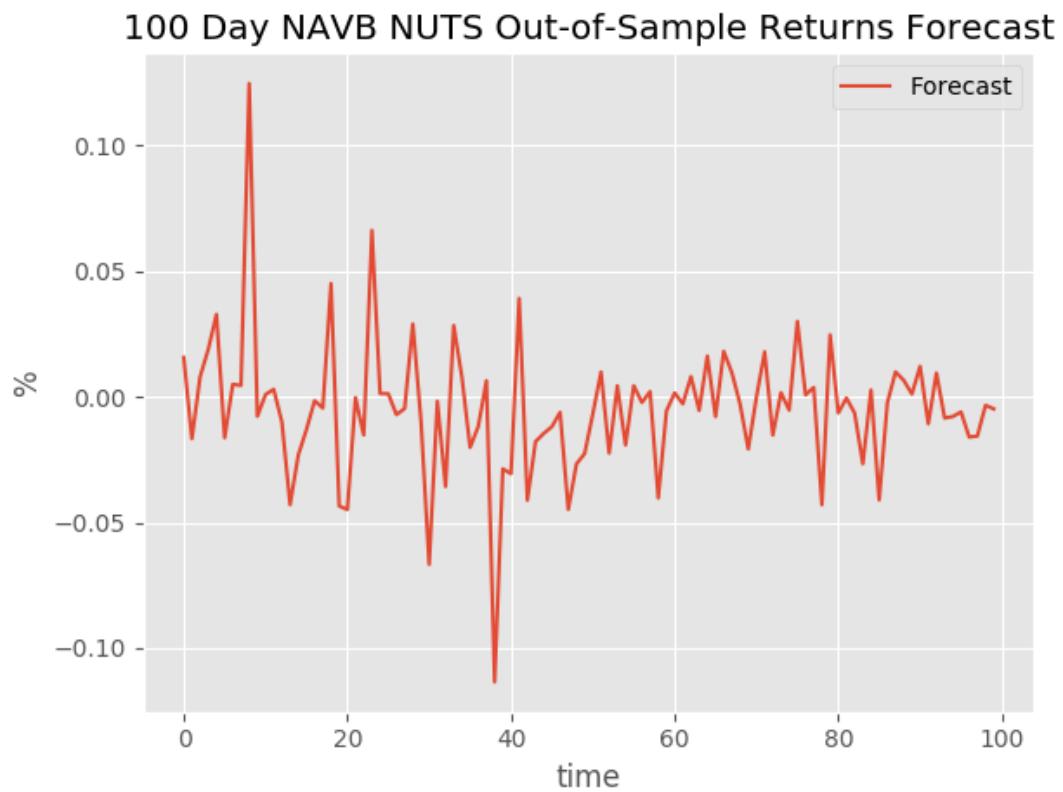


FIGURE A.158: 100 day NAVB NUTS out-of-sample forecast

HRG scored sharpe ratios of 1.362 for the original returns, and -1.812 for the predicted returns in the in-sample test.

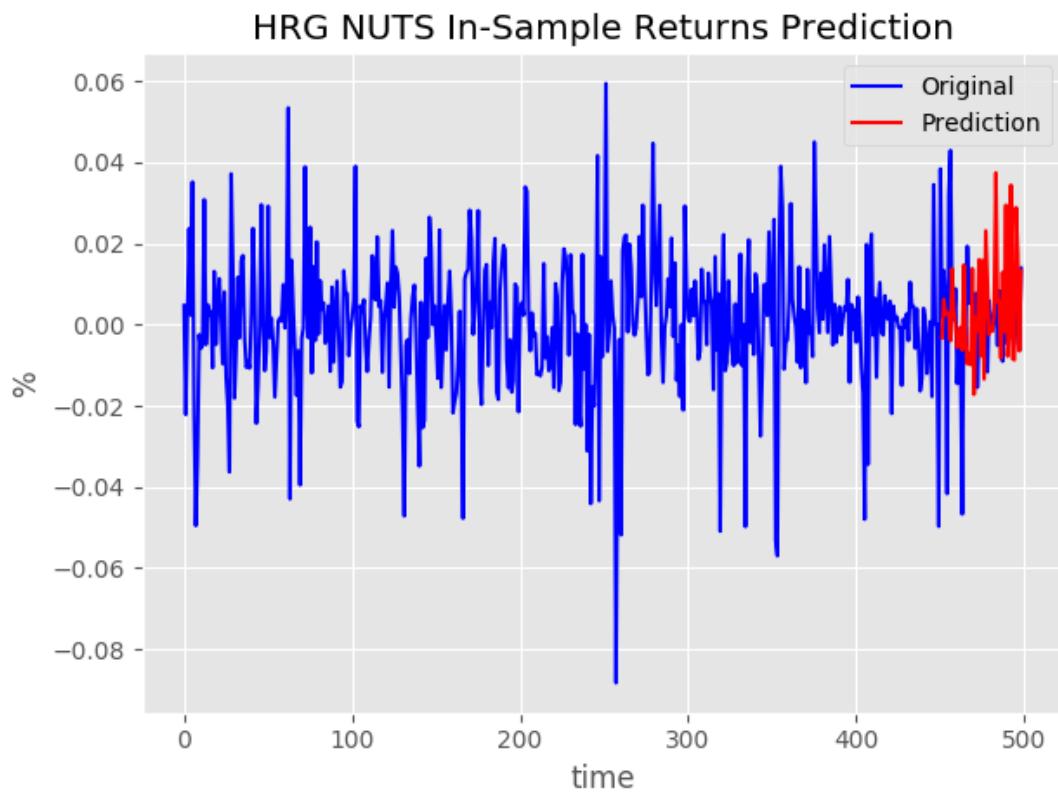


FIGURE A.159: HRG NUTS in-sample prediction

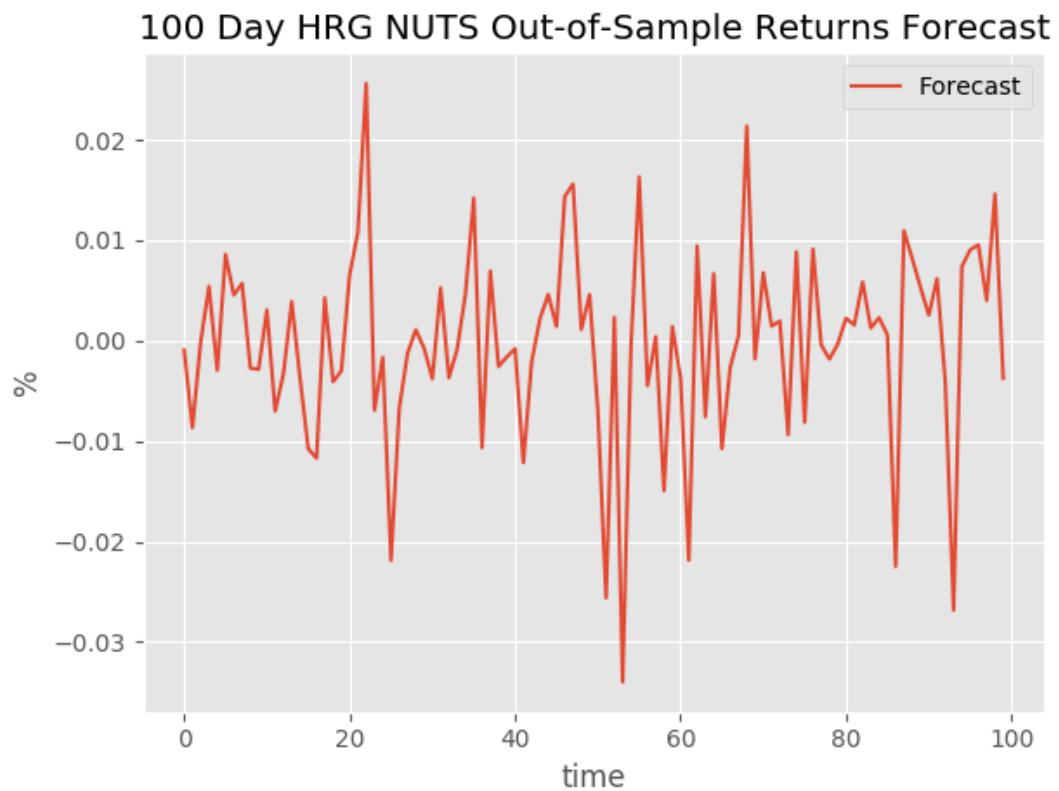


FIGURE A.160: 100 day HRG NUTS out-of-sample forecast

Bibliography

- [1] Harry Markowitz. Portfolio selection. *American Finance Association*, 1952.
- [2] Peter J. Brockwell and Richard A. Davis. *Introduction to Time Series and Forecasting*. Springer, 2016.
- [3] Tobias Moskowitz, Yao Hua Ooi, and Lasse H. Pedersen. Time series momentum. *Chicago Booth Research*, September 2011.
- [4] S. Radha and M. Thenmozhi. Forecasting short term interest rates using arma, arma-garch and arma-egarch models. *Indian Institute of Capital Markets 9th Capital Markets Conference Paper*, January 2006.
- [5] Natalia Abrosimova, Gishan Dissanaike, and Dirk Linowski. Testing weak-form efficiency of the russian stock market. In *EFA 2002 Berlin Meetings*, 2002.
- [6] Ali F. Darrat and Maosen Zhong. On testing the random walk hypothesis: A model-comparison approach, 2001.
- [7] Thomas H. Cormen and Chales E. Leiserson. *Introduction to Algorithms*. The MIT Press, 2009.
- [8] Tom M. Mitchell. *Machine Learning*. McGraw Hill, 1997.
- [9] Punniyamoorthy Murugesan and Jose Joy Thoppan. Detection of stock price manipulation using discriminant analysis, June 2012.
- [10] Manish Kumar and M. Thenmozhi. Forecasting stock index movement: A comparison of support vector machines and random forest, June 2016.
- [11] Zura Kakushadze. Mean-reversion and optimization mean-reversion and optimization. *Journal of Asset Management*, 16(1):14–40, 2015.
- [12] Germán G. Creamer and Yoav Freund. Automated trading with boosting and expert weighting. *Quantitative Finance*, Vol. 4(No. 10):pp. 401–420, April 2010.
- [13] Andrew Gelman, John B. Carlin, Hal S. Stern, David B. Dunson, Aki Vehtari, and Donald B. Rubin. *Bayesian Data Analysis*. CRC Press, 2014.

- [14] David Blanchett, Michael S. Finke, and Wade D. Pfau. Asset valuations and safe portfolio withdrawal rates, 2013.
- [15] Matthew D. Hoffman and Andrew Gelman. The no-u-turn sampler: Adaptively setting path lengths in hamiltonian monte carlo. *Journal of Machine Learning Research*, 15, 2014.
- [16] Quandl. Wiki eod stock prices, 2017.